

GNU make

GNU make 官网手册: <http://www.gnu.org/software/make/manual/make.html>

Short Table of Contents

- 1 Overview of `make`
- 2 An Introduction to Makefiles
- 3 Writing Makefiles
- 4 Writing Rules
- 5 Writing Recipes in Rules
- 6 How to Use Variables
- 7 Conditional Parts of Makefiles
- 8 Functions for Transforming Text
- 9 How to Run `make`
- 10 Using Implicit Rules
- 11 Using `make` to Update Archive Files
- 12 Extending GNU `make`
- 13 Integrating GNU `make`
- 14 Features of GNU `make`
- 15 Incompatibilities and Missing Features
- 16 Makefile Conventions
- Appendix A Quick Reference
- Appendix B Errors Generated by Make
- Appendix C Complex Makefile Example
- Appendix D GNU Free Documentation License
- Index of Concepts
- Index of Functions, Variables, & Directives

Table of Contents

- 1 Overview of `make`
 - 1.1 How to Read This Manual
 - 1.2 Problems and Bugs
- 2 An Introduction to Makefiles
 - 2.1 What a Rule Looks Like
 - 2.2 A Simple Makefile
 - 2.3 How `make` Processes a Makefile
 - 2.4 Variables Make Makefiles Simpler
 - 2.5 Letting `make` Deduce the Recipes
 - 2.6 Another Style of Makefile
 - 2.7 Rules for Cleaning the Directory
- 3 Writing Makefiles
 - 3.1 What Makefiles Contain
 - 3.1.1 Splitting Long Lines
 - 3.2 What Name to Give Your Makefile
 - 3.3 Including Other Makefiles
 - 3.4 The Variable `MAKEFILES`
 - 3.5 How Makefiles Are Remade

3.6 Overriding Part of Another Makefile

3.7 How `make` Reads a Makefile

3.8 How Makefiles Are Parsed

3.9 Secondary Expansion

4 Writing Rules

4.1 Rule Example

4.2 Rule Syntax

4.3 Types of Prerequisites

4.4 Using Wildcard Characters in File Names

4.4.1 Wildcard Examples

4.4.2 Pitfalls of Using Wildcards

4.4.3 The Function `wildcard`

4.5 Searching Directories for Prerequisites

4.5.1 `VPATH`: Search Path for All Prerequisites

4.5.2 The `vpath` Directive

4.5.3 How Directory Searches are Performed

4.5.4 Writing Recipes with Directory Search

4.5.5 Directory Search and Implicit Rules

4.5.6 Directory Search for Link Libraries

4.6 Phony Targets

4.7 Rules without Recipes or Prerequisites

4.8 Empty Target Files to Record Events

4.9 Special Built-in Target Names

4.10 Multiple Targets in a Rule

4.11 Multiple Rules for One Target

4.12 Static Pattern Rules

4.12.1 Syntax of Static Pattern Rules

4.12.2 Static Pattern Rules versus Implicit Rules

4.13 Double-Colon Rules

4.14 Generating Prerequisites Automatically

5 Writing Recipes in Rules

5.1 Recipe Syntax

5.1.1 Splitting Recipe Lines

5.1.2 Using Variables in Recipes

5.2 Recipe Echoing

5.3 Recipe Execution

5.3.1 Using One Shell

5.3.2 Choosing the Shell

5.4 Parallel Execution

5.4.1 Output During Parallel Execution

5.4.2 Input During Parallel Execution

5.5 Errors in Recipes

5.6 Interrupting or Killing `make`

5.7 Recursive Use of `make`

5.7.1 How the `MAKE` Variable Works

5.7.2 Communicating Variables to a Sub-`make`

5.7.3 Communicating Options to a Sub-`make`

5.7.4 The `'--print-directory'` Option

5.8 Defining Canned Recipes

5.9 Using Empty Recipes

6 How to Use Variables

6.1 Basics of Variable References

- 6.2 The Two Flavors of Variables
- 6.3 Advanced Features for Reference to Variables
 - 6.3.1 Substitution References
 - 6.3.2 Computed Variable Names
- 6.4 How Variables Get Their Values
- 6.5 Setting Variables
- 6.6 Appending More Text to Variables
- 6.7 The `override` Directive
- 6.8 Defining Multi-Line Variables
- 6.9 undefining Variables
- 6.10 Variables from the Environment
- 6.11 Target-specific Variable Values
- 6.12 Pattern-specific Variable Values
- 6.13 Suppressing Inheritance
- 6.14 Other Special Variables

7 Conditional Parts of Makefiles

- 7.1 Example of a Conditional
- 7.2 Syntax of Conditionals
- 7.3 Conditionals that Test Flags

8 Functions for Transforming Text

- 8.1 Function Call Syntax
- 8.2 Functions for String Substitution and Analysis
- 8.3 Functions for File Names
- 8.4 Functions for Conditionals
- 8.5 The `foreach` Function
- 8.6 The `file` Function
- 8.7 The `call` Function
- 8.8 The `value` Function
- 8.9 The `eval` Function
- 8.10 The `origin` Function
- 8.11 The `flavor` Function
- 8.12 Functions That Control Make
- 8.13 The `shell` Function
- 8.14 The `guile` Function

9 How to Run `make`

- 9.1 Arguments to Specify the Makefile
- 9.2 Arguments to Specify the Goals
- 9.3 Instead of Executing Recipes
- 9.4 Avoiding Recompilation of Some Files
- 9.5 Overriding Variables
- 9.6 Testing the Compilation of a Program
- 9.7 Summary of Options

10 Using Implicit Rules

- 10.1 Using Implicit Rules
- 10.2 Catalogue of Built-In Rules
- 10.3 Variables Used by Implicit Rules
- 10.4 Chains of Implicit Rules
- 10.5 Defining and Redefining Pattern Rules
 - 10.5.1 Introduction to Pattern Rules
 - 10.5.2 Pattern Rule Examples
 - 10.5.3 Automatic Variables
 - 10.5.4 How Patterns Match

10.5.5 Match-Anything Pattern Rules

10.5.6 Canceling Implicit Rules

10.6 Defining Last-Resort Default Rules

10.7 Old-Fashioned Suffix Rules

10.8 Implicit Rule Search Algorithm

11 Using `make` to Update Archive Files

11.1 Archive Members as Targets

11.2 Implicit Rule for Archive Member Targets

11.2.1 Updating Archive Symbol Directories

11.3 Dangers When Using Archives

11.4 Suffix Rules for Archive Files

12 Extending GNU `make`

12.1 GNU Guile Integration

12.1.1 Conversion of Guile Types

12.1.2 Interfaces from Guile to `make`

12.1.3 Example Using Guile in `make`

12.2 Loading Dynamic Objects

12.2.1 The `load` Directive

12.2.2 How Loaded Objects Are Remade

12.2.3 Loaded Object Interface

12.2.4 Example Loaded Object

13 Integrating GNU `make`

13.1 Sharing Job Slots with GNU `make`

13.1.1 POSIX Jobserver Interaction

13.1.2 Windows Jobserver Interaction

13.2 Synchronized Terminal Output

14 Features of GNU `make`

15 Incompatibilities and Missing Features

16 Makefile Conventions

16.1 General Conventions for Makefiles

16.2 Utilities in Makefiles

16.3 Variables for Specifying Commands

16.4 `DESTDIR`: Support for Staged Installs

16.5 Variables for Installation Directories

16.6 Standard Targets for Users

16.7 Install Command Categories

Appendix A Quick Reference

Appendix B Errors Generated by Make

Appendix C Complex Makefile Example

Appendix D GNU Free Documentation License

Index of Concepts

Index of Functions, Variables, & Directives

Next: Overview, Previous: (dir), Up: (dir) [Contents][Index]

GNU `make`

This file documents the GNU `make` utility, which determines automatically which pieces of a large program need to be recompiled, and issues the commands to recompile them.

Copyright © 1988, 1989, 1990, 1991, 1992, 1993, 1994, 1995, 1996, 1997, 1998, 1999, 2000, 2002, 2003, 2004, 2005, 2006, 2007, 2008, 2009, 2010, 2011, 2012, 2013, 2014, 2015, 2016, 2017, 2018, 2019, 2020 Free Software Foundation, Inc.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.3 or any later version published by the Free Software Foundation; with no Invariant Sections, with the Front-Cover Texts being “A GNU Manual,” and with the Back-Cover Texts as in (a) below. A copy of the license is included in the section entitled “GNU Free Documentation License.”

(a) The FSF’s Back-Cover Text is: “You have the freedom to copy and modify this GNU manual. Buying copies from the FSF supports it in developing GNU and promoting software freedom.”

• Overview	Overview of <code>make</code> .
• Introduction	An introduction to <code>make</code> .
• Makefiles	Makefiles tell <code>make</code> what to do.
• Rules	Rules describe when a file must be remade.
• Recipes	Recipes say how to remake a file.
• Using Variables	You can use variables to avoid repetition.
• Conditionals	Use or ignore parts of the makefile based on the values of variables.
• Functions	Many powerful ways to manipulate text.
• Invoking <code>make</code>	How to invoke <code>make</code> on the command line.
• Implicit Rules	Use implicit rules to treat many files alike, based on their file names.
• Archives	How <code>make</code> can update library archives.
• Extending <code>make</code>	Using extensions to <code>make</code> .
• Integrating <code>make</code>	Integrating <code>make</code> with other tools.
• Features	Features GNU <code>make</code> has over other <code>makes</code> .
• Missing	What GNU <code>make</code> lacks from other <code>makes</code> .
• Makefile Conventions	Conventions for writing makefiles for GNU programs.
• Quick Reference	A quick reference for experienced users.
• Error Messages	A list of common errors generated by <code>make</code> .
• Complex Makefile	A real example of a straightforward, but nontrivial, makefile.
• GNU Free Documentation License	License for copying this manual.
• Concept Index	Index of Concepts.
• Name Index	Index of Functions, Variables, & Directives.

— The Detailed Node Listing —

Overview of `make`

• Preparing	Preparing and running <code>make</code> .
• Reading	On reading this text.
• Bugs	Problems and bugs.

An Introduction to Makefiles

- [Rule Introduction](#) What a rule looks like.
- [Simple Makefile](#) A simple makefile.
- [How Make Works](#) How `make` processes this makefile.
- [Variables Simplify](#) Variables make makefiles simpler.
- [make Deduces](#) Letting `make` deduce the recipes.
- [Combine By Prerequisite](#) Another style of makefile.
- [Cleanup](#) Rules for cleaning the directory.

Writing Makefiles

- [Makefile Contents](#) What makefiles contain.
- [Makefile Names](#) How to name your makefile.
- [Include](#) How one makefile can use another makefile.
- [MAKEFILES Variable](#) The environment can specify extra makefiles.
- [Remaking Makefiles](#) How makefiles get remade.
- [Overriding Makefiles](#) How to override part of one makefile with another makefile.
- [Reading Makefiles](#) How makefiles are read in.
- [Parsing Makefiles](#) How makefiles are parsed.
- [Secondary Expansion](#) How and when secondary expansion is performed.

What Makefiles Contain

- [Splitting Lines](#) Splitting long lines in makefiles

Writing Rules

- [Rule Example](#) An example explained.
- [Rule Syntax](#) General syntax explained.
- [Prerequisite Types](#) There are two types of prerequisites.
- [Wildcards](#) Using wildcard characters such as `‘*’`.
- [Directory Search](#) Searching other directories for source files.
- [Phony Targets](#) Using a target that is not a real file’s name.
- [Force Targets](#) You can use a target without a recipe or prerequisites to mark other targets as phony.
- [Empty Targets](#) When only the date matters and the files are empty.
- [Special Targets](#) Targets with special built-in meanings.
- [Multiple Targets](#) When to make use of several targets in a rule.
- [Multiple Rules](#) How to use several rules with the same target.
- [Static Pattern](#) Static pattern rules apply to multiple targets and can vary the prerequisites according to the target name.

• Double-Colon	How to use a special kind of rule to allow several independent rules for one target.
• Automatic Prerequisites	How to automatically generate rules giving prerequisites from source files themselves.

Using Wildcard Characters in File Names

• Wildcard Examples	Several examples.
• Wildcard Pitfall	Problems to avoid.
• Wildcard Function	How to cause wildcard expansion where it does not normally take place.

Searching Directories for Prerequisites

• General Search	Specifying a search path that applies to every prerequisite.
• Selective Search	Specifying a search path for a specified class of names.
• Search Algorithm	When and how search paths are applied.
• Recipes/Search	How to write recipes that work together with search paths.
• Implicit/Search	How search paths affect implicit rules.
• Libraries/Search	Directory search for link libraries.

Static Pattern Rules

• Static Usage	The syntax of static pattern rules.
• Static versus Implicit	When are they better than implicit rules?

Writing Recipes in Rules

• Recipe Syntax	Recipe syntax features and pitfalls.
• Echoing	How to control when recipes are echoed.
• Execution	How recipes are executed.
• Parallel	How recipes can be executed in parallel.
• Errors	What happens after a recipe execution error.
• Interrupts	What happens when a recipe is interrupted.
• Recursion	Invoking <code>make</code> from makefiles.
• Canned Recipes	Defining canned recipes.
• Empty Recipes	Defining useful, do-nothing recipes.

Recipe Syntax

• Splitting Recipe Lines	Breaking long recipe lines for readability.
• Variables in Recipes	Using <code>make</code> variables in recipes.

Recipe Execution

- [One Shell](#) One shell for all lines in a recipe.
- [Choosing the Shell](#) How `make` chooses the shell used to run recipes.

Parallel Execution

- [Parallel Output](#) Handling output during parallel execution
- [Parallel Input](#) Handling input during parallel execution

Recursive Use of `make`

- [MAKE Variable](#) The special effects of using `'$(MAKE)'`.
- [Variables/Recursion](#) How to communicate variables to a `sub-make`.
- [Options/Recursion](#) How to communicate options to a `sub-make`.
- [-w Option](#) How the `'-w'` or `'--print-directory'` option helps debug use of recursive `make` commands.

How to Use Variables

- [Reference](#) How to use the value of a variable.
- [Flavors](#) Variables come in two flavors.
- [Advanced](#) Advanced features for referencing a variable.
- [Values](#) All the ways variables get their values.
- [Setting](#) How to set a variable in the makefile.
- [Appending](#) How to append more text to the old value of a variable.
- [Override Directive](#) How to set a variable in the makefile even if the user has set it with a command argument.
- [Multi-Line](#) An alternate way to set a variable to a multi-line string.
- [Undefine Directive](#) How to undefine a variable so that it appears as if it was never set.
- [Environment](#) Variable values can come from the environment.
- [Target-specific](#) Variable values can be defined on a per-target basis.
- [Pattern-specific](#) Target-specific variable values can be applied to a group of targets that match a pattern.
- [Suppressing Inheritance](#) Suppress inheritance of variables.
- [Special Variables](#) Variables with special meaning or behavior.

Advanced Features for Reference to Variables

- [Substitution Refs](#) Referencing a variable with substitutions on the value.
- [Computed Names](#) Computing the name of the variable to refer to.

Conditional Parts of Makefiles

- [Conditional](#) Example of a conditional

Example

- Conditional Syntax
- Testing Flags

The syntax of conditionals.

Conditionals that test flags.

Functions for Transforming Text

- Syntax of Functions
- Text Functions
- File Name Functions
- Conditional Functions
- Foreach Function
- File Function
- Call Function
- Value Function
- Eval Function
- Origin Function
- Flavor Function
- Make Control Functions
- Shell Function
- Guile Function

How to write a function call.

General-purpose text manipulation functions.

Functions for manipulating file names.

Functions that implement conditions.

Repeat some text with controlled variation.

Write text to a file.

Expand a user-defined function.

Return the un-expanded value of a variable.

Evaluate the arguments as makefile syntax.

Find where a variable got its value.

Find out the flavor of a variable.

Functions that control how make runs.

Substitute the output of a shell command.

Use GNU Guile embedded scripting language.

How to Run make

- Makefile Arguments
- Goals
- Instead of Execution
- Avoiding Compilation
- Overriding
- Testing
- Options Summary

How to specify which makefile to use.

How to use goal arguments to specify which parts of the makefile to use.

How to use mode flags to specify what kind of thing to do with the recipes in the makefile other than simply execute them.

How to avoid recompiling certain files.

How to override a variable to specify an alternate compiler and other things.

How to proceed past some errors, to test compilation.

Summary of Options

Using Implicit Rules

- Using Implicit
- Catalogue of Rules
- Implicit Variables
- Chained Rules
- Pattern Rules
- Last Resort

How to use an existing implicit rule to get the recipes for updating a file.

A list of built-in rules.

How to change what predefined rules do.

How to use a chain of implicit rules.

How to define new implicit rules.

How to define a recipe for rules which cannot find any.

- [Suffix Rules](#) The old-fashioned style of implicit rule.
- [Implicit Rule Search](#) The precise algorithm for applying implicit rules.

Defining and Redefining Pattern Rules

- [Pattern Intro](#) An introduction to pattern rules.
- [Pattern Examples](#) Examples of pattern rules.
- [Automatic Variables](#) How to use automatic variables in the recipe of implicit rules.
- [Pattern Match](#) How patterns match.
- [Match-Anything Rules](#) Precautions you should take prior to defining rules that can match any target file whatever.
- [Canceling Rules](#) How to override or cancel built-in rules.

Using `make` to Update Archive Files

- [Archive Members](#) Archive members as targets.
- [Archive Update](#) The implicit rule for archive member targets.
- [Archive Pitfalls](#) Dangers to watch out for when using archives.
- [Archive Suffix Rules](#) You can write a special kind of suffix rule for updating archives.

Implicit Rule for Archive Member Targets

- [Archive Symbols](#) How to update archive symbol directories.

Extending GNU `make`

- [Guile Integration](#) Using Guile as an embedded scripting language.
- [Loading Objects](#) Loading dynamic objects as extensions.

GNU Guile Integration

- [Guile Types](#) Converting Guile types to `make` strings.
- [Guile Interface](#) Invoking `make` functions from Guile.
- [Guile Example](#) Example using Guile in `make`.

Loading Dynamic Objects

- [load Directive](#) Loading dynamic objects as extensions.
- [Remaking Loaded Objects](#) How loaded objects get remade.
- [Loaded Object API](#) Programmatic interface for loaded objects.
- [Loaded Object Example](#) Example of a loaded object

- [Job Slots](#) Share job slots with GNU `make`.
- [Terminal Output](#) Control output to terminals.

Sharing Job Slots with GNU `make`

- [POSIX Jobserver](#) Using the jobserver on POSIX systems.
- [Windows Jobserver](#) Using the jobserver on Windows systems.

Next: [Introduction](#), Previous: [Top](#), Up: [Top](#) [[Contents](#)][[Index](#)]

1 Overview of `make`

The `make` utility automatically determines which pieces of a large program need to be recompiled, and issues commands to recompile them. This manual describes GNU `make`, which was implemented by Richard Stallman and Roland McGrath. Development since Version 3.76 has been handled by Paul D. Smith.

GNU `make` conforms to section 6.2 of IEEE Standard 1003.2-1992 (POSIX.2).

Our examples show C programs, since they are most common, but you can use `make` with any programming language whose compiler can be run with a shell command. Indeed, `make` is not limited to programs. You can use it to describe any task where some files must be updated automatically from others whenever the others change.

- [Preparing](#) Preparing and running `make`.
- [Reading](#) On reading this text.
- [Bugs](#) Problems and bugs.

Next: [Reading](#), Previous: [Overview](#), Up: [Overview](#) [[Contents](#)][[Index](#)]

Preparing

Preparing and Running Make

To prepare to use `make`, you must write a file called the makefile that describes the relationships among files in your program and provides commands for updating each file. In a program, typically, the executable file is updated from object files, which are in turn made by compiling source files.

Once a suitable makefile exists, each time you change some source files, this simple shell command:

```
make
```

suffices to perform all necessary recompilations. The `make` program uses the makefile data base and the last-modification times of the files to decide which of the files need to be updated. For each of those files, it issues the recipes recorded in the data base.

You can provide command line arguments to `make` to control which files should be recompiled, or how. See [How to Run `make`](#).

1.1 How to Read This Manual

If you are new to `make`, or are looking for a general introduction, read the first few sections of each chapter, skipping the later sections. In each chapter, the first few sections contain introductory or general information and the later sections contain specialized or technical information. The exception is the second chapter, [An Introduction to Makefiles](#), all of which is introductory.

If you are familiar with other `make` programs, see [Features of GNU `make`](#), which lists the enhancements GNU `make` has, and [Incompatibilities and Missing Features](#), which explains the few things GNU `make` lacks that others have.

For a quick summary, see [Options Summary](#), [Quick Reference](#), and [Special Targets](#).

Previous: [Reading](#), Up: [Overview](#) [[Contents](#)][[Index](#)]

1.2 Problems and Bugs

If you have problems with GNU `make` or think you've found a bug, please report it to the developers; we cannot promise to do anything but we might well want to fix it.

Before reporting a bug, make sure you've actually found a real bug. Carefully reread the documentation and see if it really says you can do what you're trying to do. If it's not clear whether you should be able to do something or not, report that too; it's a bug in the documentation!

Before reporting a bug or trying to fix it yourself, try to isolate it to the smallest possible makefile that reproduces the problem. Then send us the makefile and the exact results `make` gave you, including any error or warning messages. Please don't paraphrase these messages: it's best to cut and paste them into your report. When generating this small makefile, be sure to not use any non-free or unusual tools in your recipes: you can almost always emulate what such a tool would do with simple shell commands. Finally, be sure to explain what you expected to occur; this will help us decide whether the problem was really in the documentation.

Once you have a precise problem you can report it in one of two ways. Either send electronic mail to:

bug-make@gnu.org

or use our Web-based project management tool, at:

<http://savannah.gnu.org/projects/make/>

In addition to the information above, please be careful to include the version number of `make` you are using. You can get this information with the command `'make --version'`. Be sure also to include the type of machine and operating system you are using. One way to obtain this information is by looking at the final lines of output from the command `'make --help'`.

Next: [Makefiles](#), Previous: [Overview](#), Up: [Top](#) [[Contents](#)][[Index](#)]

2 An Introduction to Makefiles

You need a file called a makefile to tell `make` what to do. Most often, the makefile tells `make` how to compile and link a program.

In this chapter, we will discuss a simple makefile that describes how to compile and link a text editor which consists of eight C source files and three header files. The makefile can also tell

`make` how to run miscellaneous commands when explicitly asked (for example, to remove certain files as a clean-up operation). To see a more complex example of a makefile, see [Complex Makefile](#).

When `make` recompiles the editor, each changed C source file must be recompiled. If a header file has changed, each C source file that includes the header file must be recompiled to be safe. Each compilation produces an object file corresponding to the source file. Finally, if any source file has been recompiled, all the object files, whether newly made or saved from previous compilations, must be linked together to produce the new executable editor.

- [Rule Introduction](#) What a rule looks like.
- [Simple Makefile](#) A simple makefile.
- [How Make Works](#) How `make` processes this makefile.
- [Variables Simplify](#) Variables make makefiles simpler.
- [make Deduces](#) Letting `make` deduce the recipes.
- [Combine By Prerequisite](#) Another style of makefile.
- [Cleanup](#) Rules for cleaning the directory.

Next: [Simple Makefile](#), Previous: [Introduction](#), Up: [Introduction](#) [\[Contents\]](#)[\[Index\]](#)

2.1 What a Rule Looks Like

A simple makefile consists of “rules” with the following shape:

```
target ... : prerequisites ...
    recipe
    ...
    ...
```

A target is usually the name of a file that is generated by a program; examples of targets are executable or object files. A target can also be the name of an action to carry out, such as `'clean'` (see [Phony Targets](#)).

A prerequisite is a file that is used as input to create the target. A target often depends on several files.

A recipe is an action that `make` carries out. A recipe may have more than one command, either on the same line or each on its own line. **Please note:** you need to put a tab character at the beginning of every recipe line! This is an obscurity that catches the unwary. If you prefer to prefix your recipes with a character other than tab, you can set the `.RECIPEPREFIX` variable to an alternate character (see [Special Variables](#)).

Usually a recipe is in a rule with prerequisites and serves to create a target file if any of the prerequisites change. However, the rule that specifies a recipe for the target need not have prerequisites. For example, the rule containing the delete command associated with the target `'clean'` does not have prerequisites.

A rule, then, explains how and when to remake certain files which are the targets of the particular rule. `make` carries out the recipe on the prerequisites to create or update the target. A rule can also explain how and when to carry out an action. See [Writing Rules](#).

A makefile may contain other text besides rules, but a simple makefile need only contain rules. Rules may look somewhat more complicated than shown in this template, but all fit the pattern more or less.

Next: [How Make Works](#), Previous: [Rule Introduction](#), Up: [Introduction](#) [\[Contents\]](#)[\[Index\]](#)

2.2 A Simple Makefile

Here is a straightforward makefile that describes the way an executable file called `edit` depends on eight object files which, in turn, depend on eight C source and three header files.

In this example, all the C files include `defs.h`, but only those defining editing commands include `command.h`, and only low level files that change the editor buffer include `buffer.h`.

```
edit : main.o kbd.o command.o display.o \  
      insert.o search.o files.o utils.o  
      cc -o edit main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
main.o : main.c defs.h  
      cc -c main.c  
kbd.o : kbd.c defs.h command.h  
      cc -c kbd.c  
command.o : command.c defs.h command.h  
      cc -c command.c  
display.o : display.c defs.h buffer.h  
      cc -c display.c  
insert.o : insert.c defs.h buffer.h  
      cc -c insert.c  
search.o : search.c defs.h buffer.h  
      cc -c search.c  
files.o : files.c defs.h buffer.h command.h  
      cc -c files.c  
utils.o : utils.c defs.h  
      cc -c utils.c  
  
clean :  
      rm edit main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o
```

We split each long line into two lines using backslash/newline; this is like using one long line, but is easier to read. See [Splitting Long Lines](#).

To use this makefile to create the executable file called `edit`, type:

```
make
```

To use this makefile to delete the executable file and all the object files from the directory, type:

```
make clean
```

In the example makefile, the targets include the executable file `'edit'`, and the object files `'main.o'` and `'kbd.o'`. The prerequisites are files such as `'main.c'` and `'defs.h'`. In fact, each `'*.o'` file is both a target and a prerequisite. Recipes include `'cc -c main.c'` and `'cc -c kbd.c'`.

When a target is a file, it needs to be recompiled or relinked if any of its prerequisites change. In addition, any prerequisites that are themselves automatically generated should be updated first. In this example, `edit` depends on each of the eight object files; the object file `main.o` depends on the source file `main.c` and on the header file `defs.h`.

A recipe may follow each line that contains a target and prerequisites. These recipes say how to update the target file. A tab character (or whatever character is specified by the `.RECIPEPREFIX` variable; see [Special Variables](#)) must come at the beginning of every line in the recipe to distinguish recipes from other lines in the makefile. (Bear in mind that `make` does not know anything about how the recipes work. It is up to you to supply recipes that will update the target file properly. All `make` does is execute the recipe you have specified when the target file needs to be updated.)

The target `'clean'` is not a file, but merely the name of an action. Since you normally do not want to carry out the actions in this rule, `'clean'` is not a prerequisite of any other rule. Consequently, `make` never does anything with it unless you tell it specifically. Note that this rule not only is not a prerequisite, it also does not have any prerequisites, so the only purpose of the rule is to run the specified recipe. Targets that do not refer to files but are just actions are called phony targets. See [Phony Targets](#), for information about this kind of target. See [Errors in Recipes](#), to see how to cause `make` to ignore errors from `rm` or any other command.

2.3 How `make` Processes a Makefile

By default, `make` starts with the first target (not targets whose names start with `.'`). This is called the default goal. (Goals are the targets that `make` strives ultimately to update. You can override this behavior using the command line (see [Arguments to Specify the Goals](#)) or with the `.DEFAULT_GOAL` special variable (see [Other Special Variables](#)).

In the simple example of the previous section, the default goal is to update the executable program `edit`; therefore, we put that rule first.

Thus, when you give the command:

```
make
```

`make` reads the makefile in the current directory and begins by processing the first rule. In the example, this rule is for relinking `edit`; but before `make` can fully process this rule, it must process the rules for the files that `edit` depends on, which in this case are the object files. Each of these files is processed according to its own rule. These rules say to update each `.'o'` file by compiling its source file. The recompilation must be done if the source file, or any of the header files named as prerequisites, is more recent than the object file, or if the object file does not exist.

The other rules are processed because their targets appear as prerequisites of the goal. If some other rule is not depended on by the goal (or anything it depends on, etc.), that rule is not processed, unless you tell `make` to do so (with a command such as `make clean`).

Before recompiling an object file, `make` considers updating its prerequisites, the source file and header files. This makefile does not specify anything to be done for them—the `.'c'` and `.'h'` files are not the targets of any rules—so `make` does nothing for these files. But `make` would update automatically generated C programs, such as those made by Bison or Yacc, by their own rules at this time.

After recompiling whichever object files need it, `make` decides whether to relink `edit`. This must be done if the file `edit` does not exist, or if any of the object files are newer than it. If an object file was just recompiled, it is now newer than `edit`, so `edit` is relinked.

Thus, if we change the file `insert.c` and run `make`, `make` will compile that file to update `insert.o`, and then link `edit`. If we change the file `command.h` and run `make`, `make` will recompile the object files `kbd.o`, `command.o` and `files.o` and then link the file `edit`.

2.4 Variables Make Makefiles Simpler

In our example, we had to list all the object files twice in the rule for `edit` (repeated here):

```
edit : main.o kbd.o command.o display.o \  
      insert.o search.o files.o utils.o  
      cc -o edit main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o
```

Such duplication is error-prone; if a new object file is added to the system, we might add it to one list and forget the other. We can eliminate the risk and simplify the makefile by using a variable. Variables allow a text string to be defined once and substituted in multiple places later (see [How to Use Variables](#)).

It is standard practice for every makefile to have a variable named `objects`, `OBJECTS`, `objs`, `OBJS`, `obj`, or `OBJ` which is a list of all object file names. We would define such a variable `objects` with a line like this in the makefile:

```
objects = main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o
```

Then, each place we want to put a list of the object file names, we can substitute the variable's value by writing `'$(objects)'` (see [How to Use Variables](#)).

Here is how the complete simple makefile looks when you use a variable for the object files:

```
objects = main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
edit : $(objects)  
      cc -o edit $(objects)  
main.o : main.c defs.h  
      cc -c main.c  
kbd.o : kbd.c defs.h command.h  
      cc -c kbd.c  
command.o : command.c defs.h command.h  
      cc -c command.c  
display.o : display.c defs.h buffer.h  
      cc -c display.c  
insert.o : insert.c defs.h buffer.h  
      cc -c insert.c  
search.o : search.c defs.h buffer.h  
      cc -c search.c  
files.o : files.c defs.h buffer.h command.h  
      cc -c files.c  
utils.o : utils.c defs.h  
      cc -c utils.c  
clean :  
      rm edit $(objects)
```

Next: [Combine By Prerequisite](#), Previous: [Variables Simplify](#), Up: [Introduction](#) [[Contents](#)][[Index](#)]

2.5 Letting `make` Deduce the Recipes

It is not necessary to spell out the recipes for compiling the individual C source files, because `make` can figure them out: it has an implicit rule for updating a `'.o'` file from a correspondingly named `'.c'` file using a `'cc -c'` command. For example, it will use the recipe `'cc -c main.c -o main.o'` to compile `main.c` into `main.o`. We can therefore omit the recipes from the rules for the object files. See [Using Implicit Rules](#).

When a `'.c'` file is used automatically in this way, it is also automatically added to the list of prerequisites. We can therefore omit the `'.c'` files from the prerequisites, provided we omit the recipe.

Here is the entire example, with both of these changes, and a variable `objects` as suggested above:

```
objects = main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
edit : $(objects)  
      cc -o edit $(objects)  
  
main.o : defs.h  
kbd.o : defs.h command.h  
command.o : defs.h command.h  
display.o : defs.h buffer.h  
insert.o : defs.h buffer.h  
search.o : defs.h buffer.h  
files.o : defs.h buffer.h command.h  
utils.o : defs.h  
  
.PHONY : clean  
clean :  
      rm edit $(objects)
```

This is how we would write the makefile in actual practice. (The complications associated with `'clean'` are described elsewhere. See [Phony Targets](#), and [Errors in Recipes](#).)

Because implicit rules are so convenient, they are important. You will see them used frequently.

Next: [Cleanup](#), Previous: [make Deduces](#), Up: [Introduction](#) [[Contents](#)][[Index](#)]

2.6 Another Style of Makefile

When the objects of a makefile are created only by implicit rules, an alternative style of makefile is possible. In this style of makefile, you group entries by their prerequisites instead of by their targets. Here is what one looks like:

```
objects = main.o kbd.o command.o display.o \
        insert.o search.o files.o utils.o

edit : $(objects)
    cc -o edit $(objects)

$(objects) : defs.h
kbd.o command.o files.o : command.h
display.o insert.o search.o files.o : buffer.h
```

Here `defs.h` is given as a prerequisite of all the object files; `command.h` and `buffer.h` are prerequisites of the specific object files listed for them.

Whether this is better is a matter of taste: it is more compact, but some people dislike it because they find it clearer to put all the information about each target in one place.

Previous: [Combine By Prerequisite](#), Up: [Introduction](#) [[Contents](#)][[Index](#)]

2.7 Rules for Cleaning the Directory

Compiling a program is not the only thing you might want to write rules for. Makefiles commonly tell how to do a few other things besides compiling a program: for example, how to delete all the object files and executables so that the directory is `'clean'`.

Here is how we could write a `make` rule for cleaning our example editor:

```
clean:
    rm edit $(objects)
```

In practice, we might want to write the rule in a somewhat more complicated manner to handle unanticipated situations. We would do this:

```
.PHONY : clean
clean :
    -rm edit $(objects)
```

This prevents `make` from getting confused by an actual file called `clean` and causes it to continue in spite of errors from `rm`. (See [Phony Targets](#), and [Errors in Recipes](#).)

A rule such as this should not be placed at the beginning of the makefile, because we do not want it to run by default! Thus, in the example makefile, we want the rule for `edit`, which recompiles the editor, to remain the default goal.

Since `clean` is not a prerequisite of `edit`, this rule will not run at all if we give the command `'make'` with no arguments. In order to make the rule run, we have to type `'make clean'`. See [How to Run make](#).

Next: [Rules](#), Previous: [Introduction](#), Up: [Top](#) [[Contents](#)][[Index](#)]

3 Writing Makefiles

The information that tells `make` how to recompile a system comes from reading a data base called the makefile.

- [Makefile Contents](#) What makefiles contain.
- [Makefile Names](#) How to name your makefile.
- [Include](#) How one makefile can use another makefile.
- [MAKEFILES Variable](#) The environment can specify extra makefiles.
- [Remaking Makefiles](#) How makefiles get remade.
- [Overriding Makefiles](#) How to override part of one makefile with another makefile.
- [Reading Makefiles](#) How makefiles are read in.
- [Parsing Makefiles](#) How makefiles are parsed.
- [Secondary Expansion](#) How and when secondary expansion is performed.

Next: [Makefile Names](#), Previous: [Makefiles](#), Up: [Makefiles](#) [[Contents](#)][[Index](#)]

3.1 What Makefiles Contain

Makefiles contain five kinds of things: explicit rules, implicit rules, variable definitions, directives, and comments. Rules, variables, and directives are described at length in later chapters.

- An explicit rule says when and how to remake one or more files, called the rule's targets. It lists the other files that the targets depend on, called the prerequisites of the target, and may also give a recipe to use to create or update the targets. See [Writing Rules](#).
- An implicit rule says when and how to remake a class of files based on their names. It describes how a target may depend on a file with a name similar to the target and gives a recipe to create or update such a target. See [Using Implicit Rules](#).
- A variable definition is a line that specifies a text string value for a variable that can be substituted into the text later. The simple makefile example shows a variable definition for `objects` as a list of all object files (see [Variables Make Makefiles Simpler](#)).
- A directive is an instruction for `make` to do something special while reading the makefile. These include:
 - Reading another makefile (see [Including Other Makefiles](#)).
 - Deciding (based on the values of variables) whether to use or ignore a part of the makefile (see [Conditional Parts of Makefiles](#)).
 - Defining a variable from a verbatim string containing multiple lines (see [Defining Multi-Line Variables](#)).
- `#` in a line of a makefile starts a comment. It and the rest of the line are ignored, except that a trailing backslash not escaped by another backslash will continue the comment across multiple lines. A line containing just a comment (with perhaps spaces before it) is effectively blank, and is ignored. If you want a literal `#`, escape it with a backslash (e.g., `\#`). Comments may appear on any line in the makefile, although they are treated specially in certain situations.

You cannot use comments within variable references or function calls: any instance of `#` will be treated literally (rather than as the start of a comment) inside a variable reference or function call.

Comments within a recipe are passed to the shell, just as with any other recipe text. The shell decides how to interpret it: whether or not this is a comment is up to the shell.

Within a `define` directive, comments are not ignored during the definition of the variable, but rather kept intact in the value of the variable. When the variable is

expanded they will either be treated as `make` comments or as recipe text, depending on the context in which the variable is evaluated.

- [Splitting Lines](#) Splitting long lines in makefiles

Previous: [Makefile Contents](#), Up: [Makefile Contents](#) [[Contents](#)][[Index](#)]

3.1.1 Splitting Long Lines

Makefiles use a “line-based” syntax in which the newline character is special and marks the end of a statement. GNU `make` has no limit on the length of a statement line, up to the amount of memory in your computer.

However, it is difficult to read lines which are too long to display without wrapping or scrolling. So, you can format your makefiles for readability by adding newlines into the middle of a statement: you do this by escaping the internal newlines with a backslash (`\`) character. Where we need to make a distinction we will refer to “physical lines” as a single line ending with a newline (regardless of whether it is escaped) and a “logical line” being a complete statement including all escaped newlines up to the first non-escaped newline.

The way in which backslash/newline combinations are handled depends on whether the statement is a recipe line or a non-recipe line. Handling of backslash/newline in a recipe line is discussed later (see [Splitting Recipe Lines](#)).

Outside of recipe lines, backslash/newlines are converted into a single space character. Once that is done, all whitespace around the backslash/newline is condensed into a single space: this includes all whitespace preceding the backslash, all whitespace at the beginning of the line after the backslash/newline, and any consecutive backslash/newline combinations.

If the `.POSIX` special target is defined then backslash/newline handling is modified slightly to conform to POSIX.2: first, whitespace preceding a backslash is not removed and second, consecutive backslash/newlines are not condensed.

Splitting Without Adding Whitespace

If you need to split a line but do not want any whitespace added, you can utilize a subtle trick: replace your backslash/newline pairs with the three characters dollar sign/backslash/newline:

```
var := one$\  
      word
```

After `make` removes the backslash/newline and condenses the following line into a single space, this is equivalent to:

```
var := one$ word
```

Then `make` will perform variable expansion. The variable reference `‘$ ’` refers to a variable with the one-character name `“ ”` (space) which does not exist, and so expands to the empty string, giving a final assignment which is the equivalent of:

```
var := oneword
```

Next: [Include](#), Previous: [Makefile Contents](#), Up: [Makefiles](#) [[Contents](#)][[Index](#)]

3.2 What Name to Give Your Makefile

By default, when `make` looks for the makefile, it tries the following names, in order: `GNUmakefile`, `makefile` and `Makefile`.

Normally you should call your makefile either `makefile` or `Makefile`. (We recommend `Makefile` because it appears prominently near the beginning of a directory listing, right near other

important files such as `README`.) The first name checked, `GNUmakefile`, is not recommended for most makefiles. You should use this name if you have a makefile that is specific to GNU `make`, and will not be understood by other versions of `make`. Other `make` programs look for `makefile` and `Makefile`, but not `GNUmakefile`.

If `make` finds none of these names, it does not use any makefile. Then you must specify a goal with a command argument, and `make` will attempt to figure out how to remake it using only its built-in implicit rules. See [Using Implicit Rules](#).

If you want to use a nonstandard name for your makefile, you can specify the makefile name with the `-f` or `--file` option. The arguments `-f name` or `--file=name` tell `make` to read the file name as the makefile. If you use more than one `-f` or `--file` option, you can specify several makefiles. All the makefiles are effectively concatenated in the order specified. The default makefile names `GNUmakefile`, `makefile` and `Makefile` are not checked automatically if you specify `-f` or `--file`.

Next: [MAKEFILES Variable](#), Previous: [Makefile Names](#), Up: [Makefiles](#) [[Contents](#)][[Index](#)]

3.3 Including Other Makefiles

The `include` directive tells `make` to suspend reading the current makefile and read one or more other makefiles before continuing. The directive is a line in the makefile that looks like this:

```
include filenames...
```

`filenames` can contain shell file name patterns. If `filenames` is empty, nothing is included and no error is printed.

Extra spaces are allowed and ignored at the beginning of the line, but the first character must not be a tab (or the value of `.RECIPEPREFIX`)—if the line begins with a tab, it will be considered a recipe line. Whitespace is required between `include` and the file names, and between file names; extra whitespace is ignored there and at the end of the directive. A comment starting with `#` is allowed at the end of the line. If the file names contain any variable or function references, they are expanded. See [How to Use Variables](#).

For example, if you have three `.mk` files, `a.mk`, `b.mk`, and `c.mk`, and `$(bar)` expands to `bish bash`, then the following expression

```
include foo *.mk $(bar)
```

is equivalent to

```
include foo a.mk b.mk c.mk bish bash
```

When `make` processes an `include` directive, it suspends reading of the containing makefile and reads from each listed file in turn. When that is finished, `make` resumes reading the makefile in which the directive appears.

One occasion for using `include` directives is when several programs, handled by individual makefiles in various directories, need to use a common set of variable definitions (see [Setting Variables](#)) or pattern rules (see [Defining and Redefining Pattern Rules](#)).

Another such occasion is when you want to generate prerequisites from source files automatically; the prerequisites can be put in a file that is included by the main makefile. This practice is generally cleaner than that of somehow appending the prerequisites to the end of the main makefile as has been traditionally done with other versions of `make`. See [Automatic Prerequisites](#).

If the specified name does not start with a slash, and the file is not found in the current directory, several other directories are searched. First, any directories you have specified with the `-I` or `--include-dir` option are searched (see [Summary of Options](#)). Then the

following directories (if they exist) are searched, in this order: `prefix/include` (normally `/usr/local/include`¹) `/usr/gnu/include`, `/usr/local/include`, `/usr/include`.

If an included makefile cannot be found in any of these directories, a warning message is generated, but it is not an immediately fatal error; processing of the makefile containing the `include` continues. Once it has finished reading makefiles, `make` will try to remake any that are out of date or don't exist. See [How Makefiles Are Remade](#). Only after it has tried to find a way to remake a makefile and failed, will `make` diagnose the missing makefile as a fatal error.

If you want `make` to simply ignore a makefile which does not exist or cannot be remade, with no error message, use the `-include` directive instead of `include`, like this:

```
-include filenames...
```

This acts like `include` in every way except that there is no error (not even a warning) if any of the filenames (or any prerequisites of any of the filenames) do not exist or cannot be remade.

For compatibility with some other `make` implementations, `sinclude` is another name for `-include`.

Next: [Remaking Makefiles](#), Previous: [Include](#), Up: [Makefiles](#) [[Contents](#)][[Index](#)]

3.4 The Variable `MAKEFILES`

If the environment variable `MAKEFILES` is defined, `make` considers its value as a list of names (separated by whitespace) of additional makefiles to be read before the others. This works much like the `include` directive: various directories are searched for those files (see [Including Other Makefiles](#)). In addition, the default goal is never taken from one of these makefiles (or any makefile included by them) and it is not an error if the files listed in `MAKEFILES` are not found.

The main use of `MAKEFILES` is in communication between recursive invocations of `make` (see [Recursive Use of `make`](#)). It usually is not desirable to set the environment variable before a top-level invocation of `make`, because it is usually better not to mess with a makefile from outside. However, if you are running `make` without a specific makefile, a makefile in `MAKEFILES` can do useful things to help the built-in implicit rules work better, such as defining search paths (see [Directory Search](#)).

Some users are tempted to set `MAKEFILES` in the environment automatically on login, and program makefiles to expect this to be done. This is a very bad idea, because such makefiles will fail to work if run by anyone else. It is much better to write explicit `include` directives in the makefiles. See [Including Other Makefiles](#).

Next: [Overriding Makefiles](#), Previous: [MAKEFILES Variable](#), Up: [Makefiles](#) [[Contents](#)][[Index](#)]

3.5 How Makefiles Are Remade

Sometimes makefiles can be remade from other files, such as RCS or SCCS files. If a makefile can be remade from other files, you probably want `make` to get an up-to-date version of the makefile to read in.

To this end, after reading in all makefiles `make` will consider each as a goal target and attempt to update it. If a makefile has a rule which says how to update it (found either in that very makefile or in another one) or if an implicit rule applies to it (see [Using Implicit Rules](#)), it will be updated if necessary. After all makefiles have been checked, if any have actually been changed, `make` starts with a clean slate and reads all the makefiles over again. (It will also attempt to update each of them over again, but normally this will not change them again, since they are already up to date.) Each restart will cause the special variable `MAKE_RESTARTS` to be updated (see [Special Variables](#)).

If you know that one or more of your makefiles cannot be remade and you want to keep `make` from performing an implicit rule search on them, perhaps for efficiency reasons, you can use any normal method of preventing implicit rule look-up to do so. For example, you can write

an explicit rule with the makefile as the target, and an empty recipe (see [Using Empty Recipes](#)).

If the makefiles specify a double-colon rule to remake a file with a recipe but no prerequisites, that file will always be remade (see [Double-Colon](#)). In the case of makefiles, a makefile that has a double-colon rule with a recipe but no prerequisites will be remade every time `make` is run, and then again after `make` starts over and reads the makefiles in again. This would cause an infinite loop: `make` would constantly remake the makefile, and never do anything else. So, to avoid this, `make` will **not** attempt to remake makefiles which are specified as targets of a double-colon rule with a recipe but no prerequisites.

If you do not specify any makefiles to be read with `'-f'` or `'--file'` options, `make` will try the default makefile names; see [What Name to Give Your Makefile](#). Unlike makefiles explicitly requested with `'-f'` or `'--file'` options, `make` is not certain that these makefiles should exist. However, if a default makefile does not exist but can be created by running `make` rules, you probably want the rules to be run so that the makefile can be used.

Therefore, if none of the default makefiles exists, `make` will try to make each of them in the same order in which they are searched for (see [What Name to Give Your Makefile](#)) until it succeeds in making one, or it runs out of names to try. Note that it is not an error if `make` cannot find or make any makefile; a makefile is not always necessary.

When you use the `'-t'` or `'--touch'` option (see [Instead of Executing Recipes](#)), you would not want to use an out-of-date makefile to decide which targets to touch. So the `'-t'` option has no effect on updating makefiles; they are really updated even if `'-t'` is specified. Likewise, `'-q'` (or `'--question'`) and `'-n'` (or `'--just-print'`) do not prevent updating of makefiles, because an out-of-date makefile would result in the wrong output for other targets. Thus, `'make -f mfile -n foo'` will update `mfile`, read it in, and then print the recipe to update `foo` and its prerequisites without running it. The recipe printed for `foo` will be the one specified in the updated contents of `mfile`.

However, on occasion you might actually wish to prevent updating of even the makefiles. You can do this by specifying the makefiles as goals in the command line as well as specifying them as makefiles. When the makefile name is specified explicitly as a goal, the options `'-t'` and so on do apply to them.

Thus, `'make -f mfile -n mfile foo'` would read the makefile `mfile`, print the recipe needed to update it without actually running it, and then print the recipe needed to update `foo` without running that. The recipe for `foo` will be the one specified by the existing contents of `mfile`.

Next: [Reading Makefiles](#), Previous: [Remaking Makefiles](#), Up: [Makefiles](#) [[Contents](#)][[Index](#)]

3.6 Overriding Part of Another Makefile

Sometimes it is useful to have a makefile that is mostly just like another makefile. You can often use the `'include'` directive to include one in the other, and add more targets or variable definitions. However, it is invalid for two makefiles to give different recipes for the same target. But there is another way.

In the containing makefile (the one that wants to include the other), you can use a match-anything pattern rule to say that to remake any target that cannot be made from the information in the containing makefile, `make` should look in another makefile. See [Pattern Rules](#), for more information on pattern rules.

For example, if you have a makefile called `Makefile` that says how to make the target `'foo'` (and other targets), you can write a makefile called `GNUmakefile` that contains:

```
foo:
    frobnicate > foo

%: force
    @$(MAKE) -f Makefile $@

force: ;
```

If you say `'make foo'`, `make` will find `GNUmakefile`, read it, and see that to make `foo`, it needs to run the recipe `'frobnicate > foo'`. If you say `'make bar'`, `make` will find no way to make `bar` in `GNUmakefile`, so it will use the recipe from the pattern rule: `'make -f Makefile bar'`. If `Makefile` provides a rule for updating `bar`, `make` will apply the rule. And likewise for any other target that `GNUmakefile` does not say how to make.

The way this works is that the pattern rule has a pattern of just `'%'`, so it matches any target whatever. The rule specifies a prerequisite `force`, to guarantee that the recipe will be run even if the target file already exists. We give the `force` target an empty recipe to prevent `make` from searching for an implicit rule to build it—otherwise it would apply the same match-anything rule to `force` itself and create a prerequisite loop!

Next: [Parsing Makefiles](#), Previous: [Overriding Makefiles](#), Up: [Makefiles](#) [[Contents](#)][[Index](#)]

3.7 How `make` Reads a Makefile

GNU `make` does its work in two distinct phases. During the first phase it reads all the makefiles, included makefiles, etc. and internalizes all the variables and their values and implicit and explicit rules, and builds a dependency graph of all the targets and their prerequisites. During the second phase, `make` uses this internalized data to determine which targets need to be updated and run the recipes necessary to update them.

It's important to understand this two-phase approach because it has a direct impact on how variable and function expansion happens; this is often a source of some confusion when writing makefiles. Below is a summary of the different constructs that can be found in a makefile, and the phase in which expansion happens for each part of the construct.

We say that expansion is immediate if it happens during the first phase: `make` will expand that part of the construct as the makefile is parsed. We say that expansion is deferred if it is not immediate. Expansion of a deferred construct part is delayed until the expansion is used: either when it is referenced in an immediate context, or when it is needed during the second phase.

You may not be familiar with some of these constructs yet. You can reference this section as you become familiar with them, in later chapters.

Variable Assignment

Variable definitions are parsed as follows:

```
immediate = deferred
immediate ?= deferred
immediate := immediate
immediate ::= immediate
immediate += deferred or immediate
immediate != immediate

define immediate
    deferred
endef

define immediate =
    deferred
endef

define immediate ?=
    deferred
endef

define immediate :=
    immediate
endef

define immediate ::=
    immediate
endef

define immediate +=
    deferred or immediate
```

```
endef

define immediate !=
    immediate
endef
```

For the append operator `'+='`, the right-hand side is considered immediate if the variable was previously set as a simple variable (`':='` or `'::='`), and deferred otherwise.

For the shell assignment operator `'!='`, the right-hand side is evaluated immediately and handed to the shell. The result is stored in the variable named on the left, and that variable becomes a simple variable (and will thus be re-evaluated on each reference).

Conditional Directives

Conditional directives are parsed immediately. This means, for example, that automatic variables cannot be used in conditional directives, as automatic variables are not set until the recipe for that rule is invoked. If you need to use automatic variables in a conditional directive you must move the condition into the recipe and use shell conditional syntax instead.

Rule Definition

A rule is always expanded the same way, regardless of the form:

```
immediate : immediate ; deferred
deferred
```

That is, the target and prerequisite sections are expanded immediately, and the recipe used to build the target is always deferred. This is true for explicit rules, pattern rules, suffix rules, static pattern rules, and simple prerequisite definitions.

Next: [Secondary Expansion](#), Previous: [Reading Makefiles](#), Up: [Makefiles](#) [[Contents](#)][[Index](#)]

3.8 How Makefiles Are Parsed

GNU `make` parses makefiles line-by-line. Parsing proceeds using the following steps:

1. Read in a full logical line, including backslash-escaped lines (see [Splitting Long Lines](#)).
2. Remove comments (see [What Makefiles Contain](#)).
3. If the line begins with the recipe prefix character and we are in a rule context, add the line to the current recipe and read the next line (see [Recipe Syntax](#)).
4. Expand elements of the line which appear in an immediate expansion context (see [How make Reads a Makefile](#)).
5. Scan the line for a separator character, such as `':'` or `'='`, to determine whether the line is a macro assignment or a rule (see [Recipe Syntax](#)).
6. Internalize the resulting operation and read the next line.

An important consequence of this is that a macro can expand to an entire rule, if it is one line long. This will work:

```
myrule = target : ; echo built

$(myrule)
```

However, this will not work because `make` does not re-split lines after it has expanded them:

```
define myrule
target:
    echo built
endef

$(myrule)
```

The above makefile results in the definition of a target `'target'` with prerequisites `'echo'` and `'built'`, as if the makefile contained `target: echo built`, rather than a rule with a recipe. Newlines still present in a line after expansion is complete are ignored as normal whitespace.

In order to properly expand a multi-line macro you must use the `eval` function: this causes the `make` parser to be run on the results of the expanded macro (see [Eval Function](#)).

Previous: [Parsing Makefiles](#), Up: [Makefiles](#) [[Contents](#)][[Index](#)]

3.9 Secondary Expansion

Previously we learned that GNU `make` works in two distinct phases: a read-in phase and a target-update phase (see [How make Reads a Makefile](#)). GNU `make` also has the ability to enable a second expansion of the prerequisites (only) for some or all targets defined in the makefile. In order for this second expansion to occur, the special target `.SECONDEXPANSION` must be defined before the first prerequisite list that makes use of this feature.

If that special target is defined then in between the two phases mentioned above, right at the end of the read-in phase, all the prerequisites of the targets defined after the special target are expanded a second time. In most circumstances this secondary expansion will have no effect, since all variable and function references will have been expanded during the initial parsing of the makefiles. In order to take advantage of the secondary expansion phase of the parser, then, it's necessary to escape the variable or function reference in the makefile. In this case the first expansion merely un-escapes the reference but doesn't expand it, and expansion is left to the secondary expansion phase. For example, consider this makefile:

```
.SECONDEXPANSION:
ONEVAR = onefile
TWOVAR = twofile
myfile: $(ONEVAR) $$ (TWOVAR)
```

After the first expansion phase the prerequisites list of the `myfile` target will be `onefile` and `$(TWOVAR)`; the first (unescaped) variable reference to `ONEVAR` is expanded, while the second (escaped) variable reference is simply unescaped, without being recognized as a variable reference. Now during the secondary expansion the first word is expanded again but since it contains no variable or function references it remains the value `onefile`, while the second word is now a normal reference to the variable `TWOVAR`, which is expanded to the value `twofile`. The final result is that there are two prerequisites, `onefile` and `twofile`.

Obviously, this is not a very interesting case since the same result could more easily have been achieved simply by having both variables appear, unescaped, in the prerequisites list. One difference becomes apparent if the variables are reset; consider this example:

```
.SECONDEXPANSION:
AVAR = top
onefile: $(AVAR)
twofile: $$ (AVAR)
AVAR = bottom
```

Here the prerequisite of `onefile` will be expanded immediately, and resolve to the value `top`, while the prerequisite of `twofile` will not be full expanded until the secondary expansion and yield a value of `bottom`.

This is marginally more exciting, but the true power of this feature only becomes apparent when you discover that secondary expansions always take place within the scope of the automatic variables for that target. This means that you can use variables such as `$$`, `$$*`, etc. during the second expansion and they will have their expected values, just as in the recipe. All you have to do is defer the expansion by escaping the `$`. Also, secondary expansion occurs for both explicit and implicit (pattern) rules. Knowing this, the possible uses for this feature increase dramatically. For example:

```
.SECONDEXPANSION:
main_OBJS := main.o try.o test.o
lib_OBJS := lib.o api.o
```

```
main lib: $$($$_OJS)
```

Here, after the initial expansion the prerequisites of both the `main` and `lib` targets will be `$($$_OJS)`. During the secondary expansion, the `$_` variable is set to the name of the target and so the expansion for the `main` target will yield `$(main_OJS)`, or `main.o try.o test.o`, while the secondary expansion for the `lib` target will yield `$(lib_OJS)`, or `lib.o api.o`.

You can also mix in functions here, as long as they are properly escaped:

```
main_SRCS := main.c try.c test.c
lib_SRCS  := lib.c api.c

.SECONDEXPANSION:
main lib: $$ (patsubst %.c, %.o, $$($$_SRCS))
```

This version allows users to specify source files rather than object files, but gives the same resulting prerequisites list as the previous example.

Evaluation of automatic variables during the secondary expansion phase, especially of the target name variable `$_`, behaves similarly to evaluation within recipes. However, there are some subtle differences and “corner cases” which come into play for the different types of rule definitions that `make` understands. The subtleties of using the different automatic variables are described below.

Secondary Expansion of Explicit Rules

During the secondary expansion of explicit rules, `$_` and `$$%` evaluate, respectively, to the file name of the target and, when the target is an archive member, the target member name. The `$$<` variable evaluates to the first prerequisite in the first rule for this target. `$$^` and `$$+` evaluate to the list of all prerequisites of rules that have already appeared for the same target (`$$+` with repetitions and `$$^` without). The following example will help illustrate these behaviors:

```
.SECONDEXPANSION:

foo: foo.1 bar.1 $$< $$^ $$+    # line #1

foo: foo.2 bar.2 $$< $$^ $$+    # line #2

foo: foo.3 bar.3 $$< $$^ $$+    # line #3
```

In the first prerequisite list, all three variables (`$$<`, `$$^`, and `$$+`) expand to the empty string. In the second, they will have values `foo.1`, `foo.1 bar.1`, and `foo.1 bar.1` respectively. In the third they will have values `foo.1`, `foo.1 bar.1 foo.2 bar.2`, and `foo.1 bar.1 foo.2 bar.2 foo.1 foo.1 bar.1 foo.1 bar.1` respectively.

Rules undergo secondary expansion in makefile order, except that the rule with the recipe is always evaluated last.

The variables `$$?` and `$$*` are not available and expand to the empty string.

Secondary Expansion of Static Pattern Rules

Rules for secondary expansion of static pattern rules are identical to those for explicit rules, above, with one exception: for static pattern rules the `$$*` variable is set to the pattern stem. As with explicit rules, `$$?` is not available and expands to the empty string.

Secondary Expansion of Implicit Rules

As `make` searches for an implicit rule, it substitutes the stem and then performs secondary expansion for every rule with a matching target pattern. The value of the automatic variables is derived in the same fashion as for static pattern rules. As an example:

```
.SECONDEXPANSION:

foo: bar
```

```
foo foz: fo%: bo%

%oo: $$< $$^ $$+ $$*
```

When the implicit rule is tried for target `foo`, `$$<` expands to `bar`, `$$^` expands to `bar boo`, `$$+` also expands to `bar boo`, and `$$*` expands to `f`.

Note that the directory prefix (D), as described in [Implicit Rule Search Algorithm](#), is appended (after expansion) to all the patterns in the prerequisites list. As an example:

```
. SECONDEXPANSION:

/tmp/foo.o:

%.o: $(addsuffix /%.c,foo bar) foo.h
      @echo $^
```

The prerequisite list printed, after the secondary expansion and directory prefix reconstruction, will be `/tmp/foo/foo.c /tmp/bar/foo.c foo.h`. If you are not interested in this reconstruction, you can use `$$*` instead of `%` in the prerequisites list.

Next: [Recipes](#), Previous: [Makefiles](#), Up: [Top](#) [[Contents](#)][[Index](#)]

4 Writing Rules

A rule appears in the makefile and says when and how to remake certain files, called the rule’ s targets (most often only one per rule). It lists the other files that are the prerequisites of the target, and the recipe to use to create or update the target.

The order of rules is not significant, except for determining the default goal: the target for `make` to consider, if you do not otherwise specify one. The default goal is the target of the first rule in the first makefile. If the first rule has multiple targets, only the first target is taken as the default. There are two exceptions: a target starting with a period is not a default unless it contains one or more slashes, `’/’` , as well; and, a target that defines a pattern rule has no effect on the default goal. (See [Defining and Redefining Pattern Rules](#).)

Therefore, we usually write the makefile so that the first rule is the one for compiling the entire program or all the programs described by the makefile (often with a target called `’all’`). See [Arguments to Specify the Goals](#).

• Rule Example	An example explained.
• Rule Syntax	General syntax explained.
• Prerequisite Types	There are two types of prerequisites.
• Wildcards	Using wildcard characters such as <code>’*’</code> .
• Directory Search	Searching other directories for source files.
• Phony Targets	Using a target that is not a real file’ s name.
• Force Targets	You can use a target without a recipe or prerequisites to mark other targets as phony.
• Empty Targets	When only the date matters and the files are empty.
• Special Targets	Targets with special built-in meanings.
• Multiple Targets	When to make use of several targets in a rule.
• Multiple Rules	How to use several rules with the same target.
• Static Pattern	Static pattern rules apply to multiple targets and can vary the prerequisites according to the target name.

- [Double-Colon](#) How to use a special kind of rule to allow several independent rules for one target.
- [Automatic Prerequisites](#) How to automatically generate rules giving prerequisites from source files themselves.

Next: [Rule Syntax](#), Previous: [Rules](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.1 Rule Example

Here is an example of a rule:

```
foo.o : foo.c defs.h      # module for twiddling the frobs
      cc -c -g foo.c
```

Its target is `foo.o` and its prerequisites are `foo.c` and `defs.h`. It has one command in the recipe: `'cc -c -g foo.c'`. The recipe starts with a tab to identify it as a recipe.

This rule says two things:

- How to decide whether `foo.o` is out of date: it is out of date if it does not exist, or if either `foo.c` or `defs.h` is more recent than it.
- How to update the file `foo.o`: by running `cc` as stated. The recipe does not explicitly mention `defs.h`, but we presume that `foo.c` includes it, and that is why `defs.h` was added to the prerequisites.

Next: [Prerequisite Types](#), Previous: [Rule Example](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.2 Rule Syntax

In general, a rule looks like this:

```
targets : prerequisites
      recipe
      ...
```

or like this:

```
targets : prerequisites ; recipe
      recipe
      ...
```

The targets are file names, separated by spaces. Wildcard characters may be used (see [Using Wildcard Characters in File Names](#)) and a name of the form `a(m)` represents member `m` in archive file `a` (see [Archive Members as Targets](#)). Usually there is only one target per rule, but occasionally there is a reason to have more (see [Multiple Targets in a Rule](#)).

The recipe lines start with a tab character (or the first character in the value of the `.RECIPEPREFIX` variable; see [Special Variables](#)). The first recipe line may appear on the line after the prerequisites, with a tab character, or may appear on the same line, with a semicolon. Either way, the effect is the same. There are other differences in the syntax of recipes. See [Writing Recipes in Rules](#).

Because dollar signs are used to start `make` variable references, if you really want a dollar sign in a target or prerequisite you must write two of them, `'$$'` (see [How to Use Variables](#)). If you have enabled secondary expansion (see [Secondary Expansion](#)) and you want a literal dollar sign in the prerequisites list, you must actually write four dollar signs (`'$$$$'`).

You may split a long line by inserting a backslash followed by a newline, but this is not required, as `make` places no limit on the length of a line in a makefile.

A rule tells `make` two things: when the targets are out of date, and how to update them when necessary.

The criterion for being out of date is specified in terms of the prerequisites, which consist of file names separated by spaces. (Wildcards and archive members (see [Archives](#)) are allowed here too.) A target is out of date if it does not exist or if it is older than any of the prerequisites (by comparison of last-modification times). The idea is that the contents of the target file are computed based on information in the prerequisites, so if any of the prerequisites changes, the contents of the existing target file are no longer necessarily valid.

How to update is specified by a recipe. This is one or more lines to be executed by the shell (normally `'sh'`), but with some extra features (see [Writing Recipes in Rules](#)).

Next: [Wildcards](#), Previous: [Rule Syntax](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.3 Types of Prerequisites

There are actually two different types of prerequisites understood by GNU `make`: normal prerequisites such as described in the previous section, and order-only prerequisites. A normal prerequisite makes two statements: first, it imposes an order in which recipes will be invoked: the recipes for all prerequisites of a target will be completed before the recipe for the target is run. Second, it imposes a dependency relationship: if any prerequisite is newer than the target, then the target is considered out-of-date and must be rebuilt.

Normally, this is exactly what you want: if a target's prerequisite is updated, then the target should also be updated.

Occasionally, however, you have a situation where you want to impose a specific ordering on the rules to be invoked without forcing the target to be updated if one of those rules is executed. In that case, you want to define order-only prerequisites. Order-only prerequisites can be specified by placing a pipe symbol (`|`) in the prerequisites list: any prerequisites to the left of the pipe symbol are normal; any prerequisites to the right are order-only:

```
targets : normal-prerequisites | order-only-prerequisites
```

The normal prerequisites section may of course be empty. Also, you may still declare multiple lines of prerequisites for the same target: they are appended appropriately (normal prerequisites are appended to the list of normal prerequisites; order-only prerequisites are appended to the list of order-only prerequisites). Note that if you declare the same file to be both a normal and an order-only prerequisite, the normal prerequisite takes precedence (since they have a strict superset of the behavior of an order-only prerequisite).

Consider an example where your targets are to be placed in a separate directory, and that directory might not exist before `make` is run. In this situation, you want the directory to be created before any targets are placed into it but, because the timestamps on directories change whenever a file is added, removed, or renamed, we certainly don't want to rebuild all the targets whenever the directory's timestamp changes. One way to manage this is with order-only prerequisites: make the directory an order-only prerequisite on all the targets:

```
OBJDIR := objdir
OBJJS := $(addprefix $(OBJDIR)/,foo.o bar.o baz.o)

$(OBJDIR)/%.o : %.c
    $(COMPILE.c) $(OUTPUT_OPTION) $<

all: $(OBJJS)

$(OBJJS) : | $(OBJDIR)

$(OBJDIR) :
    mkdir $(OBJDIR)
```

Now the rule to create the `objdir` directory will be run, if needed, before any `'o'` is built, but no `'o'` will be built because the `objdir` directory timestamp changed.

Next: [Directory Search](#), Previous: [Prerequisite Types](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.4 Using Wildcard Characters in File Names

A single file name can specify many files using wildcard characters. The wildcard characters in `make` are `'*'`, `'?'` and `'[...]'`, the same as in the Bourne shell. For example, `*.c` specifies a list of all the files (in the working directory) whose names end in `'c'`.

The character `'~'` at the beginning of a file name also has special significance. If alone, or followed by a slash, it represents your home directory. For example `~/bin` expands to `/home/you/bin`. If the `'~'` is followed by a word, the string represents the home directory of the user named by that word. For example `~john/bin` expands to `/home/john/bin`. On systems which don't have a home directory for each user (such as MS-DOS or MS-Windows), this functionality can be simulated by setting the environment variable `HOME`.

Wildcard expansion is performed by `make` automatically in targets and in prerequisites. In recipes, the shell is responsible for wildcard expansion. In other contexts, wildcard expansion happens only if you request it explicitly with the `wildcard` function.

The special significance of a wildcard character can be turned off by preceding it with a backslash. Thus, `foo\*bar` would refer to a specific file whose name consists of `'foo'`, an asterisk, and `'bar'`.

- [Wildcard Examples](#) Several examples.
- [Wildcard Pitfall](#) Problems to avoid.
- [Wildcard Function](#) How to cause wildcard expansion where it does not normally take place.

Next: [Wildcard Pitfall](#), Previous: [Wildcards](#), Up: [Wildcards](#) [[Contents](#)][[Index](#)]

4.4.1 Wildcard Examples

Wildcards can be used in the recipe of a rule, where they are expanded by the shell. For example, here is a rule to delete all the object files:

```
clean:
    rm -f *.o
```

Wildcards are also useful in the prerequisites of a rule. With the following rule in the makefile, `'make print'` will print all the `'c'` files that have changed since the last time you printed them:

```
print: *.c
    lpr -p $?
    touch print
```

This rule uses `print` as an empty target file; see [Empty Target Files to Record Events](#). (The automatic variable `'$?'` is used to print only those files that have changed; see [Automatic Variables](#).)

Wildcard expansion does not happen when you define a variable. Thus, if you write this:

```
objects = *.o
```

then the value of the variable `objects` is the actual string `'*.o'`. However, if you use the value of `objects` in a target or prerequisite, wildcard expansion will take place there. If you use the value of `objects` in a recipe, the shell may perform wildcard expansion when the recipe runs. To set `objects` to the expansion, instead use:

```
objects := $(wildcard *.o)
```

See [Wildcard Function](#).

4.4.2 Pitfalls of Using Wildcards

Now here is an example of a naive way of using wildcard expansion, that does not do what you would intend. Suppose you would like to say that the executable file `foo` is made from all the object files in the directory, and you write this:

```
objects = *.o

foo : $(objects)
    cc -o foo $(CFLAGS) $(objects)
```

The value of `objects` is the actual string `'*.o'`. Wildcard expansion happens in the rule for `foo`, so that each existing `'*.o'` file becomes a prerequisite of `foo` and will be recompiled if necessary.

But what if you delete all the `'*.o'` files? When a wildcard matches no files, it is left as it is, so then `foo` will depend on the oddly-named file `*.o`. Since no such file is likely to exist, `make` will give you an error saying it cannot figure out how to make `*.o`. This is not what you want!

Actually it is possible to obtain the desired result with wildcard expansion, but you need more sophisticated techniques, including the `wildcard` function and string substitution. See [The Function `wildcard`](#).

Microsoft operating systems (MS-DOS and MS-Windows) use backslashes to separate directories in pathnames, like so:

```
c:\foo\bar\baz.c
```

This is equivalent to the Unix-style `c:/foo/bar/baz.c` (the `c:` part is the so-called drive letter). When `make` runs on these systems, it supports backslashes as well as the Unix-style forward slashes in pathnames. However, this support does not include the wildcard expansion, where backslash is a quote character. Therefore, you must use Unix-style slashes in these cases.

4.4.3 The Function `wildcard`

Wildcard expansion happens automatically in rules. But wildcard expansion does not normally take place when a variable is set, or inside the arguments of a function. If you want to do wildcard expansion in such places, you need to use the `wildcard` function, like this:

```
$(wildcard pattern...)
```

This string, used anywhere in a makefile, is replaced by a space-separated list of names of existing files that match one of the given file name patterns. If no existing file name matches a pattern, then that pattern is omitted from the output of the `wildcard` function. Note that this is different from how unmatched wildcards behave in rules, where they are used verbatim rather than ignored (see [Wildcard Pitfall](#)).

One use of the `wildcard` function is to get a list of all the C source files in a directory, like this:

```
$(wildcard *.c)
```

We can change the list of C source files into a list of object files by replacing the `'*.c'` suffix with `'*.o'` in the result, like this:

```
$(patsubst %.c, %.o, $(wildcard *.c))
```

(Here we have used another function, `patsubst`. See [Functions for String Substitution and Analysis](#).)

Thus, a makefile to compile all C source files in the directory and then link them together could be written as follows:

```
objects := $(patsubst %.c,%.o,$(wildcard *.c))

foo : $(objects)
    cc -o foo $(objects)
```

(This takes advantage of the implicit rule for compiling C programs, so there is no need to write explicit rules for compiling the files. See [The Two Flavors of Variables](#), for an explanation of `':='`, which is a variant of `'='`.)

Next: [Phony Targets](#), Previous: [Wildcards](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.5 Searching Directories for Prerequisites

For large systems, it is often desirable to put sources in a separate directory from the binaries. The directory search features of `make` facilitate this by searching several directories automatically to find a prerequisite. When you redistribute the files among directories, you do not need to change the individual rules, just the search paths.

- [General Search](#) Specifying a search path that applies to every prerequisite.
- [Selective Search](#) Specifying a search path for a specified class of names.
- [Search Algorithm](#) When and how search paths are applied.
- [Recipes/Search](#) How to write recipes that work together with search paths.
- [Implicit/Search](#) How search paths affect implicit rules.
- [Libraries/Search](#) Directory search for link libraries.

Next: [Selective Search](#), Previous: [Directory Search](#), Up: [Directory Search](#) [[Contents](#)][[Index](#)]

4.5.1 `VPATH`: Search Path for All Prerequisites

The value of the `make` variable `VPATH` specifies a list of directories that `make` should search. Most often, the directories are expected to contain prerequisite files that are not in the current directory; however, `make` uses `VPATH` as a search list for both prerequisites and targets of rules.

Thus, if a file that is listed as a target or prerequisite does not exist in the current directory, `make` searches the directories listed in `VPATH` for a file with that name. If a file is found in one of them, that file may become the prerequisite (see below). Rules may then specify the names of files in the prerequisite list as if they all existed in the current directory. See [Writing Recipes with Directory Search](#).

In the `VPATH` variable, directory names are separated by colons or blanks. The order in which directories are listed is the order followed by `make` in its search. (On MS-DOS and MS-Windows, semi-colons are used as separators of directory names in `VPATH`, since the colon can be used in the pathname itself, after the drive letter.)

For example,

```
VPATH = src:../headers
```

specifies a path containing two directories, `src` and `../headers`, which `make` searches in that order.

With this value of `VPATH`, the following rule,

```
foo.o : foo.c
```

is interpreted as if it were written like this:

```
foo.o : src/foo.c
```

assuming the file `foo.c` does not exist in the current directory but is found in the directory `src`.

Next: [Search Algorithm](#), Previous: [General Search](#), Up: [Directory Search](#) [\[Contents\]](#)[\[Index\]](#)

4.5.2 The `vpath` Directive

Similar to the `VPATH` variable, but more selective, is the `vpath` directive (note lower case), which allows you to specify a search path for a particular class of file names: those that match a particular pattern. Thus you can supply certain search directories for one class of file names and other directories (or none) for other file names.

There are three forms of the `vpath` directive:

`vpath pattern directories`

Specify the search path directories for file names that match pattern.

The search path, directories, is a list of directories to be searched, separated by colons (semi-colons on MS-DOS and MS-Windows) or blanks, just like the search path used in the `VPATH` variable.

`vpath pattern`

Clear out the search path associated with pattern.

`vpath`

Clear all search paths previously specified with `vpath` directives.

A `vpath` pattern is a string containing a `'%'` character. The string must match the file name of a prerequisite that is being searched for, the `'%'` character matching any sequence of zero or more characters (as in pattern rules; see [Defining and Redefining Pattern Rules](#)). For example, `%.h` matches files that end in `.h`. (If there is no `'%'`, the pattern must match the prerequisite exactly, which is not useful very often.)

`'%'` characters in a `vpath` directive's pattern can be quoted with preceding backslashes (`'\'`). Backslashes that would otherwise quote `'%'` characters can be quoted with more backslashes. Backslashes that quote `'%'` characters or other backslashes are removed from the pattern before it is compared to file names. Backslashes that are not in danger of quoting `'%'` characters go unmolested.

When a prerequisite fails to exist in the current directory, if the pattern in a `vpath` directive matches the name of the prerequisite file, then the directories in that directive are searched just like (and before) the directories in the `VPATH` variable.

For example,

```
vpath %.h ../headers
```

tells `make` to look for any prerequisite whose name ends in `.h` in the directory `../headers` if the file is not found in the current directory.

If several `vpath` patterns match the prerequisite file's name, then `make` processes each matching `vpath` directive one by one, searching all the directories mentioned in each directive. `make` handles multiple `vpath` directives in the order in which they appear in the makefile; multiple directives with the same pattern are independent of each other.

Thus,

```
vpath %.c foo
vpath %    blish
vpath %.c bar
```

will look for a file ending in `' .c '` in `foo`, then `blish`, then `bar`, while

```
vpath %.c foo:bar
vpath %    blish
```

will look for a file ending in `' .c '` in `foo`, then `bar`, then `blish`.

Next: [Recipes/Search](#), Previous: [Selective Search](#), Up: [Directory Search](#) [[Contents](#)][[Index](#)]

4.5.3 How Directory Searches are Performed

When a prerequisite is found through directory search, regardless of type (general or selective), the pathname located may not be the one that `make` actually provides you in the prerequisite list. Sometimes the path discovered through directory search is thrown away.

The algorithm `make` uses to decide whether to keep or abandon a path found via directory search is as follows:

1. If a target file does not exist at the path specified in the makefile, directory search is performed.
2. If the directory search is successful, that path is kept and this file is tentatively stored as the target.
3. All prerequisites of this target are examined using this same method.
4. After processing the prerequisites, the target may or may not need to be rebuilt:
 1. If the target does not need to be rebuilt, the path to the file found during directory search is used for any prerequisite lists which contain this target. In short, if `make` doesn't need to rebuild the target then you use the path found via directory search.
 2. If the target does need to be rebuilt (is out-of-date), the pathname found during directory search is thrown away, and the target is rebuilt using the file name specified in the makefile. In short, if `make` must rebuild, then the target is rebuilt locally, not in the directory found via directory search.

This algorithm may seem complex, but in practice it is quite often exactly what you want.

Other versions of `make` use a simpler algorithm: if the file does not exist, and it is found via directory search, then that pathname is always used whether or not the target needs to be built. Thus, if the target is rebuilt it is created at the pathname discovered during directory search.

If, in fact, this is the behavior you want for some or all of your directories, you can use the `GPATH` variable to indicate this to `make`.

`GPATH` has the same syntax and format as `VPATH` (that is, a space- or colon-delimited list of pathnames). If an out-of-date target is found by directory search in a directory that also appears in `GPATH`, then that pathname is not thrown away. The target is rebuilt using the expanded path.

Next: [Implicit/Search](#), Previous: [Search Algorithm](#), Up: [Directory Search](#) [[Contents](#)][[Index](#)]

4.5.4 Writing Recipes with Directory Search

When a prerequisite is found in another directory through directory search, this cannot change the recipe of the rule; they will execute as written. Therefore, you must write the recipe with care so that it will look for the prerequisite in the directory where `make` finds it.

This is done with the automatic variables such as `'$^'` (see [Automatic Variables](#)). For instance, the value of `'$^'` is a list of all the prerequisites of the rule, including the names of the directories in which they were found, and the value of `'$@'` is the target. Thus:

```
foo.o : foo.c
cc -c $(CFLAGS) $^ -o $@
```

(The variable `CFLAGS` exists so you can specify flags for C compilation by implicit rules; we use it here for consistency so it will affect all C compilations uniformly; see [Variables Used by Implicit Rules](#).)

Often the prerequisites include header files as well, which you do not want to mention in the recipe. The automatic variable `‘$<’` is just the first prerequisite:

```
VPATH = src:../headers
foo.o : foo.c defs.h hack.h
cc -c $(CFLAGS) $< -o $@
```

Next: [Libraries/Search](#), Previous: [Recipes/Search](#), Up: [Directory Search](#) [\[Contents\]](#)[\[Index\]](#)

4.5.5 Directory Search and Implicit Rules

The search through the directories specified in `VPATH` or with `vpath` also happens during consideration of implicit rules (see [Using Implicit Rules](#)).

For example, when a file `foo.o` has no explicit rule, `make` considers implicit rules, such as the built-in rule to compile `foo.c` if that file exists. If such a file is lacking in the current directory, the appropriate directories are searched for it. If `foo.c` exists (or is mentioned in the makefile) in any of the directories, the implicit rule for C compilation is applied.

The recipes of implicit rules normally use automatic variables as a matter of necessity; consequently they will use the file names found by directory search with no extra effort.

Previous: [Implicit/Search](#), Up: [Directory Search](#) [\[Contents\]](#)[\[Index\]](#)

4.5.6 Directory Search for Link Libraries

Directory search applies in a special way to libraries used with the linker. This special feature comes into play when you write a prerequisite whose name is of the form `‘-lname’`. (You can tell something strange is going on here because the prerequisite is normally the name of a file, and the file name of a library generally looks like `libname.a`, not like `‘-lname’`.)

When a prerequisite’s name has the form `‘-lname’`, `make` handles it specially by searching for the file `libname.so`, and, if it is not found, for the file `libname.a` in the current directory, in directories specified by matching `vpath` search paths and the `VPATH` search path, and then in the directories `/lib`, `/usr/lib`, and `prefix/lib` (normally `/usr/local/lib`, but MS-DOS/MS-Windows versions of `make` behave as if `prefix` is defined to be the root of the DJGPP installation tree).

For example, if there is a `/usr/lib/libcurses.a` library on your system (and no `/usr/lib/libcurses.so` file), then

```
foo : foo.c -lcurses
cc $^ -o $@
```

would cause the command `‘cc foo.c /usr/lib/libcurses.a -o foo’` to be executed when `foo` is older than `foo.c` or than `/usr/lib/libcurses.a`.

Although the default set of files to be searched for is `libname.so` and `libname.a`, this is customizable via the `.LIBPATTERNS` variable. Each word in the value of this variable is a pattern string. When a prerequisite like `‘-lname’` is seen, `make` will replace the percent in each pattern in the list with `name` and perform the above directory searches using each library file name.

The default value for `.LIBPATTERNS` is `‘lib%.so lib%.a’`, which provides the default behavior described above.

You can turn off link library expansion completely by setting this variable to an empty value.

4.6 Phony Targets

A phony target is one that is not really the name of a file; rather it is just a name for a recipe to be executed when you make an explicit request. There are two reasons to use a phony target: to avoid a conflict with a file of the same name, and to improve performance.

If you write a rule whose recipe will not create the target file, the recipe will be executed every time the target comes up for remaking. Here is an example:

```
clean:
    rm *.o temp
```

Because the `rm` command does not create a file named `clean`, probably no such file will ever exist. Therefore, the `rm` command will be executed every time you say `'make clean'`.

In this example, the `clean` target will not work properly if a file named `clean` is ever created in this directory. Since it has no prerequisites, `clean` would always be considered up to date and its recipe would not be executed. To avoid this problem you can explicitly declare the target to be phony by making it a prerequisite of the special target `.PHONY` (see [Special Built-in Target Names](#)) as follows:

```
.PHONY: clean
clean:
    rm *.o temp
```

Once this is done, `'make clean'` will run the recipe regardless of whether there is a file named `clean`.

Phony targets are also useful in conjunction with recursive invocations of `make` (see [Recursive Use of make](#)). In this situation the makefile will often contain a variable which lists a number of sub-directories to be built. A simplistic way to handle this is to define one rule with a recipe that loops over the sub-directories, like this:

```
SUBDIRS = foo bar baz

subdirs:
    for dir in $(SUBDIRS); do \
        $(MAKE) -C $$dir; \
    done
```

There are problems with this method, however. First, any error detected in a sub-make is ignored by this rule, so it will continue to build the rest of the directories even when one fails. This can be overcome by adding shell commands to note the error and exit, but then it will do so even if `make` is invoked with the `-k` option, which is unfortunate. Second, and perhaps more importantly, you cannot take advantage of `make`'s ability to build targets in parallel (see [Parallel Execution](#)), since there is only one rule.

By declaring the sub-directories as `.PHONY` targets (you must do this as the sub-directory obviously always exists; otherwise it won't be built) you can remove these problems:

```
SUBDIRS = foo bar baz

.PHONY: subdirs $(SUBDIRS)

subdirs: $(SUBDIRS)

$(SUBDIRS):
    $(MAKE) -C $@

foo: baz
```

Here we've also declared that the `foo` sub-directory cannot be built until after the `baz` sub-directory is complete; this kind of relationship declaration is particularly important when

attempting parallel builds.

The implicit rule search (see [Implicit Rules](#)) is skipped for `.PHONY` targets. This is why declaring a target as `.PHONY` is good for performance, even if you are not worried about the actual file existing.

A phony target should not be a prerequisite of a real target file; if it is, its recipe will be run every time `make` goes to update that file. As long as a phony target is never a prerequisite of a real target, the phony target recipe will be executed only when the phony target is a specified goal (see [Arguments to Specify the Goals](#)).

Phony targets can have prerequisites. When one directory contains multiple programs, it is most convenient to describe all of the programs in one makefile `./Makefile`. Since the target remade by default will be the first one in the makefile, it is common to make this a phony target named `'all'` and give it, as prerequisites, all the individual programs. For example:

```
all : prog1 prog2 prog3
.PHONY : all

prog1 : prog1.o utils.o
       cc -o prog1 prog1.o utils.o

prog2 : prog2.o
       cc -o prog2 prog2.o

prog3 : prog3.o sort.o utils.o
       cc -o prog3 prog3.o sort.o utils.o
```

Now you can say just `'make'` to remake all three programs, or specify as arguments the ones to remake (as in `'make prog1 prog3'`). Phoniness is not inherited: the prerequisites of a phony target are not themselves phony, unless explicitly declared to be so.

When one phony target is a prerequisite of another, it serves as a subroutine of the other. For example, here `'make cleanall'` will delete the object files, the difference files, and the file program:

```
.PHONY: cleanall cleanobj cleandiff

cleanall : cleanobj cleandiff
         rm program

cleanobj :
         rm *.o

cleandiff :
         rm *.diff
```

Next: [Empty Targets](#), Previous: [Phony Targets](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.7 Rules without Recipes or Prerequisites

If a rule has no prerequisites or recipe, and the target of the rule is a nonexistent file, then `make` imagines this target to have been updated whenever its rule is run. This implies that all targets depending on this one will always have their recipe run.

An example will illustrate this:

```
clean: FORCE
      rm $(objects)

FORCE:
```

Here the target `'FORCE'` satisfies the special conditions, so the target `clean` that depends on it is forced to run its recipe. There is nothing special about the name `'FORCE'`, but that is one name commonly used this way.

As you can see, using `'FORCE'` this way has the same results as using `'.PHONY: clean'`.

Using `‘.PHONY’` is more explicit and more efficient. However, other versions of `make` do not support `‘.PHONY’`; thus `‘FORCE’` appears in many makefiles. See [Phony Targets](#).

Next: [Special Targets](#), Previous: [Force Targets](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.8 Empty Target Files to Record Events

The empty target is a variant of the phony target; it is used to hold recipes for an action that you request explicitly from time to time. Unlike a phony target, this target file can really exist; but the file’s contents do not matter, and usually are empty.

The purpose of the empty target file is to record, with its last-modification time, when the rule’s recipe was last executed. It does so because one of the commands in the recipe is a `touch` command to update the target file.

The empty target file should have some prerequisites (otherwise it doesn’t make sense). When you ask to remake the empty target, the recipe is executed if any prerequisite is more recent than the target; in other words, if a prerequisite has changed since the last time you remade the target. Here is an example:

```
print: foo.c bar.c
    lpr -p $?
    touch print
```

With this rule, `‘make print’` will execute the `lpr` command if either source file has changed since the last `‘make print’`. The automatic variable `‘$?’` is used to print only those files that have changed (see [Automatic Variables](#)).

Next: [Multiple Targets](#), Previous: [Empty Targets](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.9 Special Built-in Target Names

Certain names have special meanings if they appear as targets.

`.PHONY`

The prerequisites of the special target `.PHONY` are considered to be phony targets. When it is time to consider such a target, `make` will run its recipe unconditionally, regardless of whether a file with that name exists or what its last-modification time is. See [Phony Targets](#).

`.SUFFIXES`

The prerequisites of the special target `.SUFFIXES` are the list of suffixes to be used in checking for suffix rules. See [Old-Fashioned Suffix Rules](#).

`.DEFAULT`

The recipe specified for `.DEFAULT` is used for any target for which no rules are found (either explicit rules or implicit rules). See [Last Resort](#). If a `.DEFAULT` recipe is specified, every file mentioned as a prerequisite, but not as a target in a rule, will have that recipe executed on its behalf. See [Implicit Rule Search Algorithm](#).

`.PRECIOUS`

The targets which `.PRECIOUS` depends on are given the following special treatment: if `make` is killed or interrupted during the execution of their recipes, the target is not deleted. See [Interrupting or Killing make](#). Also, if the target is an intermediate file, it will not be deleted after it is no longer needed, as is normally done. See [Chains of Implicit Rules](#). In this latter respect it overlaps with the `.SECONDARY` special target.

You can also list the target pattern of an implicit rule (such as `‘%.o’`) as a prerequisite file of the special target `.PRECIOUS` to preserve intermediate files created by rules whose target patterns match that file’s name.

`.INTERMEDIATE`

The targets which `.INTERMEDIATE` depends on are treated as intermediate files. See [Chains of Implicit Rules](#). `.INTERMEDIATE` with no prerequisites has no effect.

`.SECONDARY`

The targets which `.SECONDARY` depends on are treated as intermediate files, except that they are never automatically deleted. See [Chains of Implicit Rules](#).

`.SECONDARY` with no prerequisites causes all targets to be treated as secondary (i.e., no target is removed because it is considered intermediate).

`.SECONDEXPANSION`

If `.SECONDEXPANSION` is mentioned as a target anywhere in the makefile, then all prerequisite lists defined after it appears will be expanded a second time after all makefiles have been read in. See [Secondary Expansion](#).

`.DELETE_ON_ERROR`

If `.DELETE_ON_ERROR` is mentioned as a target anywhere in the makefile, then `make` will delete the target of a rule if it has changed and its recipe exits with a nonzero exit status, just as it does when it receives a signal. See [Errors in Recipes](#).

`.IGNORE`

If you specify prerequisites for `.IGNORE`, then `make` will ignore errors in execution of the recipe for those particular files. The recipe for `.IGNORE` (if any) is ignored.

If mentioned as a target with no prerequisites, `.IGNORE` says to ignore errors in execution of recipes for all files. This usage of `' .IGNORE '` is supported only for historical compatibility. Since this affects every recipe in the makefile, it is not very useful; we recommend you use the more selective ways to ignore errors in specific recipes. See [Errors in Recipes](#).

`.LOW_RESOLUTION_TIME`

If you specify prerequisites for `.LOW_RESOLUTION_TIME`, `make` assumes that these files are created by commands that generate low resolution time stamps. The recipe for the `.LOW_RESOLUTION_TIME` target are ignored.

The high resolution file time stamps of many modern file systems lessen the chance of `make` incorrectly concluding that a file is up to date. Unfortunately, some hosts do not provide a way to set a high resolution file time stamp, so commands like `'cp -p'` that explicitly set a file's time stamp must discard its sub-second part. If a file is created by such a command, you should list it as a prerequisite of `.LOW_RESOLUTION_TIME` so that `make` does not mistakenly conclude that the file is out of date. For example:

```
.LOW_RESOLUTION_TIME: dst
dst: src
    cp -p src dst
```

Since `'cp -p'` discards the sub-second part of `src`'s time stamp, `dst` is typically slightly older than `src` even when it is up to date. The `.LOW_RESOLUTION_TIME` line causes `make` to consider `dst` to be up to date if its time stamp is at the start of the same second that `src`'s time stamp is in.

Due to a limitation of the archive format, archive member time stamps are always low resolution. You need not list archive members as prerequisites of `.LOW_RESOLUTION_TIME`, as `make` does this automatically.

`.SILENT`

If you specify prerequisites for `.SILENT`, then `make` will not print the recipe used to remake those particular files before executing them. The recipe for `.SILENT` is ignored.

If mentioned as a target with no prerequisites, `.SILENT` says not to print any recipes before executing them. You may also use more selective ways to silence specific recipe command lines. See [Recipe Echoing](#). If you want to silence all recipes for a particular run of `make`, use the `'-s'` or `'--silent'` option (see [Options Summary](#)).

`.EXPORT_ALL_VARIABLES`

Simply by being mentioned as a target, this tells `make` to export all variables to child processes by default. See [Communicating Variables to a Sub-make](#).

.NOTPARALLEL

If .NOTPARALLEL is mentioned as a target, then this invocation of `make` will be run serially, even if the `-j` option is given. Any recursively invoked `make` command will still run recipes in parallel (unless its makefile also contains this target). Any prerequisites on this target are ignored.

.ONESHELL

If .ONESHELL is mentioned as a target, then when a target is built all lines of the recipe will be given to a single invocation of the shell rather than each line being invoked separately (see [Recipe Execution](#)).

.POSIX

If .POSIX is mentioned as a target, then the makefile will be parsed and run in POSIX-conforming mode. This does not mean that only POSIX-conforming makefiles will be accepted: all advanced GNU `make` features are still available. Rather, this target causes `make` to behave as required by POSIX in those areas where `make`'s default behavior differs.

In particular, if this target is mentioned then recipes will be invoked as if the shell had been passed the `-e` flag: the first failing command in a recipe will cause the recipe to fail immediately.

Any defined implicit rule suffix also counts as a special target if it appears as a target, and so does the concatenation of two suffixes, such as `'.c.o'`. These targets are suffix rules, an obsolete way of defining implicit rules (but a way still widely used). In principle, any target name could be special in this way if you break it in two and add both pieces to the suffix list. In practice, suffixes normally begin with `'.'`, so these special target names also begin with `'.'`. See [Old-Fashioned Suffix Rules](#).

Next: [Multiple Rules](#), Previous: [Special Targets](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.10 Multiple Targets in a Rule

When an explicit rule has multiple targets they can be treated in one of two possible ways: as independent targets or as grouped targets. The manner in which they are treated is determined by the separator that appears after the list of targets.

Rules with Independent Targets

Rules that use the standard target separator, `:`, define independent targets. This is equivalent to writing the same rule once for each target, with duplicated prerequisites and recipes. Typically, the recipe would use automatic variables such as `'$@'` to specify which target is being built.

Rules with independent targets are useful in two cases:

- You want just prerequisites, no recipe. For example:

```
kbd.o command.o files.o: command.h
```

gives an additional prerequisite to each of the three object files mentioned. It is equivalent to writing:

```
kbd.o: command.h
command.o: command.h
files.o: command.h
```

- Similar recipes work for all the targets. The automatic variable `'$@'` can be used to substitute the particular target to be remade into the commands (see [Automatic Variables](#)). For example:

```
bigoutput littleoutput : text.g
generate text.g -$(subst output,, $@) > $@
```

is equivalent to

```
bigoutput : text.g
    generate text.g -big > bigoutput
littleoutput : text.g
    generate text.g -little > littleoutput
```

Here we assume the hypothetical program `generate` makes two types of output, one if given `'-big'` and one if given `'-little'`. See [Functions for String Substitution and Analysis](#), for an explanation of the `subst` function.

Suppose you would like to vary the prerequisites according to the target, much as the variable `'$@'` allows you to vary the recipe. You cannot do this with multiple targets in an ordinary rule, but you can do it with a static pattern rule. See [Static Pattern Rules](#).

Rules with Grouped Targets

If instead of independent targets you have a recipe that generates multiple files from a single invocation, you can express that relationship by declaring your rule to use grouped targets. A grouped target rule uses the separator `&:` (the `'&'` here is used to imply “all”).

When `make` builds any one of the grouped targets, it understands that all the other targets in the group are also created as a result of the invocation of the recipe. Furthermore, if only some of the grouped targets are out of date or missing `make` will realize that running the recipe will update all of the targets.

As an example, this rule defines a grouped target:

```
foo bar biz &: baz boz
    echo $^ > foo
    echo $^ > bar
    echo $^ > biz
```

During the execution of a grouped target’ s recipe, the automatic variable `'$@'` is set to the name of the particular target in the group which triggered the rule. Caution must be used if relying on this variable in the recipe of a grouped target rule.

Unlike independent targets, a grouped target rule must include a recipe. However, targets that are members of a grouped target may also appear in independent target rule definitions that do not have recipes.

Each target may have only one recipe associated with it. If a grouped target appears in either an independent target rule or in another grouped target rule with a recipe, you will get a warning and the latter recipe will replace the former recipe. Additionally the target will be removed from the previous group and appear only in the new group.

If you would like a target to appear in multiple groups, then you must use the double-colon grouped target separator, `&::` when declaring all of the groups containing that target. Grouped double-colon targets are each considered independently, and each grouped double-colon rule’ s recipe is executed at most once, if at least one of its multiple targets requires updating.

Next: [Static Pattern](#), Previous: [Multiple Targets](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.11 Multiple Rules for One Target

One file can be the target of several rules. All the prerequisites mentioned in all the rules are merged into one list of prerequisites for the target. If the target is older than any prerequisite from any rule, the recipe is executed.

There can only be one recipe to be executed for a file. If more than one rule gives a recipe for the same file, `make` uses the last one given and prints an error message. (As a special case, if the file’ s name begins with a dot, no error message is printed. This odd behavior is only for compatibility with other implementations of `make`... you should avoid using it). Occasionally it

is useful to have the same target invoke multiple recipes which are defined in different parts of your makefile; you can use double-colon rules (see [Double-Colon](#)) for this.

An extra rule with just prerequisites can be used to give a few extra prerequisites to many files at once. For example, makefiles often have a variable, such as `objects`, containing a list of all the compiler output files in the system being made. An easy way to say that all of them must be recompiled if `config.h` changes is to write the following:

```
objects = foo.o bar.o
foo.o : defs.h
bar.o : defs.h test.h
$(objects) : config.h
```

This could be inserted or taken out without changing the rules that really specify how to make the object files, making it a convenient form to use if you wish to add the additional prerequisite intermittently.

Another wrinkle is that the additional prerequisites could be specified with a variable that you set with a command line argument to `make` (see [Overriding Variables](#)). For example,

```
extradeps=
$(objects) : $(extradeps)
```

means that the command `'make extradeps=foo.h'` will consider `foo.h` as a prerequisite of each object file, but plain `'make'` will not.

If none of the explicit rules for a target has a recipe, then `make` searches for an applicable implicit rule to find one see [Using Implicit Rules](#)).

Next: [Double-Colon](#), Previous: [Multiple Rules](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.12 Static Pattern Rules

Static pattern rules are rules which specify multiple targets and construct the prerequisite names for each target based on the target name. They are more general than ordinary rules with multiple targets because the targets do not have to have identical prerequisites. Their prerequisites must be analogous, but not necessarily identical.

- [Static Usage](#) The syntax of static pattern rules.
- [Static versus Implicit](#) When are they better than implicit rules?

Next: [Static versus Implicit](#), Previous: [Static Pattern](#), Up: [Static Pattern](#) [[Contents](#)][[Index](#)]

4.12.1 Syntax of Static Pattern Rules

Here is the syntax of a static pattern rule:

```
targets ::: target-pattern: prereq-patterns ...
      recipe
      ...
```

The targets list specifies the targets that the rule applies to. The targets can contain wildcard characters, just like the targets of ordinary rules (see [Using Wildcard Characters in File Names](#)).

The target-pattern and prereq-patterns say how to compute the prerequisites of each target. Each target is matched against the target-pattern to extract a part of the target name, called the stem. This stem is substituted into each of the prereq-patterns to make the prerequisite names (one from each prereq-pattern).

Each pattern normally contains the character `'%'` just once. When the target-pattern matches a target, the `'%'` can match any part of the target name; this part is called the stem.

The rest of the pattern must match exactly. For example, the target `foo.o` matches the pattern `'%.o'`, with `'foo'` as the stem. The targets `foo.c` and `foo.out` do not match that pattern.

The prerequisite names for each target are made by substituting the stem for the `'%'` in each prerequisite pattern. For example, if one prerequisite pattern is `%.c`, then substitution of the stem `'foo'` gives the prerequisite name `foo.c`. It is legitimate to write a prerequisite pattern that does not contain `'%'`; then this prerequisite is the same for all targets.

`'%'` characters in pattern rules can be quoted with preceding backslashes (`'\'`). Backslashes that would otherwise quote `'%'` characters can be quoted with more backslashes. Backslashes that quote `'%'` characters or other backslashes are removed from the pattern before it is compared to file names or has a stem substituted into it. Backslashes that are not in danger of quoting `'%'` characters go unmolested. For example, the pattern `the\%weird\\%pattern\\` has `'the%weird\'` preceding the operative `'%'` character, and `'pattern\\'` following it. The final two backslashes are left alone because they cannot affect any `'%'` character.

Here is an example, which compiles each of `foo.o` and `bar.o` from the corresponding `.c` file:

```
objects = foo.o bar.o

all: $(objects)

$(objects): %.o: %.c
    $(CC) -c $(CFLAGS) $< -o $@
```

Here `'$<'` is the automatic variable that holds the name of the prerequisite and `'$@'` is the automatic variable that holds the name of the target; see [Automatic Variables](#).

Each target specified must match the target pattern; a warning is issued for each target that does not. If you have a list of files, only some of which will match the pattern, you can use the `filter` function to remove non-matching file names (see [Functions for String Substitution and Analysis](#)):

```
files = foo.elc bar.o lose.o

$(filter %.o,$(files)): %.o: %.c
    $(CC) -c $(CFLAGS) $< -o $@
$(filter %.elc,$(files)): %.elc: %.el
    emacs -f batch-byte-compile $<
```

In this example the result of `'$(filter %.o,$(files))'` is `bar.o lose.o`, and the first static pattern rule causes each of these object files to be updated by compiling the corresponding C source file. The result of `'$(filter %.elc,$(files))'` is `foo.elc`, so that file is made from `foo.el`.

Another example shows how to use `$$` in static pattern rules:

```
bigoutput littleoutput : %output : text.g
    generate text.g -$$ > $@
```

When the `generate` command is run, `$$` will expand to the stem, either `'big'` or `'little'`.

Previous: [Static Usage](#), Up: [Static Pattern](#) [[Contents](#)][[Index](#)]

4.12.2 Static Pattern Rules versus Implicit Rules

A static pattern rule has much in common with an implicit rule defined as a pattern rule (see [Defining and Redefining Pattern Rules](#)). Both have a pattern for the target and patterns for constructing the names of prerequisites. The difference is in how `make` decides when the rule applies.

An implicit rule can apply to any target that matches its pattern, but it does apply only when the target has no recipe otherwise specified, and only when the prerequisites can be found. If

more than one implicit rule appears applicable, only one applies; the choice depends on the order of rules.

By contrast, a static pattern rule applies to the precise list of targets that you specify in the rule. It cannot apply to any other target and it invariably does apply to each of the targets specified. If two conflicting rules apply, and both have recipes, that's an error.

The static pattern rule can be better than an implicit rule for these reasons:

- You may wish to override the usual implicit rule for a few files whose names cannot be categorized syntactically but can be given in an explicit list.
- If you cannot be sure of the precise contents of the directories you are using, you may not be sure which other irrelevant files might lead `make` to use the wrong implicit rule. The choice might depend on the order in which the implicit rule search is done. With static pattern rules, there is no uncertainty: each rule applies to precisely the targets specified.

Next: [Automatic Prerequisites](#), Previous: [Static Pattern](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.13 Double-Colon Rules

Double-colon rules are explicit rules written with `::` instead of `:` after the target names. They are handled differently from ordinary rules when the same target appears in more than one rule. Pattern rules with double-colons have an entirely different meaning (see [Match-Anything Rules](#)).

When a target appears in multiple rules, all the rules must be the same type: all ordinary, or all double-colon. If they are double-colon, each of them is independent of the others. Each double-colon rule's recipe is executed if the target is older than any prerequisites of that rule. If there are no prerequisites for that rule, its recipe is always executed (even if the target already exists). This can result in executing none, any, or all of the double-colon rules.

Double-colon rules with the same target are in fact completely separate from one another. Each double-colon rule is processed individually, just as rules with different targets are processed.

The double-colon rules for a target are executed in the order they appear in the makefile. However, the cases where double-colon rules really make sense are those where the order of executing the recipes would not matter.

Double-colon rules are somewhat obscure and not often very useful; they provide a mechanism for cases in which the method used to update a target differs depending on which prerequisite files caused the update, and such cases are rare.

Each double-colon rule should specify a recipe; if it does not, an implicit rule will be used if one applies. See [Using Implicit Rules](#).

Previous: [Double-Colon](#), Up: [Rules](#) [[Contents](#)][[Index](#)]

4.14 Generating Prerequisites Automatically

In the makefile for a program, many of the rules you need to write often say only that some object file depends on some header file. For example, if `main.c` uses `defs.h` via an `#include`, you would write:

```
main.o: defs.h
```

You need this rule so that `make` knows that it must remake `main.o` whenever `defs.h` changes. You can see that for a large program you would have to write dozens of such rules in your makefile. And, you must always be very careful to update the makefile every time you add or remove an `#include`.

To avoid this hassle, most modern C compilers can write these rules for you, by looking at the `#include` lines in the source files. Usually this is done with the `'-M'` option to the compiler. For example, the command:

```
cc -M main.c
```

generates the output:

```
main.o : main.c defs.h
```

Thus you no longer have to write all those rules yourself. The compiler will do it for you.

Note that such a rule constitutes mentioning `main.o` in a makefile, so it can never be considered an intermediate file by implicit rule search. This means that `make` won't ever remove the file after using it; see [Chains of Implicit Rules](#).

With old `make` programs, it was traditional practice to use this compiler feature to generate prerequisites on demand with a command like `'make depend'`. That command would create a file `depend` containing all the automatically-generated prerequisites; then the makefile could use `include` to read them in (see [Include](#)).

In GNU `make`, the feature of remaking makefiles makes this practice obsolete—you need never tell `make` explicitly to regenerate the prerequisites, because it always regenerates any makefile that is out of date. See [Remaking Makefiles](#).

The practice we recommend for automatic prerequisite generation is to have one makefile corresponding to each source file. For each source file `name.c` there is a makefile `name.d` which lists what files the object file `name.o` depends on. That way only the source files that have changed need to be rescanned to produce the new prerequisites.

Here is the pattern rule to generate a file of prerequisites (i.e., a makefile) called `name.d` from a C source file called `name.c`:

```
%.d: %.c
    @set -e; rm -f $@; \
    $(CC) -M $(CPPFLAGS) $< > $@.$$$$; \
    sed 's, \($*\)\.o[ :]*, \1.o $@ : ,g' < $@.$$$$ > $@; \
    rm -f $@.$$$$
```

See [Pattern Rules](#), for information on defining pattern rules. The `'-e'` flag to the shell causes it to exit immediately if the `$(CC)` command (or any other command) fails (exits with a nonzero status).

With the GNU C compiler, you may wish to use the `'-MM'` flag instead of `'-M'`. This omits prerequisites on system header files. See [Options Controlling the Preprocessor](#) in Using GNU CC, for details.

The purpose of the `sed` command is to translate (for example):

```
main.o : main.c defs.h
```

into:

```
main.o main.d : main.c defs.h
```

This makes each `'d'` file depend on all the source and header files that the corresponding `'o'` file depends on. `make` then knows it must regenerate the prerequisites whenever any of the source or header files changes.

Once you've defined the rule to remake the `'d'` files, you then use the `include` directive to read them all in. See [Include](#). For example:

```
sources = foo.c bar.c
```

include \$(sources:.c=.d)

(This example uses a substitution variable reference to translate the list of source files `'foo.c bar.c'` into a list of prerequisite makefiles, `'foo.d bar.d'`. See [Substitution Refs](#), for full information on substitution references.) Since the `.d` files are makefiles like any others, `make` will remake them as necessary with no further work from you. See [Remaking Makefiles](#).

Note that the `.d` files contain target definitions; you should be sure to place the `include` directive after the first, default goal in your makefiles or run the risk of having a random object file become the default goal. See [How Make Works](#).

Next: [Using Variables](#), Previous: [Rules](#), Up: [Top](#) [\[Contents\]](#)[\[Index\]](#)

5 Writing Recipes in Rules

The recipe of a rule consists of one or more shell command lines to be executed, one at a time, in the order they appear. Typically, the result of executing these commands is that the target of the rule is brought up to date.

Users use many different shell programs, but recipes in makefiles are always interpreted by `/bin/sh` unless the makefile specifies otherwise. See [Recipe Execution](#).

- [Recipe Syntax](#) Recipe syntax features and pitfalls.
- [Echoing](#) How to control when recipes are echoed.
- [Execution](#) How recipes are executed.
- [Parallel](#) How recipes can be executed in parallel.
- [Errors](#) What happens after a recipe execution error.
- [Interrupts](#) What happens when a recipe is interrupted.
- [Recursion](#) Invoking `make` from makefiles.
- [Canned Recipes](#) Defining canned recipes.
- [Empty Recipes](#) Defining useful, do-nothing recipes.

Next: [Echoing](#), Previous: [Recipes](#), Up: [Recipes](#) [\[Contents\]](#)[\[Index\]](#)

5.1 Recipe Syntax

Makefiles have the unusual property that there are really two distinct syntaxes in one file. Most of the makefile uses `make` syntax (see [Writing Makefiles](#)). However, recipes are meant to be interpreted by the shell and so they are written using shell syntax. The `make` program does not try to understand shell syntax: it performs only a very few specific translations on the content of the recipe before handing it to the shell.

Each line in the recipe must start with a tab (or the first character in the value of the `.RECIPEPREFIX` variable; see [Special Variables](#)), except that the first recipe line may be attached to the target-and-prerequisites line with a semicolon in between. Any line in the makefile that begins with a tab and appears in a “rule context” (that is, after a rule has been started until another rule or variable definition) will be considered part of a recipe for that rule. Blank lines and lines of just comments may appear among the recipe lines; they are ignored.

Some consequences of these rules include:

- A blank line that begins with a tab is not blank: it’s an empty recipe (see [Empty Recipes](#)).
- A comment in a recipe is not a `make` comment; it will be passed to the shell as-is. Whether the shell treats it as a comment or not depends on your shell.
- A variable definition in a “rule context” which is indented by a tab as the first character on the line, will be considered part of a recipe, not a `make` variable definition,

and passed to the shell.

- A conditional expression (`ifdef`, `ifeq`, etc. see [Syntax of Conditionals](#)) in a “rule context” which is indented by a tab as the first character on the line, will be considered part of a recipe and be passed to the shell.

- [Splitting Recipe Lines](#) Breaking long recipe lines for readability.
- [Variables in Recipes](#) Using `make` variables in recipes.

Next: [Variables in Recipes](#), Previous: [Recipe Syntax](#), Up: [Recipe Syntax](#) [\[Contents\]](#)[\[Index\]](#)

5.1.1 Splitting Recipe Lines

One of the few ways in which `make` does interpret recipes is checking for a backslash just before the newline. As in normal makefile syntax, a single logical recipe line can be split into multiple physical lines in the makefile by placing a backslash before each newline. A sequence of lines like this is considered a single recipe line, and one instance of the shell will be invoked to run it.

However, in contrast to how they are treated in other places in a makefile (see [Splitting Long Lines](#)), backslash/newline pairs are not removed from the recipe. Both the backslash and the newline characters are preserved and passed to the shell. How the backslash/newline is interpreted depends on your shell. If the first character of the next line after the backslash/newline is the recipe prefix character (a tab by default; see [Special Variables](#)), then that character (and only that character) is removed. Whitespace is never added to the recipe.

For example, the recipe for the `all` target in this makefile:

```
all :
    @echo no\
space
    @echo no\
    space
    @echo one \
    space
    @echo one\
    space
```

consists of four separate shell commands where the output is:

```
nospace
nospace
one space
one space
```

As a more complex example, this makefile:

```
all : ; @echo 'hello \
    world' ; echo "hello \
    world"
```

will invoke one shell with a command of:

```
echo 'hello \
world' ; echo "hello \
    world"
```

which, according to shell quoting rules, will yield the following output:

```
hello \
world
hello    world
```

Notice how the backslash/newline pair was removed inside the string quoted with double quotes (`"..."`), but not from the string quoted with single quotes (`' ... '`). This is the way the

default shell (`/bin/sh`) handles backslash/newline pairs. If you specify a different shell in your makefiles it may treat them differently.

Sometimes you want to split a long line inside of single quotes, but you don't want the backslash/newline to appear in the quoted content. This is often the case when passing scripts to languages such as Perl, where extraneous backslashes inside the script can change its meaning or even be a syntax error. One simple way of handling this is to place the quoted string, or even the entire command, into a `make` variable then use the variable in the recipe. In this situation the newline quoting rules for makefiles will be used, and the backslash/newline will be removed. If we rewrite our example above using this method:

```
HELLO = 'hello \
world'

all : ; @echo $(HELLO)
```

we will get output like this:

```
hello world
```

If you like, you can also use target-specific variables (see [Target-specific Variable Values](#)) to obtain a tighter correspondence between the variable and the recipe that uses it.

Previous: [Splitting Recipe Lines](#), Up: [Recipe Syntax](#) [[Contents](#)][[Index](#)]

5.1.2 Using Variables in Recipes

The other way in which `make` processes recipes is by expanding any variable references in them (see [Basics of Variable References](#)). This occurs after `make` has finished reading all the makefiles and the target is determined to be out of date; so, the recipes for targets which are not rebuilt are never expanded.

Variable and function references in recipes have identical syntax and semantics to references elsewhere in the makefile. They also have the same quoting rules: if you want a dollar sign to appear in your recipe, you must double it (`'$$'`). For shells like the default shell, that use dollar signs to introduce variables, it's important to keep clear in your mind whether the variable you want to reference is a `make` variable (use a single dollar sign) or a shell variable (use two dollar signs). For example:

```
LIST = one two three
all:
    for i in $(LIST); do \
        echo $$i; \
    done
```

results in the following command being passed to the shell:

```
for i in one two three; do \
    echo $i; \
done
```

which generates the expected result:

```
one
two
three
```

Next: [Execution](#), Previous: [Recipe Syntax](#), Up: [Recipes](#) [[Contents](#)][[Index](#)]

5.2 Recipe Echoing

Normally `make` prints each line of the recipe before it is executed. We call this echoing because it gives the appearance that you are typing the lines yourself.

When a line starts with `@`, the echoing of that line is suppressed. The `@` is discarded before the line is passed to the shell. Typically you would use this for a command whose only effect is to print something, such as an `echo` command to indicate progress through the makefile:

```
@echo About to make distribution files
```

When `make` is given the flag `-n` or `--just-print` it only echoes most recipes, without executing them. See [Summary of Options](#). In this case even the recipe lines starting with `@` are printed. This flag is useful for finding out which recipes `make` thinks are necessary without actually doing them.

The `-s` or `--silent` flag to `make` prevents all echoing, as if all recipes started with `@`. A rule in the makefile for the special target `.SILENT` without prerequisites has the same effect (see [Special Built-in Target Names](#)).

Next: [Parallel](#), Previous: [Echoing](#), Up: [Recipes](#) [[Contents](#)][[Index](#)]

5.3 Recipe Execution

When it is time to execute recipes to update a target, they are executed by invoking a new sub-shell for each line of the recipe, unless the `.ONESHELL` special target is in effect (see [Using One Shell](#)) (In practice, `make` may take shortcuts that do not affect the results.)

Please note: this implies that setting shell variables and invoking shell commands such as `cd` that set a context local to each process will not affect the following lines in the recipe.² If you want to use `cd` to affect the next statement, put both statements in a single recipe line. Then `make` will invoke one shell to run the entire line, and the shell will execute the statements in sequence. For example:

```
foo : bar/lose
    cd $(<D) && gobble $(<F) > ../$@
```

Here we use the shell AND operator (`&&`) so that if the `cd` command fails, the script will fail without trying to invoke the `gobble` command in the wrong directory, which could cause problems (in this case it would certainly cause `../foo` to be truncated, at least).

- [One Shell](#) One shell for all lines in a recipe.
- [Choosing the Shell](#) How `make` chooses the shell used to run recipes.

Next: [Choosing the Shell](#), Previous: [Execution](#), Up: [Execution](#) [[Contents](#)][[Index](#)]

5.3.1 Using One Shell

Sometimes you would prefer that all the lines in the recipe be passed to a single invocation of the shell. There are generally two situations where this is useful: first, it can improve performance in makefiles where recipes consist of many command lines, by avoiding extra processes. Second, you might want newlines to be included in your recipe command (for example perhaps you are using a very different interpreter as your `SHELL`). If the `.ONESHELL` special target appears anywhere in the makefile then all recipe lines for each target will be provided to a single invocation of the shell. Newlines between recipe lines will be preserved. For example:

```
.ONESHELL:
foo : bar/lose
    cd $(@D)
    gobble $(@F) > ../$@
```

would now work as expected even though the commands are on different recipe lines.

If `.ONESHELL` is provided, then only the first line of the recipe will be checked for the special prefix characters (`'@'` , `'-'` , and `'+'`). Subsequent lines will include the special characters in the recipe line when the `SHELL` is invoked. If you want your recipe to start with one of these special characters you' ll need to arrange for them to not be the first characters on the first line, perhaps by adding a comment or similar. For example, this would be a syntax error in Perl because the first `'@'` is removed by `make`:

```
.ONESHELL:
SHELL = /usr/bin/perl
.SHELLFLAGS = -e
show :
    @f = qw(a b c);
    print "@f\n";
```

However, either of these alternatives would work properly:

```
.ONESHELL:
SHELL = /usr/bin/perl
.SHELLFLAGS = -e
show :
    # Make sure "@" is not the first character on the first line
    @f = qw(a b c);
    print "@f\n";
```

or

```
.ONESHELL:
SHELL = /usr/bin/perl
.SHELLFLAGS = -e
show :
    my @f = qw(a b c);
    print "@f\n";
```

As a special feature, if `SHELL` is determined to be a POSIX-style shell, the special prefix characters in “internal” recipe lines will be removed before the recipe is processed. This feature is intended to allow existing makefiles to add the `.ONESHELL` special target and still run properly without extensive modifications. Since the special prefix characters are not legal at the beginning of a line in a POSIX shell script this is not a loss in functionality. For example, this works as expected:

```
.ONESHELL:
foo : bar/lose
    @cd $(@D)
    @gobble $(@F) > ../$@
```

Even with this special feature, however, makefiles with `.ONESHELL` will behave differently in ways that could be noticeable. For example, normally if any line in the recipe fails, that causes the rule to fail and no more recipe lines are processed. Under `.ONESHELL` a failure of any but the final recipe line will not be noticed by `make`. You can modify `.SHELLFLAGS` to add the `-e` option to the shell which will cause any failure anywhere in the command line to cause the shell to fail, but this could itself cause your recipe to behave differently. Ultimately you may need to harden your recipe lines to allow them to work with `.ONESHELL`.

Previous: [One Shell](#), Up: [Execution](#) [[Contents](#)][[Index](#)]

5.3.2 Choosing the Shell

The program used as the shell is taken from the variable `SHELL`. If this variable is not set in your makefile, the program `/bin/sh` is used as the shell. The argument(s) passed to the shell are taken from the variable `.SHELLFLAGS`. The default value of `.SHELLFLAGS` is `-c` normally, or `-ec` in POSIX-conforming mode.

Unlike most variables, the variable `SHELL` is never set from the environment. This is because the `SHELL` environment variable is used to specify your personal choice of shell program for interactive use. It would be very bad for personal choices like this to affect the functioning of makefiles. See [Variables from the Environment](#).

Furthermore, when you do set `SHELL` in your makefile that value is not exported in the environment to recipe lines that `make` invokes. Instead, the value inherited from the user's environment, if any, is exported. You can override this behavior by explicitly exporting `SHELL` (see [Communicating Variables to a Sub-make](#)), forcing it to be passed in the environment to recipe lines.

However, on MS-DOS and MS-Windows the value of `SHELL` in the environment **is** used, since on those systems most users do not set this variable, and therefore it is most likely set specifically to be used by `make`. On MS-DOS, if the setting of `SHELL` is not suitable for `make`, you can set the variable `MAKESHELL` to the shell that `make` should use; if set it will be used as the shell instead of the value of `SHELL`.

Choosing a Shell in DOS and Windows

Choosing a shell in MS-DOS and MS-Windows is much more complex than on other systems.

On MS-DOS, if `SHELL` is not set, the value of the variable `COMSPEC` (which is always set) is used instead.

The processing of lines that set the variable `SHELL` in Makefiles is different on MS-DOS. The stock shell, `command.com`, is ridiculously limited in its functionality and many users of `make` tend to install a replacement shell. Therefore, on MS-DOS, `make` examines the value of `SHELL`, and changes its behavior based on whether it points to a Unix-style or DOS-style shell. This allows reasonable functionality even if `SHELL` points to `command.com`.

If `SHELL` points to a Unix-style shell, `make` on MS-DOS additionally checks whether that shell can indeed be found; if not, it ignores the line that sets `SHELL`. In MS-DOS, GNU `make` searches for the shell in the following places:

1. In the precise place pointed to by the value of `SHELL`. For example, if the makefile specifies `'SHELL = /bin/sh'`, `make` will look in the directory `/bin` on the current drive.
2. In the current directory.
3. In each of the directories in the `PATH` variable, in order.

In every directory it examines, `make` will first look for the specific file (`sh` in the example above). If this is not found, it will also look in that directory for that file with one of the known extensions which identify executable files. For example `.exe`, `.com`, `.bat`, `.btm`, `.sh`, and some others.

If any of these attempts is successful, the value of `SHELL` will be set to the full pathname of the shell as found. However, if none of these is found, the value of `SHELL` will not be changed, and thus the line that sets it will be effectively ignored. This is so `make` will only support features specific to a Unix-style shell if such a shell is actually installed on the system where `make` runs.

Note that this extended search for the shell is limited to the cases where `SHELL` is set from the Makefile; if it is set in the environment or command line, you are expected to set it to the full pathname of the shell, exactly as things are on Unix.

The effect of the above DOS-specific processing is that a Makefile that contains `'SHELL = /bin/sh'` (as many Unix makefiles do), will work on MS-DOS unaltered if you have e.g. `sh.exe` installed in some directory along your `PATH`.

Next: [Errors](#), Previous: [Execution](#), Up: [Recipes](#) [[Contents](#)][[Index](#)]

5.4 Parallel Execution

GNU `make` knows how to execute several recipes at once. Normally, `make` will execute only one recipe at a time, waiting for it to finish before executing the next. However, the `'-j'` or `'--`

`jobs` option tells `make` to execute many recipes simultaneously. You can inhibit parallelism in a particular makefile with the `.NOTPARALLEL` pseudo-target (see [Special Built-in Target Names](#)).

On MS-DOS, the `-j` option has no effect, since that system doesn't support multi-processing.

If the `-j` option is followed by an integer, this is the number of recipes to execute at once; this is called the number of job slots. If there is nothing looking like an integer after the `-j` option, there is no limit on the number of job slots. The default number of job slots is one, which means serial execution (one thing at a time).

Handling recursive `make` invocations raises issues for parallel execution. For more information on this, see [Communicating Options to a Sub-`make`](#).

If a recipe fails (is killed by a signal or exits with a nonzero status), and errors are not ignored for that recipe (see [Errors in Recipes](#)), the remaining recipe lines to remake the same target will not be run. If a recipe fails and the `-k` or `--keep-going` option was not given (see [Summary of Options](#)), `make` aborts execution. If `make` terminates for any reason (including a signal) with child processes running, it waits for them to finish before actually exiting.

When the system is heavily loaded, you will probably want to run fewer jobs than when it is lightly loaded. You can use the `-l` option to tell `make` to limit the number of jobs to run at once, based on the load average. The `-l` or `--max-load` option is followed by a floating-point number. For example,

```
-l 2.5
```

will not let `make` start more than one job if the load average is above 2.5. The `-l` option with no following number removes the load limit, if one was given with a previous `-l` option.

More precisely, when `make` goes to start up a job, and it already has at least one job running, it checks the current load average; if it is not lower than the limit given with `-l`, `make` waits until the load average goes below that limit, or until all the other jobs finish.

By default, there is no load limit.

- [Parallel Output](#) Handling output during parallel execution
- [Parallel Input](#) Handling input during parallel execution

Next: [Parallel Input](#), Previous: [Parallel](#), Up: [Parallel](#) [[Contents](#)][[Index](#)]

5.4.1 Output During Parallel Execution

When running several recipes in parallel the output from each recipe appears as soon as it is generated, with the result that messages from different recipes may be interspersed, sometimes even appearing on the same line. This can make reading the output very difficult.

To avoid this you can use the `--output-sync` (`-O`) option. This option instructs `make` to save the output from the commands it invokes and print it all once the commands are completed. Additionally, if there are multiple recursive `make` invocations running in parallel, they will communicate so that only one of them is generating output at a time.

If working directory printing is enabled (see [The `--print-directory` Option](#)), the enter/leave messages are printed around each output grouping. If you prefer not to see these messages add the `--no-print-directory` option to `MAKEFLAGS`.

There are four levels of granularity when synchronizing output, specified by giving an argument to the option (e.g., `-Oline` or `--output-sync=recurse`).

`none`

This is the default: all output is sent directly as it is generated and no synchronization is performed.

`line`

Output from each individual line of the recipe is grouped and printed as soon as that line is complete. If a recipe consists of multiple lines, they may be interspersed with lines from other recipes.

`target`

Output from the entire recipe for each target is grouped and printed once the target is complete. This is the default if the `--output-sync` or `-O` option is given with no argument.

`recurse`

Output from each recursive invocation of `make` is grouped and printed once the recursive invocation is complete.

Regardless of the mode chosen, the total build time will be the same. The only difference is in how the output appears.

The `'target'` and `'recurse'` modes both collect the output of the entire recipe of a target and display it uninterrupted when the recipe completes. The difference between them is in how recipes that contain recursive invocations of `make` are treated (see [Recursive Use of `make`](#)). For all recipes which have no recursive lines, the `'target'` and `'recurse'` modes behave identically.

If the `'recurse'` mode is chosen, recipes that contain recursive `make` invocations are treated the same as other targets: the output from the recipe, including the output from the recursive `make`, is saved and printed after the entire recipe is complete. This ensures output from all the targets built by a given recursive `make` instance are grouped together, which may make the output easier to understand. However it also leads to long periods of time during the build where no output is seen, followed by large bursts of output. If you are not watching the build as it proceeds, but instead viewing a log of the build after the fact, this may be the best option for you.

If you are watching the output, the long gaps of quiet during the build can be frustrating. The `'target'` output synchronization mode detects when `make` is going to be invoked recursively, using the standard methods, and it will not synchronize the output of those lines. The recursive `make` will perform the synchronization for its targets and the output from each will be displayed immediately when it completes. Be aware that output from recursive lines of the recipe are not synchronized (for example if the recursive line prints a message before running `make`, that message will not be synchronized).

The `'line'` mode can be useful for front-ends that are watching the output of `make` to track when recipes are started and completed.

Some programs invoked by `make` may behave differently if they determine they're writing output to a terminal versus a file (often described as "interactive" vs. "non-interactive" modes). For example, many programs that can display colored output will not do so if they determine they are not writing to a terminal. If your makefile invokes a program like this then using the output synchronization options will cause the program to believe it's running in "non-interactive" mode even though the output will ultimately go to the terminal.

Previous: [Parallel Output](#), Up: [Parallel](#) [[Contents](#)][[Index](#)]

5.4.2 Input During Parallel Execution

Two processes cannot both take input from the same device at the same time. To make sure that only one recipe tries to take input from the terminal at once, `make` will invalidate the standard input streams of all but one running recipe. If another recipe attempts to read from standard input it will usually incur a fatal error (a `'Broken pipe'` signal).

It is unpredictable which recipe will have a valid standard input stream (which will come from the terminal, or wherever you redirect the standard input of `make`). The first recipe run will always get it first, and the first recipe started after that one finishes will get it next, and so on.

We will change how this aspect of `make` works if we find a better alternative. In the mean time, you should not rely on any recipe using standard input at all if you are using the parallel

execution feature; but if you are not using this feature, then standard input works normally in all recipes.

Next: [Interrupts](#), Previous: [Parallel](#), Up: [Recipes](#) [[Contents](#)][[Index](#)]

5.5 Errors in Recipes

After each shell invocation returns, `make` looks at its exit status. If the shell completed successfully (the exit status is zero), the next line in the recipe is executed in a new shell; after the last line is finished, the rule is finished.

If there is an error (the exit status is nonzero), `make` gives up on the current rule, and perhaps on all rules.

Sometimes the failure of a certain recipe line does not indicate a problem. For example, you may use the `mkdir` command to ensure that a directory exists. If the directory already exists, `mkdir` will report an error, but you probably want `make` to continue regardless.

To ignore errors in a recipe line, write a `'-'` at the beginning of the line's text (after the initial tab). The `'-'` is discarded before the line is passed to the shell for execution.

For example,

```
clean:
    -rm -f *.o
```

This causes `make` to continue even if `rm` is unable to remove a file.

When you run `make` with the `'-i'` or `'--ignore-errors'` flag, errors are ignored in all recipes of all rules. A rule in the makefile for the special target `.IGNORE` has the same effect, if there are no prerequisites. This is less flexible but sometimes useful.

When errors are to be ignored, because of either a `'-'` or the `'-i'` flag, `make` treats an error return just like success, except that it prints out a message that tells you the status code the shell exited with, and says that the error has been ignored.

When an error happens that `make` has not been told to ignore, it implies that the current target cannot be correctly remade, and neither can any other that depends on it either directly or indirectly. No further recipes will be executed for these targets, since their preconditions have not been achieved.

Normally `make` gives up immediately in this circumstance, returning a nonzero status. However, if the `'-k'` or `'--keep-going'` flag is specified, `make` continues to consider the other prerequisites of the pending targets, remaking them if necessary, before it gives up and returns nonzero status. For example, after an error in compiling one object file, `'make -k'` will continue compiling other object files even though it already knows that linking them will be impossible. See [Summary of Options](#).

The usual behavior assumes that your purpose is to get the specified targets up to date; once `make` learns that this is impossible, it might as well report the failure immediately. The `'-k'` option says that the real purpose is to test as many of the changes made in the program as possible, perhaps to find several independent problems so that you can correct them all before the next attempt to compile. This is why Emacs' `compile` command passes the `'-k'` flag by default.

Usually when a recipe line fails, if it has changed the target file at all, the file is corrupted and cannot be used—or at least it is not completely updated. Yet the file's time stamp says that it is now up to date, so the next time `make` runs, it will not try to update that file. The situation is just the same as when the shell is killed by a signal; see [Interrupts](#). So generally the right thing to do is to delete the target file if the recipe fails after beginning to change the file. `make` will do this if `.DELETE_ON_ERROR` appears as a target. This is almost always what you want `make` to do, but it is not historical practice; so for compatibility, you must explicitly request it.

Next: [Recursion](#), Previous: [Errors](#), Up: [Recipes](#) [[Contents](#)][[Index](#)]

5.6 Interrupting or Killing `make`

If `make` gets a fatal signal while a shell is executing, it may delete the target file that the recipe was supposed to update. This is done if the target file's last-modification time has changed since `make` first checked it.

The purpose of deleting the target is to make sure that it is remade from scratch when `make` is next run. Why is this? Suppose you type `Ctrl-c` while a compiler is running, and it has begun to write an object file `foo.o`. The `Ctrl-c` kills the compiler, resulting in an incomplete file whose last-modification time is newer than the source file `foo.c`. But `make` also receives the `Ctrl-c` signal and deletes this incomplete file. If `make` did not do this, the next invocation of `make` would think that `foo.o` did not require updating—resulting in a strange error message from the linker when it tries to link an object file half of which is missing.

You can prevent the deletion of a target file in this way by making the special target `.PRECIOUS` depend on it. Before remaking a target, `make` checks to see whether it appears on the prerequisites of `.PRECIOUS`, and thereby decides whether the target should be deleted if a signal happens. Some reasons why you might do this are that the target is updated in some atomic fashion, or exists only to record a modification-time (its contents do not matter), or must exist at all times to prevent other sorts of trouble.

Although `make` does its best to clean up there are certain situations in which cleanup is impossible. For example, `make` may be killed by an uncatchable signal. Or, one of the programs `make` invokes may be killed or crash, leaving behind an up-to-date but corrupt target file: `make` will not realize that this failure requires the target to be cleaned. Or `make` itself may encounter a bug and crash.

For these reasons it's best to write defensive recipes, which won't leave behind corrupted targets even if they fail. Most commonly these recipes create temporary files rather than updating the target directly, then rename the temporary file to the final target name. Some compilers already behave this way, so that you don't need to write a defensive recipe.

Next: [Canned Recipes](#), Previous: [Interrupts](#), Up: [Recipes](#) [[Contents](#)][[Index](#)]

5.7 Recursive Use of `make`

Recursive use of `make` means using `make` as a command in a makefile. This technique is useful when you want separate makefiles for various subsystems that compose a larger system. For example, suppose you have a sub-directory `subdir` which has its own makefile, and you would like the containing directory's makefile to run `make` on the sub-directory. You can do it by writing this:

```
subsystem:
    cd subdir && $(MAKE)
```

or, equivalently, this (see [Summary of Options](#)):

```
subsystem:
    $(MAKE) -C subdir
```

You can write recursive `make` commands just by copying this example, but there are many things to know about how they work and why, and about how the sub-`make` relates to the top-level `make`. You may also find it useful to declare targets that invoke recursive `make` commands as `‘.PHONY’` (for more discussion on when this is useful, see [Phony Targets](#)).

For your convenience, when GNU `make` starts (after it has processed any `-C` options) it sets the variable `CURDIR` to the pathname of the current working directory. This value is never touched by `make` again: in particular note that if you include files from other directories the value of `CURDIR` does not change. The value has the same precedence it would have if it were set in the makefile (by default, an environment variable `CURDIR` will not override this value). Note that setting this variable has no impact on the operation of `make` (it does not cause `make` to change its working directory, for example).

- [MAKE Variable](#) The special effects of using `'$(MAKE)'` .
- [Variables/Recursion](#) How to communicate variables to a sub-`make`.
- [Options/Recursion](#) How to communicate options to a sub-`make`.
- [-w Option](#) How the `'-w'` or `'--print-directory'` option helps debug use of recursive `make` commands.

Next: [Variables/Recursion](#), Previous: [Recursion](#), Up: [Recursion](#) [[Contents](#)][[Index](#)]

5.7.1 How the `MAKE` Variable Works

Recursive `make` commands should always use the variable `MAKE`, not the explicit command name `'make'` , as shown here:

```
subsystem:
    cd subdir && $(MAKE)
```

The value of this variable is the file name with which `make` was invoked. If this file name was `/bin/make`, then the recipe executed is `'cd subdir && /bin/make'` . If you use a special version of `make` to run the top-level makefile, the same special version will be executed for recursive invocations.

As a special feature, using the variable `MAKE` in the recipe of a rule alters the effects of the `'-t'` (`'--touch'`), `'-n'` (`'--just-print'`), or `'-q'` (`'--question'`) option. Using the `MAKE` variable has the same effect as using a `'+'` character at the beginning of the recipe line. See [Instead of Executing the Recipes](#). This special feature is only enabled if the `MAKE` variable appears directly in the recipe: it does not apply if the `MAKE` variable is referenced through expansion of another variable. In the latter case you must use the `'+'` token to get these special effects.

Consider the command `'make -t'` in the above example. (The `'-t'` option marks targets as up to date without actually running any recipes; see [Instead of Execution](#).) Following the usual definition of `'-t'` , a `'make -t'` command in the example would create a file named `subsystem` and do nothing else. What you really want it to do is run `'cd subdir && make -t'` ; but that would require executing the recipe, and `'-t'` says not to execute recipes.

The special feature makes this do what you want: whenever a recipe line of a rule contains the variable `MAKE`, the flags `'-t'` , `'-n'` and `'-q'` do not apply to that line. Recipe lines containing `MAKE` are executed normally despite the presence of a flag that causes most recipes not to be run. The usual `MAKEFLAGS` mechanism passes the flags to the sub-`make` (see [Communicating Options to a Sub-`make`](#)), so your request to touch the files, or print the recipes, is propagated to the subsystem.

Next: [Options/Recursion](#), Previous: [MAKE Variable](#), Up: [Recursion](#) [[Contents](#)][[Index](#)]

5.7.2 Communicating Variables to a Sub-`make`

Variable values of the top-level `make` can be passed to the sub-`make` through the environment by explicit request. These variables are defined in the sub-`make` as defaults, but they do not override variables defined in the makefile used by the sub-`make` unless you use the `'-e'` switch (see [Summary of Options](#)).

To pass down, or export, a variable, `make` adds the variable and its value to the environment for running each line of the recipe. The sub-`make`, in turn, uses the environment to initialize its table of variable values. See [Variables from the Environment](#).

Except by explicit request, `make` exports a variable only if it is either defined in the environment initially or set on the command line, and if its name consists only of letters, numbers, and

underscores. Some shells cannot cope with environment variable names consisting of characters other than letters, numbers, and underscores.

The value of the `make` variable `SHELL` is not exported. Instead, the value of the `SHELL` variable from the invoking environment is passed to the sub-`make`. You can force `make` to export its value for `SHELL` by using the `export` directive, described below. See [Choosing the Shell](#).

The special variable `MAKEFLAGS` is always exported (unless you `unexport` it). `MAKEFILES` is exported if you set it to anything.

`make` automatically passes down variable values that were defined on the command line, by putting them in the `MAKEFLAGS` variable. See [Options/Recursion](#).

Variables are not normally passed down if they were created by default by `make` (see [Variables Used by Implicit Rules](#)). The sub-`make` will define these for itself.

If you want to export specific variables to a sub-`make`, use the `export` directive, like this:

```
export variable ...
```

If you want to prevent a variable from being exported, use the `unexport` directive, like this:

```
unexport variable ...
```

In both of these forms, the arguments to `export` and `unexport` are expanded, and so could be variables or functions which expand to a (list of) variable names to be (un)exported.

As a convenience, you can define a variable and export it at the same time by doing:

```
export variable = value
```

has the same result as:

```
variable = value
export variable
```

and

```
export variable := value
```

has the same result as:

```
variable := value
export variable
```

Likewise,

```
export variable += value
```

is just like:

```
variable += value
export variable
```

See [Appending More Text to Variables](#).

You may notice that the `export` and `unexport` directives work in `make` in the same way they work in the shell, `sh`.

If you want all variables to be exported by default, you can use `export` by itself:

```
export
```

This tells `make` that variables which are not explicitly mentioned in an `export` or `unexport` directive should be exported. Any variable given in an `unexport` directive will still not be exported. If you use `export` by itself to export variables by default, variables whose names contain characters other than alphanumerics and underscores will not be exported unless specifically mentioned in an `export` directive.

The behavior elicited by an `export` directive by itself was the default in older versions of GNU `make`. If your makefiles depend on this behavior and you want to be compatible with old versions of `make`, you can write a rule for the special target `.EXPORT_ALL_VARIABLES` instead of using the `export` directive. This will be ignored by old `makes`, while the `export` directive will cause a syntax error.

Likewise, you can use `unexport` by itself to tell `make` not to export variables by default. Since this is the default behavior, you would only need to do this if `export` had been used by itself earlier (in an included makefile, perhaps). You **cannot** use `export` and `unexport` by themselves to have variables exported for some recipes and not for others. The last `export` or `unexport` directive that appears by itself determines the behavior for the entire run of `make`.

As a special feature, the variable `MAKELEVEL` is changed when it is passed down from level to level. This variable's value is a string which is the depth of the level as a decimal number. The value is `'0'` for the top-level `make`; `'1'` for a sub-`make`, `'2'` for a sub-sub-`make`, and so on. The incrementation happens when `make` sets up the environment for a recipe.

The main use of `MAKELEVEL` is to test it in a conditional directive (see [Conditional Parts of Makefiles](#)); this way you can write a makefile that behaves one way if run recursively and another way if run directly by you.

You can use the variable `MAKEFILES` to cause all sub-`make` commands to use additional makefiles. The value of `MAKEFILES` is a whitespace-separated list of file names. This variable, if defined in the outer-level makefile, is passed down through the environment; then it serves as a list of extra makefiles for the sub-`make` to read before the usual or specified ones. See [The Variable MAKEFILES](#).

Next: [-w Option](#), Previous: [Variables/Recursion](#), Up: [Recursion](#) [[Contents](#)][[Index](#)]

5.7.3 Communicating Options to a Sub-`make`

Flags such as `'-s'` and `'-k'` are passed automatically to the sub-`make` through the variable `MAKEFLAGS`. This variable is set up automatically by `make` to contain the flag letters that `make` received. Thus, if you do `'make -ks'` then `MAKEFLAGS` gets the value `'ks'`.

As a consequence, every sub-`make` gets a value for `MAKEFLAGS` in its environment. In response, it takes the flags from that value and processes them as if they had been given as arguments. See [Summary of Options](#).

Likewise variables defined on the command line are passed to the sub-`make` through `MAKEFLAGS`. Words in the value of `MAKEFLAGS` that contain `'='`, `make` treats as variable definitions just as if they appeared on the command line. See [Overriding Variables](#).

The options `'-C'`, `'-f'`, `'-o'`, and `'-W'` are not put into `MAKEFLAGS`; these options are not passed down.

The `'-j'` option is a special case (see [Parallel Execution](#)). If you set it to some numeric value `'N'` and your operating system supports it (most any UNIX system will; others typically won't), the parent `make` and all the sub-`makes` will communicate to ensure that there are only `'N'` jobs running at the same time between them all. Note that any job that is marked recursive (see [Instead of Executing Recipes](#)) doesn't count against the total jobs (otherwise we could get `'N'` sub-`makes` running and have no slots left over for any real work!)

If your operating system doesn't support the above communication, then no `'-j'` is added to `MAKEFLAGS`, so that sub-`makes` run in non-parallel mode. If the `'-j'` option were passed down to sub-`makes` you would get many more jobs running in parallel than you asked for. If you give `'-j'` with no numeric argument, meaning to run as many jobs as possible in parallel, this is passed down, since multiple infinities are no more than one.

If you do not want to pass the other flags down, you must change the value of `MAKEFLAGS`, like this:

```
subsystem:
    cd subdir && $(MAKE) MAKEFLAGS=
```

The command line variable definitions really appear in the variable `MAKEOVERRIDES`, and `MAKEFLAGS` contains a reference to this variable. If you do want to pass flags down normally, but don't want to pass down the command line variable definitions, you can reset `MAKEOVERRIDES` to empty, like this:

```
MAKEOVERRIDES =
```

This is not usually useful to do. However, some systems have a small fixed limit on the size of the environment, and putting so much information into the value of `MAKEFLAGS` can exceed it. If you see the error message `'Arg list too long'`, this may be the problem. (For strict compliance with POSIX.2, changing `MAKEOVERRIDES` does not affect `MAKEFLAGS` if the special target ``.POSIX'` appears in the makefile. You probably do not care about this.)

A similar variable `MFLAGS` exists also, for historical compatibility. It has the same value as `MAKEFLAGS` except that it does not contain the command line variable definitions, and it always begins with a hyphen unless it is empty (`MAKEFLAGS` begins with a hyphen only when it begins with an option that has no single-letter version, such as `'--warn-undefined-variables'`). `MFLAGS` was traditionally used explicitly in the recursive `make` command, like this:

```
subsystem:
    cd subdir && $(MAKE) $(MFLAGS)
```

but now `MAKEFLAGS` makes this usage redundant. If you want your makefiles to be compatible with old `make` programs, use this technique; it will work fine with more modern `make` versions too.

The `MAKEFLAGS` variable can also be useful if you want to have certain options, such as `'-k'` (see [Summary of Options](#)), set each time you run `make`. You simply put a value for `MAKEFLAGS` in your environment. You can also set `MAKEFLAGS` in a makefile, to specify additional flags that should also be in effect for that makefile. (Note that you cannot use `MFLAGS` this way. That variable is set only for compatibility; `make` does not interpret a value you set for it in any way.)

When `make` interprets the value of `MAKEFLAGS` (either from the environment or from a makefile), it first prepends a hyphen if the value does not already begin with one. Then it chops the value into words separated by blanks, and parses these words as if they were options given on the command line (except that `'-C'`, `'-f'`, `'-h'`, `'-o'`, `'-W'`, and their long-named versions are ignored; and there is no error for an invalid option).

If you do put `MAKEFLAGS` in your environment, you should be sure not to include any options that will drastically affect the actions of `make` and undermine the purpose of makefiles and of `make` itself. For instance, the `'-t'`, `'-n'`, and `'-q'` options, if put in one of these variables, could have disastrous consequences and would certainly have at least surprising and probably annoying effects.

If you'd like to run other implementations of `make` in addition to GNU `make`, and hence do not want to add GNU `make`-specific flags to the `MAKEFLAGS` variable, you can add them to the `GNUMAKEFLAGS` variable instead. This variable is parsed just before `MAKEFLAGS`, in the same way as `MAKEFLAGS`. When `make` constructs `MAKEFLAGS` to pass to a recursive `make` it will include all flags, even those taken from `GNUMAKEFLAGS`. As a result, after parsing `GNUMAKEFLAGS` GNU `make` sets this variable to the empty string to avoid duplicating flags during recursion.

It's best to use `GNUMAKEFLAGS` only with flags which won't materially change the behavior of your makefiles. If your makefiles require GNU `make` anyway then simply use `MAKEFLAGS`. Flags such as `'--no-print-directory'` or `'--output-sync'` may be appropriate for `GNUMAKEFLAGS`.

5.7.4 The ‘--print-directory’ Option

If you use several levels of recursive `make` invocations, the ‘-w’ or ‘--print-directory’ option can make the output a lot easier to understand by showing each directory as `make` starts processing it and as `make` finishes processing it. For example, if ‘`make -w`’ is run in the directory `/u/gnu/make`, `make` will print a line of the form:

```
make: Entering directory `/u/gnu/make'.
```

before doing anything else, and a line of the form:

```
make: Leaving directory `/u/gnu/make'.
```

when processing is completed.

Normally, you do not need to specify this option because ‘`make`’ does it for you: ‘-w’ is turned on automatically when you use the ‘-C’ option, and in sub-`make`s. `make` will not automatically turn on ‘-w’ if you also use ‘-s’, which says to be silent, or if you use ‘--no-print-directory’ to explicitly disable it.

Next: [Empty Recipes](#), Previous: [Recursion](#), Up: [Recipes](#) [[Contents](#)][[Index](#)]

5.8 Defining Canned Recipes

When the same sequence of commands is useful in making various targets, you can define it as a canned sequence with the `define` directive, and refer to the canned sequence from the recipes for those targets. The canned sequence is actually a variable, so the name must not conflict with other variable names.

Here is an example of defining a canned recipe:

```
define run-yacc =
yacc $(firstword $^)
mv y.tab.c $@
endef
```

Here `run-yacc` is the name of the variable being defined; `endef` marks the end of the definition; the lines in between are the commands. The `define` directive does not expand variable references and function calls in the canned sequence; the ‘\$’ characters, parentheses, variable names, and so on, all become part of the value of the variable you are defining. See [Defining Multi-Line Variables](#), for a complete explanation of `define`.

The first command in this example runs Yacc on the first prerequisite of whichever rule uses the canned sequence. The output file from Yacc is always named `y.tab.c`. The second command moves the output to the rule’s target file name.

To use the canned sequence, substitute the variable into the recipe of a rule. You can substitute it like any other variable (see [Basics of Variable References](#)). Because variables defined by `define` are recursively expanded variables, all the variable references you wrote inside the `define` are expanded now. For example:

```
foo.c : foo.y
      $(run-yacc)
```

‘`foo.y`’ will be substituted for the variable ‘`$^`’ when it occurs in `run-yacc`’s value, and ‘`foo.c`’ for ‘`$@`’.

This is a realistic example, but this particular one is not needed in practice because `make` has an implicit rule to figure out these commands based on the file names involved (see [Using Implicit Rules](#)).

In recipe execution, each line of a canned sequence is treated just as if the line appeared on its own in the rule, preceded by a tab. In particular, `make` invokes a separate sub-shell for each

line. You can use the special prefix characters that affect command lines (`'@'` , `'-'` , and `'+'`) on each line of a canned sequence. See [Writing Recipes in Rules](#). For example, using this canned sequence:

```
define frobnicate =
@echo "frobnicating target $@"
frob-step-1 $< -o $@-step-1
frob-step-2 $@-step-1 -o $@
endef
```

`make` will not echo the first line, the `echo` command. But it will echo the following two recipe lines.

On the other hand, prefix characters on the recipe line that refers to a canned sequence apply to every line in the sequence. So the rule:

```
frob.out: frob.in
    @$(frobnicate)
```

does not echo any recipe lines. (See [Recipe Echoing](#), for a full explanation of `'@'` .)

Previous: [Canned Recipes](#), Up: [Recipes](#) [[Contents](#)][[Index](#)]

5.9 Using Empty Recipes

It is sometimes useful to define recipes which do nothing. This is done simply by giving a recipe that consists of nothing but whitespace. For example:

```
target: ;
```

defines an empty recipe for `target`. You could also use a line beginning with a recipe prefix character to define an empty recipe, but this would be confusing because such a line looks empty.

You may be wondering why you would want to define a recipe that does nothing. One reason this is useful is to prevent a target from getting implicit recipes (from implicit rules or the `.DEFAULT` special target; see [Implicit Rules](#) and see [Defining Last-Resort Default Rules](#)).

Empty recipes can also be used to avoid errors for targets that will be created as a side-effect of another recipe: if the target does not exist the empty recipe ensures that `make` won't complain that it doesn't know how to build the target, and `make` will assume the target is out of date.

You may be inclined to define empty recipes for targets that are not actual files, but only exist so that their prerequisites can be remade. However, this is not the best way to do that, because the prerequisites may not be remade properly if the target file actually does exist. See [Phony Targets](#), for a better way to do this.

Next: [Conditionals](#), Previous: [Recipes](#), Up: [Top](#) [[Contents](#)][[Index](#)]

6 How to Use Variables

A variable is a name defined in a makefile to represent a string of text, called the variable's value. These values are substituted by explicit request into targets, prerequisites, recipes, and other parts of the makefile. (In some other versions of `make`, variables are called macros.)

Variables and functions in all parts of a makefile are expanded when read, except for in recipes, the right-hand sides of variable definitions using `'='` , and the bodies of variable definitions using the `define` directive.

Variables can represent lists of file names, options to pass to compilers, programs to run, directories to look in for source files, directories to write output in, or anything else you can imagine.

A variable name may be any sequence of characters not containing `':'`, `'#'`, `'='`, or whitespace. However, variable names containing characters other than letters, numbers, and underscores should be considered carefully, as in some shells they cannot be passed through the environment to a sub-`make` (see [Communicating Variables to a Sub-`make`](#)). Variable names beginning with `'.'` and an uppercase letter may be given special meaning in future versions of `make`.

Variable names are case-sensitive. The names `'foo'`, `'FOO'`, and `'Foo'` all refer to different variables.

It is traditional to use upper case letters in variable names, but we recommend using lower case letters for variable names that serve internal purposes in the makefile, and reserving upper case for parameters that control implicit rules or for parameters that the user should override with command options (see [Overriding Variables](#)).

A few variables have names that are a single punctuation character or just a few characters. These are the automatic variables, and they have particular specialized uses. See [Automatic Variables](#).

• Reference	How to use the value of a variable.
• Flavors	Variables come in two flavors.
• Advanced	Advanced features for referencing a variable.
• Values	All the ways variables get their values.
• Setting	How to set a variable in the makefile.
• Appending	How to append more text to the old value of a variable.
• Override Directive	How to set a variable in the makefile even if the user has set it with a command argument.
• Multi-Line	An alternate way to set a variable to a multi-line string.
• Undefine Directive	How to undefine a variable so that it appears as if it was never set.
• Environment	Variable values can come from the environment.
• Target-specific	Variable values can be defined on a per-target basis.
• Pattern-specific	Target-specific variable values can be applied to a group of targets that match a pattern.
• Suppressing Inheritance	Suppress inheritance of variables.
• Special Variables	Variables with special meaning or behavior.

Next: [Flavors](#), Previous: [Using Variables](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.1 Basics of Variable References

To substitute a variable's value, write a dollar sign followed by the name of the variable in parentheses or braces: either `'$(foo)'` or `'${foo}'` is a valid reference to the variable `foo`. This special significance of `'$'` is why you must write `'$$'` to have the effect of a single dollar sign in a file name or recipe.

Variable references can be used in any context: targets, prerequisites, recipes, most directives, and new variable values. Here is an example of a common case, where a variable holds the names of all the object files in a program:

```
objects = program.o foo.o utils.o
program : $(objects)
        cc -o program $(objects)

$(objects) : defs.h
```

Variable references work by strict textual substitution. Thus, the rule

```
foo = c
prog.o : prog.$(foo)
        $(foo)$(foo) -$(foo) prog.$(foo)
```

could be used to compile a C program `prog.c`. Since spaces before the variable value are ignored in variable assignments, the value of `foo` is precisely `'c'`. (Don't actually write your makefiles this way!)

A dollar sign followed by a character other than a dollar sign, open-parenthesis or open-brace treats that single character as the variable name. Thus, you could reference the variable `x` with `'$x'`. However, this practice can lead to confusion (e.g., `'$foo'` refers to the variable `f` followed by the string `oo`) so we recommend using parentheses or braces around all variables, even single-letter variables, unless omitting them gives significant readability improvements. One place where readability is often improved is automatic variables (see [Automatic Variables](#)).

Next: [Advanced](#), Previous: [Reference](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.2 The Two Flavors of Variables

There are two ways that a variable in GNU `make` can have a value; we call them the two flavors of variables. The two flavors are distinguished in how they are defined and in what they do when expanded.

The first flavor of variable is a recursively expanded variable. Variables of this sort are defined by lines using `'='` (see [Setting Variables](#)) or by the `define` directive (see [Defining Multi-Line Variables](#)). The value you specify is installed verbatim; if it contains references to other variables, these references are expanded whenever this variable is substituted (in the course of expanding some other string). When this happens, it is called recursive expansion.

For example,

```
foo = $(bar)
bar = $(ugh)
ugh = Huh?

all:;echo $(foo)
```

will echo `'Huh?'`: `'$(foo)'` expands to `'$(bar)'` which expands to `'$(ugh)'` which finally expands to `'Huh?'`.

This flavor of variable is the only sort supported by most other versions of `make`. It has its advantages and its disadvantages. An advantage (most would say) is that:

```
CFLAGS = $(include_dirs) -O
include_dirs = -Ifoo -Ibar
```

will do what was intended: when `'CFLAGS'` is expanded in a recipe, it will expand to `'-Ifoo -Ibar -O'`. A major disadvantage is that you cannot append something on the end of a variable, as in

```
CFLAGS = $(CFLAGS) -O
```

because it will cause an infinite loop in the variable expansion. (Actually `make` detects the infinite loop and reports an error.)

Another disadvantage is that any functions (see [Functions for Transforming Text](#)) referenced in the definition will be executed every time the variable is expanded. This makes `make` run slower; worse, it causes the `wildcard` and `shell` functions to give unpredictable results because you cannot easily control when they are called, or even how many times.

To avoid all the problems and inconveniences of recursively expanded variables, there is another flavor: simply expanded variables.

Simply expanded variables are defined by lines using `:=` or `::=` (see [Setting Variables](#)). Both forms are equivalent in GNU `make`; however only the `:=` form is described by the POSIX standard (support for `::=` was added to the POSIX standard in 2012, so older versions of `make` won't accept this form either).

The value of a simply expanded variable is scanned once and for all, expanding any references to other variables and functions, when the variable is defined. The actual value of the simply expanded variable is the result of expanding the text that you write. It does not contain any references to other variables; it contains their values as of the time this variable was defined. Therefore,

```
x := foo
y := $(x) bar
x := later
```

is equivalent to

```
y := foo bar
x := later
```

When a simply expanded variable is referenced, its value is substituted verbatim.

Here is a somewhat more complicated example, illustrating the use of `:=` in conjunction with the `shell` function. (See [The `shell` Function](#).) This example also shows use of the variable `MAKELEVEL`, which is changed when it is passed down from level to level. (See [Communicating Variables to a Sub-`make`](#), for information about `MAKELEVEL`.)

```
ifeq (0,${MAKELEVEL})
whoami    := $(shell whoami)
host-type := $(shell arch)
MAKE := ${MAKE} host-type=${host-type} whoami=${whoami}
endif
```

An advantage of this use of `:=` is that a typical 'descend into a directory' recipe then looks like this:

```
${subdirs}:
    ${MAKE} -C $@ all
```

Simply expanded variables generally make complicated makefile programming more predictable because they work like variables in most programming languages. They allow you to redefine a variable using its own value (or its value processed in some way by one of the expansion functions) and to use the expansion functions much more efficiently (see [Functions for Transforming Text](#)).

You can also use them to introduce controlled leading whitespace into variable values. Leading whitespace characters are discarded from your input before substitution of variable references and function calls; this means you can include leading spaces in a variable value by protecting them with variable references, like this:

```
nullstring :=
space := $(nullstring) # end of the line
```

Here the value of the variable `space` is precisely one space. The comment `'# end of the line'` is included here just for clarity. Since trailing space characters are not stripped from variable values, just a space at the end of the line would have the same effect (but be rather hard to read). If you put whitespace at the end of a variable value, it is a good idea to put a comment like that at the end of the line to make your intent clear. Conversely, if you do not want any whitespace characters at the end of your variable value, you must remember not to put a random comment on the end of the line after some whitespace, such as this:

```
dir := /foo/bar    # directory to put the frobs in
```

Here the value of the variable `dir` is `'/foo/bar '` (with four trailing spaces), which was probably not the intention. (Imagine something like `'$(dir)/file'` with this definition!)

There is another assignment operator for variables, `'?='`. This is called a conditional variable assignment operator, because it only has an effect if the variable is not yet defined. This statement:

```
F00 ?= bar
```

is exactly equivalent to this (see [The `origin` Function](#)):

```
ifeq ($(origin F00), undefined)
  F00 = bar
endif
```

Note that a variable set to an empty value is still defined, so `'?='` will not set that variable.

Next: [Values](#), Previous: [Flavors](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.3 Advanced Features for Reference to Variables

This section describes some advanced features you can use to reference variables in more flexible ways.

- [Substitution Refs](#) Referencing a variable with substitutions on the value.
- [Computed Names](#) Computing the name of the variable to refer to.

Next: [Computed Names](#), Previous: [Advanced](#), Up: [Advanced](#) [[Contents](#)][[Index](#)]

6.3.1 Substitution References

A substitution reference substitutes the value of a variable with alterations that you specify. It has the form `'$(var:a=b)'` (or `'${var:a=b}'`) and its meaning is to take the value of the variable `var`, replace every `a` at the end of a word with `b` in that value, and substitute the resulting string.

When we say “at the end of a word”, we mean that `a` must appear either followed by whitespace or at the end of the value in order to be replaced; other occurrences of `a` in the value are unaltered. For example:

```
foo := a.o b.o l.a c.o
bar := $(foo:.o=.c)
```

sets `'bar'` to `'a.c b.c l.a c.c'`. See [Setting Variables](#).

A substitution reference is shorthand for the `patsubst` expansion function (see [Functions for String Substitution and Analysis](#)): `'$(var:a=b)'` is equivalent to `'$(patsubst %a,%b,var)'`. We provide substitution references as well as `patsubst` for compatibility with other implementations of `make`.

Another type of substitution reference lets you use the full power of the `patsubst` function. It has the same form `'$(var:a=b)'` described above, except that now `a` must contain a single `'%'` character. This case is equivalent to `'$(patsubst a,b,$(var))'`. See [Functions for String Substitution and Analysis](#), for a description of the `patsubst` function.

For example:

```
foo := a.o b.o l.a c.o
```

```
bar := $(foo:%.o=%.c)
```

sets 'bar' to 'a.c b.c l.a c.c'.

Previous: [Substitution Refs](#), Up: [Advanced](#) [[Contents](#)][[Index](#)]

6.3.2 Computed Variable Names

Computed variable names are a complicated concept needed only for sophisticated makefile programming. For most purposes you need not consider them, except to know that making a variable with a dollar sign in its name might have strange results. However, if you are the type that wants to understand everything, or you are actually interested in what they do, read on.

Variables may be referenced inside the name of a variable. This is called a computed variable name or a nested variable reference. For example,

```
x = y
y = z
a := $( $(x))
```

defines `a` as `'z'`: the `'$(x)'` inside `'$( $(x))'` expands to `'y'`, so `'$( $(x))'` expands to `'$(y)'` which in turn expands to `'z'`. Here the name of the variable to reference is not stated explicitly; it is computed by expansion of `'$(x)'`. The reference `'$(x)'` here is nested within the outer variable reference.

The previous example shows two levels of nesting, but any number of levels is possible. For example, here are three levels:

```
x = y
y = z
z = u
a := $( $( $(x)))
```

Here the innermost `'$(x)'` expands to `'y'`, so `'$( $(x))'` expands to `'$(y)'` which in turn expands to `'z'`; now we have `'$(z)'`, which becomes `'u'`.

References to recursively-expanded variables within a variable name are re-expanded in the usual fashion. For example:

```
x = $(y)
y = z
z = Hello
a := $( $(x))
```

defines `a` as `'Hello'`: `'$( $(x))'` becomes `'$( $(y))'` which becomes `'$(z)'` which becomes `'Hello'`.

Nested variable references can also contain modified references and function invocations (see [Functions for Transforming Text](#)), just like any other reference. For example, using the `subst` function (see [Functions for String Substitution and Analysis](#)):

```
x = variable1
variable2 := Hello
y = $(subst 1,2,$(x))
z = y
a := $( $( $(z)))
```

eventually defines `a` as `'Hello'`. It is doubtful that anyone would ever want to write a nested reference as convoluted as this one, but it works: `'$( $( $(z))'` expands to `'$( $(y))'` which becomes `'$( $(subst 1,2,$(x))'`. This gets the value `'variable1'` from `x` and changes it by substitution to `'variable2'`, so that the entire string becomes `'$(variable2)'`, a simple variable reference whose value is `'Hello'`.

A computed variable name need not consist entirely of a single variable reference. It can contain several variable references, as well as some invariant text. For example,

```
a_dirs := dira dirb
l_dirs := dir1 dir2

a_files := filea fileb
l_files := file1 file2

ifeq "$(use_a)" "yes"
a1 := a
else
a1 := 1
endif

ifeq "$(use_dirs)" "yes"
df := dirs
else
df := files
endif

dirs := $($a1)_$(df))
```

will give `dirs` the same value as `a_dirs`, `l_dirs`, `a_files` or `l_files` depending on the settings of `use_a` and `use_dirs`.

Computed variable names can also be used in substitution references:

```
a_objects := a.o b.o c.o
l_objects := 1.o 2.o 3.o

sources := $($a1)_objects:.o=.c)
```

defines `sources` as either `'a.c b.c c.c'` or `'1.c 2.c 3.c'`, depending on the value of `a1`.

The only restriction on this sort of use of nested variable references is that they cannot specify part of the name of a function to be called. This is because the test for a recognized function name is done before the expansion of nested references. For example,

```
ifdef do_sort
func := sort
else
func := strip
endif

bar := a d b g q c

foo := $($func) $(bar))
```

attempts to give `'foo'` the value of the variable `'sort a d b g q c'` or `'strip a d b g q c'`, rather than giving `'a d b g q c'` as the argument to either the `sort` or the `strip` function. This restriction could be removed in the future if that change is shown to be a good idea.

You can also use computed variable names in the left-hand side of a variable assignment, or in a `define` directive, as in:

```
dir = foo
$(dir)_sources := $(wildcard $(dir)/*.c)
define $(dir)_print =
lpr $($dir)_sources
endef
```

This example defines the variables `'dir'`, `'foo_sources'`, and `'foo_print'`.

Note that nested variable references are quite different from recursively expanded variables (see [The Two Flavors of Variables](#)), though both are used together in complex ways when doing makefile programming.

6.4 How Variables Get Their Values

Variables can get values in several different ways:

- You can specify an overriding value when you run `make`. See [Overriding Variables](#).
- You can specify a value in the makefile, either with an assignment (see [Setting Variables](#)) or with a verbatim definition (see [Defining Multi-Line Variables](#)).
- Variables in the environment become `make` variables. See [Variables from the Environment](#).
- Several automatic variables are given new values for each rule. Each of these has a single conventional use. See [Automatic Variables](#).
- Several variables have constant initial values. See [Variables Used by Implicit Rules](#).

6.5 Setting Variables

To set a variable from the makefile, write a line starting with the variable name followed by `'='`, `':='`, or `::='`. Whatever follows the `'='`, `':='`, or `::='` on the line becomes the value. For example,

```
objects = main.o foo.o bar.o utils.o
```

defines a variable named `objects`. Whitespace around the variable name and immediately after the `'='` is ignored.

Variables defined with `'='` are recursively expanded variables. Variables defined with `':='` or `::='` are simply expanded variables; these definitions can contain variable references which will be expanded before the definition is made. See [The Two Flavors of Variables](#).

The variable name may contain function and variable references, which are expanded when the line is read to find the actual variable name to use.

There is no limit on the length of the value of a variable except the amount of memory on the computer. You can split the value of a variable into multiple physical lines for readability (see [Splitting Long Lines](#)).

Most variable names are considered to have the empty string as a value if you have never set them. Several variables have built-in initial values that are not empty, but you can set them in the usual ways (see [Variables Used by Implicit Rules](#)). Several special variables are set automatically to a new value for each rule; these are called the automatic variables (see [Automatic Variables](#)).

If you'd like a variable to be set to a value only if it's not already set, then you can use the shorthand operator `'?='` instead of `'='`. These two settings of the variable `'FOO'` are identical (see [The `origin` Function](#)):

```
FOO ?= bar
```

and

```
ifeq ($(origin FOO), undefined)
FOO = bar
endif
```

The shell assignment operator `'!='` can be used to execute a shell script and set a variable to its output. This operator first evaluates the right-hand side, then passes that result to the shell for execution. If the result of the execution ends in a newline, that one newline is

removed; all other newlines are replaced by spaces. The resulting string is then placed into the named recursively-expanded variable. For example:

```
hash != printf '\043'
file_list != find . -name '*.*'
```

If the result of the execution could produce a `$`, and you don't intend what follows that to be interpreted as a make variable or function reference, then you must replace every `$` with `$$` as part of the execution. Alternatively, you can set a simply expanded variable to the result of running a program using the `shell` function call. See [The `shell` Function](#). For example:

```
hash := $(shell printf '\043')
var := $(shell find . -name "*.*")
```

As with the `shell` function, the exit status of the just-invoked shell script is stored in the `$.SHELLSTATUS` variable.

Next: [Override Directive](#), Previous: [Setting](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.6 Appending More Text to Variables

Often it is useful to add more text to the value of a variable already defined. You do this with a line containing `'+='`, like this:

```
objects += another.o
```

This takes the value of the variable `objects`, and adds the text `'another.o'` to it (preceded by a single space, if it has a value already). Thus:

```
objects = main.o foo.o bar.o utils.o
objects += another.o
```

sets `objects` to `'main.o foo.o bar.o utils.o another.o'`.

Using `'+='` is similar to:

```
objects = main.o foo.o bar.o utils.o
objects := $(objects) another.o
```

but differs in ways that become important when you use more complex values.

When the variable in question has not been defined before, `'+='` acts just like normal `'='`: it defines a recursively-expanded variable. However, when there is a previous definition, exactly what `'+='` does depends on what flavor of variable you defined originally. See [The Two Flavors of Variables](#), for an explanation of the two flavors of variables.

When you add to a variable's value with `'+='`, `make` acts essentially as if you had included the extra text in the initial definition of the variable. If you defined it first with `':='` or `':::='`, making it a simply-expanded variable, `'+='` adds to that simply-expanded definition, and expands the new text before appending it to the old value just as `':='` does (see [Setting Variables](#), for a full explanation of `':='` or `':::='`). In fact,

```
variable := value
variable += more
```

is exactly equivalent to:

```
variable := value
variable := $(variable) more
```

On the other hand, when you use `'+='` with a variable that you defined first to be recursively-expanded using plain `'='`, `make` does something a bit different. Recall that when you define a recursively-expanded variable, `make` does not expand the value you set for variable and function references immediately. Instead it stores the text verbatim, and saves these variable and function references to be expanded later, when you refer to the new variable (see [The Two Flavors of Variables](#)). When you use `'+='` on a recursively-expanded variable, it is this unexpanded text to which `make` appends the new text you specify.

```
variable = value
variable += more
```

is roughly equivalent to:

```
temp = value
variable = $(temp) more
```

except that of course it never defines a variable called `temp`. The importance of this comes when the variable's old value contains variable references. Take this common example:

```
CFLAGS = $(includes) -O
...
CFLAGS += -pg # enable profiling
```

The first line defines the `CFLAGS` variable with a reference to another variable, `includes`. (`CFLAGS` is used by the rules for C compilation; see [Catalogue of Built-In Rules](#).) Using `'='` for the definition makes `CFLAGS` a recursively-expanded variable, meaning `'$(includes) -O'` is not expanded when `make` processes the definition of `CFLAGS`. Thus, `includes` need not be defined yet for its value to take effect. It only has to be defined before any reference to `CFLAGS`. If we tried to append to the value of `CFLAGS` without using `'+='`, we might do it like this:

```
CFLAGS := $(CFLAGS) -pg # enable profiling
```

This is pretty close, but not quite what we want. Using `':='` redefines `CFLAGS` as a simply-expanded variable; this means `make` expands the text `'$(CFLAGS) -pg'` before setting the variable. If `includes` is not yet defined, we get `'-O -pg'`, and a later definition of `includes` will have no effect. Conversely, by using `'+='` we set `CFLAGS` to the unexpanded value `'$(includes) -O -pg'`. Thus we preserve the reference to `includes`, so if that variable gets defined at any later point, a reference like `'$(CFLAGS)'` still uses its value.

Next: [Multi-Line](#), Previous: [Appending](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.7 The `override` Directive

If a variable has been set with a command argument (see [Overriding Variables](#)), then ordinary assignments in the makefile are ignored. If you want to set the variable in the makefile even though it was set with a command argument, you can use an `override` directive, which is a line that looks like this:

```
override variable = value
```

or

```
override variable := value
```

To append more text to a variable defined on the command line, use:

```
override variable += more text
```

See [Appending More Text to Variables](#).

Variable assignments marked with the `override` flag have a higher priority than all other assignments, except another `override`. Subsequent assignments or appends to this variable which are not marked `override` will be ignored.

The `override` directive was not invented for escalation in the war between makefiles and command arguments. It was invented so you can alter and add to values that the user specifies with command arguments.

For example, suppose you always want the `'-g'` switch when you run the C compiler, but you would like to allow the user to specify the other switches with a command argument just as usual. You could use this `override` directive:

```
override CFLAGS += -g
```

You can also use `override` directives with `define` directives. This is done as you might expect:

```
override define foo =  
bar  
endef
```

See [Defining Multi-Line Variables](#).

Next: [Undefine Directive](#), Previous: [Override Directive](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.8 Defining Multi-Line Variables

Another way to set the value of a variable is to use the `define` directive. This directive has an unusual syntax which allows newline characters to be included in the value, which is convenient for defining both canned sequences of commands (see [Defining Canned Recipes](#)), and also sections of makefile syntax to use with `eval` (see [Eval Function](#)).

The `define` directive is followed on the same line by the name of the variable being defined and an (optional) assignment operator, and nothing more. The value to give the variable appears on the following lines. The end of the value is marked by a line containing just the word `endef`.

Aside from this difference in syntax, `define` works just like any other variable definition. The variable name may contain function and variable references, which are expanded when the directive is read to find the actual variable name to use.

The final newline before the `endef` is not included in the value; if you want your value to contain a trailing newline you must include a blank line. For example in order to define a variable that contains a newline character you must use two empty lines, not one:

```
define newline  
  
endef
```

You may omit the variable assignment operator if you prefer. If omitted, `make` assumes it to be `'='` and creates a recursively-expanded variable (see [The Two Flavors of Variables](#)). When using a `'+='` operator, the value is appended to the previous value as with any other append operation: with a single space separating the old and new values.

You may nest `define` directives: `make` will keep track of nested directives and report an error if they are not all properly closed with `endef`. Note that lines beginning with the recipe prefix character are considered part of a recipe, so any `define` or `endef` strings appearing on such a line will not be considered `make` directives.

```
define two-lines  
echo foo  
echo $(bar)  
endef
```

When used in a recipe, the previous example is functionally equivalent to this:

```
two-lines = echo foo; echo $(bar)
```

since two commands separated by semicolon behave much like two separate shell commands. However, note that using two separate lines means `make` will invoke the shell twice, running an independent sub-shell for each line. See [Recipe Execution](#).

If you want variable definitions made with `define` to take precedence over command-line variable definitions, you can use the `override` directive together with `define`:

```
override define two-lines =  
foo  
$(bar)  
endef
```

See [The `override` Directive](#).

Next: [Environment](#), Previous: [Multi-Line](#), Up: [Using Variables](#) [\[Contents\]](#)[\[Index\]](#)

6.9 Undefined Variables

If you want to clear a variable, setting its value to empty is usually sufficient. Expanding such a variable will yield the same result (empty string) regardless of whether it was set or not. However, if you are using the `flavor` (see [Flavor Function](#)) and `origin` (see [Origin Function](#)) functions, there is a difference between a variable that was never set and a variable with an empty value. In such situations you may want to use the `undefine` directive to make a variable appear as if it was never set. For example:

```
foo := foo  
bar = bar  
  
undefine foo  
undefine bar  
  
$(info $(origin foo))  
$(info $(flavor bar))
```

This example will print “undefined” for both variables.

If you want to undefine a command-line variable definition, you can use the `override` directive together with `undefine`, similar to how this is done for variable definitions:

```
override undefine CFLAGS
```

Next: [Target-specific](#), Previous: [Undefine Directive](#), Up: [Using Variables](#) [\[Contents\]](#)[\[Index\]](#)

6.10 Variables from the Environment

Variables in `make` can come from the environment in which `make` is run. Every environment variable that `make` sees when it starts up is transformed into a `make` variable with the same name and value. However, an explicit assignment in the makefile, or with a command argument, overrides the environment. (If the `‘-e’` flag is specified, then values from the environment override assignments in the makefile. See [Summary of Options](#). But this is not recommended practice.)

Thus, by setting the variable `CFLAGS` in your environment, you can cause all C compilations in most makefiles to use the compiler switches you prefer. This is safe for variables with standard or conventional meanings because you know that no makefile will use them for other things. (Note this is not totally reliable; some makefiles set `CFLAGS` explicitly and therefore are not affected by the value in the environment.)

When `make` runs a recipe, variables defined in the makefile are placed into the environment of each shell. This allows you to pass values to `sub-make` invocations (see [Recursive Use of `make`](#)). By default, only variables that came from the environment or the command line are passed to recursive invocations. You can use the `export` directive to pass other variables. See [Communicating Variables to a Sub-`make`](#), for full details.

Other use of variables from the environment is not recommended. It is not wise for makefiles to depend for their functioning on environment variables set up outside their control, since this would cause different users to get different results from the same makefile. This is against the whole purpose of most makefiles.

Such problems would be especially likely with the variable `SHELL`, which is normally present in the environment to specify the user's choice of interactive shell. It would be very undesirable for this choice to affect `make`; so, `make` handles the `SHELL` environment variable in a special way; see [Choosing the Shell](#).

Next: [Pattern-specific](#), Previous: [Environment](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.11 Target-specific Variable Values

Variable values in `make` are usually global; that is, they are the same regardless of where they are evaluated (unless they're reset, of course). One exception to that is automatic variables (see [Automatic Variables](#)).

The other exception is target-specific variable values. This feature allows you to define different values for the same variable, based on the target that `make` is currently building. As with automatic variables, these values are only available within the context of a target's recipe (and in other target-specific assignments).

Set a target-specific variable value like this:

```
target ... : variable-assignment
```

Target-specific variable assignments can be prefixed with any or all of the special keywords `export`, `override`, or `private`; these apply their normal behavior to this instance of the variable only.

Multiple target values create a target-specific variable value for each member of the target list individually.

The variable-assignment can be any valid form of assignment; recursive (`'='`), simple (`':='` or `::='`), appending (`+='`), or conditional (`?='`). All variables that appear within the variable-assignment are evaluated within the context of the target: thus, any previously-defined target-specific variable values will be in effect. Note that this variable is actually distinct from any "global" value: the two variables do not have to have the same flavor (recursive vs. simple).

Target-specific variables have the same priority as any other makefile variable. Variables provided on the command line (and in the environment if the `'-e'` option is in force) will take precedence. Specifying the `override` directive will allow the target-specific variable value to be preferred.

There is one more special feature of target-specific variables: when you define a target-specific variable that variable value is also in effect for all prerequisites of this target, and all their prerequisites, etc. (unless those prerequisites override that variable with their own target-specific variable value). So, for example, a statement like this:

```
prog : CFLAGS = -g
prog : prog.o foo.o bar.o
```

will set `CFLAGS` to `'-g'` in the recipe for `prog`, but it will also set `CFLAGS` to `'-g'` in the recipes that create `prog.o`, `foo.o`, and `bar.o`, and any recipes which create their prerequisites.

Be aware that a given prerequisite will only be built once per invocation of `make`, at most. If the same file is a prerequisite of multiple targets, and each of those targets has a different value for the same target-specific variable, then the first target to be built will cause that prerequisite to be built and the prerequisite will inherit the target-specific value from the first target. It will ignore the target-specific values from any other targets.

Next: [Suppressing Inheritance](#), Previous: [Target-specific](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.12 Pattern-specific Variable Values

In addition to target-specific variable values (see [Target-specific Variable Values](#)), GNU `make` supports pattern-specific variable values. In this form, the variable is defined for any target that matches the pattern specified.

Set a pattern-specific variable value like this:

```
pattern ... : variable-assignment
```

where `pattern` is a %-pattern. As with target-specific variable values, multiple pattern values create a pattern-specific variable value for each pattern individually. The variable-assignment can be any valid form of assignment. Any command line variable setting will take precedence, unless `override` is specified.

For example:

```
%.o : CFLAGS = -O
```

will assign `CFLAGS` the value of `'-O'` for all targets matching the pattern `%.o`.

If a target matches more than one pattern, the matching pattern-specific variables with longer stems are interpreted first. This results in more specific variables taking precedence over the more generic ones, for example:

```
%.o: %.c
      $(CC) -c $(CFLAGS) $(CPPFLAGS) $< -o $@

lib/%.o: CFLAGS := -fPIC -g
%.o: CFLAGS := -g

all: foo.o lib/bar.o
```

In this example the first definition of the `CFLAGS` variable will be used to update `lib/bar.o` even though the second one also applies to this target. Pattern-specific variables which result in the same stem length are considered in the order in which they were defined in the makefile.

Pattern-specific variables are searched after any target-specific variables defined explicitly for that target, and before target-specific variables defined for the parent target.

Next: [Special Variables](#), Previous: [Pattern-specific](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.13 Suppressing Inheritance

As described in previous sections, `make` variables are inherited by prerequisites. This capability allows you to modify the behavior of a prerequisite based on which targets caused it to be rebuilt. For example, you might set a target-specific variable on a `debug` target, then running `'make debug'` will cause that variable to be inherited by all prerequisites of `debug`, while just running `'make all'` (for example) would not have that assignment.

Sometimes, however, you may not want a variable to be inherited. For these situations, `make` provides the `private` modifier. Although this modifier can be used with any variable assignment, it makes the most sense with target- and pattern-specific variables. Any variable marked `private` will be visible to its local target but will not be inherited by prerequisites of

that target. A global variable marked `private` will be visible in the global scope but will not be inherited by any target, and hence will not be visible in any recipe.

As an example, consider this makefile:

```
EXTRA_CFLAGS =

prog: private EXTRA_CFLAGS = -L/usr/local/lib
prog: a.o b.o
```

Due to the `private` modifier, `a.o` and `b.o` will not inherit the `EXTRA_CFLAGS` variable assignment from the `prog` target.

Previous: [Suppressing Inheritance](#), Up: [Using Variables](#) [[Contents](#)][[Index](#)]

6.14 Other Special Variables

GNU `make` supports some variables that have special properties.

`MAKEFILE_LIST`

Contains the name of each makefile that is parsed by `make`, in the order in which it was parsed. The name is appended just before `make` begins to parse the makefile. Thus, if the first thing a makefile does is examine the last word in this variable, it will be the name of the current makefile. Once the current makefile has used `include`, however, the last word will be the just-included makefile.

If a makefile named `Makefile` has this content:

```
name1 := $(lastword $(MAKEFILE_LIST))

include inc.mk

name2 := $(lastword $(MAKEFILE_LIST))

all:
    @echo name1 = $(name1)
    @echo name2 = $(name2)
```

then you would expect to see this output:

```
name1 = Makefile
name2 = inc.mk
```

`.DEFAULT_GOAL`

Sets the default goal to be used if no targets were specified on the command line (see [Arguments to Specify the Goals](#)). The `.DEFAULT_GOAL` variable allows you to discover the current default goal, restart the default goal selection algorithm by clearing its value, or to explicitly set the default goal. The following example illustrates these cases:

```
# Query the default goal.
ifeq ($(DEFAULT_GOAL),)
    $(warning no default goal is set)
endif

.PHONY: foo
foo: ; @echo $@

$(warning default goal is $(DEFAULT_GOAL))

# Reset the default goal.
.DEFAULT_GOAL :=

.PHONY: bar
bar: ; @echo $@

$(warning default goal is $(DEFAULT_GOAL))

# Set our own.
.DEFAULT_GOAL := foo
```

This makefile prints:

```
no default goal is set
default goal is foo
default goal is bar
foo
```

Note that assigning more than one target name to `.DEFAULT_GOAL` is invalid and will result in an error.

MAKE_RESTARTS

This variable is set only if this instance of `make` has restarted (see [How Makefiles Are Remade](#)): it will contain the number of times this instance has restarted. Note this is not the same as recursion (counted by the `MAKELEVEL` variable). You should not set, modify, or export this variable.

MAKE_TERMOUT

MAKE_TERMERR

When `make` starts it will check whether stdout and stderr will show their output on a terminal. If so, it will set `MAKE_TERMOUT` and `MAKE_TERMERR`, respectively, to the name of the terminal device (or `true` if this cannot be determined). If set these variables will be marked for export. These variables will not be changed by `make` and they will not be modified if already set.

These values can be used (particularly in combination with output synchronization (see [Output During Parallel Execution](#)) to determine whether `make` itself is writing to a terminal; they can be tested to decide whether to force recipe commands to generate colored output for example.

If you invoke a sub-`make` and redirect its stdout or stderr it is your responsibility to reset or unexport these variables as well, if your makefiles rely on them.

.RECIPEPREFIX

The first character of the value of this variable is used as the character `make` assumes is introducing a recipe line. If the variable is empty (as it is by default) that character is the standard tab character. For example, this is a valid makefile:

```
.RECIPEPREFIX = >
all:
> @echo Hello, world
```

The value of `.RECIPEPREFIX` can be changed multiple times; once set it stays in effect for all rules parsed until it is modified.

.VARIABLES

Expands to a list of the names of all global variables defined so far. This includes variables which have empty values, as well as built-in variables (see [Variables Used by Implicit Rules](#)), but does not include any variables which are only defined in a target-specific context. Note that any value you assign to this variable will be ignored; it will always return its special value.

.FEATURES

Expands to a list of special features supported by this version of `make`. Possible values include, but are not limited to:

`'archives'`

Supports `ar` (archive) files using special file name syntax. See [Using `make` to Update Archive Files](#).

`'check-symlink'`

Supports the `-L` (`--check-symlink-times`) flag. See [Summary of Options](#).

`'else-if'`

Supports “else if” non-nested conditionals. See [Syntax of Conditionals](#).

`‘jobserver’`

Supports “job server” enhanced parallel builds. See [Parallel Execution](#).

`‘oneshell’`

Supports the `.ONESHELL` special target. See [Using One Shell](#).

`‘order-only’`

Supports order-only prerequisites. See [Types of Prerequisites](#).

`‘second-expansion’`

Supports secondary expansion of prerequisite lists.

`‘shortest-stem’`

Uses the “shortest stem” method of choosing which pattern, of multiple applicable options, will be used. See [How Patterns Match](#).

`‘target-specific’`

Supports target-specific and pattern-specific variable assignments. See [Target-specific Variable Values](#).

`‘undefine’`

Supports the `undefine` directive. See [Undefine Directive](#).

`‘guile’`

Has GNU Guile available as an embedded extension language. See [GNU Guile Integration](#).

`‘load’`

Supports dynamically loadable objects for creating custom extensions. See [Loading Dynamic Objects](#).

`.INCLUDE_DIRS`

Expands to a list of directories that `make` searches for included makefiles (see [Including Other Makefiles](#)).

`.EXTRA_PREREQS`

Each word in this variable is a new prerequisite which is added to targets for which it is set. These prerequisites differ from normal prerequisites in that they do not appear in any of the automatic variables (see [Automatic Variables](#)). This allows prerequisites to be defined which do not impact the recipe.

Consider a rule to link a program:

```
myprog: myprog.o file1.o file2.o
        $(CC) $(CFLAGS) $(LDFLAGS) -o $@ $^ $(LDLIBS)
```

Now suppose you want to enhance this makefile to ensure that updates to the compiler cause the program to be re-linked. You can add the compiler as a prerequisite, but you must ensure that it’s not passed as an argument to link command. You’ll need something like this:

```
myprog: myprog.o file1.o file2.o $(CC)
        $(CC) $(CFLAGS) $(LDFLAGS) -o $@ $(filter-out $(CC),$^) $(LDLIBS)
```

Then consider having multiple extra prerequisites: they would all have to be filtered out. Using `.EXTRA_PREREQS` and target-specific variables provides a simpler solution:

```
myprog: myprog.o file1.o file2.o
        $(CC) $(CFLAGS) $(LDFLAGS) -o $@ $^ $(LDLIBS)
myprog: .EXTRA_PREREQS = $(CC)
```

This feature can also be useful if you want to add prerequisites to a makefile you cannot easily modify: you can create a new file such as `extra.mk`:

```
myprog: .EXTRA_PREREQS = $(CC)
```

then invoke `make -f extra.mk -f Makefile`.

Setting `.EXTRA_PREREQS` globally will cause those prerequisites to be added to all targets (which did not themselves override it with a target-specific value). Note `make` is smart enough not to add a prerequisite listed in `.EXTRA_PREREQS` as a prerequisite to itself.

Next: [Functions](#), Previous: [Using Variables](#), Up: [Top](#) [[Contents](#)][[Index](#)]

7 Conditional Parts of Makefiles

A conditional directive causes part of a makefile to be obeyed or ignored depending on the values of variables. Conditionals can compare the value of one variable to another, or the value of a variable to a constant string. Conditionals control what `make` actually “sees” in the makefile, so they cannot be used to control recipes at the time of execution.

- [Conditional Example](#) Example of a conditional
- [Conditional Syntax](#) The syntax of conditionals.
- [Testing Flags](#) Conditionals that test flags.

Next: [Conditional Syntax](#), Previous: [Conditionals](#), Up: [Conditionals](#) [[Contents](#)][[Index](#)]

7.1 Example of a Conditional

The following example of a conditional tells `make` to use one set of libraries if the `CC` variable is `'gcc'`, and a different set of libraries otherwise. It works by controlling which of two recipe lines will be used for the rule. The result is that `'CC=gcc'` as an argument to `make` changes not only which compiler is used but also which libraries are linked.

```
libs_for_gcc = -lgnu
normal_libs =

foo: $(objects)
ifeq ($(CC), gcc)
    $(CC) -o foo $(objects) $(libs_for_gcc)
else
    $(CC) -o foo $(objects) $(normal_libs)
endif
```

This conditional uses three directives: one `ifeq`, one `else` and one `endif`.

The `ifeq` directive begins the conditional, and specifies the condition. It contains two arguments, separated by a comma and surrounded by parentheses. Variable substitution is performed on both arguments and then they are compared. The lines of the makefile following the `ifeq` are obeyed if the two arguments match; otherwise they are ignored.

The `else` directive causes the following lines to be obeyed if the previous conditional failed. In the example above, this means that the second alternative linking command is used whenever the first alternative is not used. It is optional to have an `else` in a conditional.

The `endif` directive ends the conditional. Every conditional must end with an `endif`. Unconditional makefile text follows.

As this example illustrates, conditionals work at the textual level: the lines of the conditional are treated as part of the makefile, or ignored, according to the condition. This is why the larger syntactic units of the makefile, such as rules, may cross the beginning or the end of the conditional.

When the variable `CC` has the value `'gcc'`, the above example has this effect:

```
foo: $(objects)
      $(CC) -o foo $(objects) $(libs_for_gcc)
```

When the variable `CC` has any other value, the effect is this:

```
foo: $(objects)
      $(CC) -o foo $(objects) $(normal_libs)
```

Equivalent results can be obtained in another way by conditionalizing a variable assignment and then using the variable unconditionally:

```
libs_for_gcc = -lgnu
normal_libs =

ifeq ($(CC), gcc)
    libs=$(libs_for_gcc)
else
    libs=$(normal_libs)
endif

foo: $(objects)
      $(CC) -o foo $(objects) $(libs)
```

Next: [Testing Flags](#), Previous: [Conditional Example](#), Up: [Conditionals](#) [\[Contents\]](#)[\[Index\]](#)

7.2 Syntax of Conditionals

The syntax of a simple conditional with no `else` is as follows:

```
conditional-directive
text-if-true
endif
```

The `text-if-true` may be any lines of text, to be considered as part of the makefile if the condition is true. If the condition is false, no text is used instead.

The syntax of a complex conditional is as follows:

```
conditional-directive
text-if-true
else
text-if-false
endif
```

or:

```
conditional-directive-one
text-if-one-is-true
else conditional-directive-two
text-if-two-is-true
else
text-if-one-and-two-are-false
endif
```

There can be as many “`else conditional-directive`” clauses as necessary. Once a given condition is true, `text-if-true` is used and no other clause is used; if no condition is true then `text-if-false` is used. The `text-if-true` and `text-if-false` can be any number of lines of text.

The syntax of the `conditional-directive` is the same whether the conditional is simple or complex; after an `else` or not. There are four different directives that test different conditions. Here is a table of them:

```
ifeq (arg1, arg2)
ifeq 'arg1' 'arg2'
ifeq "arg1" "arg2"
ifeq "arg1" 'arg2'
```

```
ifeq 'arg1' "arg2"
```

Expand all variable references in `arg1` and `arg2` and compare them. If they are identical, the `text-if-true` is effective; otherwise, the `text-if-false`, if any, is effective.

Often you want to test if a variable has a non-empty value. When the value results from complex expansions of variables and functions, expansions you would consider empty may actually contain whitespace characters and thus are not seen as empty. However, you can use the `strip` function (see [Text Functions](#)) to avoid interpreting whitespace as a non-empty value. For example:

```
ifeq ($(strip $(foo)),)
text-if-empty
endif
```

will evaluate `text-if-empty` even if the expansion of `$(foo)` contains whitespace characters.

```
ifneq (arg1, arg2)
ifneq 'arg1' 'arg2'
ifneq "arg1" "arg2"
ifneq "arg1" 'arg2'
ifneq 'arg1' "arg2"
```

Expand all variable references in `arg1` and `arg2` and compare them. If they are different, the `text-if-true` is effective; otherwise, the `text-if-false`, if any, is effective.

```
ifdef variable-name
```

The `ifdef` form takes the name of a variable as its argument, not a reference to a variable. If the value of that variable has a non-empty value, the `text-if-true` is effective; otherwise, the `text-if-false`, if any, is effective. Variables that have never been defined have an empty value. The text `variable-name` is expanded, so it could be a variable or function that expands to the name of a variable. For example:

```
bar = true
foo = bar
ifdef $(foo)
frobozz = yes
endif
```

The variable reference `$(foo)` is expanded, yielding `bar`, which is considered to be the name of a variable. The variable `bar` is not expanded, but its value is examined to determine if it is non-empty.

Note that `ifdef` only tests whether a variable has a value. It does not expand the variable to see if that value is nonempty. Consequently, tests using `ifdef` return true for all definitions except those like `foo =`. To test for an empty value, use `ifeq $(foo),)`. For example,

```
bar =
foo = $(bar)
ifdef foo
frobozz = yes
else
frobozz = no
endif
```

sets `'frobozz'` to `'yes'` , while:

```
foo =
ifdef foo
frobozz = yes
else
frobozz = no
endif
```

sets `'frobozz'` to `'no'` .

```
ifndef variable-name
```

If the variable `variable-name` has an empty value, the `text-if-true` is effective; otherwise, the `text-if-false`, if any, is effective. The rules for expansion and testing of `variable-name` are identical to the `ifdef` directive.

Extra spaces are allowed and ignored at the beginning of the conditional directive line, but a tab is not allowed. (If the line begins with a tab, it will be considered part of a recipe for a rule.) Aside from this, extra spaces or tabs may be inserted with no effect anywhere except within the directive name or within an argument. A comment starting with `#` may appear at the end of the line.

The other two directives that play a part in a conditional are `else` and `endif`. Each of these directives is written as one word, with no arguments. Extra spaces are allowed and ignored at the beginning of the line, and spaces or tabs at the end. A comment starting with `#` may appear at the end of the line.

Conditionals affect which lines of the makefile `make` uses. If the condition is true, `make` reads the lines of the `text-if-true` as part of the makefile; if the condition is false, `make` ignores those lines completely. It follows that syntactic units of the makefile, such as rules, may safely be split across the beginning or the end of the conditional.

`make` evaluates conditionals when it reads a makefile. Consequently, you cannot use automatic variables in the tests of conditionals because they are not defined until recipes are run (see [Automatic Variables](#)).

To prevent intolerable confusion, it is not permitted to start a conditional in one makefile and end it in another. However, you may write an `include` directive within a conditional, provided you do not attempt to terminate the conditional inside the included file.

Previous: [Conditional Syntax](#), Up: [Conditionals](#) [[Contents](#)][[Index](#)]

7.3 Conditionals that Test Flags

You can write a conditional that tests `make` command flags such as `-t` by using the variable `MAKEFLAGS` together with the `findstring` function (see [Functions for String Substitution and Analysis](#)). This is useful when `touch` is not enough to make a file appear up to date.

The `findstring` function determines whether one string appears as a substring of another. If you want to test for the `-t` flag, use `t` as the first string and the value of `MAKEFLAGS` as the other.

For example, here is how to arrange to use `ranlib -t` to finish marking an archive file up to date:

```
archive.a: ...
ifneq (,$(findstring t,$(MAKEFLAGS)))
    +touch archive.a
    +ranlib -t archive.a
else
    ranlib archive.a
endif
```

The `+` prefix marks those recipe lines as “recursive” so that they will be executed despite use of the `-t` flag. See [Recursive Use of `make`](#).

Next: [Running](#), Previous: [Conditionals](#), Up: [Top](#) [[Contents](#)][[Index](#)]

8 Functions for Transforming Text

Functions allow you to do text processing in the makefile to compute the files to operate on or the commands to use in recipes. You use a function in a function call, where you give the name of the function and some text (the arguments) for the function to operate on. The

result of the function’ s processing is substituted into the makefile at the point of the call, just as a variable might be substituted.

• Syntax of Functions	How to write a function call.
• Text Functions	General-purpose text manipulation functions.
• File Name Functions	Functions for manipulating file names.
• Conditional Functions	Functions that implement conditions.
• Foreach Function	Repeat some text with controlled variation.
• File Function	Write text to a file.
• Call Function	Expand a user-defined function.
• Value Function	Return the un-expanded value of a variable.
• Eval Function	Evaluate the arguments as makefile syntax.
• Origin Function	Find where a variable got its value.
• Flavor Function	Find out the flavor of a variable.
• Make Control Functions	Functions that control how make runs.
• Shell Function	Substitute the output of a shell command.
• Guile Function	Use GNU Guile embedded scripting language.

Next: [Text Functions](#), Previous: [Functions](#), Up: [Functions](#) [[Contents](#)][[Index](#)]

8.1 Function Call Syntax

A function call resembles a variable reference. It can appear anywhere a variable reference can appear, and it is expanded using the same rules as variable references. A function call looks like this:

```
$(function arguments)
```

or like this:

```
${function arguments}
```

Here `function` is a function name; one of a short list of names that are part of `make`. You can also essentially create your own functions by using the `call` built-in function.

The arguments are the arguments of the function. They are separated from the function name by one or more spaces or tabs, and if there is more than one argument, then they are separated by commas. Such whitespace and commas are not part of an argument’ s value. The delimiters which you use to surround the function call, whether parentheses or braces, can appear in an argument only in matching pairs; the other kind of delimiters may appear singly. If the arguments themselves contain other function calls or variable references, it is wisest to use the same kind of delimiters for all the references; write `‘$(subst a,b,$(x))’` , not `‘$(subst a,b,{x})’` . This is because it is clearer, and because only one type of delimiter is matched to find the end of the reference.

The text written for each argument is processed by substitution of variables and function calls to produce the argument value, which is the text on which the function acts. The substitution is done in the order in which the arguments appear.

Commas and unmatched parentheses or braces cannot appear in the text of an argument as written; leading spaces cannot appear in the text of the first argument as written. These characters can be put into the argument value by variable substitution. First define variables `comma` and `space` whose values are isolated comma and space characters, then substitute these variables where such characters are wanted, like this:

```
comma:= ,
empty:=
space:= $(empty) $(empty)
foo:= a b c
```

```
bar:= $(subst $(space),$(comma),$(foo))
# bar is now  'a,b,c' .
```

Here the `subst` function replaces each space with a comma, through the value of `foo`, and substitutes the result.

Next: [File Name Functions](#), Previous: [Syntax of Functions](#), Up: [Functions](#) [\[Contents\]](#)[\[Index\]](#)

8.2 Functions for String Substitution and Analysis

Here are some functions that operate on strings:

```
$(subst from,to,text)
```

Performs a textual replacement on the text `text`: each occurrence of `from` is replaced by `to`. The result is substituted for the function call. For example,

```
$(subst ee,EE,feet on the street)
```

produces the value `'fEEt on the strEEt'` .

```
$(patsubst pattern,replacement,text)
```

Finds whitespace-separated words in `text` that match `pattern` and replaces them with `replacement`. Here `pattern` may contain a `'%'` which acts as a wildcard, matching any number of any characters within a word. If `replacement` also contains a `'%'` , the `'%'` is replaced by the text that matched the `'%'` in `pattern`. Only the first `'%'` in the `pattern` and `replacement` is treated this way; any subsequent `'%'` is unchanged.

`'%'` characters in `patsubst` function invocations can be quoted with preceding backslashes (`'\'`). Backslashes that would otherwise quote `'%'` characters can be quoted with more backslashes. Backslashes that quote `'%'` characters or other backslashes are removed from the `pattern` before it is compared file names or has a stem substituted into it. Backslashes that are not in danger of quoting `'%'` characters go unmolested. For example, the `pattern` `the%\weird\\%pattern\\` has `'the%weird\'` preceding the operative `'%'` character, and `'pattern\\'` following it. The final two backslashes are left alone because they cannot affect any `'%'` character.

Whitespace between words is folded into single space characters; leading and trailing whitespace is discarded.

For example,

```
$(patsubst %.c,%.o,x.c.c bar.c)
```

produces the value `'x.c.o bar.o'` .

Substitution references (see [Substitution References](#)) are a simpler way to get the effect of the `patsubst` function:

```
$(var:pattern=replacement)
```

is equivalent to

```
$(patsubst pattern,replacement,$(var))
```

The second shorthand simplifies one of the most common uses of `patsubst`: replacing the suffix at the end of file names.

```
$(var:suffix=replacement)
```

is equivalent to

```
$(patsubst %suffix,%replacement,$(var))
```

For example, you might have a list of object files:

```
objects = foo.o bar.o baz.o
```

To get the list of corresponding source files, you could simply write:

```
$(objects:.o=.c)
```

instead of using the general form:

```
$(patsubst %.o,%.c,$(objects))
```

`$(strip string)`

Removes leading and trailing whitespace from string and replaces each internal sequence of one or more whitespace characters with a single space. Thus, `'$(strip a b c)'` results in `'a b c'`.

The function `strip` can be very useful when used in conjunction with conditionals. When comparing something with the empty string `"` using `ifeq` or `ifneq`, you usually want a string of just whitespace to match the empty string (see [Conditionals](#)).

Thus, the following may fail to have the desired results:

```
.PHONY: all
ifneq    "$$(needs_made)" ""
all:    $$(needs_made)
else
all:;@echo 'Nothing to make!'
endif
```

Replacing the variable reference `'$$(needs_made)'` with the function call `'$(strip $$(needs_made))'` in the `ifneq` directive would make it more robust.

`$(findstring find,in)`

Searches in for an occurrence of `find`. If it occurs, the value is `find`; otherwise, the value is empty. You can use this function in a conditional to test for the presence of a specific substring in a given string. Thus, the two examples,

```
$(findstring a,a b c)
$(findstring a,b c)
```

produce the values `'a'` and `"` (the empty string), respectively. See [Testing Flags](#), for a practical application of `findstring`.

`$(filter pattern...,text)`

Returns all whitespace-separated words in `text` that do match any of the pattern words, removing any words that do not match. The patterns are written using `'%'`, just like the patterns used in the `patsubst` function above.

The `filter` function can be used to separate out different types of strings (such as file names) in a variable. For example:

```
sources := foo.c bar.c baz.s ugh.h
foo:    $(sources)
        cc $(filter %.c %.s,$(sources)) -o foo
```

says that `foo` depends of `foo.c`, `bar.c`, `baz.s` and `ugh.h` but only `foo.c`, `bar.c` and `baz.s` should be specified in the command to the compiler.

`$(filter-out pattern...,text)`

Returns all whitespace-separated words in `text` that do not match any of the pattern words, removing the words that do match one or more. This is the exact opposite of the `filter` function.

For example, given:

```
objects=main1.o foo.o main2.o bar.o
mains=main1.o main2.o
```

the following generates a list which contains all the object files not in `'mains'` :

```
$(filter-out $(mains),$(objects))
```

`$(sort list)`

Sorts the words of list in lexical order, removing duplicate words. The output is a list of words separated by single spaces. Thus,

```
$(sort foo bar lose)
```

returns the value `'bar foo lose'` .

Incidentally, since `sort` removes duplicate words, you can use it for this purpose even if you don't care about the sort order.

`$(word n, text)`

Returns the *n*th word of text. The legitimate values of *n* start from 1. If *n* is bigger than the number of words in text, the value is empty. For example,

```
$(word 2, foo bar baz)
```

returns `'bar'` .

`$(wordlist s, e, text)`

Returns the list of words in text starting with word *s* and ending with word *e* (inclusive). The legitimate values of *s* start from 1; *e* may start from 0. If *s* is bigger than the number of words in text, the value is empty. If *e* is bigger than the number of words in text, words up to the end of text are returned. If *s* is greater than *e*, nothing is returned. For example,

```
$(wordlist 2, 3, foo bar baz)
```

returns `'bar baz'` .

`$(words text)`

Returns the number of words in text. Thus, the last word of text is

```
$(word $(words text), text).
```

`$(firstword names...)`

The argument *names* is regarded as a series of names, separated by whitespace. The value is the first name in the series. The rest of the names are ignored.

For example,

```
$(firstword foo bar)
```

produces the result `'foo'` . Although `$(firstword text)` is the same as `$(word 1, text)`, the `firstword` function is retained for its simplicity.

`$(lastword names...)`

The argument *names* is regarded as a series of names, separated by whitespace. The value is the last name in the series.

For example,

```
$(lastword foo bar)
```

produces the result `'bar'` . Although `$(lastword text)` is the same as `$(word $(words text), text)`, the `lastword` function was added for its simplicity and better performance.

Here is a realistic example of the use of `subst` and `patsubst`. Suppose that a makefile uses the `VPATH` variable to specify a list of directories that `make` should search for prerequisite files (see [VPATH Search Path for All Prerequisites](#)). This example shows how to tell the C compiler to search for header files in the same list of directories.

The value of `VPATH` is a list of directories separated by colons, such as `'src:../headers'`. First, the `subst` function is used to change the colons to spaces:

```
$(subst :, ,,$(VPATH))
```

This produces `'src ../headers'`. Then `patsubst` is used to turn each directory name into a `'-I'` flag. These can be added to the value of the variable `CFLAGS`, which is passed automatically to the C compiler, like this:

```
override CFLAGS += $(patsubst %,-I%,$(subst :, ,,$(VPATH)))
```

The effect is to append the text `'-Isrc -I../headers'` to the previously given value of `CFLAGS`. The `override` directive is used so that the new value is assigned even if the previous value of `CFLAGS` was specified with a command argument (see [The override Directive](#)).

Next: [Conditional Functions](#), Previous: [Text Functions](#), Up: [Functions](#) [\[Contents\]](#)[\[Index\]](#)

8.3 Functions for File Names

Several of the built-in expansion functions relate specifically to taking apart file names or lists of file names.

Each of the following functions performs a specific transformation on a file name. The argument of the function is regarded as a series of file names, separated by whitespace. (Leading and trailing whitespace is ignored.) Each file name in the series is transformed in the same way and the results are concatenated with single spaces between them.

`$(dir names...)`

Extracts the directory-part of each file name in `names`. The directory-part of the file name is everything up through (and including) the last slash in it. If the file name contains no slash, the directory part is the string `'./'`. For example,

```
$(dir src/foo.c hacks)
```

produces the result `'src/ ./'`.

`$(notdir names...)`

Extracts all but the directory-part of each file name in `names`. If the file name contains no slash, it is left unchanged. Otherwise, everything through the last slash is removed from it.

A file name that ends with a slash becomes an empty string. This is unfortunate, because it means that the result does not always have the same number of whitespace-separated file names as the argument had; but we do not see any other valid alternative.

For example,

```
$(notdir src/foo.c hacks)
```

produces the result `'foo.c hacks'`.

`$(suffix names...)`

Extracts the suffix of each file name in `names`. If the file name contains a period, the suffix is everything starting with the last period. Otherwise, the suffix is the empty string. This frequently means that the result will be empty when `names` is not, and if `names` contains multiple file names, the result may contain fewer file names.

For example,

```
$(suffix src/foo.c src-1.0/bar.c hacks)
```

produces the result `' .c .c'` .

```
$(basename names...)
```

Extracts all but the suffix of each file name in `names`. If the file name contains a period, the basename is everything starting up to (and not including) the last period. Periods in the directory part are ignored. If there is no period, the basename is the entire file name. For example,

```
$(basename src/foo.c src-1.0/bar hacks)
```

produces the result `'src/foo src-1.0/bar hacks'` .

```
$(addsuffix suffix, names...)
```

The argument `names` is regarded as a series of names, separated by whitespace; `suffix` is used as a unit. The value of `suffix` is appended to the end of each individual name and the resulting larger names are concatenated with single spaces between them. For example,

```
$(addsuffix .c, foo bar)
```

produces the result `'foo.c bar.c'` .

```
$(addprefix prefix, names...)
```

The argument `names` is regarded as a series of names, separated by whitespace; `prefix` is used as a unit. The value of `prefix` is prepended to the front of each individual name and the resulting larger names are concatenated with single spaces between them. For example,

```
$(addprefix src/, foo bar)
```

produces the result `'src/foo src/bar'` .

```
$(join list1, list2)
```

Concatenates the two arguments word by word: the two first words (one from each argument) concatenated form the first word of the result, the two second words form the second word of the result, and so on. So the *n*th word of the result comes from the *n*th word of each argument. If one argument has more words than the other, the extra words are copied unchanged into the result.

For example, `'$(join a b, .c .o)'` produces `'a.c b.o'` .

Whitespace between the words in the lists is not preserved; it is replaced with a single space.

This function can merge the results of the `dir` and `notdir` functions, to produce the original list of files which was given to those two functions.

```
$(wildcard pattern)
```

The argument `pattern` is a file name pattern, typically containing wildcard characters (as in shell file name patterns). The result of `wildcard` is a space-separated list of the names of existing files that match the pattern. See [Using Wildcard Characters in File Names](#).

```
$(realpath names...)
```

For each file name in `names` return the canonical absolute name. A canonical name does not contain any `.` or `..` components, nor any repeated path separators (`/`) or symlinks. In case of a failure the empty string is returned. Consult the `realpath(3)` documentation for a list of possible failure causes.

```
$(abspath names...)
```

For each file name in `names` return an absolute name that does not contain any `.` or `..` components, nor any repeated path separators (`/`). Note that, in contrast to `realpath`

function, `abspath` does not resolve symlinks and does not require the file names to refer to an existing file or directory. Use the `wildcard` function to test for existence.

Next: [Foreach Function](#), Previous: [File Name Functions](#), Up: [Functions](#) [\[Contents\]](#)[\[Index\]](#)

8.4 Functions for Conditionals

There are three functions that provide conditional expansion. A key aspect of these functions is that not all of the arguments are expanded initially. Only those arguments which need to be expanded, will be expanded.

```
$(if condition, then-part[, else-part])
```

The `if` function provides support for conditional expansion in a functional context (as opposed to the GNU `make` makefile conditionals such as `ifeq` (see [Syntax of Conditionals](#))).

The first argument, `condition`, first has all preceding and trailing whitespace stripped, then is expanded. If it expands to any non-empty string, then the condition is considered to be true. If it expands to an empty string, the condition is considered to be false.

If the condition is true then the second argument, `then-part`, is evaluated and this is used as the result of the evaluation of the entire `if` function.

If the condition is false then the third argument, `else-part`, is evaluated and this is the result of the `if` function. If there is no third argument, the `if` function evaluates to nothing (the empty string).

Note that only one of the `then-part` or the `else-part` will be evaluated, never both. Thus, either can contain side-effects (such as `shell` function calls, etc.)

```
$(or condition1[, condition2[, condition3...]])
```

The `or` function provides a “short-circuiting” OR operation. Each argument is expanded, in order. If an argument expands to a non-empty string the processing stops and the result of the expansion is that string. If, after all arguments are expanded, all of them are false (empty), then the result of the expansion is the empty string.

```
$(and condition1[, condition2[, condition3...]])
```

The `and` function provides a “short-circuiting” AND operation. Each argument is expanded, in order. If an argument expands to an empty string the processing stops and the result of the expansion is the empty string. If all arguments expand to a non-empty string then the result of the expansion is the expansion of the last argument.

Next: [File Function](#), Previous: [Conditional Functions](#), Up: [Functions](#) [\[Contents\]](#)[\[Index\]](#)

8.5 The `foreach` Function

The `foreach` function is very different from other functions. It causes one piece of text to be used repeatedly, each time with a different substitution performed on it. It resembles the `for` command in the shell `sh` and the `foreach` command in the C-shell `csh`.

The syntax of the `foreach` function is:

```
$(foreach var, list, text)
```

The first two arguments, `var` and `list`, are expanded before anything else is done; note that the last argument, `text`, is **not** expanded at the same time. Then for each word of the expanded value of `list`, the variable named by the expanded value of `var` is set to that word, and `text` is expanded. Presumably `text` contains references to that variable, so its expansion will be different each time.

The result is that `text` is expanded as many times as there are whitespace-separated words in `list`. The multiple expansions of `text` are concatenated, with spaces between them, to make the result of `foreach`.

This simple example sets the variable `'files'` to the list of all files in the directories in the list `'dirs'` :

```
dirs := a b c d
files := $(foreach dir,$(dirs),$(wildcard $(dir)/*))
```

Here text is `'$(wildcard $(dir)/*)'` . The first repetition finds the value `'a'` for `dir`, so it produces the same result as `'$(wildcard a/*)'` ; the second repetition produces the result of `'$(wildcard b/*)'` ; and the third, that of `'$(wildcard c/*)'` .

This example has the same result (except for setting `'dirs'`) as the following example:

```
files := $(wildcard a/* b/* c/* d/*)
```

When text is complicated, you can improve readability by giving it a name, with an additional variable:

```
find_files = $(wildcard $(dir)/*)
dirs := a b c d
files := $(foreach dir,$(dirs),$(find_files))
```

Here we use the variable `find_files` this way. We use plain `'='` to define a recursively-expanding variable, so that its value contains an actual function call to be re-expanded under the control of `foreach`; a simply-expanded variable would not do, since `wildcard` would be called only once at the time of defining `find_files`.

The `foreach` function has no permanent effect on the variable `var`; its value and flavor after the `foreach` function call are the same as they were beforehand. The other values which are taken from `list` are in effect only temporarily, during the execution of `foreach`. The variable `var` is a simply-expanded variable during the execution of `foreach`. If `var` was undefined before the `foreach` function call, it is undefined after the call. See [The Two Flavors of Variables](#).

You must take care when using complex variable expressions that result in variable names because many strange things are valid variable names, but are probably not what you intended. For example,

```
files := $(foreach Esta-escri-to-en-espanol!,b c ch,$(find_files))
```

might be useful if the value of `find_files` references the variable whose name is `'Esta-escri-to-en-espanol!'` (es un nombre bastante largo, no?), but it is more likely to be a mistake.

Next: [Call Function](#), Previous: [Foreach Function](#), Up: [Functions](#) [[Contents](#)][[Index](#)]

8.6 The `file` Function

The `file` function allows the makefile to write to or read from a file. Two modes of writing are supported: `overwrite`, where the text is written to the beginning of the file and any existing content is lost, and `append`, where the text is written to the end of the file, preserving the existing content. In both cases the file is created if it does not exist. It is a fatal error if the file cannot be opened for writing, or if the write operation fails. The `file` function expands to the empty string when writing to a file.

When reading from a file, the `file` function expands to the verbatim contents of the file, except that the final newline (if there is one) will be stripped. Attempting to read from a non-existent file expands to the empty string.

The syntax of the `file` function is:

```
$(file op filename[, text])
```

When the `file` function is evaluated all its arguments are expanded first, then the file indicated by `filename` will be opened in the mode described by `op`.

The operator `op` can be `>` to indicate the file will be overwritten with new content, `>>` to indicate the current contents of the file will be appended to, or `<` to indicate the contents of the file will be read in. The filename specifies the file to be written to or read from. There may optionally be whitespace between the operator and the file name.

When reading files, it is an error to provide a text value.

When writing files, text will be written to the file. If text does not already end in a newline a final newline will be written (even if text is the empty string). If the text argument is not given at all, nothing will be written.

For example, the `file` function can be useful if your build system has a limited command line size and your recipe runs a command that can accept arguments from a file as well. Many commands use the convention that an argument prefixed with an `@` specifies a file containing more arguments. Then you might write your recipe in this way:

```
program: $(OBJECTS)
        $(file >$.in,$^)
        $(CMD) $(CMDFLAGS) @$$.in
        @rm $.in
```

If the command required each argument to be on a separate line of the input file, you might write your recipe like this:

```
program: $(OBJECTS)
        $(file >$.in) $(foreach 0,$^,$(file >>$.in,$0))
        $(CMD) $(CMDFLAGS) @$$.in
        @rm $.in
```

Next: [Value Function](#), Previous: [File Function](#), Up: [Functions](#) [[Contents](#)][[Index](#)]

8.7 The `call` Function

The `call` function is unique in that it can be used to create new parameterized functions. You can write a complex expression as the value of a variable, then use `call` to expand it with different values.

The syntax of the `call` function is:

```
$(call variable,param,param,...)
```

When `make` expands this function, it assigns each `param` to temporary variables `$(1)`, `$(2)`, etc. The variable `$(0)` will contain `variable`. There is no maximum number of parameter arguments. There is no minimum, either, but it doesn't make sense to use `call` with no parameters.

Then `variable` is expanded as a `make` variable in the context of these temporary assignments. Thus, any reference to `$(1)` in the value of `variable` will resolve to the first `param` in the invocation of `call`.

Note that `variable` is the name of a variable, not a reference to that variable. Therefore you would not normally use a `'$'` or parentheses when writing it. (You can, however, use a variable reference in the name if you want the name not to be a constant.)

If `variable` is the name of a built-in function, the built-in function is always invoked (even if a `make` variable by that name also exists).

The `call` function expands the `param` arguments before assigning them to temporary variables. This means that `variable` values containing references to built-in functions that have special expansion rules, like `foreach` or `if`, may not work as you expect.

Some examples may make this clearer.

This macro simply reverses its arguments:

```
reverse = $(2) $(1)

foo = $(call reverse, a, b)
```

Here `foo` will contain `'b a'` .

This one is slightly more interesting: it defines a macro to search for the first instance of a program in `PATH`:

```
pathsearch = $(firstword $(wildcard $(addsuffix /$(1), $(subst :, , $(PATH)))))

LS := $(call pathsearch, ls)
```

Now the variable `LS` contains `/bin/ls` or similar.

The `call` function can be nested. Each recursive invocation gets its own local values for `$(1)`, etc. that mask the values of higher-level `call`. For example, here is an implementation of a `map` function:

```
map = $(foreach a, $(2), $(call $(1), $(a)))
```

Now you can map a function that normally takes only one argument, such as `origin`, to multiple values in one step:

```
o = $(call map, origin, o map MAKE)
```

and end up with `o` containing something like `'file file default'` .

A final caution: be careful when adding whitespace to the arguments to `call`. As with other functions, any whitespace contained in the second and subsequent arguments is kept; this can cause strange effects. It's generally safest to remove all extraneous whitespace when providing parameters to `call`.

Next: [Eval Function](#), Previous: [Call Function](#), Up: [Functions](#) [[Contents](#)][[Index](#)]

8.8 The `value` Function

The `value` function provides a way for you to use the value of a variable without having it expanded. Please note that this does not undo expansions which have already occurred; for example if you create a simply expanded variable its value is expanded during the definition; in that case the `value` function will return the same result as using the variable directly.

The syntax of the `value` function is:

```
$(value variable)
```

Note that `variable` is the name of a variable, not a reference to that variable. Therefore you would not normally use a `'$'` or parentheses when writing it. (You can, however, use a variable reference in the name if you want the name not to be a constant.)

The result of this function is a string containing the value of `variable`, without any expansion occurring. For example, in this makefile:

```
FOO = $PATH

all:
    @echo $(FOO)
    @echo $(value FOO)
```

The first output line would be `ATH`, since the `"$P"` would be expanded as a `make` variable, while the second output line would be the current value of your `$PATH` environment variable, since the `value` function avoided the expansion.

The `value` function is most often used in conjunction with the `eval` function (see [Eval Function](#)).

Next: [Origin Function](#), Previous: [Value Function](#), Up: [Functions](#) [[Contents](#)][[Index](#)]

8.9 The `eval` Function

The `eval` function is very special: it allows you to define new makefile constructs that are not constant; which are the result of evaluating other variables and functions. The argument to the `eval` function is expanded, then the results of that expansion are parsed as makefile syntax. The expanded results can define new `make` variables, targets, implicit or explicit rules, etc.

The result of the `eval` function is always the empty string; thus, it can be placed virtually anywhere in a makefile without causing syntax errors.

It's important to realize that the `eval` argument is expanded twice; first by the `eval` function, then the results of that expansion are expanded again when they are parsed as makefile syntax. This means you may need to provide extra levels of escaping for "\$" characters when using `eval`. The `value` function (see [Value Function](#)) can sometimes be useful in these situations, to circumvent unwanted expansions.

Here is an example of how `eval` can be used; this example combines a number of concepts and other functions. Although it might seem overly complex to use `eval` in this example, rather than just writing out the rules, consider two things: first, the template definition (in `PROGRAM_template`) could need to be much more complex than it is here; and second, you might put the complex, "generic" part of this example into another makefile, then include it in all the individual makefiles. Now your individual makefiles are quite straightforward.

```
PROGRAMS      = server client

server_OBJS = server.o server_priv.o server_access.o
server_LIBS = priv protocol

client_OBJS = client.o client_api.o client_mem.o
client_LIBS = protocol

# Everything after this is generic

.PHONY: all
all: $(PROGRAMS)

define PROGRAM_template =
$(1): $$($1)_OBJS $$($1)_LIBS:%=-1%
ALL_OBJS += $$($1)_OBJS
endef

$(foreach prog,$(PROGRAMS),$(eval $(call PROGRAM_template,$(prog))))

$(PROGRAMS):
$(LINK.o) $^ $(LDLIBS) -o $@

clean:
rm -f $(ALL_OBJS) $(PROGRAMS)
```

Next: [Flavor Function](#), Previous: [Eval Function](#), Up: [Functions](#) [[Contents](#)][[Index](#)]

8.10 The `origin` Function

The `origin` function is unlike most other functions in that it does not operate on the values of variables; it tells you something about a variable. Specifically, it tells you where it came from.

The syntax of the `origin` function is:

```
$(origin variable)
```

Note that `variable` is the name of a variable to inquire about, not a reference to that variable. Therefore you would not normally use a `'$'` or parentheses when writing it. (You can, however, use a variable reference in the name if you want the name not to be a constant.)

The result of this function is a string telling you how the variable `variable` was defined:

`'undefined'`

if variable was never defined.

`'default'`

if variable has a default definition, as is usual with `CC` and so on. See [Variables Used by Implicit Rules](#). Note that if you have redefined a default variable, the `origin` function will return the origin of the later definition.

`'environment'`

if variable was inherited from the environment provided to `make`.

`'environment override'`

if variable was inherited from the environment provided to `make`, and is overriding a setting for variable in the makefile as a result of the `'-e'` option (see [Summary of Options](#)).

`'file'`

if variable was defined in a makefile.

`'command line'`

if variable was defined on the command line.

`'override'`

if variable was defined with an `override` directive in a makefile (see [The override Directive](#)).

`'automatic'`

if variable is an automatic variable defined for the execution of the recipe for each rule (see [Automatic Variables](#)).

This information is primarily useful (other than for your curiosity) to determine if you want to believe the value of a variable. For example, suppose you have a makefile `foo` that includes another makefile `bar`. You want a variable `bletch` to be defined in `bar` if you run the command `'make -f bar'`, even if the environment contains a definition of `bletch`. However, if `foo` defined `bletch` before including `bar`, you do not want to override that definition. This could be done by using an `override` directive in `foo`, giving that definition precedence over the later definition in `bar`; unfortunately, the `override` directive would also override any command line definitions. So, `bar` could include:

```
ifdef bletch
ifeq "$(origin bletch)" "environment"
bletch = barf, gag, etc.
endif
endif
```

If `bletch` has been defined from the environment, this will redefine it.

If you want to override a previous definition of `bletch` if it came from the environment, even under `'-e'`, you could instead write:

```
ifneq "$(findstring environment,$(origin bletch))" ""
bletch = barf, gag, etc.
endif
```

Here the redefinition takes place if `'$(origin bletch)'` returns either `'environment'` or `'environment override'`. See [Functions for String Substitution and Analysis](#).

8.11 The `flavor` Function

The `flavor` function, like the `origin` function, does not operate on the values of variables but rather it tells you something about a variable. Specifically, it tells you the flavor of a variable (see [The Two Flavors of Variables](#)).

The syntax of the `flavor` function is:

```
$(flavor variable)
```

Note that `variable` is the name of a variable to inquire about, not a reference to that variable. Therefore you would not normally use a `'$'` or parentheses when writing it. (You can, however, use a variable reference in the name if you want the name not to be a constant.)

The result of this function is a string that identifies the flavor of the variable `variable`:

```
'undefined'
```

if variable was never defined.

```
'recursive'
```

if variable is a recursively expanded variable.

```
'simple'
```

if variable is a simply expanded variable.

8.12 Functions That Control Make

These functions control the way `make` runs. Generally, they are used to provide information to the user of the makefile or to cause `make` to stop if some sort of environmental error is detected.

```
$(error text...)
```

Generates a fatal error where the message is `text`. Note that the error is generated whenever this function is evaluated. So, if you put it inside a recipe or on the right side of a recursive variable assignment, it won't be evaluated until later. The text will be expanded before the error is generated.

For example,

```
ifdef ERROR1
$(error error is $(ERROR1))
endif
```

will generate a fatal error during the read of the makefile if the `make` variable `ERROR1` is defined. Or,

```
ERR = $(error found an error!)

.PHONY: err
err: ; $(ERR)
```

will generate a fatal error while `make` is running, if the `err` target is invoked.

```
$(warning text...)
```

This function works similarly to the `error` function, above, except that `make` doesn't exit. Instead, text is expanded and the resulting message is displayed, but processing of the makefile continues.

The result of the expansion of this function is the empty string.

`$(info text...)`

This function does nothing more than print its (expanded) argument(s) to standard output. No makefile name or line number is added. The result of the expansion of this function is the empty string.

Next: [Guile Function](#), Previous: [Make Control Functions](#), Up: [Functions](#) [[Contents](#)][[Index](#)]

8.13 The `shell` Function

The `shell` function is unlike any other function other than the `wildcard` function (see [The Function `wildcard`](#)) in that it communicates with the world outside of `make`.

The `shell` function performs the same function that backquotes (```) perform in most shells: it does command expansion. This means that it takes as an argument a shell command and evaluates to the output of the command. The only processing `make` does on the result is to convert each newline (or carriage-return / newline pair) to a single space. If there is a trailing (carriage-return and) newline it will simply be removed.

The commands run by calls to the `shell` function are run when the function calls are expanded (see [How `make` Reads a Makefile](#)). Because this function involves spawning a new shell, you should carefully consider the performance implications of using the `shell` function within recursively expanded variables vs. simply expanded variables (see [The Two Flavors of Variables](#)).

After the `shell` function or `!=` assignment operator is used, its exit status is placed in the `.SHELLSTATUS` variable.

Here are some examples of the use of the `shell` function:

```
contents := $(shell cat foo)
```

sets `contents` to the contents of the file `foo`, with a space (rather than a newline) separating each line.

```
files := $(shell echo *.c)
```

sets `files` to the expansion of `*.c`. Unless `make` is using a very strange shell, this has the same result as `$(wildcard *.c)` (as long as at least one `.c` file exists).

Previous: [Shell Function](#), Up: [Functions](#) [[Contents](#)][[Index](#)]

8.14 The `guile` Function

If GNU `make` is built with support for GNU Guile as an embedded extension language then the `guile` function will be available. The `guile` function takes one argument which is first expanded by `make` in the normal fashion, then passed to the GNU Guile evaluator. The result of the evaluator is converted into a string and used as the expansion of the `guile` function in the makefile. See [GNU Guile Integration](#) for details on writing extensions to `make` in Guile.

You can determine whether GNU Guile support is available by checking the `.FEATURES` variable for the word `guile`.

Next: [Implicit Rules](#), Previous: [Functions](#), Up: [Top](#) [[Contents](#)][[Index](#)]

9 How to Run `make`

A makefile that says how to recompile a program can be used in more than one way. The simplest use is to recompile every file that is out of date. Usually, makefiles are written so that

if you run `make` with no arguments, it does just that.

But you might want to update only some of the files; you might want to use a different compiler or different compiler options; you might want just to find out which files are out of date without changing them.

By giving arguments when you run `make`, you can do any of these things and many others.

The exit status of `make` is always one of three values:

0

The exit status is zero if `make` is successful.

2

The exit status is two if `make` encounters any errors. It will print messages describing the particular errors.

1

The exit status is one if you use the `'-q'` flag and `make` determines that some target is not already up to date. See [Instead of Executing Recipes](#).

- [Makefile Arguments](#) How to specify which makefile to use.
- [Goals](#) How to use goal arguments to specify which parts of the makefile to use.
- [Instead of Execution](#) How to use mode flags to specify what kind of thing to do with the recipes in the makefile other than simply execute them.
- [Avoiding Compilation](#) How to avoid recompiling certain files.
- [Overriding](#) How to override a variable to specify an alternate compiler and other things.
- [Testing](#) How to proceed past some errors, to test compilation.
- [Options Summary](#) Summary of Options

Next: [Goals](#), Previous: [Running](#), Up: [Running](#) [\[Contents\]](#)[\[Index\]](#)

9.1 Arguments to Specify the Makefile

The way to specify the name of the makefile is with the `'-f'` or `'--file'` option (`'--makefile'` also works). For example, `'-f altmake'` says to use the file `altmake` as the makefile.

If you use the `'-f'` flag several times and follow each `'-f'` with an argument, all the specified files are used jointly as makefiles.

If you do not use the `'-f'` or `'--file'` flag, the default is to try `GNUmakefile`, `makefile`, and `Makefile`, in that order, and use the first of these three which exists or can be made (see [Writing Makefiles](#)).

Next: [Instead of Execution](#), Previous: [Makefile Arguments](#), Up: [Running](#) [\[Contents\]](#)[\[Index\]](#)

9.2 Arguments to Specify the Goals

The goals are the targets that `make` should strive ultimately to update. Other targets are updated as well if they appear as prerequisites of goals, or prerequisites of prerequisites of goals, etc.

By default, the goal is the first target in the makefile (not counting targets that start with a period). Therefore, makefiles are usually written so that the first target is for compiling the entire program or programs they describe. If the first rule in the makefile has several targets, only the first target in the rule becomes the default goal, not the whole list. You can manage

the selection of the default goal from within your makefile using the `.DEFAULT_GOAL` variable (see [Other Special Variables](#)).

You can also specify a different goal or goals with command line arguments to `make`. Use the name of the goal as an argument. If you specify several goals, `make` processes each of them in turn, in the order you name them.

Any target in the makefile may be specified as a goal (unless it starts with `'-'` or contains an `'='`, in which case it will be parsed as a switch or variable definition, respectively). Even targets not in the makefile may be specified, if `make` can find implicit rules that say how to make them.

`Make` will set the special variable `MAKECMDGOALS` to the list of goals you specified on the command line. If no goals were given on the command line, this variable is empty. Note that this variable should be used only in special circumstances.

An example of appropriate use is to avoid including `.d` files during `clean` rules (see [Automatic Prerequisites](#)), so `make` won't create them only to immediately remove them again:

```
sources = foo.c bar.c

ifneq ($(MAKECMDGOALS), clean)
include $(sources:.c=.d)
endif
```

One use of specifying a goal is if you want to compile only a part of the program, or only one of several programs. Specify as a goal each file that you wish to remake. For example, consider a directory containing several programs, with a makefile that starts like this:

```
.PHONY: all
all: size nm ld ar as
```

If you are working on the program `size`, you might want to say `'make size'` so that only the files of that program are recompiled.

Another use of specifying a goal is to make files that are not normally made. For example, there may be a file of debugging output, or a version of the program that is compiled specially for testing, which has a rule in the makefile but is not a prerequisite of the default goal.

Another use of specifying a goal is to run the recipe associated with a phony target (see [Phony Targets](#)) or empty target (see [Empty Target Files to Record Events](#)). Many makefiles contain a phony target named `clean` which deletes everything except source files. Naturally, this is done only if you request it explicitly with `'make clean'`. Following is a list of typical phony and empty target names. See [Standard Targets](#), for a detailed list of all the standard target names which GNU software packages use.

`all`

Make all the top-level targets the makefile knows about.

`clean`

Delete all files that are normally created by running `make`.

`mostlyclean`

Like `'clean'`, but may refrain from deleting a few files that people normally don't want to recompile. For example, the `'mostlyclean'` target for GCC does not delete `libgcc.a`, because recompiling it is rarely necessary and takes a lot of time.

`distclean`

`realclean`

`clobber`

Any of these targets might be defined to delete more files than `'clean'` does. For example, this would delete configuration files or links that you would normally create as

preparation for compilation, even if the makefile itself cannot create these files.

`install`

Copy the executable file into a directory that users typically search for commands; copy any auxiliary files that the executable uses into the directories where it will look for them.

`print`

Print listings of the source files that have changed.

`tar`

Create a tar file of the source files.

`shar`

Create a shell archive (shar file) of the source files.

`dist`

Create a distribution file of the source files. This might be a tar file, or a shar file, or a compressed version of one of the above, or even more than one of the above.

`TAGS`

Update a tags table for this program.

`check`

`test`

Perform self tests on the program this makefile builds.

Next: [Avoiding Compilation](#), Previous: [Goals](#), Up: [Running](#) [\[Contents\]](#)[\[Index\]](#)

9.3 Instead of Executing Recipes

The makefile tells `make` how to tell whether a target is up to date, and how to update each target. But updating the targets is not always what you want. Certain options specify other activities for `make`.

`'-n'`

`'--just-print'`

`'--dry-run'`

`'--recon'`

“No-op” . Causes `make` to print the recipes that are needed to make the targets up to date, but not actually execute them. Note that some recipes are still executed, even with this flag (see [How the MAKE Variable Works](#)). Also any recipes needed to update included makefiles are still executed (see [How Makefiles Are Remade](#)).

`'-t'`

`'--touch'`

“Touch” . Marks targets as up to date without actually changing them. In other words, `make` pretends to update the targets but does not really change their contents; instead only their modified times are updated.

`'-q'`

`'--question'`

“Question” . Silently check whether the targets are up to date, but do not execute recipes; the exit code shows whether any updates are needed.

`'-W file'`

`'--what-if=file'`

`'--assume-new=file'`

`'--new-file=file'`

“What if” . Each `'-w'` flag is followed by a file name. The given files' modification times are recorded by `make` as being the present time, although the actual modification times remain the same. You can use the `'-w'` flag in conjunction with the `'-n'` flag to see what would happen if you were to modify specific files.

With the `'-n'` flag, `make` prints the recipe that it would normally execute but usually does not execute it.

With the `'-t'` flag, `make` ignores the recipes in the rules and uses (in effect) the command `touch` for each target that needs to be remade. The `touch` command is also printed, unless `'-s'` or `.SILENT` is used. For speed, `make` does not actually invoke the program `touch`. It does the work directly.

With the `'-q'` flag, `make` prints nothing and executes no recipes, but the exit status code it returns is zero if and only if the targets to be considered are already up to date. If the exit status is one, then some updating needs to be done. If `make` encounters an error, the exit status is two, so you can distinguish an error from a target that is not up to date.

It is an error to use more than one of these three flags in the same invocation of `make`.

The `'-n'` , `'-t'` , and `'-q'` options do not affect recipe lines that begin with `'+'` characters or contain the strings `'$(MAKE)'` or `'${MAKE}'` . Note that only the line containing the `'+'` character or the strings `'$(MAKE)'` or `'${MAKE}'` is run regardless of these options. Other lines in the same rule are not run unless they too begin with `'+'` or contain `'$(MAKE)'` or `'${MAKE}'` (See [How the MAKE Variable Works](#).)

The `'-t'` flag prevents phony targets (see [Phony Targets](#)) from being updated, unless there are recipe lines beginning with `'+'` or containing `'$(MAKE)'` or `'${MAKE}'` .

The `'-w'` flag provides two features:

- If you also use the `'-n'` or `'-q'` flag, you can see what `make` would do if you were to modify some files.
- Without the `'-n'` or `'-q'` flag, when `make` is actually executing recipes, the `'-w'` flag can direct `make` to act as if some files had been modified, without actually running the recipes for those files.

Note that the options `'-p'` and `'-v'` allow you to obtain other information about `make` or about the makefiles in use (see [Summary of Options](#)).

Next: [Overriding](#), Previous: [Instead of Execution](#), Up: [Running](#) [[Contents](#)][[Index](#)]

9.4 Avoiding Recompilation of Some Files

Sometimes you may have changed a source file but you do not want to recompile all the files that depend on it. For example, suppose you add a macro or a declaration to a header file that many other files depend on. Being conservative, `make` assumes that any change in the header file requires recompilation of all dependent files, but you know that they do not need to be recompiled and you would rather not waste the time waiting for them to compile.

If you anticipate the problem before changing the header file, you can use the `'-t'` flag. This flag tells `make` not to run the recipes in the rules, but rather to mark the target up to date by changing its last-modification date. You would follow this procedure:

1. Use the command `'make'` to recompile the source files that really need recompilation, ensuring that the object files are up-to-date before you begin.
2. Make the changes in the header files.
3. Use the command `'make -t'` to mark all the object files as up to date. The next time you run `make`, the changes in the header files will not cause any recompilation.

If you have already changed the header file at a time when some files do need recompilation, it is too late to do this. Instead, you can use the `'-o file'` flag, which marks a specified file as “old” (see [Summary of Options](#)). This means that the file itself will not be remade, and nothing else will be remade on its account. Follow this procedure:

1. Recompile the source files that need compilation for reasons independent of the particular header file, with `'make -o headerfile'`. If several header files are involved, use a separate `'-o'` option for each header file.
2. Touch all the object files with `'make -t'`.

Next: [Testing](#), Previous: [Avoiding Compilation](#), Up: [Running](#) [[Contents](#)][[Index](#)]

9.5 Overriding Variables

An argument that contains `'='` specifies the value of a variable: `'v=x'` sets the value of the variable `v` to `x`. If you specify a value in this way, all ordinary assignments of the same variable in the makefile are ignored; we say they have been overridden by the command line argument.

The most common way to use this facility is to pass extra flags to compilers. For example, in a properly written makefile, the variable `CFLAGS` is included in each recipe that runs the C compiler, so a file `foo.c` would be compiled something like this:

```
cc -c $(CFLAGS) foo.c
```

Thus, whatever value you set for `CFLAGS` affects each compilation that occurs. The makefile probably specifies the usual value for `CFLAGS`, like this:

```
CFLAGS=-g
```

Each time you run `make`, you can override this value if you wish. For example, if you say `'make CFLAGS='-g -O''`, each C compilation will be done with `'cc -c -g -O'`. (This also illustrates how you can use quoting in the shell to enclose spaces and other special characters in the value of a variable when you override it.)

The variable `CFLAGS` is only one of many standard variables that exist just so that you can change them this way. See [Variables Used by Implicit Rules](#), for a complete list.

You can also program the makefile to look at additional variables of your own, giving the user the ability to control other aspects of how the makefile works by changing the variables.

When you override a variable with a command line argument, you can define either a recursively-expanded variable or a simply-expanded variable. The examples shown above make a recursively-expanded variable; to make a simply-expanded variable, write `':='` or `::='` instead of `'='`. But, unless you want to include a variable reference or function call in the value that you specify, it makes no difference which kind of variable you create.

There is one way that the makefile can change a variable that you have overridden. This is to use the `override` directive, which is a line that looks like this: `'override variable = value'` (see [The `override` Directive](#)).

Next: [Options Summary](#), Previous: [Overriding](#), Up: [Running](#) [[Contents](#)][[Index](#)]

9.6 Testing the Compilation of a Program

Normally, when an error happens in executing a shell command, `make` gives up immediately, returning a nonzero status. No further recipes are executed for any target. The error implies that the goal cannot be correctly remade, and `make` reports this as soon as it knows.

When you are compiling a program that you have just changed, this is not what you want. Instead, you would rather that `make` try compiling every file that can be tried, to show you as

many compilation errors as possible.

On these occasions, you should use the `'-k'` or `'--keep-going'` flag. This tells `make` to continue to consider the other prerequisites of the pending targets, remaking them if necessary, before it gives up and returns nonzero status. For example, after an error in compiling one object file, `'make -k'` will continue compiling other object files even though it already knows that linking them will be impossible. In addition to continuing after failed shell commands, `'make -k'` will continue as much as possible after discovering that it does not know how to make a target or prerequisite file. This will always cause an error message, but without `'-k'`, it is a fatal error (see [Summary of Options](#)).

The usual behavior of `make` assumes that your purpose is to get the goals up to date; once `make` learns that this is impossible, it might as well report the failure immediately. The `'-k'` flag says that the real purpose is to test as much as possible of the changes made in the program, perhaps to find several independent problems so that you can correct them all before the next attempt to compile. This is why Emacs' `M-x compile` command passes the `'-k'` flag by default.

Previous: [Testing](#), Up: [Running](#) [[Contents](#)][[Index](#)]

9.7 Summary of Options

Here is a table of all the options `make` understands:

`'-b'`

`'-m'`

These options are ignored for compatibility with other versions of `make`.

`'-B'`

`'--always-make'`

Consider all targets out-of-date. GNU `make` proceeds to consider targets and their prerequisites using the normal algorithms; however, all targets so considered are always remade regardless of the status of their prerequisites. To avoid infinite recursion, if `MAKE_RESTARTS` (see [Other Special Variables](#)) is set to a number greater than 0 this option is disabled when considering whether to remake makefiles (see [How Makefiles Are Remade](#)).

`'-C dir'`

`'--directory=dir'`

Change to directory `dir` before reading the makefiles. If multiple `'-C'` options are specified, each is interpreted relative to the previous one: `'-C / -C etc'` is equivalent to `'-C /etc'`. This is typically used with recursive invocations of `make` (see [Recursive Use of make](#)).

`'-d'`

Print debugging information in addition to normal processing. The debugging information says which files are being considered for remaking, which file-times are being compared and with what results, which files actually need to be remade, which implicit rules are considered and which are applied—everything interesting about how `make` decides what to do. The `-d` option is equivalent to `'--debug=a'` (see below).

`'--debug[=options]'`

Print debugging information in addition to normal processing. Various levels and types of output can be chosen. With no arguments, print the “basic” level of debugging. Possible arguments are below; only the first character is considered, and values must be comma- or space-separated.

a (all)

All types of debugging output are enabled. This is equivalent to using `'-d'`.

`b` (`basic`)

Basic debugging prints each target that was found to be out-of-date, and whether the build was successful or not.

`v` (`verbose`)

A level above `'basic'` ; includes messages about which makefiles were parsed, prerequisites that did not need to be rebuilt, etc. This option also enables `'basic'` messages.

`i` (`implicit`)

Prints messages describing the implicit rule searches for each target. This option also enables `'basic'` messages.

`j` (`jobs`)

Prints messages giving details on the invocation of specific sub-commands.

`m` (`makefile`)

By default, the above messages are not enabled while trying to remake the makefiles. This option enables messages while rebuilding makefiles, too. Note that the `'all'` option does enable this option. This option also enables `'basic'` messages.

`n` (`none`)

Disable all debugging currently enabled. If additional debugging flags are encountered after this they will still take effect.

`'-e'`

`'--environment-overrides'`

Give variables taken from the environment precedence over variables from makefiles. See [Variables from the Environment](#).

`'-E string'`

`'--eval=string'`

Evaluate string as makefile syntax. This is a command-line version of the `eval` function (see [Eval Function](#)). The evaluation is performed after the default rules and variables have been defined, but before any makefiles are read.

`'-f file'`

`'--file=file'`

`'--makefile=file'`

Read the file named file as a makefile. See [Writing Makefiles](#).

`'-h'`

`'--help'`

Remind you of the options that `make` understands and then exit.

`'-i'`

`'--ignore-errors'`

Ignore all errors in recipes executed to remake files. See [Errors in Recipes](#).

`'-I dir'`

`'--include-dir=dir'`

Specifies a directory dir to search for included makefiles. See [Including Other Makefiles](#). If several `'-I'` options are used to specify several directories, the directories are searched in the order specified.

`'-j [jobs]'`

`'--jobs[=jobs]'`

Specifies the number of recipes (jobs) to run simultaneously. With no argument, `make` runs as many recipes simultaneously as possible. If there is more than one `'-j'` option, the last one is effective. See [Parallel Execution](#), for more information on how recipes are run. Note that this option is ignored on MS-DOS.

`'-k'`

`'--keep-going'`

Continue as much as possible after an error. While the target that failed, and those that depend on it, cannot be remade, the other prerequisites of these targets can be processed all the same. See [Testing the Compilation of a Program](#).

`'-l [load]'`

`'--load-average[=load]'`

`'--max-load[=load]'`

Specifies that no new recipes should be started if there are other recipes running and the load average is at least load (a floating-point number). With no argument, removes a previous load limit. See [Parallel Execution](#).

`'-L'`

`'--check-symlink-times'`

On systems that support symbolic links, this option causes `make` to consider the timestamps on any symbolic links in addition to the timestamp on the file referenced by those links. When this option is provided, the most recent timestamp among the file and the symbolic links is taken as the modification time for this target file.

`'-n'`

`'--just-print'`

`'--dry-run'`

`'--recon'`

Print the recipe that would be executed, but do not execute it (except in certain circumstances). See [Instead of Executing Recipes](#).

`'-o file'`

`'--old-file=file'`

`'--assume-old=file'`

Do not remake the file `file` even if it is older than its prerequisites, and do not remake anything on account of changes in file. Essentially the file is treated as very old and its rules are ignored. See [Avoiding Recompilation of Some Files](#).

`'-O[type]'`

`'--output-sync[=type]'`

Ensure that the complete output from each recipe is printed in one uninterrupted sequence. This option is only useful when using the `--jobs` option to run multiple recipes simultaneously (see [Parallel Execution](#)) Without this option output will be displayed as it is generated by the recipes.

With no type or the type `'target'`, output from the entire recipe of each target is grouped together. With the type `'line'`, output from each line in the recipe is grouped together. With the type `'recurse'`, the output from an entire recursive make is grouped together. With the type `'none'`, no output synchronization is performed. See [Output During Parallel Execution](#).

`'-p'`

`'--print-data-base'`

Print the data base (rules and variable values) that results from reading the makefiles; then execute as usual or as otherwise specified. This also prints the version information given by the `'-v'` switch (see below). To print the data base without trying to remake any files, use `'make -qp'`. To print the data base of predefined rules and variables, use `'make -p -f /dev/null'`. The data base output contains file name and line number information for recipe and variable definitions, so it can be a useful debugging tool in complex environments.

`'-q'`

`'--question'`

“Question mode”. Do not run any recipes, or print anything; just return an exit status that is zero if the specified targets are already up to date, one if any remaking is required, or two if an error is encountered. See [Instead of Executing Recipes](#).

`'-r'`

`'--no-builtin-rules'`

Eliminate use of the built-in implicit rules (see [Using Implicit Rules](#)). You can still define your own by writing pattern rules (see [Defining and Redefining Pattern Rules](#)). The `'-r'` option also clears out the default list of suffixes for suffix rules (see [Old-Fashioned Suffix Rules](#)). But you can still define your own suffixes with a rule for `.SUFFIXES`, and then define your own suffix rules. Note that only rules are affected by the `-r` option; default variables remain in effect (see [Variables Used by Implicit Rules](#)); see the `'-R'` option below.

`'-R'`

`'--no-builtin-variables'`

Eliminate use of the built-in rule-specific variables (see [Variables Used by Implicit Rules](#)). You can still define your own, of course. The `'-R'` option also automatically enables the `'-r'` option (see above), since it doesn't make sense to have implicit rules without any definitions for the variables that they use.

`'-s'`

`'--silent'`

`'--quiet'`

Silent operation; do not print the recipes as they are executed. See [Recipe Echoing](#).

`'-S'`

`'--no-keep-going'`

`'--stop'`

Cancel the effect of the `'-k'` option. This is never necessary except in a recursive `make` where `'-k'` might be inherited from the top-level `make` via `MAKEFLAGS` (see [Recursive Use of make](#)) or if you set `'-k'` in `MAKEFLAGS` in your environment.

`'-t'`

`'--touch'`

Touch files (mark them up to date without really changing them) instead of running their recipes. This is used to pretend that the recipes were done, in order to fool future invocations of `make`. See [Instead of Executing Recipes](#).

`'--trace'`

Show tracing information for `make` execution. Prints the entire recipe to be executed, even for recipes that are normally silent (due to `.SILENT` or `'@'`). Also prints the makefile

name and line number where the recipe was defined, and information on why the target is being rebuilt.

`'-v'`

`'--version'`

Print the version of the `make` program plus a copyright, a list of authors, and a notice that there is no warranty; then exit.

`'-w'`

`'--print-directory'`

Print a message containing the working directory both before and after executing the makefile. This may be useful for tracking down errors from complicated nests of recursive `make` commands. See [Recursive Use of `make`](#). (In practice, you rarely need to specify this option since `'make'` does it for you; see [The `'--print-directory'` Option](#).)

`'--no-print-directory'`

Disable printing of the working directory under `-w`. This option is useful when `-w` is turned on automatically, but you do not want to see the extra messages. See [The `'--print-directory'` Option](#).

`'-W file'`

`'--what-if=file'`

`'--new-file=file'`

`'--assume-new=file'`

Pretend that the target file has just been modified. When used with the `'-n'` flag, this shows you what would happen if you were to modify that file. Without `'-n'`, it is almost the same as running a `touch` command on the given file before running `make`, except that the modification time is changed only in the imagination of `make`. See [Instead of Executing Recipes](#).

`'--warn-undefined-variables'`

Issue a warning message whenever `make` sees a reference to an undefined variable. This can be helpful when you are trying to debug makefiles which use variables in complex ways.

Next: [Archives](#), Previous: [Running](#), Up: [Top](#) [[Contents](#)][[Index](#)]

10 Using Implicit Rules

Certain standard ways of remaking target files are used very often. For example, one customary way to make an object file is from a C source file using the C compiler, `cc`.

Implicit rules tell `make` how to use customary techniques so that you do not have to specify them in detail when you want to use them. For example, there is an implicit rule for C compilation. File names determine which implicit rules are run. For example, C compilation typically takes a `.c` file and makes a `.o` file. So `make` applies the implicit rule for C compilation when it sees this combination of file name endings.

A chain of implicit rules can apply in sequence; for example, `make` will remake a `.o` file from a `.y` file by way of a `.c` file.

The built-in implicit rules use several variables in their recipes so that, by changing the values of the variables, you can change the way the implicit rule works. For example, the variable `CFLAGS` controls the flags given to the C compiler by the implicit rule for C compilation.

You can define your own implicit rules by writing pattern rules.

Suffix rules are a more limited way to define implicit rules. Pattern rules are more general and clearer, but suffix rules are retained for compatibility.

• Using Implicit	How to use an existing implicit rule to get the recipes for updating a file.
• Catalogue of Rules	A list of built-in rules.
• Implicit Variables	How to change what predefined rules do.
• Chained Rules	How to use a chain of implicit rules.
• Pattern Rules	How to define new implicit rules.
• Last Resort	How to define a recipe for rules which cannot find any.
• Suffix Rules	The old-fashioned style of implicit rule.
• Implicit Rule Search	The precise algorithm for applying implicit rules.

Next: [Catalogue of Rules](#), Previous: [Implicit Rules](#), Up: [Implicit Rules](#) [\[Contents\]](#)[\[Index\]](#)

10.1 Using Implicit Rules

To allow `make` to find a customary method for updating a target file, all you have to do is refrain from specifying recipes yourself. Either write a rule with no recipe, or don't write a rule at all. Then `make` will figure out which implicit rule to use based on which kind of source file exists or can be made.

For example, suppose the makefile looks like this:

```
foo : foo.o bar.o
    cc -o foo foo.o bar.o $(CFLAGS) $(LDFLAGS)
```

Because you mention `foo.o` but do not give a rule for it, `make` will automatically look for an implicit rule that tells how to update it. This happens whether or not the file `foo.o` currently exists.

If an implicit rule is found, it can supply both a recipe and one or more prerequisites (the source files). You would want to write a rule for `foo.o` with no recipe if you need to specify additional prerequisites, such as header files, that the implicit rule cannot supply.

Each implicit rule has a target pattern and prerequisite patterns. There may be many implicit rules with the same target pattern. For example, numerous rules make `'.o'` files: one, from a `'.c'` file with the C compiler; another, from a `'.p'` file with the Pascal compiler; and so on. The rule that actually applies is the one whose prerequisites exist or can be made. So, if you have a file `foo.c`, `make` will run the C compiler; otherwise, if you have a file `foo.p`, `make` will run the Pascal compiler; and so on.

Of course, when you write the makefile, you know which implicit rule you want `make` to use, and you know it will choose that one because you know which possible prerequisite files are supposed to exist. See [Catalogue of Built-In Rules](#), for a catalogue of all the predefined implicit rules.

Above, we said an implicit rule applies if the required prerequisites "exist or can be made". A file "can be made" if it is mentioned explicitly in the makefile as a target or a prerequisite, or if an implicit rule can be recursively found for how to make it. When an implicit prerequisite is the result of another implicit rule, we say that chaining is occurring. See [Chains of Implicit Rules](#).

In general, `make` searches for an implicit rule for each target, and for each double-colon rule, that has no recipe. A file that is mentioned only as a prerequisite is considered a target whose rule specifies nothing, so implicit rule search happens for it. See [Implicit Rule Search Algorithm](#), for the details of how the search is done.

Note that explicit prerequisites do not influence implicit rule search. For example, consider this explicit rule:

```
foo.o: foo.p
```

The prerequisite on `foo.p` does not necessarily mean that `make` will remake `foo.o` according to the implicit rule to make an object file, a `.o` file, from a Pascal source file, a `.p` file. For example, if `foo.c` also exists, the implicit rule to make an object file from a C source file is used instead, because it appears before the Pascal rule in the list of predefined implicit rules (see [Catalogue of Built-In Rules](#)).

If you do not want an implicit rule to be used for a target that has no recipe, you can give that target an empty recipe by writing a semicolon (see [Defining Empty Recipes](#)).

Next: [Implicit Variables](#), Previous: [Using Implicit](#), Up: [Implicit Rules](#) [[Contents](#)][[Index](#)]

10.2 Catalogue of Built-In Rules

Here is a catalogue of predefined implicit rules which are always available unless the makefile explicitly overrides or cancels them. See [Canceling Implicit Rules](#), for information on canceling or overriding an implicit rule. The `‘-r’` or `‘--no-builtin-rules’` option cancels all predefined rules.

This manual only documents the default rules available on POSIX-based operating systems. Other operating systems, such as VMS, Windows, OS/2, etc. may have different sets of default rules. To see the full list of default rules and variables available in your version of GNU `make`, run `‘make -p’` in a directory with no makefile.

Not all of these rules will always be defined, even when the `‘-r’` option is not given. Many of the predefined implicit rules are implemented in `make` as suffix rules, so which ones will be defined depends on the suffix list (the list of prerequisites of the special target `.SUFFIXES`). The default suffix list is: `.out, .a, .ln, .o, .c, .cc, .C, .cpp, .p, .f, .F, .m, .r, .y, .l, .ym, .lm, .s, .S, .mod, .sym, .def, .h, .info, .dvi, .tex, .texinfo, .texi, .txinfo, .w, .ch, .web, .sh, .elc, .el`. All of the implicit rules described below whose prerequisites have one of these suffixes are actually suffix rules. If you modify the suffix list, the only predefined suffix rules in effect will be those named by one or two of the suffixes that are on the list you specify; rules whose suffixes fail to be on the list are disabled. See [Old-Fashioned Suffix Rules](#), for full details on suffix rules.

Compiling C programs

`n.o` is made automatically from `n.c` with a recipe of the form `‘$(CC) $(CPPFLAGS) $(CFLAGS) -c’`.

Compiling C++ programs

`n.o` is made automatically from `n.cc`, `n.cpp`, or `n.C` with a recipe of the form `‘$(CXX) $(CPPFLAGS) $(CXXFLAGS) -c’`. We encourage you to use the suffix `‘.cc’` for C++ source files instead of `‘.C’`.

Compiling Pascal programs

`n.o` is made automatically from `n.p` with the recipe `‘$(PC) $(PFLAGS) -c’`.

Compiling Fortran and Ratfor programs

`n.o` is made automatically from `n.r`, `n.F` or `n.f` by running the Fortran compiler. The precise recipe used is as follows:

```
‘.f’
‘$(FC) $(FFLAGS) -c’
‘.F’
‘$(FC) $(FFLAGS) $(CPPFLAGS) -c’
‘.r’
```

`'$(FC) $(FFLAGS) $(RFLAGS) -c'` .

Preprocessing Fortran and Ratfor programs

`n.f` is made automatically from `n.r` or `n.F`. This rule runs just the preprocessor to convert a Ratfor or preprocessable Fortran program into a strict Fortran program. The precise recipe used is as follows:

`'.F'`

`'$(FC) $(CPPFLAGS) $(FFLAGS) -F'` .

`'.r'`

`'$(FC) $(FFLAGS) $(RFLAGS) -F'` .

Compiling Modula-2 programs

`n.sym` is made from `n.def` with a recipe of the form `'$(M2C) $(M2FLAGS) $(DEFFLAGS)'` . `n.o` is made from `n.mod`; the form is: `'$(M2C) $(M2FLAGS) $(MODFLAGS)'` .

Assembling and preprocessing assembler programs

`n.o` is made automatically from `n.s` by running the assembler, `as`. The precise recipe is `'$(AS) $(ASFLAGS)'` .

`n.s` is made automatically from `n.S` by running the C preprocessor, `cpp`. The precise recipe is `'$(CPP) $(CPPFLAGS)'` .

Linking a single object file

`n` is made automatically from `n.o` by running the linker (usually called `ld`) via the C compiler. The precise recipe used is `'$(CC) $(LDFLAGS) n.o $(LOADLIBES) $(LDLIBS)'` .

This rule does the right thing for a simple program with only one source file. It will also do the right thing if there are multiple object files (presumably coming from various other source files), one of which has a name matching that of the executable file. Thus,

```
x: y.o z.o
```

when `x.c`, `y.c` and `z.c` all exist will execute:

```
cc -c x.c -o x.o
cc -c y.c -o y.o
cc -c z.c -o z.o
cc x.o y.o z.o -o x
rm -f x.o
rm -f y.o
rm -f z.o
```

In more complicated cases, such as when there is no object file whose name derives from the executable file name, you must write an explicit recipe for linking.

Each kind of file automatically made into `'.o'` object files will be automatically linked by using the compiler (`'$(CC)'` , `'$(FC)'` or `'$(PC)'` ; the C compiler `'$(CC)'` is used to assemble `'.s'` files) without the `'-c'` option. This could be done by using the `'.o'` object files as intermediates, but it is faster to do the compiling and linking in one step, so that's how it's done.

Yacc for C programs

`n.c` is made automatically from `n.y` by running Yacc with the recipe `'$(YACC) $(YFLAGS)'` .

Lex for C programs

`n.c` is made automatically from `n.l` by running Lex. The actual recipe is `'$(LEX) $(LFLAGS)'` .

Lex for Ratfor programs

`n.r` is made automatically from `n.l` by running Lex. The actual recipe is `'$(LEX) $(LFLAGS)'` .

The convention of using the same suffix `'.1'` for all Lex files regardless of whether they produce C code or Ratfor code makes it impossible for `make` to determine automatically which of the two languages you are using in any particular case. If `make` is called upon to remake an object file from a `'.1'` file, it must guess which compiler to use. It will guess the C compiler, because that is more common. If you are using Ratfor, make sure `make` knows this by mentioning `n.r` in the makefile. Or, if you are using Ratfor exclusively, with no C files, remove `'.c'` from the list of implicit rule suffixes with:

```
.SUFFIXES:
.SUFFIXES: .o .r .f .l ...
```

Making Lint Libraries from C, Yacc, or Lex programs

`n.ln` is made from `n.c` by running `lint`. The precise recipe is

`'$(LINT) $(LINTFLAGS) $(CPPFLAGS) -i'`. The same recipe is used on the C code produced from `n.y` or `n.l`.

TeX and Web

`n.dvi` is made from `n.tex` with the recipe `'$(TEX)'`. `n.tex` is made from `n.web` with

`'$(WEAVE)'`, or from `n.w` (and from `n.ch` if it exists or can be made) with `'$(CWEAVE)'`. `n.p` is made from `n.web` with `'$(TANGLE)'` and `n.c` is made from `n.w` (and from `n.ch` if it exists or can be made) with `'$(CTANGLE)'`.

Texinfo and Info

`n.dvi` is made from `n.texinfo`, `n.texi`, or `n.txinfo`, with the recipe

`'$(TEXI2DVI) $(TEXI2DVI_FLAGS)'`. `n.info` is made from `n.texinfo`, `n.texi`, or `n.txinfo`, with the recipe `'$(MAKEINFO) $(MAKEINFO_FLAGS)'`.

RCS

Any file `n` is extracted if necessary from an RCS file named either `n,v` or `RCS/n,v`. The precise recipe used is `'$(CO) $(COFLAGS)'`. `n` will not be extracted from RCS if it already exists, even if the RCS file is newer. The rules for RCS are terminal (see [Match-Anything Pattern Rules](#)), so RCS files cannot be generated from another source; they must actually exist.

SCCS

Any file `n` is extracted if necessary from an SCCS file named either `s,n` or `SCCS/s,n`. The precise recipe used is `'$(GET) $(GFLAGS)'`. The rules for SCCS are terminal (see [Match-Anything Pattern Rules](#)), so SCCS files cannot be generated from another source; they must actually exist.

For the benefit of SCCS, a file `n` is copied from `n.sh` and made executable (by everyone). This is for shell scripts that are checked into SCCS. Since RCS preserves the execution permission of a file, you do not need to use this feature with RCS.

We recommend that you avoid using of SCCS. RCS is widely held to be superior, and is also free. By choosing free software in place of comparable (or inferior) proprietary software, you support the free software movement.

Usually, you want to change only the variables listed in the table above, which are documented in the following section.

However, the recipes in built-in implicit rules actually use variables such as `COMPILE.c`, `LINK.p`, and `PREPROCESS.S`, whose values contain the recipes listed above.

`make` follows the convention that the rule to compile a `.x` source file uses the variable `COMPILE.x`. Similarly, the rule to produce an executable from a `.x` file uses `LINK.x`; and the rule to preprocess a `.x` file uses `PREPROCESS.x`.

Every rule that produces an object file uses the variable `OUTPUT_OPTION`. `make` defines this variable either to contain `'-o $@'`, or to be empty, depending on a compile-time option. You need the `'-o'` option to ensure that the output goes into the right file when the source file is in a different directory, as when using `VPATH` (see [Directory Search](#)). However, compilers on some systems do not accept a `'-o'` switch for object files. If you use such a system, and use `VPATH`,

some compilations will put their output in the wrong place. A possible workaround for this problem is to give `OUTPUT_OPTION` the value `' ; mv $*.o $@ '`.

Next: [Chained Rules](#), Previous: [Catalogue of Rules](#), Up: [Implicit Rules](#) [[Contents](#)][[Index](#)]

10.3 Variables Used by Implicit Rules

The recipes in built-in implicit rules make liberal use of certain predefined variables. You can alter the values of these variables in the makefile, with arguments to `make`, or in the environment to alter how the implicit rules work without redefining the rules themselves. You can cancel all variables used by implicit rules with the `'-R'` or `'--no-builtin-variables'` option.

For example, the recipe used to compile a C source file actually says `'$(CC) -c $(CFLAGS) $(CPPFLAGS)'`. The default values of the variables used are `'cc'` and nothing, resulting in the command `'cc -c'`. By redefining `'CC'` to `'ncc'`, you could cause `'ncc'` to be used for all C compilations performed by the implicit rule. By redefining `'CFLAGS'` to be `'-g'`, you could pass the `'-g'` option to each compilation. All implicit rules that do C compilation use `'$(CC)'` to get the program name for the compiler and all include `'$(CFLAGS)'` among the arguments given to the compiler.

The variables used in implicit rules fall into two classes: those that are names of programs (like `CC`) and those that contain arguments for the programs (like `CFLAGS`). (The “name of a program” may also contain some command arguments, but it must start with an actual executable program name.) If a variable value contains more than one argument, separate them with spaces.

The following tables describe of some of the more commonly-used predefined variables. This list is not exhaustive, and the default values shown here may not be what `make` selects for your environment. To see the complete list of predefined variables for your instance of GNU `make` you can run `'make -p'` in a directory with no makefiles.

Here is a table of some of the more common variables used as names of programs in built-in rules:

AR	Archive-maintaining program; default <code>'ar'</code> .
AS	Program for compiling assembly files; default <code>'as'</code> .
CC	Program for compiling C programs; default <code>'cc'</code> .
CXX	Program for compiling C++ programs; default <code>'g++'</code> .
CPP	Program for running the C preprocessor, with results to standard output; default <code>'\$(CC) -E'</code> .
FC	Program for compiling or preprocessing Fortran and Ratfor programs; default <code>'f77'</code> .
M2C	Program to use to compile Modula-2 source code; default <code>'m2c'</code> .
PC	Program for compiling Pascal programs; default <code>'pc'</code> .
CO	Program for extracting a file from RCS; default <code>'co'</code> .
GET	

Program for extracting a file from SCCS; default `'get'` .

LEX

Program to use to turn Lex grammars into source code; default `'lex'` .

YACC

Program to use to turn Yacc grammars into source code; default `'yacc'` .

LINT

Program to use to run lint on source code; default `'lint'` .

MAKEINFO

Program to convert a Texinfo source file into an Info file; default `'makeinfo'` .

TEX

Program to make TeX DVI files from TeX source; default `'tex'` .

TEXI2DVI

Program to make TeX DVI files from Texinfo source; default `'texi2dvi'` .

WEAVE

Program to translate Web into TeX; default `'weave'` .

CWEAVE

Program to translate C Web into TeX; default `'cweave'` .

TANGLE

Program to translate Web into Pascal; default `'tangle'` .

CTANGLE

Program to translate C Web into C; default `'ctangle'` .

RM

Command to remove a file; default `'rm -f'` .

Here is a table of variables whose values are additional arguments for the programs above. The default values for all of these is the empty string, unless otherwise noted.

ARFLAGS

Flags to give the archive-maintaining program; default `'rv'` .

ASFLAGS

Extra flags to give to the assembler (when explicitly invoked on a `'s'` or `'.s'` file).

CFLAGS

Extra flags to give to the C compiler.

CXXFLAGS

Extra flags to give to the C++ compiler.

COFLAGS

Extra flags to give to the RCS `co` program.

CPPFLAGS

Extra flags to give to the C preprocessor and programs that use it (the C and Fortran compilers).

FFLAGS

Extra flags to give to the Fortran compiler.

GFLAGS

Extra flags to give to the SCCS `get` program.

LDFLAGS

Extra flags to give to compilers when they are supposed to invoke the linker, `'ld'` , such as `-L`. Libraries (`-lfoo`) should be added to the `LDLIBS` variable instead.

LDLIBS

Library flags or names given to compilers when they are supposed to invoke the linker, `'ld' . LOADLIBES` is a deprecated (but still supported) alternative to `LDLIBS`. Non-library linker flags, such as `-L`, should go in the `LD_FLAGS` variable.

LFLAGS

Extra flags to give to Lex.

YFLAGS

Extra flags to give to Yacc.

PFLAGS

Extra flags to give to the Pascal compiler.

RFLAGS

Extra flags to give to the Fortran compiler for Ratfor programs.

LINTFLAGS

Extra flags to give to lint.

Next: [Pattern Rules](#), Previous: [Implicit Variables](#), Up: [Implicit Rules](#) [[Contents](#)][[Index](#)]

10.4 Chains of Implicit Rules

Sometimes a file can be made by a sequence of implicit rules. For example, a file `n.o` could be made from `n.y` by running first Yacc and then `cc`. Such a sequence is called a chain.

If the file `n.c` exists, or is mentioned in the makefile, no special searching is required: `make` finds that the object file can be made by C compilation from `n.c`; later on, when considering how to make `n.c`, the rule for running Yacc is used. Ultimately both `n.c` and `n.o` are updated.

However, even if `n.c` does not exist and is not mentioned, `make` knows how to envision it as the missing link between `n.o` and `n.y`! In this case, `n.c` is called an intermediate file. Once `make` has decided to use the intermediate file, it is entered in the data base as if it had been mentioned in the makefile, along with the implicit rule that says how to create it.

Intermediate files are remade using their rules just like all other files. But intermediate files are treated differently in two ways.

The first difference is what happens if the intermediate file does not exist. If an ordinary file `b` does not exist, and `make` considers a target that depends on `b`, it invariably creates `b` and then updates the target from `b`. But if `b` is an intermediate file, then `make` can leave well enough alone. It won't bother updating `b`, or the ultimate target, unless some prerequisite of `b` is newer than that target or there is some other reason to update that target.

The second difference is that if `make` does create `b` in order to update something else, it deletes `b` later on after it is no longer needed. Therefore, an intermediate file which did not exist before `make` also does not exist after `make`. `make` reports the deletion to you by printing a `'rm -f'` command showing which file it is deleting.

Ordinarily, a file cannot be intermediate if it is mentioned in the makefile as a target or prerequisite. However, you can explicitly mark a file as intermediate by listing it as a prerequisite of the special target `.INTERMEDIATE`. This takes effect even if the file is mentioned explicitly in some other way.

You can prevent automatic deletion of an intermediate file by marking it as a secondary file. To do this, list it as a prerequisite of the special target `.SECONDARY`. When a file is secondary, `make` will not create the file merely because it does not already exist, but `make` does not automatically delete the file. Marking a file as secondary also marks it as intermediate.

You can list the target pattern of an implicit rule (such as `'%.o'`) as a prerequisite of the special target `.PRECIOUS` to preserve intermediate files made by implicit rules whose target patterns match that file's name; see [Interrupts](#).

A chain can involve more than two implicit rules. For example, it is possible to make a file `foo` from `RCS/foo.y,v` by running `RCS`, `Yacc` and `cc`. Then both `foo.y` and `foo.c` are intermediate files that are deleted at the end.

No single implicit rule can appear more than once in a chain. This means that `make` will not even consider such a ridiculous thing as making `foo` from `foo.o.o` by running the linker twice. This constraint has the added benefit of preventing any infinite loop in the search for an implicit rule chain.

There are some special implicit rules to optimize certain cases that would otherwise be handled by rule chains. For example, making `foo` from `foo.c` could be handled by compiling and linking with separate chained rules, using `foo.o` as an intermediate file. But what actually happens is that a special rule for this case does the compilation and linking with a single `cc` command. The optimized rule is used in preference to the step-by-step chain because it comes earlier in the ordering of rules.

Finally, for performance reasons `make` will not consider non-terminal match-anything rules (i.e., `‘%:’`) when searching for a rule to build a prerequisite of an implicit rule (see [Match-Anything Rules](#)).

Next: [Last Resort](#), Previous: [Chained Rules](#), Up: [Implicit Rules](#) [[Contents](#)][[Index](#)]

10.5 Defining and Redefining Pattern Rules

You define an implicit rule by writing a pattern rule. A pattern rule looks like an ordinary rule, except that its target contains the character `‘%’` (exactly one of them). The target is considered a pattern for matching file names; the `‘%’` can match any nonempty substring, while other characters match only themselves. The prerequisites likewise use `‘%’` to show how their names relate to the target name.

Thus, a pattern rule `‘%.o : %.c’` says how to make any file `stem.o` from another file `stem.c`.

Note that expansion using `‘%’` in pattern rules occurs **after** any variable or function expansions, which take place when the makefile is read. See [How to Use Variables](#), and [Functions for Transforming Text](#).

• Pattern Intro	An introduction to pattern rules.
• Pattern Examples	Examples of pattern rules.
• Automatic Variables	How to use automatic variables in the recipe of implicit rules.
• Pattern Match	How patterns match.
• Match-Anything Rules	Precautions you should take prior to defining rules that can match any target file whatever.
• Canceling Rules	How to override or cancel built-in rules.

Next: [Pattern Examples](#), Previous: [Pattern Rules](#), Up: [Pattern Rules](#) [[Contents](#)][[Index](#)]

10.5.1 Introduction to Pattern Rules

A pattern rule contains the character `‘%’` (exactly one of them) in the target; otherwise, it looks exactly like an ordinary rule. The target is a pattern for matching file names; the `‘%’` matches any nonempty substring, while other characters match only themselves.

For example, `‘%.c’` as a pattern matches any file name that ends in `‘.c’`. `‘s.%.c’` as a pattern matches any file name that starts with `‘s.’`, ends in `‘.c’` and is at least five characters long. (There must be at least one character to match the `‘%’`.) The substring that the `‘%’` matches is called the stem.

`'%` in a prerequisite of a pattern rule stands for the same stem that was matched by the `'%` in the target. In order for the pattern rule to apply, its target pattern must match the file name under consideration and all of its prerequisites (after pattern substitution) must name files that exist or can be made. These files become prerequisites of the target.

Thus, a rule of the form

```
% .o : % .c ; recipe...
```

specifies how to make a file `n. o`, with another file `n. c` as its prerequisite, provided that `n. c` exists or can be made.

There may also be prerequisites that do not use `'%` ; such a prerequisite attaches to every file made by this pattern rule. These unvarying prerequisites are useful occasionally.

A pattern rule need not have any prerequisites that contain `'%` , or in fact any prerequisites at all. Such a rule is effectively a general wildcard. It provides a way to make any file that matches the target pattern. See [Last Resort](#).

More than one pattern rule may match a target. In this case `make` will choose the “best fit” rule. See [How Patterns Match](#).

Pattern rules may have more than one target; however, every target must contain a `%` character. Pattern rules are always treated as grouped targets (see [Multiple Targets in a Rule](#)) regardless of whether they use the `:` or `&:` separator.

Next: [Automatic Variables](#), Previous: [Pattern Intro](#), Up: [Pattern Rules](#) [[Contents](#)][[Index](#)]

10.5.2 Pattern Rule Examples

Here are some examples of pattern rules actually predefined in `make`. First, the rule that compiles `' .c '` files into `' .o '` files:

```
% .o : % .c
      $(CC) -c $(CFLAGS) $(CPPFLAGS) $< -o $@
```

defines a rule that can make any file `x. o` from `x. c`. The recipe uses the automatic variables `'$@'` and `'$<'` to substitute the names of the target file and the source file in each case where the rule applies (see [Automatic Variables](#)).

Here is a second built-in rule:

```
% :: RCS/% ,v
      $(CO) $(COFLAGS) $<
```

defines a rule that can make any file `x` whatsoever from a corresponding file `x, v` in the sub-directory `RCS`. Since the target is `'%` , this rule will apply to any file whatever, provided the appropriate prerequisite file exists. The double colon makes the rule terminal, which means that its prerequisite may not be an intermediate file (see [Match-Anything Pattern Rules](#)).

This pattern rule has two targets:

```
% .tab.c %.tab.h : %.y
      bison -d $<
```

This tells `make` that the recipe `'bison -d x.y'` will make both `x.tab.c` and `x.tab.h`. If the file `foo` depends on the files `parse.tab.o` and `scan.o` and the file `scan.o` depends on the file `parse.tab.h`, when `parse.y` is changed, the recipe `'bison -d parse.y'` will be executed only once, and the prerequisites of both `parse.tab.o` and `scan.o` will be satisfied. (Presumably the file `parse.tab.o` will be recompiled from `parse.tab.c` and the file `scan.o` from `scan.c`, while `foo` is linked from `parse.tab.o`, `scan.o`, and its other prerequisites, and it will execute happily ever after.)

10.5.3 Automatic Variables

Suppose you are writing a pattern rule to compile a `.c` file into a `.o` file: how do you write the `cc` command so that it operates on the right source file name? You cannot write the name in the recipe, because the name is different each time the implicit rule is applied.

What you do is use a special feature of `make`, the automatic variables. These variables have values computed afresh for each rule that is executed, based on the target and prerequisites of the rule. In this example, you would use `$(@)` for the object file name and `$(<)` for the source file name.

It's very important that you recognize the limited scope in which automatic variable values are available: they only have values within the recipe. In particular, you cannot use them anywhere within the target list of a rule; they have no value there and will expand to the empty string. Also, they cannot be accessed directly within the prerequisite list of a rule. A common mistake is attempting to use `$(@)` within the prerequisites list; this will not work. However, there is a special feature of GNU `make`, secondary expansion (see [Secondary Expansion](#)), which will allow automatic variable values to be used in prerequisite lists.

Here is a table of automatic variables:

`$(@)`

The file name of the target of the rule. If the target is an archive member, then `$(@)` is the name of the archive file. In a pattern rule that has multiple targets (see [Introduction to Pattern Rules](#)), `$(@)` is the name of whichever target caused the rule's recipe to be run.

`$(%)`

The target member name, when the target is an archive member. See [Archives](#). For example, if the target is `foo.a(bar.o)` then `$(%)` is `bar.o` and `$(@)` is `foo.a`. `$(%)` is empty when the target is not an archive member.

`$(<)`

The name of the first prerequisite. If the target got its recipe from an implicit rule, this will be the first prerequisite added by the implicit rule (see [Implicit Rules](#)).

`$(?)`

The names of all the prerequisites that are newer than the target, with spaces between them. If the target does not exist, all prerequisites will be included. For prerequisites which are archive members, only the named member is used (see [Archives](#)).

`$(^)`

The names of all the prerequisites, with spaces between them. For prerequisites which are archive members, only the named member is used (see [Archives](#)). A target has only one prerequisite on each other file it depends on, no matter how many times each file is listed as a prerequisite. So if you list a prerequisite more than once for a target, the value of `$(^)` contains just one copy of the name. This list does **not** contain any of the order-only prerequisites; for those see the `$(|)` variable, below.

`$(+)`

This is like `$(^)`, but prerequisites listed more than once are duplicated in the order they were listed in the makefile. This is primarily useful for use in linking commands where it is meaningful to repeat library file names in a particular order.

`$(|)`

The names of all the order-only prerequisites, with spaces between them.

`$(*)`

The stem with which an implicit rule matches (see [How Patterns Match](#)). If the target is `dir/a.foo.b` and the target pattern is `a.%b` then the stem is `dir/foo`. The stem is useful for constructing names of related files.

In a static pattern rule, the stem is part of the file name that matched the `'%'` in the target pattern.

In an explicit rule, there is no stem; so `'$*'` cannot be determined in that way. Instead, if the target name ends with a recognized suffix (see [Old-Fashioned Suffix Rules](#)), `'$*'` is set to the target name minus the suffix. For example, if the target name is `'foo.c'`, then `'$*'` is set to `'foo'`, since `'.c'` is a suffix. GNU `make` does this bizarre thing only for compatibility with other implementations of `make`. You should generally avoid using `'$*'` except in implicit rules or static pattern rules.

If the target name in an explicit rule does not end with a recognized suffix, `'$*'` is set to the empty string for that rule.

`'$?'` is useful even in explicit rules when you wish to operate on only the prerequisites that have changed. For example, suppose that an archive named `lib` is supposed to contain copies of several object files. This rule copies just the changed object files into the archive:

```
lib: foo.o bar.o lose.o win.o
    ar r lib $?
```

Of the variables listed above, four have values that are single file names, and three have values that are lists of file names. These seven have variants that get just the file's directory name or just the file name within the directory. The variant variables' names are formed by appending `'D'` or `'F'`, respectively. The functions `dir` and `notdir` can be used to obtain a similar effect (see [Functions for File Names](#)). Note, however, that the `'D'` variants all omit the trailing slash which always appears in the output of the `dir` function. Here is a table of the variants:

`'$(@D)'`

The directory part of the file name of the target, with the trailing slash removed. If the value of `'$@'` is `dir/foo.o` then `'$(@D)'` is `dir`. This value is `.` if `'$@'` does not contain a slash.

`'$(@F)'`

The file-within-directory part of the file name of the target. If the value of `'$@'` is `dir/foo.o` then `'$(@F)'` is `foo.o`. `'$(@F)'` is equivalent to `'$(notdir $@)'`.

`'$(*D)'`

`'$(*F)'`

The directory part and the file-within-directory part of the stem; `dir` and `foo` in this example.

`'$(%D)'`

`'$(%F)'`

The directory part and the file-within-directory part of the target archive member name. This makes sense only for archive member targets of the form `archive(member)` and is useful only when member may contain a directory name. (See [Archive Members as Targets](#).)

`'$(<D)'`

`'$(<F)'`

The directory part and the file-within-directory part of the first prerequisite.

`'$(^D)'`

`'$(^F)'`

Lists of the directory parts and the file-within-directory parts of all prerequisites.

`'$(+D)'`

`'$(+F)'`

Lists of the directory parts and the file-within-directory parts of all prerequisites, including multiple instances of duplicated prerequisites.

`'$(?D)'`

`'$(?F)'`

Lists of the directory parts and the file-within-directory parts of all prerequisites that are newer than the target.

Note that we use a special stylistic convention when we talk about these automatic variables; we write “the value of `'$<'`” , rather than “the variable `<`” as we would write for ordinary variables such as `objects` and `CFLAGS`. We think this convention looks more natural in this special case. Please do not assume it has a deep significance; `'$<'` refers to the variable named `<` just as `'$(CFLAGS)'` refers to the variable named `CFLAGS`. You could just as well use `'$(<)'` in place of `'$<'` .

Next: [Match-Anything Rules](#), Previous: [Automatic Variables](#), Up: [Pattern Rules](#) [[Contents](#)][[Index](#)]

10.5.4 How Patterns Match

A target pattern is composed of a `'%'` between a prefix and a suffix, either or both of which may be empty. The pattern matches a file name only if the file name starts with the prefix and ends with the suffix, without overlap. The text between the prefix and the suffix is called the stem. Thus, when the pattern `'%.o'` matches the file name `test.o`, the stem is `'test'` . The pattern rule prerequisites are turned into actual file names by substituting the stem for the character `'%'` . Thus, if in the same example one of the prerequisites is written as `'%.c'` , it expands to `'test.c'` .

When the target pattern does not contain a slash (and it usually does not), directory names in the file names are removed from the file name before it is compared with the target prefix and suffix. After the comparison of the file name to the target pattern, the directory names, along with the slash that ends them, are added on to the prerequisite file names generated from the pattern rule’ s prerequisite patterns and the file name. The directories are ignored only for the purpose of finding an implicit rule to use, not in the application of that rule. Thus, `'e%t'` matches the file name `src/eat`, with `'src/a'` as the stem. When prerequisites are turned into file names, the directories from the stem are added at the front, while the rest of the stem is substituted for the `'%'` . The stem `'src/a'` with a prerequisite pattern `'c%r'` gives the file name `src/car`.

A pattern rule can be used to build a given file only if there is a target pattern that matches the file name, and all prerequisites in that rule either exist or can be built. The rules you write take precedence over those that are built in. Note however, that a rule whose prerequisites actually exist or are mentioned always takes priority over a rule with prerequisites that must be made by chaining other implicit rules.

It is possible that more than one pattern rule will meet these criteria. In that case, `make` will choose the rule with the shortest stem (that is, the pattern that matches most specifically). If more than one pattern rule has the shortest stem, `make` will choose the first one found in the makefile.

This algorithm results in more specific rules being preferred over more generic ones; for example:

```
%.o: %.c
    $(CC) -c $(CFLAGS) $(CPPFLAGS) $< -o $@

%.o : %.f
    $(COMPILE.F) $(OUTPUT_OPTION) $<

lib/%.o: lib/%.c
    $(CC) -fPIC -c $(CFLAGS) $(CPPFLAGS) $< -o $@
```

Given these rules and asked to build `bar.o` where both `bar.c` and `bar.f` exist, `make` will choose the first rule and compile `bar.c` into `bar.o`. In the same situation where `bar.c` does not exist, then

`make` will choose the second rule and compile `bar.f` into `bar.o`.

If `make` is asked to build `lib/bar.o` and both `lib/bar.c` and `lib/bar.f` exist, then the third rule will be chosen since the stem for this rule (`'bar'`) is shorter than the stem for the first rule (`'lib/bar'`). If `lib/bar.c` does not exist then the third rule is not eligible and the second rule will be used, even though the stem is longer.

Next: [Canceling Rules](#), Previous: [Pattern Match](#), Up: [Pattern Rules](#) [\[Contents\]](#)[\[Index\]](#)

10.5.5 Match-Anything Pattern Rules

When a pattern rule's target is just `'%'`, it matches any file name whatever. We call these rules match-anything rules. They are very useful, but it can take a lot of time for `make` to think about them, because it must consider every such rule for each file name listed either as a target or as a prerequisite.

Suppose the makefile mentions `foo.c`. For this target, `make` would have to consider making it by linking an object file `foo.c.o`, or by C compilation-and-linking in one step from `foo.c.c`, or by Pascal compilation-and-linking from `foo.c.p`, and many other possibilities.

We know these possibilities are ridiculous since `foo.c` is a C source file, not an executable. If `make` did consider these possibilities, it would ultimately reject them, because files such as `foo.c.o` and `foo.c.p` would not exist. But these possibilities are so numerous that `make` would run very slowly if it had to consider them.

To gain speed, we have put various constraints on the way `make` considers match-anything rules. There are two different constraints that can be applied, and each time you define a match-anything rule you must choose one or the other for that rule.

One choice is to mark the match-anything rule as terminal by defining it with a double colon. When a rule is terminal, it does not apply unless its prerequisites actually exist. Prerequisites that could be made with other implicit rules are not good enough. In other words, no further chaining is allowed beyond a terminal rule.

For example, the built-in implicit rules for extracting sources from RCS and SCCS files are terminal; as a result, if the file `foo.c.v` does not exist, `make` will not even consider trying to make it as an intermediate file from `foo.c.v.o` or from `RCS/SCCS/s.foo.c.v`. RCS and SCCS files are generally ultimate source files, which should not be remade from any other files; therefore, `make` can save time by not looking for ways to remake them.

If you do not mark the match-anything rule as terminal, then it is non-terminal. A non-terminal match-anything rule cannot apply to a prerequisite of an implicit rule, or to a file name that indicates a specific type of data. A file name indicates a specific type of data if some non-match-anything implicit rule target matches it.

For example, the file name `foo.c` matches the target for the pattern rule `'%.c : %.y'` (the rule to run Yacc). Regardless of whether this rule is actually applicable (which happens only if there is a file `foo.y`), the fact that its target matches is enough to prevent consideration of any non-terminal match-anything rules for the file `foo.c`. Thus, `make` will not even consider trying to make `foo.c` as an executable file from `foo.c.o`, `foo.c.c`, `foo.c.p`, etc.

The motivation for this constraint is that non-terminal match-anything rules are used for making files containing specific types of data (such as executable files) and a file name with a recognized suffix indicates some other specific type of data (such as a C source file).

Special built-in dummy pattern rules are provided solely to recognize certain file names so that non-terminal match-anything rules will not be considered. These dummy rules have no prerequisites and no recipes, and they are ignored for all other purposes. For example, the built-in implicit rule

`%.p :`

exists to make sure that Pascal source files such as `foo.p` match a specific target pattern and thereby prevent time from being wasted looking for `foo.p.o` or `foo.p.c`.

Dummy pattern rules such as the one for `'%.p'` are made for every suffix listed as valid for use in suffix rules (see [Old-Fashioned Suffix Rules](#)).

Previous: [Match-Anything Rules](#), Up: [Pattern Rules](#) [[Contents](#)][[Index](#)]

10.5.6 Canceling Implicit Rules

You can override a built-in implicit rule (or one you have defined yourself) by defining a new pattern rule with the same target and prerequisites, but a different recipe. When the new rule is defined, the built-in one is replaced. The new rule's position in the sequence of implicit rules is determined by where you write the new rule.

You can cancel a built-in implicit rule by defining a pattern rule with the same target and prerequisites, but no recipe. For example, the following would cancel the rule that runs the assembler:

```
%.o : %.s
```

Next: [Suffix Rules](#), Previous: [Pattern Rules](#), Up: [Implicit Rules](#) [[Contents](#)][[Index](#)]

10.6 Defining Last-Resort Default Rules

You can define a last-resort implicit rule by writing a terminal match-anything pattern rule with no prerequisites (see [Match-Anything Rules](#)). This is just like any other pattern rule; the only thing special about it is that it will match any target. So such a rule's recipe is used for all targets and prerequisites that have no recipe of their own and for which no other implicit rule applies.

For example, when testing a makefile, you might not care if the source files contain real data, only that they exist. Then you might do this:

```
%. :  
    touch $@
```

to cause all the source files needed (as prerequisites) to be created automatically.

You can instead define a recipe to be used for targets for which there are no rules at all, even ones which don't specify recipes. You do this by writing a rule for the target `.DEFAULT`. Such a rule's recipe is used for all prerequisites which do not appear as targets in any explicit rule, and for which no implicit rule applies. Naturally, there is no `.DEFAULT` rule unless you write one.

If you use `.DEFAULT` with no recipe or prerequisites:

```
.DEFAULT:
```

the recipe previously stored for `.DEFAULT` is cleared. Then `make` acts as if you had never defined `.DEFAULT` at all.

If you do not want a target to get the recipe from a match-anything pattern rule or `.DEFAULT`, but you also do not want any recipe to be run for the target, you can give it an empty recipe (see [Defining Empty Recipes](#)).

You can use a last-resort rule to override part of another makefile. See [Overriding Part of Another Makefile](#).

Next: [Implicit Rule Search](#), Previous: [Last Resort](#), Up: [Implicit Rules](#) [[Contents](#)][[Index](#)]

10.7 Old-Fashioned Suffix Rules

Suffix rules are the old-fashioned way of defining implicit rules for `make`. Suffix rules are obsolete because pattern rules are more general and clearer. They are supported in GNU `make` for compatibility with old makefiles. They come in two kinds: double-suffix and single-suffix.

A double-suffix rule is defined by a pair of suffixes: the target suffix and the source suffix. It matches any file whose name ends with the target suffix. The corresponding implicit prerequisite is made by replacing the target suffix with the source suffix in the file name. A two-suffix rule `‘.c.o’` (whose target and source suffixes are `‘.o’` and `‘.c’`) is equivalent to the pattern rule `‘%.o : %.c’`.

A single-suffix rule is defined by a single suffix, which is the source suffix. It matches any file name, and the corresponding implicit prerequisite name is made by appending the source suffix. A single-suffix rule whose source suffix is `‘.c’` is equivalent to the pattern rule `‘% : %.c’`.

Suffix rule definitions are recognized by comparing each rule’s target against a defined list of known suffixes. When `make` sees a rule whose target is a known suffix, this rule is considered a single-suffix rule. When `make` sees a rule whose target is two known suffixes concatenated, this rule is taken as a double-suffix rule.

For example, `‘.c’` and `‘.o’` are both on the default list of known suffixes. Therefore, if you define a rule whose target is `‘.c.o’`, `make` takes it to be a double-suffix rule with source suffix `‘.c’` and target suffix `‘.o’`. Here is the old-fashioned way to define the rule for compiling a C source file:

```
.c.o:
    $(CC) -c $(CFLAGS) $(CPPFLAGS) -o $@ $<
```

Suffix rules cannot have any prerequisites of their own. If they have any, they are treated as normal files with funny names, not as suffix rules. Thus, the rule:

```
.c.o: foo.h
    $(CC) -c $(CFLAGS) $(CPPFLAGS) -o $@ $<
```

tells how to make the file `.c.o` from the prerequisite file `foo.h`, and is not at all like the pattern rule:

```
%.o: %.c foo.h
    $(CC) -c $(CFLAGS) $(CPPFLAGS) -o $@ $<
```

which tells how to make `‘.o’` files from `‘.c’` files, and makes all `‘.o’` files using this pattern rule also depend on `foo.h`.

Suffix rules with no recipe are also meaningless. They do not remove previous rules as do pattern rules with no recipe (see [Canceling Implicit Rules](#)). They simply enter the suffix or pair of suffixes concatenated as a target in the data base.

The known suffixes are simply the names of the prerequisites of the special target `.SUFFIXES`. You can add your own suffixes by writing a rule for `.SUFFIXES` that adds more prerequisites, as in:

```
.SUFFIXES: .hack .win
```

which adds `‘.hack’` and `‘.win’` to the end of the list of suffixes.

If you wish to eliminate the default known suffixes instead of just adding to them, write a rule for `.SUFFIXES` with no prerequisites. By special dispensation, this eliminates all existing prerequisites of `.SUFFIXES`. You can then write another rule to add the suffixes you want. For example,

```
.SUFFIXES:          # Delete the default suffixes
.SUFFIXES: .c .o .h # Define our suffix list
```

The `'-r'` or `'--no-builtin-rules'` flag causes the default list of suffixes to be empty.

The variable `SUFFIXES` is defined to the default list of suffixes before `make` reads any makefiles. You can change the list of suffixes with a rule for the special target `.SUFFIXES`, but that does not alter this variable.

Previous: [Suffix Rules](#), Up: [Implicit Rules](#) [[Contents](#)][[Index](#)]

10.8 Implicit Rule Search Algorithm

Here is the procedure `make` uses for searching for an implicit rule for a target `t`. This procedure is followed for each double-colon rule with no recipe, for each target of ordinary rules none of which have a recipe, and for each prerequisite that is not the target of any rule. It is also followed recursively for prerequisites that come from implicit rules, in the search for a chain of rules.

Suffix rules are not mentioned in this algorithm because suffix rules are converted to equivalent pattern rules once the makefiles have been read in.

For an archive member target of the form `'archive(member)'`, the following algorithm is run twice, first using the entire target name `t`, and second using `'(member)'` as the target `t` if the first run found no rule.

1. Split `t` into a directory part, called `d`, and the rest, called `n`. For example, if `t` is `'src/foo.o'`, then `d` is `'src/'` and `n` is `'foo.o'`.
2. Make a list of all the pattern rules one of whose targets matches `t` or `n`. If the target pattern contains a slash, it is matched against `t`; otherwise, against `n`.
3. If any rule in that list is not a match-anything rule, or if `t` is a prerequisite of an implicit rule, then remove all non-terminal match-anything rules from the list.
4. Remove from the list all rules with no recipe.
5. For each pattern rule in the list:
 1. Find the stem `s`, which is the nonempty part of `t` or `n` matched by the `'%'` in the target pattern.
 2. Compute the prerequisite names by substituting `s` for `'%'`; if the target pattern does not contain a slash, append `d` to the front of each prerequisite name.
 3. Test whether all the prerequisites exist or ought to exist. (If a file name is mentioned in the makefile as a target or as an explicit prerequisite, then we say it ought to exist.)

If all prerequisites exist or ought to exist, or there are no prerequisites, then this rule applies.
6. If no pattern rule has been found so far, try harder. For each pattern rule in the list:
 1. If the rule is terminal, ignore it and go on to the next rule.
 2. Compute the prerequisite names as before.
 3. Test whether all the prerequisites exist or ought to exist.
 4. For each prerequisite that does not exist, follow this algorithm recursively to see if the prerequisite can be made by an implicit rule.
 5. If all prerequisites exist, ought to exist, or can be made by implicit rules, then this rule applies.
7. If no implicit rule applies, the rule for `.DEFAULT`, if any, applies. In that case, give `t` the same recipe that `.DEFAULT` has. Otherwise, there is no recipe for `t`.

Once a rule that applies has been found, for each target pattern of the rule other than the one that matched `t` or `n`, the `'%'` in the pattern is replaced with `s` and the resultant file name

is stored until the recipe to remake the target file `t` is executed. After the recipe is executed, each of these stored file names are entered into the data base and marked as having been updated and having the same update status as the file `t`.

When the recipe of a pattern rule is executed for `t`, the automatic variables are set corresponding to the target and prerequisites. See [Automatic Variables](#).

Next: [Extending make](#), Previous: [Implicit Rules](#), Up: [Top](#) [\[Contents\]](#)[\[Index\]](#)

11 Using `make` to Update Archive Files

Archive files are files containing named sub-files called members; they are maintained with the program `ar` and their main use is as subroutine libraries for linking.

- [Archive Members](#) Archive members as targets.
- [Archive Update](#) The implicit rule for archive member targets.
- [Archive Pitfalls](#) Dangers to watch out for when using archives.
- [Archive Suffix Rules](#) You can write a special kind of suffix rule for updating archives.

Next: [Archive Update](#), Previous: [Archives](#), Up: [Archives](#) [\[Contents\]](#)[\[Index\]](#)

11.1 Archive Members as Targets

An individual member of an archive file can be used as a target or prerequisite in `make`. You specify the member named `member` in archive file `archive` as follows:

```
archive(member)
```

This construct is available only in targets and prerequisites, not in recipes! Most programs that you might use in recipes do not support this syntax and cannot act directly on archive members. Only `ar` and other programs specifically designed to operate on archives can do so. Therefore, valid recipes to update an archive member target probably must use `ar`. For example, this rule says to create a member `hack.o` in archive `foolib` by copying the file `hack.o`:

```
foolib(hack.o) : hack.o
ar cr foolib hack.o
```

In fact, nearly all archive member targets are updated in just this way and there is an implicit rule to do it for you. **Please note:** The `'c'` flag to `ar` is required if the archive file does not already exist.

To specify several members in the same archive, you can write all the member names together between the parentheses. For example:

```
foolib(hack.o kludge.o)
```

is equivalent to:

```
foolib(hack.o) foolib(kludge.o)
```

You can also use shell-style wildcards in an archive member reference. See [Using Wildcard Characters in File Names](#). For example, `'foolib(*.o)'` expands to all existing members of the `foolib` archive whose names end in `'o'`; perhaps `'foolib(hack.o) foolib(kludge.o)'`.

11.2 Implicit Rule for Archive Member Targets

Recall that a target that looks like `a(m)` stands for the member named `m` in the archive file `a`.

When `make` looks for an implicit rule for such a target, as a special feature it considers implicit rules that match `(m)`, as well as those that match the actual target `a(m)`.

This causes one special rule whose target is `(%)` to match. This rule updates the target `a(m)` by copying the file `m` into the archive. For example, it will update the archive member target `foo.a(bar.o)` by copying the file `bar.o` into the archive `foo.a` as a member named `bar.o`.

When this rule is chained with others, the result is very powerful. Thus, `'make "foo.a(bar.o)'"` (the quotes are needed to protect the `'('` and `')'` from being interpreted specially by the shell) in the presence of a file `bar.c` is enough to cause the following recipe to be run, even without a makefile:

```
cc -c bar.c -o bar.o
ar r foo.a bar.o
rm -f bar.o
```

Here `make` has envisioned the file `bar.o` as an intermediate file. See [Chains of Implicit Rules](#).

Implicit rules such as this one are written using the automatic variable `'$%'`. See [Automatic Variables](#).

An archive member name in an archive cannot contain a directory name, but it may be useful in a makefile to pretend that it does. If you write an archive member target `foo.a(dir/file.o)`, `make` will perform automatic updating with this recipe:

```
ar r foo.a dir/file.o
```

which has the effect of copying the file `dir/file.o` into a member named `file.o`. In connection with such usage, the automatic variables `%D` and `%F` may be useful.

- [Archive Symbols](#) How to update archive symbol directories.

11.2.1 Updating Archive Symbol Directories

An archive file that is used as a library usually contains a special member named `__.SYMDEF` that contains a directory of the external symbol names defined by all the other members. After you update any other members, you need to update `__.SYMDEF` so that it will summarize the other members properly. This is done by running the `ranlib` program:

```
ranlib archivefile
```

Normally you would put this command in the rule for the archive file, and make all the members of the archive file prerequisites of that rule. For example,

```
libfoo.a: libfoo.a(x.o) libfoo.a(y.o) ...
    ranlib libfoo.a
```

The effect of this is to update archive members `x.o`, `y.o`, etc., and then update the symbol directory member `__.SYMDEF` by running `ranlib`. The rules for updating the members are not shown here; most likely you can omit them and use the implicit rule which copies files into the archive, as described in the preceding section.

This is not necessary when using the GNU `ar` program, which updates the `__SYMDEF` member automatically.

Next: [Archive Suffix Rules](#), Previous: [Archive Update](#), Up: [Archives](#) [\[Contents\]](#)[\[Index\]](#)

11.3 Dangers When Using Archives

It is important to be careful when using parallel execution (the `-j` switch; see [Parallel Execution](#)) and archives. If multiple `ar` commands run at the same time on the same archive file, they will not know about each other and can corrupt the file.

Possibly a future version of `make` will provide a mechanism to circumvent this problem by serializing all recipes that operate on the same archive file. But for the time being, you must either write your makefiles to avoid this problem in some other way, or not use `-j`.

Previous: [Archive Pitfalls](#), Up: [Archives](#) [\[Contents\]](#)[\[Index\]](#)

11.4 Suffix Rules for Archive Files

You can write a special kind of suffix rule for dealing with archive files. See [Suffix Rules](#), for a full explanation of suffix rules. Archive suffix rules are obsolete in GNU `make`, because pattern rules for archives are a more general mechanism (see [Archive Update](#)). But they are retained for compatibility with other `makes`.

To write a suffix rule for archives, you simply write a suffix rule using the target suffix `'.a'` (the usual suffix for archive files). For example, here is the old-fashioned suffix rule to update a library archive from C source files:

```
.c.a:
    $(CC) $(CFLAGS) $(CPPFLAGS) -c $< -o $*.o
    $(AR) r $@ $*.o
    $(RM) $*.o
```

This works just as if you had written the pattern rule:

```
(%.o): %.c
    $(CC) $(CFLAGS) $(CPPFLAGS) -c $< -o $*.o
    $(AR) r $@ $*.o
    $(RM) $*.o
```

In fact, this is just what `make` does when it sees a suffix rule with `'.a'` as the target suffix. Any double-suffix rule `'.x.a'` is converted to a pattern rule with the target pattern `'(%.o)'` and a prerequisite pattern of `'%.x'`.

Since you might want to use `'.a'` as the suffix for some other kind of file, `make` also converts archive suffix rules to pattern rules in the normal way (see [Suffix Rules](#)). Thus a double-suffix rule `'.x.a'` produces two pattern rules: `'(%.o): %.x'` and `'%.a: %.x'`.

Next: [Integrating make](#), Previous: [Archives](#), Up: [Top](#) [\[Contents\]](#)[\[Index\]](#)

12 Extending GNU `make`

GNU `make` provides many advanced capabilities, including many useful functions. However, it does not contain a complete programming language and so it has limitations. Sometimes these limitations can be overcome through use of the `shell` function to invoke a separate program, although this can be inefficient.

In cases where the built-in capabilities of GNU `make` are insufficient to your requirements there are two options for extending `make`. On systems where it's provided, you can utilize GNU Guile as an embedded scripting language (see [GNU Guile Integration](#)). On systems which

support dynamically loadable objects, you can write your own extension in any language (which can be compiled into such an object) and load it to provide extended capabilities (see [The `load` Directive](#)).

- [Guile Integration](#) Using Guile as an embedded scripting language.
- [Loading Objects](#) Loading dynamic objects as extensions.

Next: [Loading Objects](#), Previous: [Extending make](#), Up: [Extending make](#) [\[Contents\]](#)[\[Index\]](#)

12.1 GNU Guile Integration

GNU `make` may be built with support for GNU Guile as an embedded extension language. Guile implements the Scheme language. A review of GNU Guile and the Scheme language and its features is beyond the scope of this manual: see the documentation for GNU Guile and Scheme.

You can determine if `make` contains support for Guile by examining the `.FEATURES` variable; it will contain the word `guile` if Guile support is available.

The Guile integration provides one new `make` function: `guile`. The `guile` function takes one argument which is first expanded by `make` in the normal fashion, then passed to the GNU Guile evaluator. The result of the evaluator is converted into a string and used as the expansion of the `guile` function in the makefile.

In addition, GNU `make` exposes Guile procedures for use in Guile scripts.

- [Guile Types](#) Converting Guile types to `make` strings.
- [Guile Interface](#) Invoking `make` functions from Guile.
- [Guile Example](#) Example using Guile in `make`.

Next: [Guile Interface](#), Previous: [Guile Integration](#), Up: [Guile Integration](#) [\[Contents\]](#)[\[Index\]](#)

12.1.1 Conversion of Guile Types

There is only one “data type” in `make`: a string. GNU Guile, on the other hand, provides a rich variety of different data types. An important aspect of the interface between `make` and GNU Guile is the conversion of Guile data types into `make` strings.

This conversion is relevant in two places: when a makefile invokes the `guile` function to evaluate a Guile expression, the result of that evaluation must be converted into a `make` string so it can be further evaluated by `make`. And secondly, when a Guile script invokes one of the procedures exported by `make` the argument provided to the procedure must be converted into a string.

The conversion of Guile types into `make` strings is as below:

`#f`

False is converted into the empty string: in `make` conditionals the empty string is considered false.

`#t`

True is converted to the string `'#t'`: in `make` conditionals any non-empty string is considered true.

`symbol`

`number`

A symbol or number is converted into the string representation of that symbol or number.

`character`

A printable character is converted to the same character.

`string`

A string containing only printable characters is converted to the same string.

`list`

A list is converted recursively according to the above rules. This implies that any structured list will be flattened (that is, a result of `'(a b (c d) e)'` will be converted to the `make string 'a b c d e'`).

`other`

Any other Guile type results in an error. In future versions of `make`, other Guile types may be converted.

The translation of `'#f'` (to the empty string) and `'#t'` (to the non-empty string `'#t'`) is designed to allow you to use Guile boolean results directly as `make` boolean conditions. For example:

```
$(if $(guile (access? "myfile" R_OK)),$(info myfile exists))
```

As a consequence of these conversion rules you must consider the result of your Guile script, as that result will be converted into a string and parsed by `make`. If there is no natural result for the script (that is, the script exists solely for its side-effects), you should add `'#f'` as the final expression in order to avoid syntax errors in your makefile.

Next: [Guile Example](#), Previous: [Guile Types](#), Up: [Guile Integration](#) [[Contents](#)][[Index](#)]

12.1.2 Interfaces from Guile to `make`

In addition to the `guile` function available in makefiles, `make` exposes some procedures for use in your Guile scripts. At startup `make` creates a new Guile module, `gnu make`, and exports these procedures as public interfaces from that module:

`gmkn-expand`

This procedure takes a single argument which is converted into a string. The string is expanded by `make` using normal `make` expansion rules. The result of the expansion is converted into a Guile string and provided as the result of the procedure.

`gmkn-eval`

This procedure takes a single argument which is converted into a string. The string is evaluated by `make` as if it were a makefile. This is the same capability available via the `eval` function (see [Eval Function](#)). The result of the `gmkn-eval` procedure is always the empty string.

Note that `gmkn-eval` is not quite the same as using `gmkn-expand` with the `eval` function: in the latter case the evaluated string will be expanded twice; first by `gmkn-expand`, then again by the `eval` function.

Previous: [Guile Interface](#), Up: [Guile Integration](#) [[Contents](#)][[Index](#)]

12.1.3 Example Using Guile in `make`

Here is a very simple example using GNU Guile to manage writing to a file. These Guile procedures simply open a file, allow writing to the file (one string per line), and close the file. Note that because we cannot store complex values such as Guile ports in `make` variables, we'll keep the port as a global variable in the Guile interpreter.

You can create Guile functions easily using `define/undef` to create a Guile script, then use the `guile` function to internalize it:

```
define GUILEIO
;; A simple Guile IO library for GNU make

(define MKPORT #f)
```

```

(define (mkopen name mode)
  (set! MKPORT (open-file name mode))
  #f)

(define (mkwrite s)
  (display s MKPORT)
  (newline MKPORT)
  #f)

(define (mkclose)
  (close-port MKPORT)
  #f)

#f
endif

# Internalize the Guile IO functions
$(guile $(GUILEIO))

```

If you have a significant amount of Guile support code, you might consider keeping it in a different file (e.g., `guileio.scm`) and then loading it in your makefile using the `guile` function:

```
$(guile (load "guileio.scm"))
```

An advantage to this method is that when editing `guileio.scm`, your editor will understand that this file contains Scheme syntax rather than makefile syntax.

Now you can use these Guile functions to create files. Suppose you need to operate on a very large list, which cannot fit on the command line, but the utility you’re using accepts the list as input as well:

```

prog: $(PREREQS)
    @$(guile (mkopen "tmp.out" "w")) \
      $(foreach X,$^,$(guile (mkwrite "$ (X) "))) \
      $(guile (mkclose))
    $(LINK) < tmp.out

```

A more comprehensive suite of file manipulation procedures is possible of course. You could, for example, maintain multiple output files at the same time by choosing a symbol for each one and using it as the key to a hash table, where the value is a port, then returning the symbol to be stored in a `make` variable.

Previous: [Guile Integration](#), Up: [Extending make](#) [[Contents](#)][[Index](#)]

12.2 Loading Dynamic Objects

Warning: The `load` directive and extension capability is considered a “technology preview” in this release of GNU `make`. We encourage you to experiment with this feature and we appreciate any feedback on it. However we cannot guarantee to maintain backward-compatibility in the next release. Consider using GNU Guile instead for extending GNU `make` (see [The `guile` Function](#)).

Many operating systems provide a facility for dynamically loading compiled objects. If your system provides this facility, GNU `make` can make use of it to load dynamic objects at runtime, providing new capabilities which may then be invoked by your makefile.

The `load` directive is used to load a dynamic object. Once the object is loaded, a “setup” function will be invoked to allow the object to initialize itself and register new facilities with GNU `make`. A dynamic object might include new `make` functions, for example, and the “setup” function would register them with GNU `make`’s function handling system.

- [load Directive](#) Loading dynamic objects as extensions.

- [Remaking Loaded Objects](#) How loaded objects get remade.
- [Loaded Object API](#) Programmatic interface for loaded objects.
- [Loaded Object Example](#) Example of a loaded object

Next: [Remaking Loaded Objects](#), Previous: [Loading Objects](#), Up: [Loading Objects](#) [\[Contents\]](#)[\[Index\]](#)

12.2.1 The `load` Directive

Objects are loaded into GNU `make` by placing the `load` directive into your makefile. The syntax of the `load` directive is as follows:

```
load object-file ...
```

or:

```
load object-file(symbol-name) ...
```

The file `object-file` is dynamically loaded by GNU `make`. If `object-file` does not include a directory path then it is first looked for in the current directory. If it is not found there, or a directory path is included, then system-specific paths will be searched. If the load fails for any reason, `make` will print a message and exit.

If the load succeeds `make` will invoke an initializing function.

If `symbol-name` is provided, it will be used as the name of the initializing function.

If no `symbol-name` is provided, the initializing function name is created by taking the base file name of `object-file`, up to the first character which is not a valid symbol name character (alphanumerics and underscores are valid symbol name characters). To this prefix will be appended the suffix `_gmk_setup`.

More than one object file may be loaded with a single `load` directive, and both forms of `load` arguments may be used in the same directive.

The initializing function will be provided the file name and line number of the invocation of the `load` operation. It should return a value of type `int`, which must be 0 on failure and non-0 on success. If the return value is -1, then GNU `make` will not attempt to rebuild the object file (see [How Loaded Objects Are Remade](#)).

For example:

```
load ../mk_funcs.so
```

will load the dynamic object `../mk_funcs.so`. After the object is loaded, `make` will invoke the function (assumed to be defined by the shared object) `mk_funcs_gmk_setup`.

On the other hand:

```
load ../mk_funcs.so(init_mk_func)
```

will load the dynamic object `../mk_funcs.so`. After the object is loaded, `make` will invoke the function `init_mk_func`.

Regardless of how many times an object file appears in a `load` directive, it will only be loaded (and its setup function will only be invoked) once.

After an object has been successfully loaded, its file name is appended to the `.LOADED` variable.

If you would prefer that failure to load a dynamic object not be reported as an error, you can use the `-load` directive instead of `load`. GNU `make` will not fail and no message will be generated if an object fails to load. The failed object is not added to the `.LOADED` variable, which can then be consulted to determine if the load was successful.

12.2.2 How Loaded Objects Are Remade

Loaded objects undergo the same re-make procedure as makefiles (see [How Makefiles Are Remade](#)). If any loaded object is recreated, then `make` will start from scratch and re-read all the makefiles, and reload the object files again. It is not necessary for the loaded object to do anything special to support this.

It's up to the makefile author to provide the rules needed for rebuilding the loaded object.

12.2.3 Loaded Object Interface

Warning: For this feature to be useful your extensions will need to invoke various functions internal to GNU `make`. The programming interfaces provided in this release should not be considered stable: functions may be added, removed, or change calling signatures or implementations in future versions of GNU `make`.

To be useful, loaded objects must be able to interact with GNU `make`. This interaction includes both interfaces the loaded object provides to makefiles and also interfaces `make` provides to the loaded object to manipulate `make`'s operation.

The interface between loaded objects and `make` is defined by the `gnumake.h` C header file. All loaded objects written in C should include this header file. Any loaded object not written in C will need to implement the interface defined in this header file.

Typically, a loaded object will register one or more new GNU `make` functions using the `gmk_add_function` routine from within its setup function. The implementations of these `make` functions may make use of the `gmk_expand` and `gmk_eval` routines to perform their tasks, then optionally return a string as the result of the function expansion.

Loaded Object Licensing

Every dynamic extension should define the global symbol `plugin_is_GPL_compatible` to assert that it has been licensed under a GPL-compatible license. If this symbol does not exist, `make` emits a fatal error and exits when it tries to load your extension.

The declared type of the symbol should be `int`. It does not need to be in any allocated section, though. The code merely asserts that the symbol exists in the global scope. Something like this is enough:

```
int plugin_is_GPL_compatible;
```

Data Structures

`gmk_floc`

This structure represents a filename/location pair. It is provided when defining items, so GNU `make` can inform the user later where the definition occurred if necessary.

Registering Functions

There is currently one way for makefiles to invoke operations provided by the loaded object: through the `make` function call interface. A loaded object can register one or more new functions which may then be invoked from within the makefile in the same way as any other function.

Use `gmk_add_function` to create a new `make` function. Its arguments are as follows:

`name`

The function name. This is what the makefile should use to invoke the function. The name must be between 1 and 255 characters long and it may only contain alphanumeric, period (`'.'`), dash (`'-'`), and underscore (`'_'`) characters. It may not begin with a period.

`func_ptr`

A pointer to a function that `make` will invoke when it expands the function in a makefile. This function must be defined by the loaded object.

`min_args`

The minimum number of arguments the function will accept. Must be between 0 and 255. GNU `make` will check this and fail before invoking `func_ptr` if the function was invoked with too few arguments.

`max_args`

The maximum number of arguments the function will accept. Must be between 0 and 255. GNU `make` will check this and fail before invoking `func_ptr` if the function was invoked with too few arguments. If the value is 0, then any number of arguments is accepted. If the value is greater than 0, then it must be greater than or equal to `min_args`.

`flags`

Flags that specify how this function will operate; the desired flags should be OR' d together. If the `GMK_FUNC_NOEXPAND` flag is given then the function arguments will not be expanded before the function is called; otherwise they will be expanded first.

Registered Function Interface

A function registered with `make` must match the `gm_k_func_ptr` type. It will be invoked with three parameters: `name` (the name of the function), `argc` (the number of arguments to the function), and `argv` (an array of pointers to arguments to the function). The last pointer (that is, `argv[argc]`) will be null (0).

The return value of the function is the result of expanding the function. If the function expands to nothing the return value may be null. Otherwise, it must be a pointer to a string created with `gm_k_alloc`. Once the function returns, `make` owns this string and will free it when appropriate; it cannot be accessed by the loaded object.

GNU `make` Facilities

There are some facilities exported by GNU `make` for use by loaded objects. Typically these would be run from within the setup function and/or the functions registered via `gm_k_add_function`, to retrieve or modify the data `make` works with.

`gm_k_expand`

This function takes a string and expands it using `make` expansion rules. The result of the expansion is returned in a nil-terminated string buffer. The caller is responsible for calling `gm_k_free` with a pointer to the returned buffer when done.

`gm_k_eval`

This function takes a buffer and evaluates it as a segment of makefile syntax. This function can be used to define new variables, new rules, etc. It is equivalent to using the `eval make` function.

Note that there is a difference between `gm_k_eval` and calling `gm_k_expand` with a string using the `eval` function: in the latter case the string will be expanded twice; once by `gm_k_expand` and then again by the `eval` function. Using `gm_k_eval` the buffer is only expanded once, at most (as it' s read by the `make` parser).

Memory Management

Some systems allow for different memory management schemes. Thus you should never pass memory that you' ve allocated directly to any `make` function, nor should you attempt to directly free any memory returned to you by any `make` function. Instead, use the `gm_k_alloc` and `gm_k_free` functions.

In particular, the string returned to `make` by a function registered using `gmk_add_function` must be allocated using `gmk_alloc`, and the string returned from the `make gmk_expand` function must be freed (when no longer needed) using `gmk_free`.

`gmk_alloc`

Return a pointer to a newly-allocated buffer. This function will always return a valid pointer; if not enough memory is available `make` will exit.

`gmk_free`

Free a buffer returned to you by `make`. Once the `gmk_free` function returns the string will no longer be valid.

Previous: [Loaded Object API](#), Up: [Loading Objects](#) [\[Contents\]](#)[\[Index\]](#)

12.2.4 Example Loaded Object

Let’ s suppose we wanted to write a new GNU `make` function that would create a temporary file and return its name. We would like our function to take a prefix as an argument. First we can write the function in a file `mk_temp.c`:

```
#include <stdlib.h>
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <errno.h>

#include <gnumake.h>

int plugin_is_GPL_compatible;

char *
gen_tmpfile(const char *nm, int argc, char **argv)
{
    int fd;

    /* Compute the size of the filename and allocate space for it.  */
    int len = strlen (argv[0]) + 6 + 1;
    char *buf = gmk_alloc (len);

    strcpy (buf, argv[0]);
    strcat (buf, "XXXXXX");

    fd = mkstemp(buf);
    if (fd >= 0)
    {
        /* Don't leak the file descriptor.  */
        close (fd);
        return buf;
    }

    /* Failure.  */
    fprintf (stderr, "mkstemp(%s) failed: %s\n", buf, strerror (errno));
    gmk_free (buf);
    return NULL;
}

int
mk_temp_gmk_setup ()
{
    /* Register the function with make name "mk-temp".  */
    gmk_add_function ("mk-temp", gen_tmpfile, 1, 1, 1);
    return 1;
}
```

Next, we will write a makefile that can build this shared object, load it, and use it:

```
all:
    @echo Temporary file: $(mk-temp tmpfile.)

load mk_temp.so
```

```
mk_temp.so: mk_temp.c
$(CC) -shared -fPIC -o $ $<
```

On MS-Windows, due to peculiarities of how shared objects are produced, the compiler needs to scan the import library produced when building `make`, typically called `libgnumake-version.dll.a`, where `version` is the version of the load object API. So the recipe to produce a shared object will look on Windows like this (assuming the API version is 1):

```
mk_temp.dll: mk_temp.c
$(CC) -shared -o $ $< -lgnumake-1
```

Now when you run `make` you'll see something like:

```
$ make
cc -shared -fPIC -o mk_temp.so mk_temp.c
Temporary filename: tmpfile.A7JEwd
```

Next: [Features](#), Previous: [Extending make](#), Up: [Top](#) [[Contents](#)][[Index](#)]

13 Integrating GNU `make`

GNU `make` is often one component in a larger system of tools, including integrated development environments, compiler toolchains, and others. The role of `make` is to start commands and determine whether they succeeded or not: no special integration is needed to accomplish that. However, sometimes it is convenient to bind `make` more tightly with other parts of the system, both higher-level (tools that invoke `make`) and lower-level (tools that `make` invokes).

- [Job Slots](#) Share job slots with GNU `make`.
- [Terminal Output](#) Control output to terminals.

Next: [Terminal Output](#), Previous: [Integrating make](#), Up: [Integrating make](#) [[Contents](#)][[Index](#)]

13.1 Sharing Job Slots with GNU `make`

GNU `make` has the ability to run multiple recipes in parallel (see [Parallel Execution](#)) and to cap the total number of parallel jobs even across recursive invocations of `make` (see [Communicating Options to a Sub-`make`](#)). Tools that `make` invokes which are also able to run multiple operations in parallel, either using multiple threads or multiple processes, can be enhanced to participate in GNU `make`'s job management facility to ensure that the total number of active threads/processes running on the system does not exceed the maximum number of slots provided to GNU `make`.

GNU `make` uses a method called the "jobserver" to control the number of active jobs across recursive invocations. The actual implementation of the jobserver varies across different operating systems, but some fundamental aspects are always true.

First, only command lines that `make` understands to be recursive invocations of `make` (see [How the MAKE Variable Works](#)) will have access to the jobserver. When writing makefiles you must be sure to mark the command as recursive (most commonly by prefixing the command line with the `+` indicator (see [Recursive Use of `make`](#))).

Second, `make` will provide information necessary for accessing the jobserver through the environment to its children, in the `MAKEFLAGS` environment variable. Tools which want to participate in the jobserver protocol will need to parse this environment variable, as described in subsequent sections.

Third, every command `make` starts has one implicit job slot reserved for it before it starts. Any tool which wants to participate in the jobserver protocol should assume it can always run one

job without having to contact the jobserver at all.

Finally, it's critical that tools that participate in the jobserver protocol return the exact number of slots they obtained from the jobserver back to the jobserver before they exit, even under error conditions. Remember that the implicit job slot should **not** be returned to the jobserver! Returning too few slots means that those slots will be lost for the rest of the build process; returning too many slots means that extra slots will be available. The top-level `make` command will print an error message at the end of the build if it detects an incorrect number of slots available in the jobserver.

As an example, suppose you are implementing a linker which provides for multithreaded operation. You would like to enhance the linker so that if it is invoked by GNU `make` it can participate in the jobserver protocol to control how many threads are used during link. First you will need to modify the linker to determine if the `MAKEFLAGS` environment variable is set. Next you will need to parse the value of that variable to determine if the jobserver is available, and how to access it. If it is available then you can access it to obtain job slots controlling how much parallelism your tool can use. Once done your tool must return those job slots back to the jobserver.

- [POSIX Jobserver](#) Using the jobserver on POSIX systems.
- [Windows Jobserver](#) Using the jobserver on Windows systems.

Next: [Windows Jobserver](#), Previous: [Job Slots](#), Up: [Job Slots](#) [\[Contents\]](#)[\[Index\]](#)

13.1.1 POSIX Jobserver Interaction

On POSIX systems the jobserver is implemented as a simple UNIX pipe. The pipe will be pre-loaded with one single-character token for each available job. To obtain an extra slot you must read a single character from the jobserver pipe; to release a slot you must write a single character back into the jobserver pipe. Note that the read side of the jobserver pipe is set to “blocking” mode.

To access the pipe you must parse the `MAKEFLAGS` variable and look for the argument string `--jobserver-auth=R,W` where ‘R’ and ‘W’ are non-negative integers representing file descriptors: ‘R’ is the read file descriptor and ‘W’ is the write file descriptor.

It's important that when you release the job slot, you write back the same character you read from the pipe for that slot. Don't assume that all tokens are the same character; different characters may have different meanings to GNU `make`. The order is not important, since `make` has no idea in what order jobs will complete anyway.

There are various error conditions you must consider to ensure your implementation is robust:

- Usually you will have a command-line argument controlling the parallel operation of your tool. Consider whether your tool should detect situations where both the jobserver and the command-line argument are specified, and how it should react.
- If your tool determines that the `--jobserver-auth` option is available in `MAKEFLAGS` but that the file descriptors specified are closed, this means that the calling `make` process did not think that your tool was a recursive `make` invocation (e.g., the command line was not prefixed with a `+` character). You should notify your users of this situation.
- Your tool should also examine the first word of the `MAKEFLAGS` variable and look for the character `n`. If this character is present then `make` was invoked with the `‘-n’` option and your tool should stop without performing any operations.
- Your tool should be sure to write back the tokens it read, even under error conditions. This includes not only errors in your tool but also outside influences such as interrupts (`SIGINT`), etc. You may want to install signal handlers to manage this write-back.

Previous: [POSIX Jobserver](#), Up: [Job Slots](#) [\[Contents\]](#)[\[Index\]](#)

13.1.2 Windows Jobserver Interaction

On Windows systems the jobserver is implemented as a named semaphore. The semaphore will be set with an initial count equal to the number of available slots; to obtain a slot you must wait on the semaphore (with or without a timeout). To release a slot, release the semaphore.

To access the semaphore you must parse the `MAKEFLAGS` variable and look for the argument string `--jobserver-auth=NAME` where `'NAME'` is the name of the named semaphore. Use this name with `OpenSemaphore` to create a handle to the semaphore.

There are various error conditions you must consider to ensure your implementation is robust:

- Usually you will have a command-line argument controlling the parallel operation of your tool. Consider whether your tool should detect situations where both the jobserver and the command-line argument are specified, and how it should react.
- Your tool should be sure to release the semaphore for the tokens it read, even under error conditions. This includes not only errors in your tool but also outside influences such as interrupts (`SIGINT`), etc. You may want to install signal handlers to manage this write-back.

Previous: [Job Slots](#), Up: [Integrating make](#) [[Contents](#)][[Index](#)]

13.2 Synchronized Terminal Output

Normally GNU `make` will invoke all commands with access to the same standard and error outputs that `make` itself was started with. A number of tools will detect whether the output is a terminal or not-a-terminal, and use this information to change the output style. For example if the output goes to a terminal the tool may add control characters that set color, or even change the location of the cursor. If the output is not going to a terminal then these special control characters are not emitted so that they don't corrupt log files, etc.

The `--output-sync` (see [Output During Parallel Output](#)) option will defeat the terminal detection. When output synchronization is enabled GNU `make` arranges for all command output to be written to a file, so that its output can be written as a block without interference from other commands. This means that all tools invoked by `make` will believe that their output is not going to be displayed on a terminal, even when it will be (because `make` will display it there after the command is completed).

In order to facilitate tools which would like to determine whether or not their output will be displayed on a terminal, GNU `make` will set the `MAKE_TERMOUT` and `MAKE_TERMERR` environment variables before invoking any commands. Tools which would like to determine whether standard or error output (respectively) will be displayed on a terminal can check these environment variables to determine if they exist and contain a non-empty value. If so the tool can assume that the output will (eventually) be displayed on a terminal. If the variables are not set or have an empty value, then the tool should fall back to its normal methods of detecting whether output is going to a terminal or not.

The content of the variables can be parsed to determine the type of terminal which will be used to display the output.

Similarly, environments which invoke `make` and would like to capture the output and eventually display it on a terminal (or some display which can interpret terminal control characters) can set these variables before invoking `make`. GNU `make` will not modify these environment variables if they already exist when it starts.

Next: [Missing](#), Previous: [Integrating make](#), Up: [Top](#) [[Contents](#)][[Index](#)]

14 Features of GNU `make`

Here is a summary of the features of GNU `make`, for comparison with and credit to other versions of `make`. We consider the features of `make` in 4.2 BSD systems as a baseline. If you are

concerned with writing portable makefiles, you should not use the features of `make` listed here, nor the ones in [Missing](#).

Many features come from the version of `make` in System V.

- The `VPATH` variable and its special meaning. See [Searching Directories for Prerequisites](#). This feature exists in System V `make`, but is undocumented. It is documented in 4.3 BSD `make` (which says it mimics System V's `VPATH` feature).
- Included makefiles. See [Including Other Makefiles](#). Allowing multiple files to be included with a single directive is a GNU extension.
- Variables are read from and communicated via the environment. See [Variables from the Environment](#).
- Options passed through the variable `MAKEFLAGS` to recursive invocations of `make`. See [Communicating Options to a Sub-`make`](#).
- The automatic variable `$$` is set to the member name in an archive reference. See [Automatic Variables](#).
- The automatic variables `$@`, `$*`, `$<`, `$%`, and `$?` have corresponding forms like `$(@F)` and `$(@D)`. We have generalized this to `$$` as an obvious extension. See [Automatic Variables](#).
- Substitution variable references. See [Basics of Variable References](#).
- The command line options `'-b'` and `'-m'`, accepted and ignored. In System V `make`, these options actually do something.
- Execution of recursive commands to run `make` via the variable `MAKE` even if `'-n'`, `'-q'` or `'-t'` is specified. See [Recursive Use of `make`](#).
- Support for suffix `'.a'` in suffix rules. See [Archive Suffix Rules](#). This feature is obsolete in GNU `make`, because the general feature of rule chaining (see [Chains of Implicit Rules](#)) allows one pattern rule for installing members in an archive (see [Archive Update](#)) to be sufficient.
- The arrangement of lines and backslash/newline combinations in recipes is retained when the recipes are printed, so they appear as they do in the makefile, except for the stripping of initial whitespace.

The following features were inspired by various other versions of `make`. In some cases it is unclear exactly which versions inspired which others.

- Pattern rules using `'%'`. This has been implemented in several versions of `make`. We're not sure who invented it first, but it's been spread around a bit. See [Defining and Redefining Pattern Rules](#).
- Rule chaining and implicit intermediate files. This was implemented by Stu Feldman in his version of `make` for AT&T Eighth Edition Research Unix, and later by Andrew Hume of AT&T Bell Labs in his `mk` program (where he terms it "transitive closure"). We do not really know if we got this from either of them or thought it up ourselves at the same time. See [Chains of Implicit Rules](#).
- The automatic variable `$$` containing a list of all prerequisites of the current target. We did not invent this, but we have no idea who did. See [Automatic Variables](#). The automatic variable `$$` is a simple extension of `$$`.
- The "what if" flag (`'-W'` in GNU `make`) was (as far as we know) invented by Andrew Hume in `mk`. See [Instead of Executing Recipes](#).
- The concept of doing several things at once (parallelism) exists in many incarnations of `make` and similar programs, though not in the System V or BSD implementations. See [Recipe Execution](#).
- A number of different build tools that support parallelism also support collecting output and displaying as a single block. See [Output During Parallel Execution](#).
- Modified variable references using pattern substitution come from SunOS 4. See [Basics of Variable References](#). This functionality was provided in GNU `make` by the `patsubst` function before the alternate syntax was implemented for compatibility with SunOS 4. It is not altogether clear who inspired whom, since GNU `make` had `patsubst` before SunOS 4 was released.

- The special significance of '+' characters preceding recipe lines (see [Instead of Executing Recipes](#)) is mandated by IEEE Standard 1003.2-1992 (POSIX.2).
- The '+=' syntax to append to the value of a variable comes from SunOS 4 `make`. See [Appending More Text to Variables](#).
- The syntax 'archive(mem1 mem2...)' to list multiple members in a single archive file comes from SunOS 4 `make`. See [Archive Members](#).
- The `-include` directive to include makefiles with no error for a nonexistent file comes from SunOS 4 `make`. (But note that SunOS 4 `make` does not allow multiple makefiles to be specified in one `-include` directive.) The same feature appears with the name `sinclude` in SGI `make` and perhaps others.
- The `!=` shell assignment operator exists in many BSD of `make` and is purposefully implemented here to behave identically to those implementations.
- Various build management tools are implemented using scripting languages such as Perl or Python and thus provide a natural embedded scripting language, similar to GNU `make`'s integration of GNU Guile.

The remaining features are inventions new in GNU `make`:

- Use the '-v' or '--version' option to print version and copyright information.
- Use the '-h' or '--help' option to summarize the options to `make`.
- Simply-expanded variables. See [The Two Flavors of Variables](#).
- Pass command line variable assignments automatically through the variable `MAKE` to recursive `make` invocations. See [Recursive Use of make](#).
- Use the '-C' or '--directory' command option to change directory. See [Summary of Options](#).
- Make verbatim variable definitions with `define`. See [Defining Multi-Line Variables](#).
- Declare phony targets with the special target `.PHONY`.
Andrew Hume of AT&T Bell Labs implemented a similar feature with a different syntax in his `mk` program. This seems to be a case of parallel discovery. See [Phony Targets](#).
- Manipulate text by calling functions. See [Functions for Transforming Text](#).
- Use the '-o' or '--old-file' option to pretend a file's modification-time is old. See [Avoiding Recompilation of Some Files](#).
- Conditional execution.
This feature has been implemented numerous times in various versions of `make`; it seems a natural extension derived from the features of the C preprocessor and similar macro languages and is not a revolutionary concept. See [Conditional Parts of Makefiles](#).
- Specify a search path for included makefiles. See [Including Other Makefiles](#).
- Specify extra makefiles to read with an environment variable. See [The Variable MAKEFILES](#).
- Strip leading sequences of './' from file names, so that `./file` and `file` are considered to be the same file.
- Use a special search method for library prerequisites written in the form '-lname'. See [Directory Search for Link Libraries](#).
- Allow suffixes for suffix rules (see [Old-Fashioned Suffix Rules](#)) to contain any characters. In other versions of `make`, they must begin with '.' and not contain any '/' characters.
- Keep track of the current level of `make` recursion using the variable `MAKELEVEL`. See [Recursive Use of make](#).
- Provide any goals given on the command line in the variable `MAKECMDGOALS`. See [Arguments to Specify the Goals](#).
- Specify static pattern rules. See [Static Pattern Rules](#).
- Provide selective `vpath` search. See [Searching Directories for Prerequisites](#).
- Provide computed variable references. See [Basics of Variable References](#).

- Update makefiles. See [How Makefiles Are Remade](#). System V `make` has a very, very limited form of this functionality in that it will check out SCCS files for makefiles.
- Various new built-in implicit rules. See [Catalogue of Built-In Rules](#).
- Load dynamic objects which can modify the behavior of `make`. See [Loading Dynamic Objects](#).

Next: [Makefile Conventions](#), Previous: [Features](#), Up: [Top](#) [[Contents](#)][[Index](#)]

15 Incompatibilities and Missing Features

The `make` programs in various other systems support a few features that are not implemented in GNU `make`. The POSIX.2 standard (IEEE Standard 1003.2-1992) which specifies `make` does not require any of these features.

- A target of the form `'file((entry))'` stands for a member of archive file `file`. The member is chosen, not by name, but by being an object file which defines the linker symbol entry.

This feature was not put into GNU `make` because of the non-modularity of putting knowledge into `make` of the internal format of archive file symbol tables. See [Updating Archive Symbol Directories](#).

- Suffixes (used in suffix rules) that end with the character `'~'` have a special meaning to System V `make`; they refer to the SCCS file that corresponds to the file one would get without the `'~'`. For example, the suffix rule `'c~.o'` would make the file `n.o` from the SCCS file `s,n.c`. For complete coverage, a whole series of such suffix rules is required. See [Old-Fashioned Suffix Rules](#).

In GNU `make`, this entire series of cases is handled by two pattern rules for extraction from SCCS, in combination with the general feature of rule chaining. See [Chains of Implicit Rules](#).

- In System V and 4.3 BSD `make`, files found by `VPATH` search (see [Searching Directories for Prerequisites](#)) have their names changed inside recipes. We feel it is much cleaner to always use automatic variables and thus make this feature unnecessary.
- In some Unix `make`s, the automatic variable `$(*)` appearing in the prerequisites of a rule has the amazingly strange “feature” of expanding to the full name of the target of that rule. We cannot imagine what went on in the minds of Unix `make` developers to do this; it is utterly inconsistent with the normal definition of `$(*)`.
- In some Unix `make`s, implicit rule search (see [Using Implicit Rules](#)) is apparently done for all targets, not just those without recipes. This means you can do:

```
foo.o:
    cc -c foo.c
```

and Unix `make` will intuit that `foo.o` depends on `foo.c`.

We feel that such usage is broken. The prerequisite properties of `make` are well-defined (for GNU `make`, at least), and doing such a thing simply does not fit the model.

- GNU `make` does not include any built-in implicit rules for compiling or preprocessing EFL programs. If we hear of anyone who is using EFL, we will gladly add them.
- It appears that in SVR4 `make`, a suffix rule can be specified with no recipe, and it is treated as if it had an empty recipe (see [Empty Recipes](#)). For example:

```
.c.a:
```

will override the built-in `.c.a` suffix rule.

We feel that it is cleaner for a rule without a recipe to always simply add to the prerequisite list for the target. The above example can be easily rewritten to get the

desired behavior in GNU `make`:

```
.c.a: ;
```

- Some versions of `make` invoke the shell with the `'-e'` flag, except under `'-k'` (see [Testing the Compilation of a Program](#)). The `'-e'` flag tells the shell to exit as soon as any program it runs returns a nonzero status. We feel it is cleaner to write each line of the recipe to stand on its own and not require this special treatment.

Next: [Quick Reference](#), Previous: [Missing](#), Up: [Top](#) [\[Contents\]](#)[\[Index\]](#)

16 Makefile Conventions

This describes conventions for writing the Makefiles for GNU programs. Using Automake will help you write a Makefile that follows these conventions. For more information on portable Makefiles, see POSIX and [Portable Make Programming](#) in Autoconf.

• Makefile Basics	General conventions for Makefiles.
• Utilities in Makefiles	Utilities to be used in Makefiles.
• Command Variables	Variables for specifying commands.
• DESTDIR	Supporting staged installs.
• Directory Variables	Variables for installation directories.
• Standard Targets	Standard targets for users.
• Install Command Categories	Three categories of commands in the <code>'install'</code> rule: normal, pre-install and post-install.

Next: [Utilities in Makefiles](#), Up: [Makefile Conventions](#) [\[Contents\]](#)[\[Index\]](#)

16.1 General Conventions for Makefiles

Every Makefile should contain this line:

```
SHELL = /bin/sh
```

to avoid trouble on systems where the `SHELL` variable might be inherited from the environment. (This is never a problem with GNU `make`.)

Different `make` programs have incompatible suffix lists and implicit rules, and this sometimes creates confusion or misbehavior. So it is a good idea to set the suffix list explicitly using only the suffixes you need in the particular Makefile, like this:

```
.SUFFIXES:
.SUFFIXES: .c .o
```

The first line clears out the suffix list, the second introduces all suffixes which may be subject to implicit rules in this Makefile.

Don't assume that `.` is in the path for command execution. When you need to run programs that are a part of your package during the make, please make sure that it uses `./` if the program is built as part of the make or `$(srcdir)/` if the file is an unchanging part of the source code. Without one of these prefixes, the current search path is used.

The distinction between `./` (the build directory) and `$(srcdir)/` (the source directory) is important because users can build in a separate directory using the `'--srcdir'` option to `configure`. A rule of the form:

```
foo.1 : foo.man sedscript
      sed -f sedscript foo.man > foo.1
```

will fail when the build directory is not the source directory, because `foo.man` and `sedscript` are in the source directory.

When using GNU `make`, relying on `'VPATH'` to find the source file will work in the case where there is a single dependency file, since the `make` automatic variable `'$<'` will represent the source file wherever it is. (Many versions of `make` set `'$<'` only in implicit rules.) A Makefile target like

```
foo.o : bar.c
      $(CC) -I. -I$(srcdir) $(CFLAGS) -c bar.c -o foo.o
```

should instead be written as

```
foo.o : bar.c
      $(CC) -I. -I$(srcdir) $(CFLAGS) -c $< -o $@
```

in order to allow `'VPATH'` to work correctly. When the target has multiple dependencies, using an explicit `'$(srcdir)'` is the easiest way to make the rule work well. For example, the target above for `foo.1` is best written as:

```
foo.1 : foo.man sedscript
      sed -f $(srcdir)/sedscript $(srcdir)/foo.man > $@
```

GNU distributions usually contain some files which are not source files—for example, Info files, and the output from Autoconf, Automake, Bison or Flex. Since these files normally appear in the source directory, they should always appear in the source directory, not in the build directory. So Makefile rules to update them should put the updated files in the source directory.

However, if a file does not appear in the distribution, then the Makefile should not put it in the source directory, because building a program in ordinary circumstances should not modify the source directory in any way.

Try to make the build and installation targets, at least (and all their subtargets) work correctly with a parallel `make`.

Next: [Command Variables](#), Previous: [Makefile Basics](#), Up: [Makefile Conventions](#) [[Contents](#)][[Index](#)]

16.2 Utilities in Makefiles

Write the Makefile commands (and any shell scripts, such as `configure`) to run under `sh` (both the traditional Bourne shell and the POSIX shell), not `csh`. Don't use any special features of `ksh` or `bash`, or POSIX features not widely supported in traditional Bourne `sh`.

The `configure` script and the Makefile rules for building and installation should not use any utilities directly except these:

```
awk cat cmp cp diff echo egrep expr false grep install-info ln ls
mkdir mv printf pwd rm rmdir sed sleep sort tar test touch tr true
```

Compression programs such as `gzip` can be used in the `dist` rule.

Generally, stick to the widely-supported (usually POSIX-specified) options and features of these programs. For example, don't use `'mkdir -p'`, convenient as it may be, because a few systems don't support it at all and with others, it is not safe for parallel execution. For a list of known incompatibilities, see [Portable Shell Programming](#) in Autoconf.

It is a good idea to avoid creating symbolic links in makefiles, since a few file systems don't support them.

The Makefile rules for building and installation can also use compilers and related programs, but should do so via `make` variables so that the user can substitute alternatives. Here are some of the programs we mean:

```
ar bison cc flex install ld ldconfig lex
make makeinfo ranlib texi2dvi yacc
```

Use the following `make` variables to run those programs:

```
$(AR) $(BISON) $(CC) $(FLEX) $(INSTALL) $(LD) $(LDCONFIG) $(LEX)
$(MAKE) $(MAKEINFO) $(RANLIB) $(TEXI2DVI) $(YACC)
```

When you use `ranlib` or `ldconfig`, you should make sure nothing bad happens if the system does not have the program in question. Arrange to ignore an error from that command, and print a message before the command to tell the user that failure of this command does not mean a problem. (The Autoconf `'AC_PROG_RANLIB'` macro can help with this.)

If you use symbolic links, you should implement a fallback for systems that don't have symbolic links.

Additional utilities that can be used via Make variables are:

```
chgrp chmod chown mknod
```

It is ok to use other utilities in Makefile portions (or scripts) intended only for particular systems where you know those utilities exist.

Next: [DESTDIR](#), Previous: [Utilities in Makefiles](#), Up: [Makefile Conventions](#) [[Contents](#)][[Index](#)]

16.3 Variables for Specifying Commands

Makefiles should provide variables for overriding certain commands, options, and so on.

In particular, you should run most utility programs via variables. Thus, if you use Bison, have a variable named `BISON` whose default value is set with `'BISON = bison'`, and refer to it with `$(BISON)` whenever you need to use Bison.

File management utilities such as `ln`, `rm`, `mv`, and so on, need not be referred to through variables in this way, since users don't need to replace them with other programs.

Each program-name variable should come with an options variable that is used to supply options to the program. Append `'FLAGS'` to the program-name variable name to get the options variable name—for example, `BISONFLAGS`. (The names `CFLAGS` for the C compiler, `YFLAGS` for `yacc`, and `LFLAGS` for `lex`, are exceptions to this rule, but we keep them because they are standard.) Use `CPPFLAGS` in any compilation command that runs the preprocessor, and use `LDFLAGS` in any compilation command that does linking as well as in any direct use of `ld`.

If there are C compiler options that must be used for proper compilation of certain files, do not include them in `CFLAGS`. Users expect to be able to specify `CFLAGS` freely themselves. Instead, arrange to pass the necessary options to the C compiler independently of `CFLAGS`, by writing them explicitly in the compilation commands or by defining an implicit rule, like this:

```
CFLAGS = -g
ALL_CFLAGS = -I. $(CFLAGS)
.c.o:
    $(CC) -c $(CPPFLAGS) $(ALL_CFLAGS) $<
```

Do include the `'-g'` option in `CFLAGS`, because that is not required for proper compilation. You can consider it a default that is only recommended. If the package is set up so that it is compiled with GCC by default, then you might as well include `'-O'` in the default value of `CFLAGS` as well.

Put `CFLAGS` last in the compilation command, after other variables containing compiler options, so the user can use `CFLAGS` to override the others.

`CFLAGS` should be used in every invocation of the C compiler, both those which do compilation and those which do linking.

Every Makefile should define the variable `INSTALL`, which is the basic command for installing a file into the system.

Every Makefile should also define the variables `INSTALL_PROGRAM` and `INSTALL_DATA`. (The default for `INSTALL_PROGRAM` should be `$(INSTALL)`; the default for `INSTALL_DATA` should be `${INSTALL} -m 644`.) Then it should use those variables as the commands for actual installation, for executables and non-executables respectively. Minimal use of these variables is as follows:

```
$(INSTALL_PROGRAM) foo $(bindir)/foo
$(INSTALL_DATA) libfoo.a $(libdir)/libfoo.a
```

However, it is preferable to support a `DESTDIR` prefix on the target files, as explained in the next section.

It is acceptable, but not required, to install multiple files in one command, with the final argument being a directory, as in:

```
$(INSTALL_PROGRAM) foo bar baz $(bindir)
```

Next: [Directory Variables](#), Previous: [Command Variables](#), Up: [Makefile Conventions](#) [[Contents](#)][[Index](#)]

16.4 `DESTDIR`: Support for Staged Installs

`DESTDIR` is a variable prepended to each installed target file, like this:

```
$(INSTALL_PROGRAM) foo $(DESTDIR)$ (bindir)/foo
$(INSTALL_DATA) libfoo.a $(DESTDIR)$ (libdir)/libfoo.a
```

The `DESTDIR` variable is specified by the user on the `make` command line as an absolute file name. For example:

```
make DESTDIR=/tmp/stage install
```

`DESTDIR` should be supported only in the `install*` and `uninstall*` targets, as those are the only targets where it is useful.

If your installation step would normally install `/usr/local/bin/foo` and `/usr/local/lib/libfoo.a`, then an installation invoked as in the example above would install `/tmp/stage/usr/local/bin/foo` and `/tmp/stage/usr/local/lib/libfoo.a` instead.

Prepending the variable `DESTDIR` to each target in this way provides for staged installs, where the installed files are not placed directly into their expected location but are instead copied into a temporary location (`DESTDIR`). However, installed files maintain their relative directory structure and any embedded file names will not be modified.

You should not set the value of `DESTDIR` in your `Makefile` at all; then the files are installed into their expected locations by default. Also, specifying `DESTDIR` should not change the operation of the software in any way, so its value should not be included in any file contents.

`DESTDIR` support is commonly used in package creation. It is also helpful to users who want to understand what a given package will install where, and to allow users who don't normally have permissions to install into protected areas to build and install before gaining those permissions. Finally, it can be useful with tools such as `stow`, where code is installed in one place but made to appear to be installed somewhere else using symbolic links or special mount operations. So, we strongly recommend GNU packages support `DESTDIR`, though it is not an absolute requirement.

Next: [Standard Targets](#), Previous: [DESTDIR](#), Up: [Makefile Conventions](#) [[Contents](#)][[Index](#)]

16.5 Variables for Installation Directories

Installation directories should always be named by variables, so it is easy to install in a nonstandard place. The standard names for these variables and the values they should have in GNU packages are described below. They are based on a standard file system layout; variants of it are used in GNU/Linux and other modern operating systems.

Installers are expected to override these values when calling `make` (e.g., `make prefix=/usr install`) or `configure` (e.g., `configure --prefix=/usr`). GNU packages should not try to guess which value should be appropriate for these variables on the system they are being installed onto: use the default settings specified here so that all GNU packages behave identically, allowing the installer to achieve any desired layout.

All installation directories, and their parent directories, should be created (if necessary) before they are installed into.

These first two variables set the root for the installation. All the other installation directories should be subdirectories of one of these two, and nothing should be directly installed into these two directories.

`prefix`

A prefix used in constructing the default values of the variables listed below. The default value of `prefix` should be `/usr/local`. When building the complete GNU system, the prefix will be empty and `/usr` will be a symbolic link to `/`. (If you are using Autoconf, write it as `'@prefix@'` .)

Running `'make install'` with a different value of `prefix` from the one used to build the program should not recompile the program.

`exec_prefix`

A prefix used in constructing the default values of some of the variables listed below. The default value of `exec_prefix` should be `$(prefix)`. (If you are using Autoconf, write it as `'@exec_prefix@'` .)

Generally, `$(exec_prefix)` is used for directories that contain machine-specific files (such as executables and subroutine libraries), while `$(prefix)` is used directly for other directories.

Running `'make install'` with a different value of `exec_prefix` from the one used to build the program should not recompile the program.

Executable programs are installed in one of the following directories.

`bindir`

The directory for installing executable programs that users can run. This should normally be `/usr/local/bin`, but write it as `$(exec_prefix)/bin`. (If you are using Autoconf, write it as `'@bindir@'` .)

`sbindir`

The directory for installing executable programs that can be run from the shell, but are only generally useful to system administrators. This should normally be `/usr/local/sbin`, but write it as `$(exec_prefix)/sbin`. (If you are using Autoconf, write it as `'@sbindir@'` .)

`libexecdir`

The directory for installing executable programs to be run by other programs rather than by users. This directory should normally be `/usr/local/libexec`, but write it as `$(exec_prefix)/libexec`. (If you are using Autoconf, write it as `'@libexecdir@'` .)

The definition of `'libexecdir'` is the same for all packages, so you should install your data in a subdirectory thereof. Most packages install their data under `$(libexecdir)/package-name/`, possibly within additional subdirectories thereof, such as `$(libexecdir)/package-name/machine/version`.

Data files used by the program during its execution are divided into categories in two ways.

- Some files are normally modified by programs; others are never normally modified (though users may edit some of these).

- Some files are architecture-independent and can be shared by all machines at a site; some are architecture-dependent and can be shared only by machines of the same kind and operating system; others may never be shared between two machines.

This makes for six different possibilities. However, we want to discourage the use of architecture-dependent files, aside from object files and libraries. It is much cleaner to make other data files architecture-independent, and it is generally not hard.

Here are the variables Makefiles should use to specify directories to put these various kinds of files in:

`'datarootdir'`

The root of the directory tree for read-only architecture-independent data files. This should normally be `/usr/local/share`, but write it as `$(prefix)/share`. (If you are using Autoconf, write it as `'@datarootdir@'`.) `'datadir'`'s default value is based on this variable; so are `'infodir'`, `'mandir'`, and others.

`'datadir'`

The directory for installing idiosyncratic read-only architecture-independent data files for this program. This is usually the same place as `'datarootdir'`, but we use the two separate variables so that you can move these program-specific files without altering the location for Info files, man pages, etc.

This should normally be `/usr/local/share`, but write it as `$(datarootdir)`. (If you are using Autoconf, write it as `'@datadir@'`.)

The definition of `'datadir'` is the same for all packages, so you should install your data in a subdirectory thereof. Most packages install their data under `$(datadir)/package-name/`.

`'sysconfdir'`

The directory for installing read-only data files that pertain to a single machine—that is to say, files for configuring a host. Mailer and network configuration files, `/etc/passwd`, and so forth belong here. All the files in this directory should be ordinary ASCII text files. This directory should normally be `/usr/local/etc`, but write it as `$(prefix)/etc`. (If you are using Autoconf, write it as `'@sysconfdir@'`.)

Do not install executables here in this directory (they probably belong in `$(libexecdir)` or `$(sbindir)`). Also do not install files that are modified in the normal course of their use (programs whose purpose is to change the configuration of the system excluded). Those probably belong in `$(localstatedir)`.

`'sharedstatedir'`

The directory for installing architecture-independent data files which the programs modify while they run. This should normally be `/usr/local/com`, but write it as `$(prefix)/com`. (If you are using Autoconf, write it as `'@sharedstatedir@'`.)

`'localstatedir'`

The directory for installing data files which the programs modify while they run, and that pertain to one specific machine. Users should never need to modify files in this directory to configure the package's operation; put such configuration information in separate files that go in `$(datadir)` or `$(sysconfdir)`. `$(localstatedir)` should normally be `/usr/local/var`, but write it as `$(prefix)/var`. (If you are using Autoconf, write it as `'@localstatedir@'`.)

`'runstatedir'`

The directory for installing data files which the programs modify while they run, that pertain to one specific machine, and which need not persist longer than the execution of the program—which is generally long-lived, for example, until the next reboot. PID files for system daemons are a typical use. In addition, this directory should not be cleaned except perhaps at reboot, while the general `/tmp` (`TMPDIR`) may be cleaned arbitrarily. This should normally be `/var/run`, but write it as `$(localstatedir)/run`. Having it as a separate variable allows the use of `/run` if desired, for example. (If you are using Autoconf 2.70 or later, write it as `'@runstatedir@'`.)

These variables specify the directory for installing certain specific types of files, if your program has them. Every GNU package should have Info files, so every program needs `'infodir'`, but not all need `'libdir'` or `'lispdir'`.

`'includedir'`

The directory for installing header files to be included by user programs with the C `'#include'` preprocessor directive. This should normally be `/usr/local/include`, but write it as `$(prefix)/include`. (If you are using Autoconf, write it as `'@includedir@'`.)

Most compilers other than GCC do not look for header files in directory `/usr/local/include`. So installing the header files this way is only useful with GCC.

Sometimes this is not a problem because some libraries are only really intended to work with GCC. But some libraries are intended to work with other compilers. They should install their header files in two places, one specified by `includedir` and one specified by `oldincludedir`.

`'oldincludedir'`

The directory for installing `'#include'` header files for use with compilers other than GCC. This should normally be `/usr/include`. (If you are using Autoconf, you can write it as `'@oldincludedir@'`.)

The Makefile commands should check whether the value of `oldincludedir` is empty. If it is, they should not try to use it; they should cancel the second installation of the header files.

A package should not replace an existing header in this directory unless the header came from the same package. Thus, if your Foo package provides a header file `foo.h`, then it should install the header file in the `oldincludedir` directory if either (1) there is no `foo.h` there or (2) the `foo.h` that exists came from the Foo package.

To tell whether `foo.h` came from the Foo package, put a magic string in the file—part of a comment—and `grep` for that string.

`'docdir'`

The directory for installing documentation files (other than Info) for this package. By default, it should be `/usr/local/share/doc/yourpkg`, but it should be written as `$(datarootdir)/doc/yourpkg`. (If you are using Autoconf, write it as `'@docdir@'`.) The `yourpkg` subdirectory, which may include a version number, prevents collisions among files with common names, such as `README`.

`'infodir'`

The directory for installing the Info files for this package. By default, it should be `/usr/local/share/info`, but it should be written as `$(datarootdir)/info`. (If you are using Autoconf, write it as `'@infodir@'`.) `infodir` is separate from `docdir` for compatibility with existing practice.

`'htmldir'`

`'dvidir'`

`'pdfdir'`

`'psdir'`

Directories for installing documentation files in the particular format. They should all be set to `$(docdir)` by default. (If you are using Autoconf, write them as `'@htmldir@'`, `'@dvidir@'`, etc.) Packages which supply several translations of their documentation should install them in `'$(htmldir)/' ll`, `'$(pdfdir)/' ll`, etc. where `ll` is a locale abbreviation such as `'en'` or `'pt_BR'`.

`'libdir'`

The directory for object files and libraries of object code. Do not install executables here, they probably ought to go in `$(libexecdir)` instead. The value of `libdir` should normally be `/usr/local/lib`, but write it as `$(exec_prefix)/lib`. (If you are using Autoconf, write it as `'@libdir@'`.)

`'lispdir'`

The directory for installing any Emacs Lisp files in this package. By default, it should be `/usr/local/share/emacs/site-lisp`, but it should be written as `$(datarootdir)/emacs/site-lisp`.

If you are using Autoconf, write the default as `'@lispdir@'`. In order to make `'@lispdir@'` work, you need the following lines in your `configure.ac` file:

```
lispdir='${datarootdir}/emacs/site-lisp'
AC_SUBST(lispdir)
```

`'localedir'`

The directory for installing locale-specific message catalogs for this package. By default, it should be `/usr/local/share/locale`, but it should be written as `$(datarootdir)/locale`. (If you are using Autoconf, write it as `'@localedir@'`.) This directory usually has a subdirectory per locale.

Unix-style man pages are installed in one of the following:

`'mandir'`

The top-level directory for installing the man pages (if any) for this package. It will normally be `/usr/local/share/man`, but you should write it as `$(datarootdir)/man`. (If you are using Autoconf, write it as `'@mandir@'`.)

`'man1dir'`

The directory for installing section 1 man pages. Write it as `$(mandir)/man1`.

`'man2dir'`

The directory for installing section 2 man pages. Write it as `$(mandir)/man2`

`'...'`

Don't make the primary documentation for any GNU software be a man page. Write a manual in Texinfo instead. Man pages are just for the sake of people running GNU software on Unix, which is a secondary application only.

`'manext'`

The file name extension for the installed man page. This should contain a period followed by the appropriate digit; it should normally be `'.1'`.

`'man1ext'`

The file name extension for installed section 1 man pages.

`'man2ext'`

The file name extension for installed section 2 man pages.

`'...'`

Use these names instead of `'manext'` if the package needs to install man pages in more than one section of the manual.

And finally, you should set the following variable:

`'srcdir'`

The directory for the sources being compiled. The value of this variable is normally inserted by the `configure` shell script. (If you are using Autoconf, use `'srcdir = @srcdir@'`.)

For example:

```
# Common prefix for installation directories.
# NOTE: This directory must exist when you start the install.
prefix = /usr/local
datarootdir = $(prefix)/share
datadir = $(datarootdir)
```

```
exec_prefix = $(prefix)
# Where to put the executable for the command 'gcc'.
bindir = $(exec_prefix)/bin
# Where to put the directories used by the compiler.
libexecdir = $(exec_prefix)/libexec
# Where to put the Info files.
infodir = $(datarootdir)/info
```

If your program installs a large number of files into one of the standard user-specified directories, it might be useful to group them into a subdirectory particular to that program. If you do this, you should write the `install` rule to create these subdirectories.

Do not expect the user to include the subdirectory name in the value of any of the variables listed above. The idea of having a uniform set of variable names for installation directories is to enable the user to specify the exact same values for several different GNU packages. In order for this to be useful, all the packages must be designed so that they will work sensibly when the user does so.

At times, not all of these variables may be implemented in the current release of Autoconf and/or Automake; but as of Autoconf 2.60, we believe all of them are. When any are missing, the descriptions here serve as specifications for what Autoconf will implement. As a programmer, you can either use a development version of Autoconf or avoid using these variables until a stable release is made which supports them.

Next: [Install Command Categories](#), Previous: [Directory Variables](#), Up: [Makefile Conventions](#) [[Contents](#)][[Index](#)]

16.6 Standard Targets for Users

All GNU programs should have the following targets in their Makefiles:

`'all'`

Compile the entire program. This should be the default target. This target need not rebuild any documentation files; Info files should normally be included in the distribution, and DVI (and other documentation format) files should be made only when explicitly asked for.

By default, the Make rules should compile and link with `'-g'`, so that executable programs have debugging symbols. Otherwise, you are essentially helpless in the face of a crash, and it is often far from easy to reproduce with a fresh build.

`'install'`

Compile the program and copy the executables, libraries, and so on to the file names where they should reside for actual use. If there is a simple test to verify that a program is properly installed, this target should run that test.

Do not strip executables when installing them. This helps eventual debugging that may be needed later, and nowadays disk space is cheap and dynamic loaders typically ensure debug sections are not loaded during normal execution. Users that need stripped binaries may invoke the `install-strip` target to do that.

If possible, write the `install` target rule so that it does not modify anything in the directory where the program was built, provided `'make all'` has just been done. This is convenient for building the program under one user name and installing it under another.

The commands should create all the directories in which files are to be installed, if they don't already exist. This includes the directories specified as the values of the variables `prefix` and `exec_prefix`, as well as all subdirectories that are needed. One way to do this is by means of an `installdirs` target as described below.

Use `'-'` before any command for installing a man page, so that `make` will ignore any errors. This is in case there are systems that don't have the Unix man page documentation system installed.

The way to install Info files is to copy them into `$(infodir)` with `$(INSTALL_DATA)` (see [Command Variables](#)), and then run the `install-info` program if it is present. `install-info` is

a program that edits the `Info dir` file to add or update the menu entry for the given Info file; it is part of the Texinfo package.

Here is a sample rule to install an Info file that also tries to handle some additional situations, such as `install-info` not being present.

```
do-install-info: foo.info installdirs
    $(NORMAL_INSTALL)
# Prefer an info file in . to one in srcdir.
    if test -f foo.info; then d=. ; \
    else d="$(srcdir)"; fi; \
    $(INSTALL_DATA) $$d/foo.info \
    "$(DESTDIR)$(infodir)/foo.info"
# Run install-info only if it exists.
# Use 'if' instead of just prepending '-' to the
# line so we notice real errors from install-info.
# Use '$(SHELL) -c' because some shells do not
# fail gracefully when there is an unknown command.
    $(POST_INSTALL)
    if $(SHELL) -c 'install-info --version' \
    >/dev/null 2>&1; then \
    install-info --dir-file="$(DESTDIR)$(infodir)/dir" \
    "$(DESTDIR)$(infodir)/foo.info"; \
    else true; fi
```

When writing the `install` target, you must classify all the commands into three categories: normal ones, pre-installation commands and post-installation commands. See [Install Command Categories](#).

```
'install-html'
```

```
'install-dvi'
```

```
'install-pdf'
```

```
'install-ps'
```

These targets install documentation in formats other than Info; they're intended to be called explicitly by the person installing the package, if that format is desired. GNU prefers Info files, so these must be installed by the `install` target.

When you have many documentation files to install, we recommend that you avoid collisions and clutter by arranging for these targets to install in subdirectories of the appropriate installation directory, such as `htmldir`. As one example, if your package has multiple manuals, and you wish to install HTML documentation with many files (such as the "split" mode output by `makeinfo --html`), you'll certainly want to use subdirectories, or two nodes with the same name in different manuals will overwrite each other.

Please make these `install-format` targets invoke the commands for the format target, for example, by making format a dependency.

```
'uninstall'
```

Delete all the installed files—the copies that the `'install'` and `'install-*` targets create.

This rule should not modify the directories where compilation is done, only the directories where files are installed.

The uninstallation commands are divided into three categories, just like the installation commands. See [Install Command Categories](#).

```
'install-strip'
```

Like `install`, but strip the executable files while installing them. In simple cases, this target can use the `install` target in a simple way:

```
install-strip:
    $(MAKE) INSTALL_PROGRAM='$(INSTALL_PROGRAM) -s' \
    install
```

But if the package installs scripts as well as real executables, the `install-strip` target can't just refer to the `install` target; it has to strip the executables but not the scripts.

`install-strip` should not strip the executables in the build directory which are being copied for installation. It should only strip the copies that are installed.

Normally we do not recommend stripping an executable unless you are sure the program has no bugs. However, it can be reasonable to install a stripped executable for actual execution while saving the unstripped executable elsewhere in case there is a bug.

`'clean'`

Delete all files in the current directory that are normally created by building the program. Also delete files in other directories if they are created by this makefile. However, don't delete the files that record the configuration. Also preserve files that could be made by building, but normally aren't because the distribution comes with them. There is no need to delete parent directories that were created with `'mkdir -p'`, since they could have existed anyway.

Delete `.dvi` files here if they are not part of the distribution.

`'distclean'`

Delete all files in the current directory (or created by this makefile) that are created by configuring or building the program. If you have unpacked the source and built the program without creating any other files, `'make distclean'` should leave only the files that were in the distribution. However, there is no need to delete parent directories that were created with `'mkdir -p'`, since they could have existed anyway.

`'mostlyclean'`

Like `'clean'`, but may refrain from deleting a few files that people normally don't want to recompile. For example, the `'mostlyclean'` target for GCC does not delete `libgcc.a`, because recompiling it is rarely necessary and takes a lot of time.

`'maintainer-clean'`

Delete almost everything that can be reconstructed with this Makefile. This typically includes everything deleted by `distclean`, plus more: C source files produced by Bison, tags tables, Info files, and so on.

The reason we say "almost everything" is that running the command `'make maintainer-clean'` should not delete `configure` even if `configure` can be remade using a rule in the Makefile. More generally, `'make maintainer-clean'` should not delete anything that needs to exist in order to run `configure` and then begin to build the program. Also, there is no need to delete parent directories that were created with `'mkdir -p'`, since they could have existed anyway. These are the only exceptions; `maintainer-clean` should delete everything else that can be rebuilt.

The `'maintainer-clean'` target is intended to be used by a maintainer of the package, not by ordinary users. You may need special tools to reconstruct some of the files that `'make maintainer-clean'` deletes. Since these files are normally included in the distribution, we don't take care to make them easy to reconstruct. If you find you need to unpack the full distribution again, don't blame us.

To help make users aware of this, the commands for the special `maintainer-clean` target should start with these two:

```
@echo 'This command is intended for maintainers to use; it'
@echo 'deletes files that may need special tools to rebuild.'
```

`'TAGS'`

Update a tags table for this program.

`'info'`

Generate any Info files needed. The best way to write the rules is as follows:

```
info: foo.info
```

```
foo.info: foo.texi chap1.texi chap2.texi
$(MAKEINFO) $(srcdir)/foo.texi
```

You must define the variable `MAKEINFO` in the Makefile. It should run the `makeinfo` program, which is part of the Texinfo distribution.

Normally a GNU distribution comes with Info files, and that means the Info files are present in the source directory. Therefore, the Make rule for an info file should update it in the source directory. When users build the package, ordinarily Make will not update the Info files because they will already be up to date.

```
‘dvi’
```

```
‘html’
```

```
‘pdf’
```

```
‘ps’
```

Generate documentation files in the given format. These targets should always exist, but any or all can be a no-op if the given output format cannot be generated. These targets should not be dependencies of the `all` target; the user must manually invoke them.

Here’s an example rule for generating DVI files from Texinfo:

```
dvi: foo.dvi

foo.dvi: foo.texi chap1.texi chap2.texi
$(TEXI2DVI) $(srcdir)/foo.texi
```

You must define the variable `TEXI2DVI` in the Makefile. It should run the program `texi2dvi`, which is part of the Texinfo distribution. (`texi2dvi` uses TeX to do the real work of formatting. TeX is not distributed with Texinfo.) Alternatively, write only the dependencies, and allow GNU `make` to provide the command.

Here’s another example, this one for generating HTML from Texinfo:

```
html: foo.html

foo.html: foo.texi chap1.texi chap2.texi
$(TEXI2HTML) $(srcdir)/foo.texi
```

Again, you would define the variable `TEXI2HTML` in the Makefile; for example, it might run `makeinfo --no-split --html` (`makeinfo` is part of the Texinfo distribution).

```
‘dist’
```

Create a distribution tar file for this program. The tar file should be set up so that the file names in the tar file start with a subdirectory name which is the name of the package it is a distribution for. This name can include the version number.

For example, the distribution tar file of GCC version 1.40 unpacks into a subdirectory named `gcc-1.40`.

The easiest way to do this is to create a subdirectory appropriately named, use `ln` or `cp` to install the proper files in it, and then `tar` that subdirectory.

Compress the tar file with `gzip`. For example, the actual distribution file for GCC version 1.40 is called `gcc-1.40.tar.gz`. It is ok to support other free compression formats as well.

The `dist` target should explicitly depend on all non-source files that are in the distribution, to make sure they are up to date in the distribution. See [Making Releases](#) in GNU Coding Standards.

```
‘check’
```

Perform self-tests (if any). The user must build the program before running the tests, but need not install the program; you should write the self-tests so that they work when the program is built but not installed.

The following targets are suggested as conventional names, for programs in which they are useful.

installcheck

Perform installation tests (if any). The user must build and install the program before running the tests. You should not assume that `$(bindir)` is in the search path.

installdirs

It's useful to add a target named `'installdirs'` to create the directories where files are installed, and their parent directories. There is a script called `mkinstalldirs` which is convenient for this; you can find it in the Gnulib package. You can use a rule like this:

```
# Make sure all installation directories (e.g. $(bindir))
# actually exist by making them if necessary.
installdirs: mkinstalldirs
    $(srcdir)/mkinstalldirs $(bindir) $(datadir) \
                                $(libdir) $(infodir) \
                                $(mandir)
```

or, if you wish to support `DESTDIR` (strongly encouraged),

```
# Make sure all installation directories (e.g. $(bindir))
# actually exist by making them if necessary.
installdirs: mkinstalldirs
    $(srcdir)/mkinstalldirs \
        $(DESTDIR)$(bindir) $(DESTDIR)$(datadir) \
        $(DESTDIR)$(libdir) $(DESTDIR)$(infodir) \
        $(DESTDIR)$(mandir)
```

This rule should not modify the directories where compilation is done. It should do nothing but create installation directories.

Previous: [Standard Targets](#), Up: [Makefile Conventions](#) [[Contents](#)][[Index](#)]

16.7 Install Command Categories

When writing the `install` target, you must classify all the commands into three categories: normal ones, pre-installation commands and post-installation commands.

Normal commands move files into their proper places, and set their modes. They may not alter any files except the ones that come entirely from the package they belong to.

Pre-installation and post-installation commands may alter other files; in particular, they can edit global configuration files or data bases.

Pre-installation commands are typically executed before the normal commands, and post-installation commands are typically run after the normal commands.

The most common use for a post-installation command is to run `install-info`. This cannot be done with a normal command, since it alters a file (the Info directory) which does not come entirely and solely from the package being installed. It is a post-installation command because it needs to be done after the normal command which installs the package's Info files.

Most programs don't need any pre-installation commands, but we have the feature just in case it is needed.

To classify the commands in the `install` rule into these three categories, insert category lines among them. A category line specifies the category for the commands that follow.

A category line consists of a tab and a reference to a special Make variable, plus an optional comment at the end. There are three variables you can use, one for each category; the variable name specifies the category. Category lines are no-ops in ordinary execution because these three Make variables are normally undefined (and you should not define them in the makefile).

Here are the three possible category lines, each with a comment that explains what it means:

```
$(PRE_INSTALL)      # Pre-install commands follow.
$(POST_INSTALL)     # Post-install commands follow.
$(NORMAL_INSTALL)   # Normal commands follow.
```

If you don’ t use a category line at the beginning of the `install` rule, all the commands are classified as normal until the first category line. If you don’ t use any category lines, all the commands are classified as normal.

These are the category lines for `uninstall`:

```
$(PRE_UNINSTALL)    # Pre-uninstall commands follow.
$(POST_UNINSTALL)   # Post-uninstall commands follow.
$(NORMAL_UNINSTALL) # Normal commands follow.
```

Typically, a pre-uninstall command would be used for deleting entries from the Info directory.

If the `install` or `uninstall` target has any dependencies which act as subroutines of installation, then you should start each dependency’ s commands with a category line, and start the main target’ s commands with a category line also. This way, you can ensure that each command is placed in the right category regardless of which of the dependencies actually run.

Pre-installation and post-installation commands should not run any programs except for these:

```
[ basename bash cat chgrp chmod chown cmp cp dd diff echo
egrep expand expr false fgrep find getopt grep gunzip gzip
hostname install install-info kill ldconfig ln ls md5sum
mkdir mkfifo mknod mv printenv pwd rm rmdir sed sort tee
test touch true uname xargs yes
```

The reason for distinguishing the commands in this way is for the sake of making binary packages. Typically a binary package contains all the executables and other files that need to be installed, and has its own method of installing them—so it does not need to run the normal installation commands. But installing the binary package does need to execute the pre-installation and post-installation commands.

Programs to build binary packages work by extracting the pre-installation and post-installation commands. Here is one way of extracting the pre-installation commands (the `-s` option to `make` is needed to silence messages about entering subdirectories):

```
make -s -n install -o all \
    PRE_INSTALL=pre-install \
    POST_INSTALL=post-install \
    NORMAL_INSTALL=normal-install \
    | gawk -f pre-install.awk
```

where the file `pre-install.awk` could contain this:

```
$0 ~ /^ (normal-install|post-install) [ \t]*$/ {on = 0}
on {print $0}
$0 ~ /^pre-install [ \t]*$/ {on = 1}
```

Next: [Error Messages](#), Previous: [Makefile Conventions](#), Up: [Top](#) [\[Contents\]](#)[\[Index\]](#)

Appendix A Quick Reference

This appendix summarizes the directives, text manipulation functions, and special variables which GNU `make` understands. See [Special Targets](#), [Catalogue of Built-In Rules](#), and [Summary of Options](#), for other summaries.

Here is a summary of the directives GNU `make` recognizes:

```
define variable
define variable =
```

```
define variable :=
define variable ::=
define variable +=
define variable ?=
endif
```

Define multi-line variables.
See [Multi-Line](#).

```
undefine variable
```

Undefining variables.
See [Undefine Directive](#).

```
ifdef variable
ifndef variable
ifeq (a,b)
ifeq "a" "b"
ifeq 'a' 'b'
ifneq (a,b)
ifneq "a" "b"
ifneq 'a' 'b'
else
endif
```

Conditionally evaluate part of the makefile.
See [Conditionals](#).

```
include file
-include file
sinclude file
```

Include another makefile.
See [Including Other Makefiles](#).

```
override variable-assignment
```

Define a variable, overriding any previous definition, even one from the command line.
See [The `override` Directive](#).

```
export
```

Tell `make` to export all variables to child processes by default.
See [Communicating Variables to a Sub-`make`](#).

```
export variable
export variable-assignment
unexport variable
```

Tell `make` whether or not to export a particular variable to child processes.
See [Communicating Variables to a Sub-`make`](#).

```
private variable-assignment
```

Do not allow this variable assignment to be inherited by prerequisites.
See [Suppressing Inheritance](#).

```
vpath pattern path
```

Specify a search path for files matching a `'%'` pattern.
See [The `vpath` Directive](#).

```
vpath pattern
```

Remove all search paths previously specified for pattern.

```
vpath
```

Remove all search paths previously specified in any `vpath` directive.

Here is a summary of the built-in functions (see [Functions](#)):

```
$(subst from,to,text)
```

Replace from with to in text.

See [Functions for String Substitution and Analysis](#).

```
$(patsubst pattern,replacement,text)
```

Replace words matching pattern with replacement in text.

See [Functions for String Substitution and Analysis](#).

```
$(strip string)
```

Remove excess whitespace characters from string.

See [Functions for String Substitution and Analysis](#).

```
$(findstring find,text)
```

Locate find in text.

See [Functions for String Substitution and Analysis](#).

```
$(filter pattern...,text)
```

Select words in text that match one of the pattern words.

See [Functions for String Substitution and Analysis](#).

```
$(filter-out pattern...,text)
```

Select words in text that do not match any of the pattern words.

See [Functions for String Substitution and Analysis](#).

```
$(sort list)
```

Sort the words in list lexicographically, removing duplicates.

See [Functions for String Substitution and Analysis](#).

```
$(word n,text)
```

Extract the nth word (one-origin) of text.

See [Functions for String Substitution and Analysis](#).

```
$(words text)
```

Count the number of words in text.

See [Functions for String Substitution and Analysis](#).

```
$(wordlist s,e,text)
```

Returns the list of words in text from s to e.

See [Functions for String Substitution and Analysis](#).

```
$(firstword names...)
```

Extract the first word of names.

See [Functions for String Substitution and Analysis](#).

```
$(lastword names...)
```

Extract the last word of names.

See [Functions for String Substitution and Analysis](#).

```
$(dir names...)
```

Extract the directory part of each file name.

See [Functions for File Names](#).

```
$(notdir names...)
```

Extract the non-directory part of each file name.

See [Functions for File Names](#).

```
$(suffix names...)
```

Extract the suffix (the last ' .' and following characters) of each file name.

See [Functions for File Names](#).

```
$(basename names...)
```

Extract the base name (name without suffix) of each file name.

See [Functions for File Names](#).

```
$(addsuffix suffix,names...)
```

Append suffix to each word in names.

See [Functions for File Names](#).

`$(addprefix prefix, names...)`

Prepend prefix to each word in names.

See [Functions for File Names](#).

`$(join list1, list2)`

Join two parallel lists of words.

See [Functions for File Names](#).

`$(wildcard pattern...)`

Find file names matching a shell file name pattern (not a `'%'` pattern).

See [The Function](#) `wildcard`.

`$(realpath names...)`

For each file name in names, expand to an absolute name that does not contain any `.`, `..`, nor symlinks.

See [Functions for File Names](#).

`$(abspath names...)`

For each file name in names, expand to an absolute name that does not contain any `.` or `..` components, but preserves symlinks.

See [Functions for File Names](#).

`$(error text...)`

When this function is evaluated, `make` generates a fatal error with the message text.

See [Functions That Control Make](#).

`$(warning text...)`

When this function is evaluated, `make` generates a warning with the message text.

See [Functions That Control Make](#).

`$(shell command)`

Execute a shell command and return its output.

See [The `shell` Function](#).

`$(origin variable)`

Return a string describing how the `make` variable variable was defined.

See [The `origin` Function](#).

`$(flavor variable)`

Return a string describing the flavor of the `make` variable variable.

See [The `flavor` Function](#).

`$(foreach var, words, text)`

Evaluate text with var bound to each word in words, and concatenate the results.

See [The `foreach` Function](#).

`$(if condition, then-part[, else-part])`

Evaluate the condition condition; if it's non-empty substitute the expansion of the then-part otherwise substitute the expansion of the else-part.

See [Functions for Conditionals](#).

`$(or condition1[, condition2[, condition3...]])`

Evaluate each condition conditionN one at a time; substitute the first non-empty expansion. If all expansions are empty, substitute the empty string.

See [Functions for Conditionals](#).

`$(and condition1[, condition2[, condition3...]])`

Evaluate each condition conditionN one at a time; if any expansion results in the empty string substitute the empty string. If all expansions result in a non-empty string, substitute the expansion of the last condition.

See [Functions for Conditionals](#).

`$(call var, param, ...)`

Evaluate the variable var replacing any references to `$(1)`, `$(2)` with the first, second, etc. param values.

See [The `call` Function](#).

`$(eval text)`

Evaluate text then read the results as makefile commands. Expands to the empty string.

See [The `eval` Function](#).

`$(file op filename, text)`

Expand the arguments, then open the file filename using mode op and write text to that file.

See [The `file` Function](#).

`$(value var)`

Evaluates to the contents of the variable var, with no expansion performed on it.

See [The `value` Function](#).

Here is a summary of the automatic variables. See [Automatic Variables](#), for full information.

`$@`

The file name of the target.

`$%`

The target member name, when the target is an archive member.

`$<`

The name of the first prerequisite.

`$?`

The names of all the prerequisites that are newer than the target, with spaces between them. For prerequisites which are archive members, only the named member is used (see [Archives](#)).

`^`

`$+`

The names of all the prerequisites, with spaces between them. For prerequisites which are archive members, only the named member is used (see [Archives](#)). The value of `^` omits duplicate prerequisites, while `+` retains them and preserves their order.

`$*`

The stem with which an implicit rule matches (see [How Patterns Match](#)).

`$(@D)`

`$(@F)`

The directory part and the file-within-directory part of `$@`.

`$(*D)`

`$(*F)`

The directory part and the file-within-directory part of `$*`.

`$(%D)`

`$(%F)`

The directory part and the file-within-directory part of `$%`.

`$(<D)`

`$(<F)`

The directory part and the file-within-directory part of `$<`.

`$(^D)`

`$(^F)`

The directory part and the file-within-directory part of `^`.

`$(+D)`

`$(+F)`

The directory part and the file-within-directory part of `$+`.

`$(?D)`

`$(?F)`

The directory part and the file-within-directory part of `$(?)`.

These variables are used specially by GNU `make`:

`MAKEFILES`

Makefiles to be read on every invocation of `make`.

See [The Variable `MAKEFILES`](#).

`VPATH`

Directory search path for files not found in the current directory.

See [VPATH Search Path for All Prerequisites](#).

`SHELL`

The name of the system default command interpreter, usually `/bin/sh`. You can set `SHELL` in the makefile to change the shell used to run recipes. See [Recipe Execution](#). The `SHELL` variable is handled specially when importing from and exporting to the environment.

See [Choosing the Shell](#).

`MAKESHELL`

On MS-DOS only, the name of the command interpreter that is to be used by `make`. This value takes precedence over the value of `SHELL`. See [MAKESHELL variable](#).

`MAKE`

The name with which `make` was invoked. Using this variable in recipes has special meaning. See [How the `MAKE` Variable Works](#).

`MAKE_VERSION`

The built-in variable `'MAKE_VERSION'` expands to the version number of the GNU `make` program.

`MAKE_HOST`

The built-in variable `'MAKE_HOST'` expands to a string representing the host that GNU `make` was built to run on.

`MAKELEVEL`

The number of levels of recursion (sub-`makes`).

See [Variables/Recursion](#).

`MAKEFLAGS`

The flags given to `make`. You can set this in the environment or a makefile to set flags.

See [Communicating Options to a Sub-`make`](#).

It is never appropriate to use `MAKEFLAGS` directly in a recipe line: its contents may not be quoted correctly for use in the shell. Always allow recursive `make`'s to obtain these values through the environment from its parent.

`GNUMAKEFLAGS`

Other flags parsed by `make`. You can set this in the environment or a makefile to set `make` command-line flags. GNU `make` never sets this variable itself. This variable is only needed if you'd like to set GNU `make`-specific flags in a POSIX-compliant makefile. This variable will be seen by GNU `make` and ignored by other `make` implementations. It's not needed if you only use GNU `make`; just use `MAKEFLAGS` directly. See [Communicating Options to a Sub-`make`](#).

`MAKECMDGOALS`

The targets given to `make` on the command line. Setting this variable has no effect on the operation of `make`.

See [Arguments to Specify the Goals](#).

`CURDIR`

Set to the absolute pathname of the current working directory (after all `-C` options are processed, if any). Setting this variable has no effect on the operation of `make`.

See [Recursive Use of `make`](#).

`SUFFIXES`

The default list of suffixes before `make` reads any makefiles.

`.LIBPATTERNS`

Defines the naming of the libraries `make` searches for, and their order.
See [Directory Search for Link Libraries](#).

Next: [Complex Makefile](#), Previous: [Quick Reference](#), Up: [Top](#) [[Contents](#)][[Index](#)]

Appendix B Errors Generated by Make

Here is a list of the more common errors you might see generated by `make`, and some information about what they mean and how to fix them.

Sometimes `make` errors are not fatal, especially in the presence of a `-` prefix on a recipe line, or the `-k` command line option. Errors that are fatal are prefixed with the string `***`.

Error messages are all either prefixed with the name of the program (usually `'make'`), or, if the error is found in a makefile, the name of the file and line number containing the problem.

In the table below, these common prefixes are left off.

`'[foo] Error NN'`

`'[foo] signal description'`

These errors are not really `make` errors at all. They mean that a program that `make` invoked as part of a recipe returned a non-0 error code (`'Error NN'`), which `make` interprets as failure, or it exited in some other abnormal fashion (with a signal of some type). See [Errors in Recipes](#).

If no `***` is attached to the message, then the sub-process failed but the rule in the makefile was prefixed with the `-` special character, so `make` ignored the error.

`'missing separator. Stop.'`

`'missing separator (did you mean TAB instead of 8 spaces?). Stop.'`

This means that `make` could not understand much of anything about the makefile line it just read. GNU `make` looks for various separators (`:`, `=`, recipe prefix characters, etc.) to indicate what kind of line it's parsing. This message means it couldn't find a valid one.

One of the most common reasons for this message is that you (or perhaps your oh-so-helpful editor, as is the case with many MS-Windows editors) have attempted to indent your recipe lines with spaces instead of a tab character. In this case, `make` will use the second form of the error above. Remember that every line in the recipe must begin with a tab character (unless you set `.RECIPEPREFIX`; see [Special Variables](#)). Eight spaces do not count. See [Rule Syntax](#).

`'recipe commences before first target. Stop.'`

`'missing rule before recipe. Stop.'`

This means the first thing in the makefile seems to be part of a recipe: it begins with a recipe prefix character and doesn't appear to be a legal `make` directive (such as a variable assignment). Recipes must always be associated with a target.

The second form is generated if the line has a semicolon as the first non-whitespace character; `make` interprets this to mean you left out the "target: prerequisite" section of a rule. See [Rule Syntax](#).

`'No rule to make target `xxx'.'`

`'No rule to make target `xxx', needed by `yyy'.'`

This means that `make` decided it needed to build a target, but then couldn't find any instructions in the makefile on how to do that, either explicit or implicit (including in the default rules database).

If you want that file to be built, you will need to add a rule to your makefile describing how that target can be built. Other possible sources of this problem are typos in the

makefile (if that file name is wrong) or a corrupted source tree (if that file is not supposed to be built, but rather only a prerequisite).

```
'No targets specified and no makefile found. Stop.'
```

```
'No targets. Stop.'
```

The former means that you didn't provide any targets to be built on the command line, and `make` couldn't find any makefiles to read in. The latter means that some makefile was found, but it didn't contain any default goal and none was given on the command line. GNU `make` has nothing to do in these situations. See [Arguments to Specify the Makefile](#).

```
'Makefile `xxx` was not found.'
```

```
'Included makefile `xxx` was not found.'
```

A makefile specified on the command line (first form) or included (second form) was not found.

```
'warning: overriding recipe for target `xxx`'
```

```
'warning: ignoring old recipe for target `xxx`'
```

GNU `make` allows only one recipe to be specified per target (except for double-colon rules). If you give a recipe for a target which already has been defined to have one, this warning is issued and the second recipe will overwrite the first. See [Multiple Rules for One Target](#).

```
'Circular xxx <- yyy dependency dropped.'
```

This means that `make` detected a loop in the dependency graph: after tracing the prerequisite `yyy` of target `xxx`, and its prerequisites, etc., one of them depended on `xxx` again.

```
'Recursive variable `xxx` references itself (eventually). Stop.'
```

This means you've defined a normal (recursive) `make` variable `xxx` that, when it's expanded, will refer to itself (`xxx`). This is not allowed; either use simply-expanded variables (`:=` or `::=`) or use the append operator (`+=`). See [How to Use Variables](#).

```
'Unterminated variable reference. Stop.'
```

This means you forgot to provide the proper closing parenthesis or brace in your variable or function reference.

```
'insufficient arguments to function `xxx`. Stop.'
```

This means you haven't provided the requisite number of arguments for this function. See the documentation of the function for a description of its arguments. See [Functions for Transforming Text](#).

```
'missing target pattern. Stop.'
```

```
'multiple target patterns. Stop.'
```

```
'target pattern contains no `%'. Stop.'
```

```
'mixed implicit and static pattern rules. Stop.'
```

These errors are generated for malformed static pattern rules (see [Syntax of Static Pattern Rules](#)). The first means the target-pattern part of the rule is empty; the second means there are multiple pattern characters (%) in the target-pattern part; the third means there are no pattern characters in the target-pattern part; and the fourth means that all three parts of the static pattern rule contain pattern characters (%)—the first part should not contain pattern characters.

If you see these errors and you aren't trying to create a static pattern rule, check the value of any variables in your target and prerequisite lists to be sure they do not contain colons.

```
'warning: -jN forced in submake: disabling jobserver mode.'
```

This warning and the next are generated if `make` detects error conditions related to parallel processing on systems where `sub-makes` can communicate (see [Communicating Options to a Sub-make](#)). This warning is generated if a recursive invocation of a `make` process is forced to have `‘-jN’` in its argument list (where N is greater than one). This could happen, for example, if you set the `MAKE` environment variable to `‘make -j2’`. In this case, the `sub-make` doesn’t communicate with other `make` processes and will simply pretend it has two jobs of its own.

```
‘warning: jobserver unavailable: using -j1. Add `+’ to parent make rule.’
```

In order for `make` processes to communicate, the parent will pass information to the child. Since this could result in problems if the child process isn’t actually a `make`, the parent will only do this if it thinks the child is a `make`. The parent uses the normal algorithms to determine this (see [How the MAKE Variable Works](#)). If the makefile is constructed such that the parent doesn’t know the child is a `make` process, then the child will receive only part of the information necessary. In this case, the child will generate this warning message and proceed with its build in a sequential manner.

```
‘warning: ignoring prerequisites on suffix rule definition’
```

According to POSIX, a suffix rule cannot contain prerequisites. If a rule that could be a suffix rule has prerequisites it is interpreted as a simple explicit rule, with an odd target name. This requirement is obeyed when POSIX-conforming mode is enabled (the `.POSIX` target is defined). In versions of GNU `make` prior to 4.3, no warning was emitted and a suffix rule was created, however all prerequisites were ignored and were not part of the suffix rule. Starting with GNU `make` 4.3 the behavior is the same, and in addition this warning is generated. In a future version the POSIX-conforming behavior will be the only behavior: no rule with a prerequisite can be suffix rule and this warning will be removed.

Next: [GNU Free Documentation License](#), Previous: [Error Messages](#), Up: [Top](#) [[Contents](#)][[Index](#)]

Appendix C Complex Makefile Example

Here is the makefile for the GNU `tar` program. This is a moderately complex makefile. The first line uses a `#!` setting to allow the makefile to be executed directly.

Because it is the first target, the default goal is `‘all’`. An interesting feature of this makefile is that `testpad.h` is a source file automatically created by the `testpad` program, itself compiled from `testpad.c`.

If you type `‘make’` or `‘make all’`, then `make` creates the `tar` executable, the `rmt` daemon that provides remote tape access, and the `tar.info` Info file.

If you type `‘make install’`, then `make` not only creates `tar`, `rmt`, and `tar.info`, but also installs them.

If you type `‘make clean’`, then `make` removes the `‘.o’` files, and the `tar`, `rmt`, `testpad`, `testpad.h`, and `core` files.

If you type `‘make distclean’`, then `make` not only removes the same files as does `‘make clean’` but also the `TAGS`, `Makefile`, and `config.status` files. (Although it is not evident, this makefile (and `config.status`) is generated by the user with the `configure` program, which is provided in the `tar` distribution, but is not shown here.)

If you type `‘make realclean’`, then `make` removes the same files as does `‘make distclean’` and also removes the Info files generated from `tar.texinfo`.

In addition, there are targets `shar` and `dist` that create distribution kits.

```
#!/usr/bin/make -f
# Generated automatically from Makefile.in by configure.
# Un*x Makefile for GNU tar program.
# Copyright (C) 1991 Free Software Foundation, Inc.
```

```

# This program is free software; you can redistribute
# it and/or modify it under the terms of the GNU
# General Public License ...
...
...

SHELL = /bin/sh

#### Start of system configuration section. ####

srcdir = .

# If you use gcc, you should either run the
# fixincludes script that comes with it or else use
# gcc with the -traditional option. Otherwise ioctl
# calls will be compiled incorrectly on some systems.
CC = gcc -O
YACC = bison -y
INSTALL = /usr/local/bin/install -c
INSTALLDATA = /usr/local/bin/install -c -m 644

# Things you might add to DEFS:
# -DSTDC_HEADERS      If you have ANSI C headers and
#                      libraries.
# -DPOSIX             If you have POSIX.1 headers and
#                      libraries.
# -DBSD42             If you have sys/dir.h (unless
#                      you use -DPOSIX), sys/file.h,
#                      and st_blocks in `struct stat'.
# -DUSG               If you have System V/ANSI C
#                      string and memory functions
#                      and headers, sys/sysmacros.h,
#                      fcntl.h, getcwd, no valloc,
#                      and ndir.h (unless
#                      you use -DDIRENT).
# -DNO_MEMORY_H       If USG or STDC_HEADERS but do not
#                      include memory.h.
# -DDIRENT             If USG and you have dirent.h
#                      instead of ndir.h.
# -DSIGTYPE=int       If your signal handlers
#                      return int, not void.
# -DNO_MTIO           If you lack sys/mtio.h
#                      (magtape ioctls).
# -DNO_REMOTE          If you do not have a remote shell
#                      or rexec.
# -DUSE_REXEC          To use rexec for remote tape
#                      operations instead of
#                      forking rsh or remsh.
# -DVPRINTF_MISSING   If you lack vprintf function
#                      (but have _doprnt).
# -DDOPRNT_MISSING    If you lack _doprnt function.
#                      Also need to define
#                      -DVPRINTF_MISSING.
# -DFTIME_MISSING      If you lack ftime system call.
# -DSTRSTR_MISSING     If you lack strstr function.
# -DVALLOC_MISSING     If you lack valloc function.
# -DMKDIR_MISSING      If you lack mkdir and
#                      rmdir system calls.
# -DRENAME_MISSING     If you lack rename system call.
# -DFTRUNCATE_MISSING  If you lack ftruncate
#                      system call.
# -DV7                 On Version 7 Unix (not
#                      tested in a long time).
# -DEMUL_OPEN3         If you lack a 3-argument version
#                      of open, and want to emulate it
#                      with system calls you do have.
# -DNO_OPEN3           If you lack the 3-argument open
#                      and want to disable the tar -k
#                      option instead of emulating open.
# -DXENIX              If you have sys/inode.h
#                      and need it 94 to be included.

DEFS = -DSIGTYPE=int -DDIRENT -DSTRSTR_MISSING \
       -DVPRINTF_MISSING -DBSD42

# Set this to rtapelib.o unless you defined NO_REMOTE,
# in which case make it empty.
RTAPELIB = rtapelib.o
LIBS =

```

```

DEF_AR_FILE = /dev/rmt8
DEFBLOCKING = 20

CDEBUG = -g
CFLAGS = $(CDEBUG) -I. -I$(srcdir) $(DEFS) \
        -DDEF_AR_FILE=\"$(DEF_AR_FILE)\" \
        -DDEFBLOCKING=$(DEFBLOCKING)
LDFLAGS = -g

prefix = /usr/local
# Prefix for each installed program,
# normally empty or `g'.
binprefix =

# The directory to install tar in.
bindir = $(prefix)/bin

# The directory to install the info files in.
infodir = $(prefix)/info

#### End of system configuration section. ####

SRCS_C = tar.c create.c extract.c buffer.c \
        getoldopt.c update.c gnu.c mangle.c \
        version.c list.c names.c diffarch.c \
        port.c wildmat.c getopt.c getopt1.c \
        regex.c
SRCS_Y = getdate.y
SRCS = $(SRCS_C) $(SRCS_Y)
OBJS = $(SRCS_C:.c=.o) $(SRCS_Y:.y=.o) $(RTAPELIB)
AUX = README COPYING ChangeLog Makefile.in \
        makefile.pc configure configure.in \
        tar.texinfo tar.info* texinfo.tex \
        tar.h port.h open3.h getopt.h regex.h \
        rmt.h rmt.c rtapelib.c alloca.c \
        msd_dir.h msd_dir.c tcexparg.c \
        level-0 level-1 backup-specs testpad.c

.PHONY: all
all: tar rmt tar.info

tar: $(OBJS)
$(CC) $(LDFLAGS) -o $@ $(OBJS) $(LIBS)

rmt: rmt.c
$(CC) $(CFLAGS) $(LDFLAGS) -o $@ rmt.c

tar.info: tar.texinfo
        makeinfo tar.texinfo

.PHONY: install
install: all
$(INSTALL) tar $(bindir)/$(binprefix)tar
-test ! -f rmt || $(INSTALL) rmt /etc/rmt
$(INSTALLDATA) $(srcdir)/tar.info* $(infodir)

$(OBJS): tar.h port.h testpad.h
regex.o buffer.o tar.o: regex.h
# getdate.y has 8 shift/reduce conflicts.

testpad.h: testpad
        ./testpad

testpad: testpad.o
$(CC) -o $@ testpad.o

TAGS: $(SRCS)
        etags $(SRCS)

.PHONY: clean
clean:
        rm -f *.o tar rmt testpad testpad.h core

```

```

.PHONY: distclean
distclean: clean
    rm -f TAGS Makefile config.status

.PHONY: realclean
realclean: distclean
    rm -f tar.info*

.PHONY: shar
shar: $(SRCS) $(AUX)
    shar $(SRCS) $(AUX) | compress \
    > tar-`sed -e '/version_string/!d' \
    -e 's/[^0-9.]*\([0-9.]*\).*\/1/' \
    -e q
    version.c`.shar.Z

.PHONY: dist
dist: $(SRCS) $(AUX)
    echo tar-`sed \
    -e '/version_string/!d' \
    -e 's/[^0-9.]*\([0-9.]*\).*\/1/' \
    -e q
    version.c` > .fname
    -rm -rf `cat .fname`
    mkdir `cat .fname`
    ln $(SRCS) $(AUX) `cat .fname`
    tar chZf `cat .fname`.tar.Z `cat .fname`
    -rm -rf `cat .fname`.fname

tar.zoo: $(SRCS) $(AUX)
    -rm -rf tmp.dir
    -mkdir tmp.dir
    -rm tar.zoo
    for X in $(SRCS) $(AUX) ; do \
        echo $$X ; \
        sed 's/$$/^M/' $$X \
        > tmp.dir/$$X ; done
    cd tmp.dir ; zoo aM ../tar.zoo *
    -rm -rf tmp.dir

```

Next: [Concept Index](#), Previous: [Complex Makefile](#), Up: [Top](#) [[Contents](#)][[Index](#)]

Appendix D GNU Free Documentation License

Version 1.3, 3 November 2008

Copyright © 2000, 2001, 2002, 2007, 2008 Free Software Foundation, Inc.

<https://fsf.org/>

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

0. PREAMBLE

The purpose of this License is to make a manual, textbook, or other functional and useful document free in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or noncommercially. Secondly, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of “copyleft”, which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to

software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

1. APPLICABILITY AND DEFINITIONS

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The “Document” , below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as “you” . You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A “Modified Version” of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A “Secondary Section” is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document’ s overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The “Invariant Sections” are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The “Cover Texts” are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A “Transparent” copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not “Transparent” is called “Opaque” .

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The “Title Page” means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, “Title Page” means the text near the most prominent appearance of the work’ s title, preceding the beginning of the body of the text.

The “publisher” means any person or entity that distributes copies of the Document to the public.

A section “Entitled XYZ” means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ stands for a specific section name mentioned below, such as “Acknowledgements”, “Dedications”, “Endorsements”, or “History” .) To “Preserve the Title” of such a section when you modify the Document means that it remains a section “Entitled XYZ” according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document’s license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version

under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

1. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any, be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.
2. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.
3. State on the Title page the name of the publisher of the Modified Version, as the publisher.
4. Preserve all the copyright notices of the Document.
5. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
6. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
7. Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document' s license notice.
8. Include an unaltered copy of this License.
9. Preserve the section Entitled "History" , Preserve its Title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
10. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.
11. For any section Entitled "Acknowledgements" or "Dedications" , Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.
12. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.
13. Delete any section Entitled "Endorsements" . Such a section may not be included in the Modified Version.
14. Do not retitle any existing section to be Entitled "Endorsements" or to conflict in title with any Invariant Section.
15. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their titles to the list of Invariant Sections in the Modified Version' s license notice. These titles must be distinct from any other section titles.

You may add a section Entitled “Endorsements” , provided it contains nothing but endorsements of your Modified Version by various parties—for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled “History” in the various original documents, forming one section Entitled “History” ; likewise combine any sections Entitled “Acknowledgements” , and any sections Entitled “Dedications” . You must delete all sections Entitled “Endorsements.”

6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an “aggregate” if the copyright resulting from the compilation is not used to limit the legal rights of the compilation’s users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document’s Cover Texts may be placed on covers that bracket the Document within

the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled “Acknowledgements” , “Dedications” , or “History” , the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense, or distribute it is void, and will automatically terminate your rights under this License.

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, receipt of a copy of some or all of the same material does not give you any rights to use it.

10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See <https://www.gnu.org/copyleft/>.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License “or any later version” applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation. If the Document specifies that a proxy can decide which future versions of this License can be used, that proxy’s public statement of acceptance of a version permanently authorizes you to choose that version for the Document.

11. RELICENSING

“Massive Multiauthor Collaboration Site” (or “MMC Site”) means any World Wide Web server that publishes copyrightable works and also provides prominent facilities for anybody to edit those works. A public wiki that anybody can edit is an example of such a server. A “Massive Multiauthor Collaboration” (or “MMC”) contained in the site means any set of copyrightable works thus published on the MMC site.

“CC-BY-SA” means the Creative Commons Attribution-Share Alike 3.0 license published by Creative Commons Corporation, a not-for-profit corporation with a principal place of business in San Francisco, California, as well as future copyleft versions of that license published by that same organization.

“Incorporate” means to publish or republish a Document, in whole or in part, as part of another Document.

An MMC is “eligible for relicensing” if it is licensed under this License, and if all works that were first published under this License somewhere other than this MMC, and subsequently incorporated in whole or in part into the MMC, (1) had no cover texts or invariant sections, and (2) were thus incorporated prior to November 1, 2008.

The operator of an MMC Site may republish an MMC contained in the site under CC-BY-SA on the same site at any time before August 1, 2009, provided the MMC is eligible for relicensing.

ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

```
Copyright (C)  year  your name.
Permission is granted to copy, distribute and/or modify this document
under the terms of the GNU Free Documentation License, Version 1.3
or any later version published by the Free Software Foundation;
with no Invariant Sections, no Front-Cover Texts, and no Back-Cover
Texts.  A copy of the license is included in the section entitled ``GNU
Free Documentation License''.
```

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the “with... Texts.” line with this:

```
with the Invariant Sections being list their titles, with
the Front-Cover Texts being list, and with the Back-Cover Texts
being list.
```

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.

Next: [Name Index](#), Previous: [GNU Free Documentation License](#), Up: [Top](#) [[Contents](#)][[Index](#)]

Index of Concepts

Jump to: [!](#) <#> [\\$](#) [%](#) [*](#) [+](#) [,](#) [-](#) [.](#) [:](#) [=](#) [?](#) [@](#) [\[](#) [\](#) [_](#) [~](#)
[A](#) [B](#) [C](#) [D](#) [E](#) [F](#) [G](#) [H](#) [I](#) [J](#) [K](#) [L](#) [M](#) [N](#) [O](#) [P](#) [Q](#) [R](#) [S](#) [T](#) [U](#) [V](#) [W](#) [Y](#)

Index Entry	Section
!	
!=:	Setting
!=, expansion:	Reading Makefiles
#	
# (comments), in makefile:	Makefile Contents
# (comments), in recipes:	Recipe Syntax
#include:	Automatic Prerequisites

\$

\$, in function call:	Syntax of Functions
\$, in rules:	Rule Syntax
\$, in variable name:	Computed Names
\$, in variable reference:	Reference

%

%, in pattern rules:	Pattern Intro
%, quoting in <code>patsubst</code>:	Text Functions
%, quoting in static pattern:	Static Usage
%, quoting in <code>vpath</code>:	Selective Search
%, quoting with <code>\</code> (backslash):	Selective Search
%, quoting with <code>\</code> (backslash):	Static Usage
%, quoting with <code>\</code> (backslash):	Text Functions

*

* (wildcard character):	Wildcards
---	---------------------------

+

+, and <code>define</code>:	Canned Recipes
+, and recipe execution:	Instead of Execution
+, and recipes:	MAKE Variable
+=:	Appending
+=, expansion:	Reading Makefiles
+=, expansion:	Reading Makefiles

,

, <code>v</code> (RCS file extension):	Catalogue of Rules
--	------------------------------------

-

- (in recipes):	Errors
-, and <code>define</code>:	Canned Recipes
--always-make:	Options Summary
--assume-new:	Instead of Execution
--assume-new:	Options Summary
--assume-new, and recursion:	Options/Recursion
--assume-old:	Avoiding Compilation
--assume-old:	Options Summary
--assume-old, and recursion:	Options/Recursion
--check-symlink-times:	Options Summary
--debug:	Options Summary
--directory:	Recursion
--directory:	Options Summary
--directory, and <code>--print-directory</code>:	-w Option
--directory, and recursion:	Options/Recursion
--dry-run:	Echoing
--dry-run:	Instead of Execution

`--dry-run:`

`--environment-overrides:`

`--eval:`

`--file:`

`--file:`

`--file:`

`--file, and recursion:`

`--help:`

`--ignore-errors:`

`--ignore-errors:`

`--include-dir:`

`--include-dir:`

`--jobs:`

`--jobs:`

`--jobs, and recursion:`

`--just-print:`

`--just-print:`

`--just-print:`

`--keep-going:`

`--keep-going:`

`--keep-going:`

`--load-average:`

`--load-average:`

`--makefile:`

`--makefile:`

`--makefile:`

`--max-load:`

`--max-load:`

`--new-file:`

`--new-file:`

`--new-file, and recursion:`

`--no-builtin-rules:`

`--no-builtin-variables:`

`--no-keep-going:`

`--no-print-directory:`

`--no-print-directory:`

`--old-file:`

`--old-file:`

`--old-file, and recursion:`

`--output-sync:`

`--output-sync:`

`--print-data-base:`

`--print-directory:`

`--print-directory, and --directory:`

`--print-directory, and recursion:`

`--print-directory, disabling:`

`--question:`

`--question:`

`--quiet:`

`--quiet:`

Options Summary

Options Summary

Options Summary

Makefile Names

Makefile Arguments

Options Summary

Options/Recursion

Options Summary

Errors

Options Summary

Include

Options Summary

Parallel

Options Summary

Options/Recursion

Echoing

Instead of Execution

Options Summary

Errors

Testing

Options Summary

Parallel

Options Summary

Makefile Names

Makefile Arguments

Options Summary

Parallel

Options Summary

Instead of Execution

Options Summary

Options/Recursion

Options Summary

Options Summary

-w Option

Options Summary

Avoiding Compilation

Options Summary

Options/Recursion

Parallel Output

Options Summary

Options Summary

Options Summary

-w Option

-w Option

-w Option

Instead of Execution

Options Summary

Echoing

Options Summary

<code>--recon:</code>	Echoing
<code>--recon:</code>	Instead of Execution
<code>--recon:</code>	Options Summary
<code>--silent:</code>	Echoing
<code>--silent:</code>	Options Summary
<code>--stop:</code>	Options Summary
<code>--touch:</code>	Instead of Execution
<code>--touch:</code>	Options Summary
<code>--touch, and recursion:</code>	MAKE Variable
<code>--trace:</code>	Options Summary
<code>--version:</code>	Options Summary
<code>--warn-undefined-variables:</code>	Options Summary
<code>--what-if:</code>	Instead of Execution
<code>--what-if:</code>	Options Summary
<code>-b:</code>	Options Summary
<code>-B:</code>	Options Summary
<code>-C:</code>	Recursion
<code>-C:</code>	Options Summary
<code>-C, and -w:</code>	-w Option
<code>-C, and recursion:</code>	Options/Recursion
<code>-d:</code>	Options Summary
<code>-e:</code>	Options Summary
<code>-E:</code>	Options Summary
<code>-e (shell flag):</code>	Automatic Prerequisites
<code>-f:</code>	Makefile Names
<code>-f:</code>	Makefile Arguments
<code>-f:</code>	Options Summary
<code>-f, and recursion:</code>	Options/Recursion
<code>-h:</code>	Options Summary
<code>-I:</code>	Include
<code>-i:</code>	Errors
<code>-i:</code>	Options Summary
<code>-I:</code>	Options Summary
<code>-j:</code>	Parallel
<code>-j:</code>	Options Summary
<code>-j, and archive update:</code>	Archive Pitfalls
<code>-j, and recursion:</code>	Options/Recursion
<code>-k:</code>	Errors
<code>-k:</code>	Testing
<code>-k:</code>	Options Summary
<code>-l:</code>	Options Summary
<code>-L:</code>	Options Summary
<code>-l (library search):</code>	Libraries/Search
<code>-l (load average):</code>	Parallel
<code>-m:</code>	Options Summary
<code>-M (to compiler):</code>	Automatic Prerequisites
<code>-MM (to GNU compiler):</code>	Automatic Prerequisites
<code>-n:</code>	Echoing
<code>-n:</code>	Instead of Execution
<code>-n:</code>	Options Summary

<code>-O:</code>	Parallel Output
<code>-o:</code>	Avoiding Compilation
<code>-o:</code>	Options Summary
<code>-O:</code>	Options Summary
<code>-o, and recursion:</code>	Options/Recursion
<code>-p:</code>	Options Summary
<code>-q:</code>	Instead of Execution
<code>-q:</code>	Options Summary
<code>-r:</code>	Options Summary
<code>-R:</code>	Options Summary
<code>-s:</code>	Echoing
<code>-s:</code>	Options Summary
<code>-S:</code>	Options Summary
<code>-t:</code>	Instead of Execution
<code>-t:</code>	Options Summary
<code>-t, and recursion:</code>	MAKE Variable
<code>-v:</code>	Options Summary
<code>-W:</code>	Instead of Execution
<code>-w:</code>	Options Summary
<code>-W:</code>	Options Summary
<code>-w, and -C:</code>	-w Option
<code>-W, and recursion:</code>	Options/Recursion
<code>-w, and recursion:</code>	-w Option
<code>-w, disabling:</code>	-w Option
.	
<code>. a (archives):</code>	Archive Suffix Rules
<code>. c:</code>	Catalogue of Rules
<code>. C:</code>	Catalogue of Rules
<code>. cc:</code>	Catalogue of Rules
<code>. ch:</code>	Catalogue of Rules
<code>. cpp:</code>	Catalogue of Rules
<code>. d:</code>	Automatic Prerequisites
<code>. def:</code>	Catalogue of Rules
<code>. dvi:</code>	Catalogue of Rules
<code>. f:</code>	Catalogue of Rules
<code>. F:</code>	Catalogue of Rules
<code>. info:</code>	Catalogue of Rules
<code>. l:</code>	Catalogue of Rules
<code>. LIBPATTERNS, and link libraries:</code>	Libraries/Search
<code>. ln:</code>	Catalogue of Rules
<code>. mod:</code>	Catalogue of Rules
<code>. o:</code>	Catalogue of Rules
<code>. o:</code>	Catalogue of Rules
<code>. ONESHELL, use of:</code>	One Shell
<code>. p:</code>	Catalogue of Rules
<code>. PRECIOUS intermediate files:</code>	Chained Rules
<code>. r:</code>	Catalogue of Rules
<code>. s:</code>	Catalogue of Rules
<code>. S:</code>	Catalogue of Rules

.sh:	Catalogue of Rules
.SHELLFLAGS, value of:	Choosing the Shell
.sym:	Catalogue of Rules
.tex:	Catalogue of Rules
.texi:	Catalogue of Rules
.texinfo:	Catalogue of Rules
.txinfo:	Catalogue of Rules
.w:	Catalogue of Rules
.web:	Catalogue of Rules
.y:	Catalogue of Rules
:	
:: rules (double-colon):	Double-Colon
::=:	Flavors
::=:	Setting
:=:	Flavors
:=:	Setting
=	
=:	Flavors
=:	Setting
=, expansion:	Reading Makefiles
?	
? (wildcard character):	Wildcards
?=:	Flavors
?=:	Setting
?=?, expansion:	Reading Makefiles
@	
@ (in recipes):	Echoing
@, and define:	Canned Recipes
[	
[...] (wildcard characters):	Wildcards
\	
\ (backslash), for continuation lines:	Simple Makefile
\ (backslash), in recipes:	Splitting Recipe Lines
\ (backslash), to quote %:	Selective Search
\ (backslash), to quote %:	Static Usage
\ (backslash), to quote %:	Text Functions
-	
___.SYMDEF:	Archive Symbols
~	

A

abspath:	File Name Functions
algorithm for directory search:	Search Algorithm
all (standard target):	Goals
appending to variables:	Appending
ar:	Implicit Variables
archive:	Archives
archive member targets:	Archive Members
archive symbol directory updating:	Archive Symbols
archive, and -j:	Archive Pitfalls
archive, and parallel execution:	Archive Pitfalls
archive, suffix rule for:	Archive Suffix Rules
Arg list too long:	Options/Recursion
arguments of functions:	Syntax of Functions
as:	Catalogue of Rules
as:	Implicit Variables
assembly, rule to compile:	Catalogue of Rules
automatic generation of prerequisites:	Include
automatic generation of prerequisites:	Automatic Prerequisites
automatic variables:	Automatic Variables
automatic variables in prerequisites:	Automatic Variables

B

backquotes:	Shell Function
backslash (\), for continuation lines:	Simple Makefile
backslash (\), in recipes:	Splitting Recipe Lines
backslash (\), to quote %:	Selective Search
backslash (\), to quote %:	Static Usage
backslash (\), to quote %:	Text Functions
backslash (\), to quote newlines:	Splitting Lines
backslashes in pathnames and wildcard expansion:	Wildcard Pitfall
basename:	File Name Functions
binary packages:	Install Command Categories
broken pipe:	Parallel Input
bugs, reporting:	Bugs
built-in special targets:	Special Targets

C

C++, rule to compile:	Catalogue of Rules
C, rule to compile:	Catalogue of Rules
canned recipes:	Canned Recipes
cc:	Catalogue of Rules
cc:	Implicit Variables
cd (shell command):	Execution
cd (shell command):	MAKE Variable
chains of rules:	Chained Rules
check (standard target):	Goals

<code>clean</code> (standard target):	Goals
<code>clean</code> target:	Simple Makefile
<code>clean</code> target:	Cleanup
cleaning up:	Cleanup
<code>clobber</code> (standard target):	Goals
<code>co</code> :	Catalogue of Rules
<code>co</code> :	Implicit Variables
combining rules by prerequisite:	Combine By Prerequisite
command expansion:	Shell Function
command line variable definitions, and recursion:	Options/Recursion
command line variables:	Overriding
commands, sequences of:	Canned Recipes
comments, in makefile:	Makefile Contents
comments, in recipes:	Recipe Syntax
compatibility:	Features
compatibility in exporting:	Variables/Recursion
compilation, testing:	Testing
computed variable name:	Computed Names
conditional expansion:	Conditional Functions
conditional variable assignment:	Flavors
conditionals:	Conditionals
continuation lines:	Simple Makefile
controlling make:	Make Control Functions
conventions for makefiles:	Makefile Conventions
convert guile types:	Guile Types
<code>ctangle</code> :	Catalogue of Rules
<code>ctangle</code> :	Implicit Variables
<code>cweave</code> :	Catalogue of Rules
<code>cweave</code> :	Implicit Variables

D	
data base of <code>make</code> rules:	Options Summary
deducing recipes (implicit rules):	make Deduces
default directories for included makefiles:	Include
default goal:	How Make Works
default goal:	Rules
default makefile name:	Makefile Names
default rules, last-resort:	Last Resort
define, expansion:	Reading Makefiles
defining variables verbatim:	Multi-Line
deletion of target files:	Errors
deletion of target files:	Interrupts
directive:	Makefile Contents
directories, creating installation:	Directory Variables
directories, printing them:	-w Option
directories, updating archive symbol:	Archive Symbols
directory part:	File Name Functions
directory search (<code>VPATH</code>):	Directory Search
directory search (<code>VPATH</code>), and implicit rules:	Implicit/Search
directory search (<code>VPATH</code>), and link libraries:	Libraries/Search

directory search (<code>VPATH</code>), and recipes:	Recipes/Search
directory search algorithm:	Search Algorithm
directory search, traditional (<code>GPATH</code>):	Search Algorithm
<code>dist</code> (standard target):	Goals
<code>distclean</code> (standard target):	Goals
dollar sign (<code>\$</code>), in function call:	Syntax of Functions
dollar sign (<code>\$</code>), in rules:	Rule Syntax
dollar sign (<code>\$</code>), in variable name:	Computed Names
dollar sign (<code>\$</code>), in variable reference:	Reference
DOS, choosing a shell in:	Choosing the Shell
double-colon rules:	Double-Colon
duplicate words, removing:	Text Functions
E	
E2BIG:	Options/Recursion
echoing of recipes:	Echoing
editor:	Introduction
Emacs (<code>M-x compile</code>):	Errors
empty recipes:	Empty Recipes
empty targets:	Empty Targets
environment:	Environment
environment, and recursion:	Variables/Recursion
environment, <code>SHELL</code> in:	Choosing the Shell
error, stopping on:	Make Control Functions
errors (in recipes):	Errors
errors with wildcards:	Wildcard Pitfall
evaluating makefile syntax:	Eval Function
example of loaded objects:	Loaded Object Example
example using Guile:	Guile Example
execution, in parallel:	Parallel
execution, instead of:	Instead of Execution
execution, of recipes:	Execution
exit status (errors):	Errors
exit status of make:	Running
expansion, secondary:	Secondary Expansion
explicit rule, definition of:	Makefile Contents
explicit rule, expansion:	Reading Makefiles
explicit rules, secondary expansion of:	Secondary Expansion
exporting variables:	Variables/Recursion
extensions, Guile:	Guile Integration
extensions, load directive:	load Directive
extensions, loading:	Loading Objects
F	
<code>f77</code> :	Catalogue of Rules
<code>f77</code> :	Implicit Variables
FDL, GNU Free Documentation License:	GNU Free Documentation License
features of GNU <code>make</code> :	Features

features, missing:	Missing
file name functions:	File Name Functions
file name of makefile:	Makefile Names
file name of makefile, how to specify:	Makefile Names
file name prefix, adding:	File Name Functions
file name suffix:	File Name Functions
file name suffix, adding:	File Name Functions
file name with wildcards:	Wildcards
file name, abspath of:	File Name Functions
file name, basename of:	File Name Functions
file name, directory part:	File Name Functions
file name, nondirectory part:	File Name Functions
file name, realpath of:	File Name Functions
file, reading from:	File Function
file, writing to:	File Function
files, assuming new:	Instead of Execution
files, assuming old:	Avoiding Compilation
files, avoiding recompilation of:	Avoiding Compilation
files, intermediate:	Chained Rules
filtering out words:	Text Functions
filtering words:	Text Functions
finding strings:	Text Functions
flags:	Options Summary
flags for compilers:	Implicit Variables
flavor of variable:	Flavor Function
flavors of variables:	Flavors
FORCE:	Force Targets
force targets:	Force Targets
Fortran, rule to compile:	Catalogue of Rules
functions:	Functions
functions, for controlling make:	Make Control Functions
functions, for file names:	File Name Functions
functions, for text:	Text Functions
functions, syntax of:	Syntax of Functions
functions, user defined:	Call Function

G

g++:	Catalogue of Rules
g++:	Implicit Variables
gcc:	Catalogue of Rules
generating prerequisites automatically:	Include
generating prerequisites automatically:	Automatic Prerequisites
get:	Catalogue of Rules
get:	Implicit Variables
globbing (wildcards):	Wildcards
goal:	How Make Works
goal, default:	How Make Works
goal, default:	Rules
goal, how to specify:	Goals
grouped targets:	Multiple Targets

Guile: Guile: Guile example: guile, conversion of types:	Guile Function Guile Integration Guile Example Guile Types
H	
home directory:	Wildcards
I	
IEEE Standard 1003.2: ifdef, expansion: ifeq, expansion: ifndef, expansion: ifneq, expansion: implicit rule: implicit rule, and directory search: implicit rule, and <code>VPATH</code>: implicit rule, definition of: implicit rule, expansion: implicit rule, how to use: implicit rule, introduction to: implicit rule, predefined: implicit rule, search algorithm: implicit rules, secondary expansion of: included makefiles, default directories: including (<code>MAKEFILES</code> variable): including (<code>MAKEFILE_LIST</code> variable): including other makefiles: incompatibilities: independent targets: Info, rule to format: inheritance, suppressing: input during parallel execution: install (standard target): installation directories, creating: installations, staged: interface for loaded objects: intermediate files: intermediate files, preserving: intermediate targets, explicit: interrupt:	Overview Reading Makefiles Reading Makefiles Reading Makefiles Reading Makefiles Implicit Rules Implicit/Search Implicit/Search Makefile Contents Reading Makefiles Using Implicit make Deduces Catalogue of Rules Implicit Rule Search Secondary Expansion Include MAKEFILES Variable Special Variables Include Missing Multiple Targets Catalogue of Rules Suppressing Inheritance Parallel Input Goals Directory Variables DESTDIR Loaded Object API Chained Rules Chained Rules Special Targets Interrupts
J	
job slots: job slots, and recursion: job slots, sharing: jobs, limiting based on load: jobserver: jobserver on POSIX:	Parallel Options/Recursion Job Slots Parallel Job Slots POSIX Jobserver

jobserver on Windows: joining lists of words:	Windows Jobserver File Name Functions
K	
killing (interruption):	Interrupts
L	
last-resort default rules: ld: lex: lex: Lex, rule to run: libraries for linking, directory search: library archive, suffix rule for: limiting jobs based on load: link libraries, and directory search: link libraries, patterns matching: linking, predefined rule for: lint: lint: lint, rule to run: list of all prerequisites: list of changed prerequisites: load average: load directive: loaded object API: loaded object example: loaded object licensing: loaded objects: loaded objects, remaking of: long lines, splitting: loops in variable expansion: lpr (shell command): lpr (shell command):	Last Resort Catalogue of Rules Catalogue of Rules Implicit Variables Catalogue of Rules Libraries/Search Archive Suffix Rules Parallel Libraries/Search Libraries/Search Catalogue of Rules Catalogue of Rules Implicit Variables Catalogue of Rules Automatic Variables Automatic Variables Parallel load Directive Loaded Object API Loaded Object Example Loaded Object API Loading Objects Remaking Loaded Objects Splitting Lines Flavors Wildcard Examples Empty Targets
M	
m2c: m2c: macro: make depend: make extensions: make integration: make interface to guile: make procedures in guile: makefile: makefile name: makefile name, how to specify: makefile rule parts: makefile syntax, evaluating:	Catalogue of Rules Implicit Variables Using Variables Automatic Prerequisites Extending make Integrating make Guile Interface Guile Interface Introduction Makefile Names Makefile Names Rule Introduction Eval Function

makefile, and <code>MAKEFILES</code> variable:	MAKEFILES Variable
makefile, conventions for:	Makefile Conventions
makefile, how <code>make</code> processes:	How Make Works
makefile, how to write:	Makefiles
makefile, including:	Include
makefile, overriding:	Overriding Makefiles
makefile, reading:	Reading Makefiles
makefile, remaking of:	Remaking Makefiles
makefile, simple:	Simple Makefile
makefiles, and <code>MAKEFILE_LIST</code> variable:	Special Variables
makefiles, and special variables:	Special Variables
makefiles, parsing:	Parsing Makefiles
<code>makeinfo</code> :	Catalogue of Rules
<code>makeinfo</code> :	Implicit Variables
match-anything rule:	Match-Anything Rules
match-anything rule, used to override:	Overriding Makefiles
missing features:	Missing
mistakes with wildcards:	Wildcard Pitfall
modified variable reference:	Substitution Refs
Modula-2, rule to compile:	Catalogue of Rules
<code>mostlyclean</code> (standard target):	Goals
multi-line variable definition:	Multi-Line
multiple rules for one target:	Multiple Rules
multiple rules for one target (<code>::</code>):	Double-Colon
multiple targets:	Multiple Targets
multiple targets, in pattern rule:	Pattern Intro

N

name of makefile:	Makefile Names
name of makefile, how to specify:	Makefile Names
nested variable reference:	Computed Names
newline, quoting, in makefile:	Simple Makefile
newline, quoting, in recipes:	Splitting Recipe Lines
nondirectory part:	File Name Functions
normal prerequisites:	Prerequisite Types

O

<code>obj</code> :	Variables Simplify
<code>OBJ</code> :	Variables Simplify
<code>objects</code> :	Variables Simplify
<code>OBJECTS</code> :	Variables Simplify
objects, loaded:	Loading Objects
<code>objs</code> :	Variables Simplify
<code>OBJS</code> :	Variables Simplify
old-fashioned suffix rules:	Suffix Rules
options:	Options Summary
options, and recursion:	Options/Recursion
options, setting from environment:	Options/Recursion
options, setting in makefiles:	Options/Recursion

order of pattern rules:	Pattern Match
order-only prerequisites:	Prerequisite Types
origin of variable:	Origin Function
output during parallel execution:	Parallel Output
output during parallel execution:	Options Summary
overriding makefiles:	Overriding Makefiles
overriding variables with arguments:	Overriding
overriding with <code>override</code> :	Override Directive

P

parallel execution:	Parallel
parallel execution, and archive update:	Archive Pitfalls
parallel execution, input during:	Parallel Input
parallel execution, output during:	Parallel Output
parallel execution, output during:	Options Summary
parallel execution, overriding:	Special Targets
parallel output to terminal:	Terminal Output
parsing makefiles:	Parsing Makefiles
parts of makefile rule:	Rule Introduction
Pascal, rule to compile:	Catalogue of Rules
pattern rule:	Pattern Intro
pattern rule, expansion:	Reading Makefiles
pattern rules, order of:	Pattern Match
pattern rules, static (not implicit):	Static Pattern
pattern rules, static, syntax of:	Static Usage
pattern-specific variables:	Pattern-specific
<code>pc</code> :	Catalogue of Rules
<code>pc</code> :	Implicit Variables
phony targets:	Phony Targets
phony targets and recipe execution:	Instead of Execution
pitfalls of wildcards:	Wildcard Pitfall
<code>plugin_is_GPL_compatible</code> :	Loaded Object API
portability:	Features
POSIX:	Overview
POSIX:	Options/Recursion
POSIX-conforming mode, setting:	Special Targets
post-installation commands:	Install Command Categories
pre-installation commands:	Install Command Categories
precious targets:	Special Targets
predefined rules and variables, printing:	Options Summary
prefix, adding:	File Name Functions
prerequisite:	Rules
prerequisite pattern, implicit:	Pattern Intro
prerequisite pattern, static (not implicit):	Static Usage
prerequisite types:	Prerequisite Types
prerequisite, expansion:	Reading Makefiles
prerequisites:	Rule Syntax
prerequisites, and automatic variables:	Automatic Variables
prerequisites, automatic generation:	Include
prerequisites, automatic generation:	Automatic Prerequisites

prerequisites, introduction to:	Rule Introduction
prerequisites, list of all:	Automatic Variables
prerequisites, list of changed:	Automatic Variables
prerequisites, normal:	Prerequisite Types
prerequisites, order-only:	Prerequisite Types
prerequisites, varying (static pattern):	Static Pattern
preserving intermediate files:	Chained Rules
preserving with <code>.PRECIOUS</code> :	Special Targets
preserving with <code>.PRECIOUS</code> :	Chained Rules
preserving with <code>.SECONDARY</code> :	Special Targets
<code>print</code> (standard target):	Goals
<code>print</code> target:	Wildcard Examples
<code>print</code> target:	Empty Targets
printing directories:	-w Option
printing messages:	Make Control Functions
printing of recipes:	Echoing
printing user warnings:	Make Control Functions
problems and bugs, reporting:	Bugs
problems with wildcards:	Wildcard Pitfall
processing a makefile:	How Make Works
Q	
question mode:	Instead of Execution
quoting <code>%</code> , in <code>patsubst</code> :	Text Functions
quoting <code>%</code> , in static pattern:	Static Usage
quoting <code>%</code> , in <code>vpath</code> :	Selective Search
quoting newline, in makefile:	Simple Makefile
quoting newline, in recipes:	Splitting Recipe Lines
R	
Ratfor, rule to compile:	Catalogue of Rules
RCS, rule to extract from:	Catalogue of Rules
reading from a file:	File Function
reading makefiles:	Reading Makefiles
<code>README</code> :	Makefile Names
<code>realclean</code> (standard target):	Goals
<code>realpath</code> :	File Name Functions
recipe:	Simple Makefile
recipe execution, single invocation:	Special Targets
recipe lines, single shell:	One Shell
recipe syntax:	Recipe Syntax
recipe, execution:	Execution
recipes:	Rule Syntax
recipes:	Recipes
recipes setting shell variables:	Execution
recipes, and directory search:	Recipes/Search
recipes, backslash (<code>\</code>) in:	Splitting Recipe Lines
recipes, canned:	Canned Recipes
recipes, comments in:	Recipe Syntax

recipes, echoing:	Echoing
recipes, empty:	Empty Recipes
recipes, errors in:	Errors
recipes, execution in parallel:	Parallel
recipes, how to write:	Recipes
recipes, instead of executing:	Instead of Execution
recipes, introduction to:	Rule Introduction
recipes, quoting newlines in:	Splitting Recipe Lines
recipes, splitting:	Splitting Recipe Lines
recipes, using variables in:	Variables in Recipes
recompilation:	Introduction
recompilation, avoiding:	Avoiding Compilation
recording events with empty targets:	Empty Targets
recursion:	Recursion
recursion, and <code>-C</code> :	Options/Recursion
recursion, and <code>-f</code> :	Options/Recursion
recursion, and <code>-j</code> :	Options/Recursion
recursion, and <code>-o</code> :	Options/Recursion
recursion, and <code>-t</code> :	MAKE Variable
recursion, and <code>-W</code> :	Options/Recursion
recursion, and <code>-w</code> :	-w Option
recursion, and command line variable definitions:	Options/Recursion
recursion, and environment:	Variables/Recursion
recursion, and <code>MAKE</code> variable:	MAKE Variable
recursion, and <code>MAKEFILES</code> variable:	MAKEFILES Variable
recursion, and options:	Options/Recursion
recursion, and printing directories:	-w Option
recursion, and variables:	Variables/Recursion
recursion, level of:	Variables/Recursion
recursive variable expansion:	Using Variables
recursive variable expansion:	Flavors
recursively expanded variables:	Flavors
reference to variables:	Reference
reference to variables:	Advanced
relinking:	How Make Works
remaking loaded objects:	Remaking Loaded Objects
remaking makefiles:	Remaking Makefiles
removal of target files:	Errors
removal of target files:	Interrupts
removing duplicate words:	Text Functions
removing targets on failure:	Special Targets
removing whitespace from split lines:	Splitting Lines
removing, to clean up:	Cleanup
reporting bugs:	Bugs
<code>rm</code> :	Implicit Variables
<code>rm</code> (shell command):	Simple Makefile
<code>rm</code> (shell command):	Wildcard Examples
<code>rm</code> (shell command):	Phony Targets
<code>rm</code> (shell command):	Errors
rule prerequisites:	Rule Syntax

rule syntax:	Rule Syntax
rule targets:	Rule Syntax
rule, double-colon (::):	Double-Colon
rule, explicit, definition of:	Makefile Contents
rule, how to write:	Rules
rule, implicit:	Implicit Rules
rule, implicit, and directory search:	Implicit/Search
rule, implicit, and <code>VPATH</code> :	Implicit/Search
rule, implicit, chains of:	Chained Rules
rule, implicit, definition of:	Makefile Contents
rule, implicit, how to use:	Using Implicit
rule, implicit, introduction to:	make Deduces
rule, implicit, predefined:	Catalogue of Rules
rule, introduction to:	Rule Introduction
rule, multiple for one target:	Multiple Rules
rule, no recipe or prerequisites:	Force Targets
rule, pattern:	Pattern Intro
rule, static pattern:	Static Pattern
rule, static pattern versus implicit:	Static versus Implicit
rule, with multiple targets:	Multiple Targets
rules, and \$:	Rule Syntax

S

<code>s.</code> (SCCS file prefix):	Catalogue of Rules
SCCS, rule to extract from:	Catalogue of Rules
search algorithm, implicit rule:	Implicit Rule Search
search path for prerequisites (<code>VPATH</code>):	Directory Search
search path for prerequisites (<code>VPATH</code>), and implicit rules:	Implicit/Search
search path for prerequisites (<code>VPATH</code>), and link libraries:	Libraries/Search
searching for strings:	Text Functions
secondary expansion:	Secondary Expansion
secondary expansion and explicit rules:	Secondary Expansion
secondary expansion and implicit rules:	Secondary Expansion
secondary expansion and static pattern rules:	Secondary Expansion
secondary files:	Chained Rules
secondary targets:	Special Targets
<code>sed</code> (shell command):	Automatic Prerequisites
selecting a word:	Text Functions
selecting word lists:	Text Functions
sequences of commands:	Canned Recipes
setting options from environment:	Options/Recursion
setting options in makefiles:	Options/Recursion
setting variables:	Setting
several rules for one target:	Multiple Rules
several targets in a rule:	Multiple Targets
<code>shar</code> (standard target):	Goals
shell command, function for:	Shell Function
shell file name pattern (in <code>include</code>):	Include

shell variables, setting in recipes:	Execution
shell wildcards (in <code>include</code>):	Include
shell, choosing the:	Choosing the Shell
SHELL, exported value:	Variables/Recursion
SHELL, import from environment:	Environment
shell, in DOS and Windows:	Choosing the Shell
SHELL, MS-DOS specifics:	Choosing the Shell
SHELL, value of:	Choosing the Shell
signal:	Interrupts
silent operation:	Echoing
simple makefile:	Simple Makefile
simple variable expansion:	Using Variables
simplifying with variables:	Variables Simplify
simply expanded variables:	Flavors
sorting words:	Text Functions
spaces, in variable values:	Flavors
spaces, stripping:	Text Functions
special targets:	Special Targets
special variables:	Special Variables
specifying makefile name:	Makefile Names
splitting long lines:	Splitting Lines
splitting recipes:	Splitting Recipe Lines
staged installs:	DESTDIR
standard input:	Parallel Input
standards conformance:	Overview
standards for makefiles:	Makefile Conventions
static pattern rule:	Static Pattern
static pattern rule, syntax of:	Static Usage
static pattern rule, versus implicit:	Static versus Implicit
static pattern rules, secondary expansion of:	Secondary Expansion
stem:	Static Usage
stem:	Pattern Match
stem, shortest:	Pattern Match
stem, variable for:	Automatic Variables
stopping make:	Make Control Functions
strings, searching for:	Text Functions
stripping whitespace:	Text Functions
sub-make:	Variables/Recursion
subdirectories, recursion for:	Recursion
substitution variable reference:	Substitution Refs
suffix rule:	Suffix Rules
suffix rule, for archive:	Archive Suffix Rules
suffix, adding:	File Name Functions
suffix, function to find:	File Name Functions
suffix, substituting in variables:	Substitution Refs
suppressing inheritance:	Suppressing Inheritance
switches:	Options Summary
symbol directories, updating archive:	Archive Symbols
syntax of recipe:	Recipe Syntax
syntax of rules:	Rule Syntax

T

tab character (in commands):	Rule Syntax
tabs in rules:	Rule Introduction
TAGS (standard target):	Goals
tangle:	Catalogue of Rules
tangle:	Implicit Variables
tar (standard target):	Goals
target:	Rules
target pattern, implicit:	Pattern Intro
target pattern, static (not implicit):	Static Usage
target, deleting on error:	Errors
target, deleting on interrupt:	Interrupts
target, expansion:	Reading Makefiles
target, multiple in pattern rule:	Pattern Intro
target, multiple rules for one:	Multiple Rules
target, touching:	Instead of Execution
target-specific variables:	Target-specific
targets:	Rule Syntax
targets without a file:	Phony Targets
targets, built-in special:	Special Targets
targets, empty:	Empty Targets
targets, force:	Force Targets
targets, grouped:	Multiple Targets
targets, independent:	Multiple Targets
targets, introduction to:	Rule Introduction
targets, multiple:	Multiple Targets
targets, phony:	Phony Targets
terminal rule:	Match-Anything Rules
terminal, output to:	Terminal Output
test (standard target):	Goals
testing compilation:	Testing
tex:	Catalogue of Rules
tex:	Implicit Variables
TeX, rule to run:	Catalogue of Rules
texi2dvi:	Catalogue of Rules
texi2dvi:	Implicit Variables
Texinfo, rule to format:	Catalogue of Rules
tilde (~):	Wildcards
tools, sharing job slots:	Job Slots
touch (shell command):	Wildcard Examples
touch (shell command):	Empty Targets
touching files:	Instead of Execution
traditional directory search (GPATH):	Search Algorithm
types of prerequisites:	Prerequisite Types
types, conversion of:	Guile Types

U

undefined variables, warning message:	Options Summary
undefining variable:	Undefine Directive

updating archive symbol directories:	Archive Symbols
updating loaded objects:	Remaking Loaded Objects
updating makefiles:	Remaking Makefiles
user defined functions:	Call Function

V	
value:	Using Variables
value, how a variable gets it:	Values
variable:	Using Variables
variable definition:	Makefile Contents
variable references in recipes:	Variables in Recipes
variables:	Variables Simplify
variables, '\$' in name:	Computed Names
variables, and implicit rule:	Automatic Variables
variables, appending to:	Appending
variables, automatic:	Automatic Variables
variables, command line:	Overriding
variables, command line, and recursion:	Options/Recursion
variables, computed names:	Computed Names
variables, conditional assignment:	Flavors
variables, defining verbatim:	Multi-Line
variables, environment:	Variables/Recursion
variables, environment:	Environment
variables, exporting:	Variables/Recursion
variables, flavor of:	Flavor Function
variables, flavors:	Flavors
variables, how they get their values:	Values
variables, how to reference:	Reference
variables, loops in expansion:	Flavors
variables, modified reference:	Substitution Refs
variables, multi-line:	Multi-Line
variables, nested references:	Computed Names
variables, origin of:	Origin Function
variables, overriding:	Override Directive
variables, overriding with arguments:	Overriding
variables, pattern-specific:	Pattern-specific
variables, recursively expanded:	Flavors
variables, setting:	Setting
variables, simply expanded:	Flavors
variables, spaces in values:	Flavors
variables, substituting suffix in:	Substitution Refs
variables, substitution reference:	Substitution Refs
variables, target-specific:	Target-specific
variables, unexpanded value:	Value Function
variables, warning for undefined:	Options Summary
varying prerequisites:	Static Pattern
verbatim variable definition:	Multi-Line
vpath:	Directory Search
VPATH, and implicit rules:	Implicit/Search
VPATH, and link libraries:	Libraries/Search

W

warnings, printing:	Make Control Functions
weave:	Catalogue of Rules
weave:	Implicit Variables
Web, rule to run:	Catalogue of Rules
what if:	Instead of Execution
whitespace, avoiding on line split:	Splitting Lines
whitespace, in variable values:	Flavors
whitespace, stripping:	Text Functions
wildcard:	Wildcards
wildcard pitfalls:	Wildcard Pitfall
wildcard, function:	File Name Functions
wildcard, in archive member:	Archive Members
wildcard, in <code>include</code> :	Include
wildcards and MS-DOS/MS-Windows backslashes:	Wildcard Pitfall
Windows, choosing a shell in:	Choosing the Shell
word, selecting a:	Text Functions
words, extracting first:	Text Functions
words, extracting last:	Text Functions
words, filtering:	Text Functions
words, filtering out:	Text Functions
words, finding number:	Text Functions
words, iterating over:	Foreach Function
words, joining lists:	File Name Functions
words, removing duplicates:	Text Functions
words, selecting lists of:	Text Functions
writing recipes:	Recipes
writing rules:	Rules
writing to a file:	File Function

Y

yacc:	Catalogue of Rules
yacc:	Implicit Variables
yacc:	Canned Recipes
Yacc, rule to run:	Catalogue of Rules

Jump to: ! # \$ % * + , - . : = ? @ [\ _ ~
 A B C D E F G H I J K L M N O P Q R S T U V W Y

Previous: [Concept Index](#), Up: [Top](#) [\[Contents\]](#)[\[Index\]](#)

Index of Functions, Variables, & Directives

Jump to: \$ % * + - . / < ? @ ^ |
 A B C D E F G I J L M N O P R S T U V W Y

\$

\$%:	Automatic Variables
\$ (%D) :	Automatic Variables
\$ (%F) :	Automatic Variables
\$ (*D) :	Automatic Variables
\$ (*F) :	Automatic Variables
\$ (+D) :	Automatic Variables
\$ (+F) :	Automatic Variables
\$ (<D) :	Automatic Variables
\$ (<F) :	Automatic Variables
\$ (?D) :	Automatic Variables
\$ (?F) :	Automatic Variables
\$ (@D) :	Automatic Variables
\$ (@F) :	Automatic Variables
\$ (^D) :	Automatic Variables
\$ (^F) :	Automatic Variables
\$*:	Automatic Variables
\$*, and static pattern:	Static Usage
\$+:	Automatic Variables
\$<:	Automatic Variables
\$?:	Automatic Variables
\$@:	Automatic Variables
\$^:	Automatic Variables
\$:	Automatic Variables

%

% (automatic variable):	Automatic Variables
%D (automatic variable):	Automatic Variables
%F (automatic variable):	Automatic Variables

*

* (automatic variable):	Automatic Variables
* (automatic variable), unsupported bizarre usage:	Missing
*D (automatic variable):	Automatic Variables
*F (automatic variable):	Automatic Variables

+

+ (automatic variable):	Automatic Variables
+D (automatic variable):	Automatic Variables
+F (automatic variable):	Automatic Variables

—

`-load:` load Directive

•

. DEFAULT:	Special Targets
. DEFAULT:	Last Resort
. DEFAULT, and empty recipes:	Empty Recipes

<code>.DEFAULT_GOAL</code> (define default goal):	Special Variables
<code>.DELETE_ON_ERROR</code> :	Special Targets
<code>.DELETE_ON_ERROR</code> :	Errors
<code>.EXPORT_ALL_VARIABLES</code> :	Special Targets
<code>.EXPORT_ALL_VARIABLES</code> :	Variables/Recursion
<code>.EXTRA_PREREQS</code> (prerequisites not added to automatic variables):	Special Variables
<code>.FEATURES</code> (list of supported features):	Special Variables
<code>.IGNORE</code> :	Special Targets
<code>.IGNORE</code> :	Errors
<code>.INCLUDE_DIRS</code> (list of include directories):	Special Variables
<code>.INTERMEDIATE</code> :	Special Targets
<code>.LIBPATTERNS</code> :	Libraries/Search
<code>.LOADED</code> :	load Directive
<code>.LOW_RESOLUTION_TIME</code> :	Special Targets
<code>.NOTPARALLEL</code> :	Special Targets
<code>.ONESHELL</code> :	Special Targets
<code>.ONESHELL</code> :	One Shell
<code>.PHONY</code> :	Phony Targets
<code>.PHONY</code> :	Special Targets
<code>.POSIX</code> :	Special Targets
<code>.POSIX</code> :	Options/Recursion
<code>.PRECIOUS</code> :	Special Targets
<code>.PRECIOUS</code> :	Interrupts
<code>.RECIPEPREFIX</code> (change the recipe prefix character):	Special Variables
<code>.SECONDARY</code> :	Special Targets
<code>.SECONDEXPANSION</code> :	Secondary Expansion
<code>.SECONDEXPANSION</code> :	Special Targets
<code>.SHELLFLAGS</code> :	Choosing the Shell
<code>.SHELLFLAGS</code> :	Choosing the Shell
<code>.SHELLSTATUS</code> :	Shell Function
<code>.SILENT</code> :	Special Targets
<code>.SILENT</code> :	Echoing
<code>.SUFFIXES</code> :	Special Targets
<code>.SUFFIXES</code> :	Suffix Rules
<code>.VARIABLES</code> (list of variables):	Special Variables
<code>/</code>	
<code>/usr/gnu/include</code> :	Include
<code>/usr/include</code> :	Include
<code>/usr/local/include</code> :	Include
<code><</code>	
<code><</code> (automatic variable):	Automatic Variables
<code><D</code> (automatic variable):	Automatic Variables
<code><F</code> (automatic variable):	Automatic Variables
<code>?</code>	
<code>?</code> (automatic variable):	Automatic Variables

?D (automatic variable):	Automatic Variables
?F (automatic variable):	Automatic Variables
@	
@ (automatic variable):	Automatic Variables
@D (automatic variable):	Automatic Variables
@F (automatic variable):	Automatic Variables
^	
^ (automatic variable):	Automatic Variables
^D (automatic variable):	Automatic Variables
^F (automatic variable):	Automatic Variables
 	
 (automatic variable):	Automatic Variables
A	
abspath:	File Name Functions
addprefix:	File Name Functions
addsuffix:	File Name Functions
and:	Conditional Functions
AR:	Implicit Variables
ARFLAGS:	Implicit Variables
AS:	Implicit Variables
ASFLAGS:	Implicit Variables
B	
basename:	File Name Functions
bindir:	Directory Variables
C	
call:	Call Function
CC:	Implicit Variables
CFLAGS:	Implicit Variables
CO:	Implicit Variables
COFLAGS:	Implicit Variables
COMSPEC:	Choosing the Shell
CPP:	Implicit Variables
CPPFLAGS:	Implicit Variables
CTANGLE:	Implicit Variables
CURDIR:	Recursion
CWEAVE:	Implicit Variables
CXX:	Implicit Variables
CXXFLAGS:	Implicit Variables
D	
define:	Multi-Line

DESTDIR:	DESTDIR
dir:	File Name Functions
E	
else:	Conditional Syntax
endif:	Multi-Line
endif:	Conditional Syntax
error:	Make Control Functions
eval:	Eval Function
exec_prefix:	Directory Variables
export:	Variables/Recursion
F	
FC:	Implicit Variables
FFLAGS:	Implicit Variables
file:	File Function
filter:	Text Functions
filter-out:	Text Functions
findstring:	Text Functions
firstword:	Text Functions
flavor:	Flavor Function
foreach:	Foreach Function
G	
GET:	Implicit Variables
GFLAGS:	Implicit Variables
gmk-eval:	Guile Interface
gmk-expand:	Guile Interface
gmk_add_function:	Loaded Object API
gmk_alloc:	Loaded Object API
gmk_eval:	Loaded Object API
gmk_expand:	Loaded Object API
gmk_free:	Loaded Object API
gmk_func_ptr:	Loaded Object API
GNUMakefile:	Makefile Names
GPATH:	Search Algorithm
guile:	Guile Function
I	
if:	Conditional Functions
if:	Conditional Functions
ifdef:	Conditional Syntax
ifeq:	Conditional Syntax
ifndef:	Conditional Syntax
ifneq:	Conditional Syntax
include:	Include
info:	Make Control Functions

J	
<code>join:</code>	File Name Functions
L	
<code>lastword:</code>	Text Functions
<code>LDFLAGS:</code>	Implicit Variables
<code>LDLIBS:</code>	Implicit Variables
<code>LEX:</code>	Implicit Variables
<code>LFLAGS:</code>	Implicit Variables
<code>libexecdir:</code>	Directory Variables
<code>LINT:</code>	Implicit Variables
<code>LINTFLAGS:</code>	Implicit Variables
<code>load:</code>	load Directive
<code>LOADLIBES:</code>	Implicit Variables
M	
<code>M2C:</code>	Implicit Variables
<code>MAKE:</code>	MAKE Variable
<code>MAKE:</code>	Flavors
<code>MAKECMDGOALS:</code>	Goals
<code>Makefile:</code>	Makefile Names
<code>makefile:</code>	Makefile Names
<code>MAKEFILES:</code>	MAKEFILES Variable
<code>MAKEFILES:</code>	Variables/Recursion
<code>MAKEFILE_LIST</code> (list of parsed makefiles):	Special Variables
<code>MAKEFLAGS:</code>	Options/Recursion
<code>MAKEINFO:</code>	Implicit Variables
<code>MAKELEVEL:</code>	Variables/Recursion
<code>MAKELEVEL:</code>	Flavors
<code>MAKEOVERRIDES:</code>	Options/Recursion
<code>MAKESHELL</code> (MS-DOS alternative to <code>SHELL</code>):	Choosing the Shell
<code>MAKE_HOST:</code>	Quick Reference
<code>MAKE_RESTARTS</code> (number of times <code>make</code> has restarted):	Special Variables
<code>MAKE_TERMERR</code> (whether <code>stderr</code> is a terminal):	Special Variables
<code>MAKE_TERMOUT</code> (whether <code>stdout</code> is a terminal):	Special Variables
<code>MAKE_VERSION:</code>	Quick Reference
<code>MFLAGS:</code>	Options/Recursion
N	
<code>notdir:</code>	File Name Functions
O	
<code>or:</code>	Conditional Functions
<code>origin:</code>	Origin Function
<code>OUTPUT_OPTION:</code>	Catalogue of Rules
<code>override:</code>	Override Directive

P

<code>patsubst:</code>	Substitution Refs
<code>patsubst:</code>	Text Functions
<code>PC:</code>	Implicit Variables
<code>PFLAGS:</code>	Implicit Variables
<code>prefix:</code>	Directory Variables
<code>private:</code>	Suppressing Inheritance

R

<code>realpath:</code>	File Name Functions
<code>RFLAGS:</code>	Implicit Variables
<code>RM:</code>	Implicit Variables

S

<code>sbindir:</code>	Directory Variables
<code>SHELL:</code>	Choosing the Shell
<code>SHELL:</code>	Choosing the Shell
<code>shell:</code>	Shell Function
<code>SHELL (recipe execution):</code>	Execution
<code>sort:</code>	Text Functions
<code>strip:</code>	Text Functions
<code>subst:</code>	Multiple Targets
<code>subst:</code>	Text Functions
<code>suffix:</code>	File Name Functions
<code>SUFFIXES:</code>	Suffix Rules

T

<code>TANGLE:</code>	Implicit Variables
<code>TEX:</code>	Implicit Variables
<code>TEXI2DVI:</code>	Implicit Variables

U

<code>undefine:</code>	Undefine Directive
<code>unexport:</code>	Variables/Recursion

V

<code>value:</code>	Value Function
<code>VPATH:</code>	Directory Search
<code>VPATH:</code>	General Search
<code>vpath:</code>	Directory Search
<code>vpath:</code>	Selective Search

W

<code>warning:</code>	Make Control Functions
<code>WEAVE:</code>	Implicit Variables
<code>wildcard:</code>	Wildcard Function

wildcard:
word:
wordlist:
words:

File Name Functions
Text Functions
Text Functions
Text Functions

Y

YACC:
YFLAGS:

Implicit Variables
Implicit Variables

Jump to: \$ % * + - . / < ? @ ^ |
 A B C D E F G I J L M N O P R S T U V W Y

Footnotes

(1)

GNU Make compiled for MS-DOS and MS-Windows behaves as if prefix has been defined to be the root of the DJGPP tree hierarchy.

(2)

On MS-DOS, the value of current working directory is **global**, so changing it will affect the following recipe lines on those systems.